Package 'bigutilsr'

October 12, 2022

Title Utility Functions for Large-scale Data

Version 0.3.4

Date 2021-04-08

Description Utility functions for large-scale data. For now, package 'bigutilsr'
mainly includes functions for outlier detection and unbiased PCA projection.

License GPL-3

Encoding UTF-8

Language en-US

ByteCompile TRUE

RoxygenNote 7.1.1

URL <https://github.com/privefl/bigutilsr>

BugReports <https://github.com/privefl/bigutilsr/issues>

LinkingTo Rcpp

Imports bigassertr (>= 0.1.1), bigparallelr (>= 0.2.3), nabor, Rcpp,
robustbase, RSpectra, stats

Suggests covr, Gmedian, mvtnorm, rrcov, spelling, testthat (>= 2.1.0)

NeedsCompilation yes

Author Florian Privé [aut, cre]

Maintainer Florian Privé <florian.prive.21@gmail.com>

Repository CRAN

Date/Publication 2021-04-13 22:00:11 UTC

R topics documented:

as_model_matrix   *Transform a data frame*

Description

Transform a data frame into a matrix using one hot encoding.

Usage

```
as_model_matrix(df, intercept = FALSE)
```

Arguments

df   A data frame.

intercept   Whether to have a column with all 1s. Default is FALSE.

Value

A matrix.

Examples

```
mat <- as_model_matrix(iris)
str(mat)
```

---

covRob	*Deprecated*

Description

Deprecated

Usage

```
covRob(data, ...)
```

Arguments

data	A matrix.
...	Not used.

See Also

covrob_ogk() dist_ogk()

---

covrob_ogk	*Robust Location and Scatter Estimation - Ortogonalized Gnanadesikan-Kettenring (OGK)*

Description

Computes a robust multivariate location and scatter estimate with a high breakdown point, using the pairwise algorithm proposed by Marona and Zamar (2002) which in turn is based on the pairwise robust estimator proposed by Gnanadesikan-Kettenring (1972).

Usage

```
covrob_ogk(U, niter = 2, beta = 0.9)

dist_ogk(U, niter = 2, beta = 0.9)
```

Arguments

U	A matrix with no missing values and at least 2 columns.
niter	Number of number of iterations for the first step of the algorithm, usually 1 or 2 since iterations beyond the second do not lead to improvement.
beta	Coverage parameter for the final reweighted estimate. Default is `0.9`.

Details

The method proposed by Marona and Zamar (2002) allowes to obtain positive-definite and almost affine equivariant robust scatter matrices starting from any pairwise robust scatter matrix. The default robust estimate of covariance between two random vectors used is the one proposed by Gnanadesikan and Kettenring (1972) but the user can choose any other method by redefining the function in slot `vrob` of the control object `CovControlOgk`. Similarly, the function for computing the robust univariate location and dispersion used is the `tau scale` defined in Yohai and Zamar (1998) but it can be redefined in the control object.

The estimates obtained by the OGK method, similarly as in `CovMcd` are returned as 'raw' estimates. To improve the estimates a reweighting step is performed using the coverage parameter `beta` and these reweighted estimates are returned as 'final' estimates.

Value

`covrob_ogk()`: list of robust estimates, `$cov` and `$center`.

`dist_ogk()`: vector of robust Mahalanobis (squared) distances.

References

Maronna, R.A. and Zamar, R.H. (2002) Robust estimates of location and dispersion of high-dimensional datasets; *Technometrics* **44**(4), 307–317.

Yohai, R.A. and Zamar, R.H. (1998) High breakdown point estimates of regression by means of the minimization of efficient scale *JASA* **86**, 403–413.

Gnanadesikan, R. and John R. Kettenring (1972) Robust estimates, residuals, and outlier detection with multiresponse data. *Biometrics* **28**, 81–124.

Todorov V & Filzmoser P (2009), An Object Oriented Framework for Robust Multivariate Analysis. *Journal of Statistical Software*, **32**(3), 1–47. URL https://www.jstatsoft.org/v32/i03/.

See Also

rrcov::CovOgk()

stats::mahalanobis()

Examples

```
X <- readRDS(system.file("testdata", "three-pops.rds", package = "bigutilsr"))
svd <- svds(scale(X), k = 5)

U <- svd$u
dist <- dist_ogk(U)
str(dist)
```

---

geometric_median   *Geometric median*

Description

Compute the geometric median, i.e. the point that minimizes the sum of all Euclidean distances to the observations (rows of U).

Usage

```
geometric_median(U, tol = 1e-10, maxiter = 1000, by_grp = NULL)
```

Arguments

U	A matrix (e.g. PC scores).
tol	Convergence criterion. Default is 1e-10.
maxiter	Maximum number of iterations. Default is 1000.
by_grp	Possibly a vector for splitting rows of U into groups before computing the geometric mean for each group. Default is NULL (ignored).

Value

The geometric median of all rows of U, a vector of the same size as ncol(U). If providing by_grp, then a matrix with rows being the geometric median within each group.

Examples

```
X <- readRDS(system.file("testdata", "three-pops.rds", package = "bigutilsr"))
pop <- rep(1:3, c(143, 167, 207))

svd <- svds(scale(X), k = 5)
U <- sweep(svd$u, 2, svd$d, '*')
plot(U, col = pop, pch = 20)

med_all <- geometric_median(U)
points(t(med_all), pch = 20, col = "blue", cex = 4)

med_pop <- geometric_median(U, by_grp = pop)
points(med_pop, pch = 20, col = "blue", cex = 2)
```

---

hist_out	*Outlier detection (histogram)*

Description

Outlier detection based on departure from histogram. Suitable for compact values (need a space between main values and outliers).

Usage

```
hist_out(x, breaks = nclass.scottRob, pmax_out = 0.2, nboot = NULL)
```

Arguments

x	Numeric vector (with compact values).
breaks	Same parameter as for hist(). Default uses a robust version of Scott's rule. You can also use "FD" or nclass.FD for a bit more bins.
pmax_out	Percentage at each side that can be considered outliers at each step. Default is 0.2.
nboot	Number of bootstrap replicates to estimate limits more robustly. Default is NULL (no bootstrap, even if **I would recommend to use it**).

Value

A list with

- x: the initial vector, whose outliers have been removed,
- lim: lower and upper limits for outlier removal,
- all_lim: all bootstrap replicates for lim (if nboot not NULL).

Examples

```
set.seed(1)
x <- rnorm(1000)
str(hist_out(x))

# Easy to separate
x2 <- c(x, rnorm(50, mean = 7))
hist(x2, breaks = nclass.scottRob)
str(hist_out(x2))

# More difficult to separate
x3 <- c(x, rnorm(50, mean = 6))
hist(x3, breaks = nclass.scottRob)
str(hist_out(x3))
str(hist_out(x3, nboot = 999))
```

---

knn_parallel   *Find K nearest neighbours for multiple query points*

Description

Find K nearest neighbours for multiple query points

Usage

```
knn_parallel(data, query = data, k, ..., ncores = bigparallelr::nb_cores())
```

Arguments

data	Mxd matrix of M target points with dimension d
query	Nxd matrix of N query points with dimension d (nb data and query must have same dimension). If missing defaults to data i.e. a self-query.
k	an integer number of nearest neighbours to find
...	Arguments passed on to nabor::knn
	eps An approximate error bound. The default of 0 implies exact matching.
	searchtype A character vector or integer indicating the search type. The default value of 1L is equivalent to "auto". See details.
	radius Maximum radius search bound. The default of 0 implies no radius bound.
ncores	Number of cores to use. Default uses bigparallelr::nb_cores().

Value

A list with elements nn.idx (1-indexed indices) and nn.dists (distances), both of which are N x k matrices. See details for the results obtained with1 invalid inputs.

Examples

```
## Not run: knn_parallel(matrix(1:4, 2), k = 2, ncores = 2)
```

LOF   *Local Outlier Factor (LOF)*

Description

LOF: Identifying Density-Based Local Outliers.

Usage

```
LOF(
  U,
  seq_k = c(4, 10, 30),
  combine = max,
  robMaha = FALSE,
  log = TRUE,
  ncores = 1
)
```

Arguments

U	A matrix, from which to detect outliers (rows). E.g. PC scores.
seq_k	Sequence of numbers of nearest neighbors to use. If multiple k are provided, this returns the combination of statistics. Default is c(4, 10, 30) and use max to combine (see combine).
combine	How to combine results for multiple k? Default uses max.
robMaha	Whether to use a robust Mahalanobis distance instead of the normal euclidean distance? Default is FALSE, meaning using euclidean.
log	Whether to return the logarithm of LOFs? Default is TRUE.
ncores	Number of cores to use. Default is 1.

References

Breunig, Markus M., et al. "LOF: identifying density-based local outliers." ACM sigmod record. Vol. 29. No. 2. ACM, 2000.

See Also

prob_dist()

Examples

```
X <- readRDS(system.file("testdata", "three-pops.rds", package = "bigutilsr"))
svd <- svds(scale(X), k = 10)

llof <- LOF(svd$u)
hist(llof, breaks = nclass.scottRob)
```

```
tukey_mc_up(llof)

llof_maha <- LOF(svd$u, robMaha = TRUE)
hist(llof_maha, breaks = nclass.scottRob)
tukey_mc_up(llof_maha)

lof <- LOF(svd$u, log = FALSE)
hist(lof, breaks = nclass.scottRob)
str(hist_out(lof))
str(hist_out(lof, nboot = 100))
str(hist_out(lof, nboot = 100, breaks = "FD"))
```

---

maha_trans   *Transform matrix*

Description

Transform matrix to use Mahalanobis distance instead of Euclidean one.

Usage

```
maha_trans(U, estim = covrob_ogk(U))
```

Arguments

U	A matrix (e.g. PC scores).
estim	List of location and scatter estimates, $cov and $center.

Value

U, transformed.

Examples

```
X <- readRDS(system.file("testdata", "three-pops.rds", package = "bigutilsr"))
svd <- svds(scale(X), k = 5)

U <- svd$u
dist1 <- dist_ogk(U)

U.maha <- maha_trans(U)
dist2 <- rowSums(U.maha^2)
all.equal(dist2, dist1)
```

---

nclass.scottRob   *Compute the Number of Classes for a Histogram*

Description

Compute the Number of Classes for a Histogram

Usage

```
nclass.scottRob(x)
```

Arguments

x   a data vector.

Value

The suggested number of classes.

References

Scott, D. W. (1979). On optimal and data-based histograms. Biometrika, 66, 605–610. doi:
10.2307/2335182.

Examples

```
x <- rnorm(1000)
hist(x, breaks = nclass.scott)
hist(x, breaks = nclass.scottRob)

x2 <- c(x, rnorm(50, mean = 50))
hist(x2, breaks = nclass.scott)
hist(x2, breaks = nclass.scott,    xlim = c(-5, 5))
hist(x2, breaks = nclass.scottRob, xlim = c(-5, 5))
```

---

pca_nspike   *Number of spikes in PCA*

Description

Estimate the number of distant spikes based on the histogram of eigenvalues.

Usage

```
pca_nspike(eigval, breaks = "FD", nboot = 100)
```

Arguments

`eigval`	Eigenvalues (squared singular values).
`breaks`	Same parameter as for `hist()`. Default uses a robust version of Scott's rule. You can also use `"FD"` or `nclass.FD` for a bit more bins.
`nboot`	Number of bootstrap replicates to estimate limits more robustly. Default is `100`.

Value

The estimated number of distant spikes.

Examples

```
N <- 400; M <- 2000; K <- 8
U <- matrix(0, N, K); U[] <- rnorm(length(U))
V <- matrix(0, M, K); V[] <- rnorm(length(V))
# X = U V^T + E
X <- tcrossprod(U, V) + 15 * rnorm(N * M)
pca <- prcomp(X)
eigval <- pca$sdev^2
plot(head(eigval, -1), log = "xy", pch = 20)
pca_nspike(eigval)
```

`pca_OADP_proj`	*OADP projection*

Description

Online Augmentation, Decomposition, and Procrustes (OADP) projection of PC loadings onto some study data `X`.

Usage

```
pca_OADP_proj(X, loadings, sval)

pca_OADP_proj2(XV, X_norm, sval)
```

Arguments

`X`	Data to get PC loadings into.
`loadings`	PC loadings of the reference PCA to project.
`sval`	Singular values of the reference PCA (sqrt of the eigen values). Only the `ncol(loadings)` first ones will be used.
`XV`	`X %*% loadings`
`X_norm`	Vector of sums of squared rows (e.g. `rowSums(X^2)`).

Value

- pca_OADP_proj(): A list with the simple projection X %*% loadings and the projection based on OADP.

- pca_OADP_proj2(): The projection based on OADP only (a matrix of same size of XV).

Examples

```
X <- readRDS(system.file("testdata", "three-pops.rds", package = "bigutilsr"))
N <- 400; M <- ncol(X)
ind <- sample(nrow(X), N)

# Compute SVD using one part of samples
svd <- svds(X[ind, ], k = 5)
U <- sweep(svd$u, 2, svd$d, '*')
col <- 2:3
plot(U[, col])
points(cbind(0, 0), pch = 8, col = "green", cex = 2)

# Projecting other samples
proj <- pca_OADP_proj(X = X[-ind, ], loadings = svd$v, sval = svd$d)
points(proj$simple_proj[, col], col = "red", pch = 20)     # shrunk towards 0
points(proj$OADP_proj[, col], col = "blue", pch = 20)      # unshrunk
```

---

predict.Procrustes   *Predict method*

Description

Predict method for class Procrustes.

Usage

```
## S3 method for class 'Procrustes'
predict(object, X, ...)
```

Arguments

object	Object of class Procrustes.
X	New matrix to transform.
...	Not used.

Value

X, transformed.

See Also

procrustes().

---

prob_dist   *Probabilistic set distance*

Description

Probabilistic set distance

Usage

```
prob_dist(U, kNN = 5, robMaha = FALSE, ncores = 1)
```

Arguments

U	A matrix, from which to detect outliers (rows). E.g. PC scores.
kNN	Number of nearest neighbors to use. Default is 5.
robMaha	Whether to use a robust Mahalanobis distance instead of the normal euclidean distance? Default is FALSE, meaning using euclidean.
ncores	Number of cores to use. Default is 1.

References

Kriegel, Hans-Peter, et al. "LoOP: local outlier probabilities." Proceedings of the 18th ACM conference on Information and knowledge management. ACM, 2009.

See Also

[LOF()](#)

Examples

```
X <- readRDS(system.file("testdata", "three-pops.rds", package = "bigutilsr"))
svd <- svds(scale(X), k = 10)
U <- svd$u

test <- prob_dist(U)
plof <- test$dist.self / test$dist.nn
plof_ish <- test$dist.self / sqrt(test$dist.nn)
plot(U[, 1:2], col = (plof_ish > tukey_mc_up(plof_ish)) + 1, pch = 20)
plot(U[, 3:4], col = (plof_ish > tukey_mc_up(plof_ish)) + 1, pch = 20)
plot(U[, 5:6], col = (plof_ish > tukey_mc_up(plof_ish)) + 1, pch = 20)
```

---

procrustes   *Procrustes transform*

Description

Procrustes transform Y = pXR (after centering), where p is a scaling coefficient and R is a rotation matrix that minimize ‖Y - pXR‖_F.

Usage

```
procrustes(Y, X, n_iter_max = 1000, epsilon_min = 1e-07)
```

Arguments

Y	Reference matrix.
X	Matrix to transform (ncol(X) >= ncol(Y)).
n_iter_max	Maximum number of iterations. Default is 1000.
epsilon_min	Convergence criterion. Default is 1e-7.

Value

Object of class "procrustes", a list with the following elements:

- $R: the rotation matrix to apply to X,
- $rho: the scaling coefficient to apply to X,
- $c: the column centering to apply to the resulting matrix,
- $diff: the average difference between Y and X transformed.

You can use method predict() to apply this transformation to other data.

Examples

```
A <- matrix(rnorm(200), ncol = 20)
B <- matrix(rnorm(length(A)), nrow = nrow(A))

proc <- procrustes(B, A)
str(proc)
plot(B, predict(proc, A)); abline(0, 1, col = "red")
```

---

rollmean *Gaussian smoothing*

Description

Gaussian smoothing

Usage

```
rollmean(x, size)
```

Arguments

x   Numeric vector.

size   Radius of the smoothing (smaller than half of the length of x). If using `size = 0`, it returns x.

Value

Numeric vector of the same length as x, smoothed.

Examples

```
(x <- rnorm(10))
rollmean(x, 3)
```

---

tukey_mc_up *Outlier detection threshold (upper)*

Description

Outlier detection threshold (upper) based on Tukey's rule, corrected for skewness using the 'med-couple', and possibly corrected for multiple testing.

Usage

```
tukey_mc_up(x, coef = NULL, alpha = 0.05, a = -4, b = 3)
```

Arguments

x	Numeric vector. Should be somewhat normally distributed.
coef	number determining how far 'whiskers' extend out from the box. If NULL (default), this is computed to get an type-I error of alpha, after adjusting for multiple testing. A standard value to use is 1.5.
alpha	See coef. Default is 0.05.
a	scaling factors multiplied by the medcouple [mc]() to determine outlyer boundaries; see the references.
b	scaling factors multiplied by the medcouple [mc]() to determine outlyer boundaries; see the references.

References

Hubert, M. and Vandervieren, E. (2008). An adjusted boxplot for skewed distributions, *Computational Statistics and Data Analysis* **52**, 5186–5201.

See Also

robustbase::adjbox()

Examples

```
hist(x <- c(rnorm(3, m = 6), rnorm(1e4, m = 0)))
(q <- tukey_mc_up(x))
abline(v = q, col = "red")
which(x > q)
```

Index