Package 'extremeStat'

January 13, 2024

Type Package

Title Extreme Value Statistics and Quantile Estimation

Version 1.5.9

Date 2024-01-12

Depends R (>= 2.10)

Imports lmomco (>= 2.2.5), berryFunctions (>= 1.15.6), pbapply,
RColorBrewer, grDevices, graphics, methods, stats, utils, evir,
ismev, fExtremes, extRemes, evd, Renext

Author Berry Boessenkool

Maintainer Berry Boessenkool <berry-b@gmx.de>

Description Fit, plot and compare several (extreme value) distribution functions.
Compute (truncated) distribution quantile estimates and plot return periods on a linear scale.
On the fitting method, see Asquith (2011): Distributional Analysis with L-moment Statistics [...] ISBN 1463508417.

License GPL (>= 2)

URL https://github.com/brry/extremeStat

RoxygenNote 7.3.0

Encoding UTF-8

Suggests testthat, knitr, rmarkdown

VignetteBuilder knitr

BugReports https://github.com/brry/extremeStat

NeedsCompilation no

Repository CRAN

Date/Publication 2024-01-13 13:00:10 UTC

R topics documented:

---

annMax   *annual discharge maxima (streamflow)*

Description

Annual discharge maxima of a stream in Austria called Griesler or Fuschler Ache, at the measurement station (gauge) near St. Lorenz, catchment area ca 100 km^2. Extracted from the time series 1976-2010 with a resolution of 15 Minutes.

Format

num [1:35] 61.5 77 37 69.3 75.6 74.9 43.7 50.8 55.6 84.1 ...

Source

Hydrographische Dienste Oberoesterreich und Salzburg, analyzed by package author (<berry-b@gmx.de>)

Examples

```
data(annMax)
str(annMax)
str(annMax)
plot(1976:2010, annMax, type="l", las=1, main="annMax dataset from Austria")
# Moving Average with different window widths:
berryFunctions::movAvLines(annMax, x=1976:2010, lwd=3, alpha=0.7)
```

---

distLexBoot	*Bootstrapping uncertainty intervals for return periods*

Description

Calculates and plots bootstrap uncertainty intervals for `plotLextreme`.

Usage

```
distLexBoot(
  dlf,
  nbest = 3,
  selection = NULL,
  n = 100,
  prop = 0.8,
  conf.lev = 0.95,
  replace = FALSE,
  RPs = NULL,
  log = TRUE,
  progbars = TRUE,
  quiet = FALSE
)
```

Arguments

dlf	dlf object, as returned by `distLextreme`
nbest	Number of best fitted distribution functions in dlf for which bootstrapping is to be done. Overridden by `selection`. DEFAULT: 3
selection	Character vector with distribution function names to be used. Suggested to keep this low. DEFAULT: NULL
n	Number of subsamples to be processed (computing time increases extraordinarily). DEFAULT: 100
prop	Proportion of sample to be used in each run. DEFAULT: 0.8
conf.lev	Confidence level (Proportion of subsamples within 'confidence interval'). Quantiles extracted from this value are passed to `quantileMean`. DEFAULT: 0.95
replace	Logical: replace in each `sample`? DEFAULT: FALSE
RPs	Return Period vector, by default calculated internally based on value of `log`. DEFAULT: NULL
log	RPs suitable for plot on a logarithmic axis? DEFAULT: TRUE
progbars	Show progress bar for Monte Carlo simulation? DEFAULT: TRUE
quiet	Logical: suppress messages? See `distLquantile`. DEFAULT: FALSE

Details

Has not been thoroughly tested yet. Bootstrapping defaults can probably be improved.

Value

invisible dlf object, see [printL](#). Additional elements are: exBootCL (confidence level), exBootRPs (x values for plot) exBootSim (all simulation results) and exBootCI (aggregated into CI band). The last two are each a list with a matrix (return levels)

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Sept 2015 + Dec 2016

See Also

[plotLexBoot](#), [distLextreme](#)

Examples

```
data(annMax)
dlf <- distLextreme(annMax, selection=c("gum","gev","wak","nor"))
dlfB <- distLexBoot(dlf, nbest=4, conf.lev=0.5, n=10) # n low for quick example tests
plotLexBoot(dlfB)
plotLexBoot(dlfB, selection=c("nor","gev"))
plotLexBoot(dlfB, selection=c("gum","gev","wak","nor"), order=FALSE)
```

---

distLextreme   *Extreme value stats*

Description

Extreme value statistics for flood risk estimation. Input: vector with annual discharge maxima (or all observations for POT approach). Output: discharge estimates for given return periods, parameters of several distributions (fit based on L-moments), quality of fits, plot with linear/logarithmic axis. (plotting positions by Weibull and Gringorton).

Usage

```
distLextreme(
  dat = NULL,
  dlf = NULL,
  RPs = c(2, 5, 10, 20, 50),
  npy = 1,
  truncate = 0,
  quiet = FALSE,
  ...
)
```

Arguments

dat	Vector with *either* (for Block Maxima Approach) extreme values like annual discharge maxima *or* (for Peak Over Threshold approach) all values in time-series. Ignored if dlf is given. DEFAULT: NULL
dlf	List as returned by [distLfit](). See also [distLquantile](). Overrides dat! DEFAULT: NULL
RPs	Return Periods (in years) for which discharge is estimated. DEFAULT: c(2,5,10,20,50)
npy	Number of observations per year. Leave npy=1 if you use annual block maxima (and leave truncate at 0). If you use a POT approach (see [vignette]() and examples below) e.g. on daily data, use npy=365.24. DEFAULT: 1
truncate	Truncated proportion to determine POT threshold, see [distLquantile](). DEFAULT: 0
quiet	Suppress notes and progbars? DEFAULT: FALSE
...	Further arguments passed to [distLquantile]() like truncate, selection, time, progbars

Details

[plotLextreme]() adds weibull and gringorton plotting positions to the distribution lines, which are estimated from the L-moments of the data itself.
I personally believe that if you have, say, 35 values in dat, the highest return period should be around 36 years (Weibull) and not 60 (Gringorton).
The plotting positions don't affect the distribution parameter estimation, so this dispute is not really important. But if you care, go ahead and google "weibull vs gringorton plotting positions".

Plotting positions are not used for fitting distributions, but for plotting only. The ranks of ascendingly sorted extreme values are used to compute the probability of non-exceedance Pn:
Pn_w <- Rank /(n+1) # Weibull
Pn_g <- (Rank-0.44)/(n+0.12) # Gringorton (taken from lmom:::evplot.default)
Finally: RP = Return period = recurrence interval = 1/P_exceedance = 1/(1-P_nonexc.), thus:
RPweibull = 1/(1-Pn_w) and analogous for gringorton.

Value

invisible dlf object, see [printL](). The added element is returnlev, a data.frame with the return level (discharge) for all given RPs and for each distribution. Note that this differs from [distLquantile]() (matrix output, not data.frame)

Note

This function replaces berryFunctions::extremeStatLmom

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, 2012 (first draft) - 2014 & 2015 (main updates)

References

https://RclickHandbuch.wordpress.com Chapter 15 (German)
Christoph Mudersbach: Untersuchungen zur Ermittlung von hydrologischen Bemessungsgroessen mit Verfahren der instationaeren Extremwertstatistik

See Also

distLfit. distLexBoot for confidence interval from Bootstrapping. fevd in the package extRemes.

Examples

```
# Basic examples
# BM vs POT
# Plotting options
# weighted mean based on Goodness of fit (GOF)
# Effect of data proportion used to estimate GOF
# compare extremeStat with other packages

library(lmomco)
library(berryFunctions)

data(annMax) # annual streamflow maxima in river in Austria

# Basic examples ---------------------------------------------------------
dlf <- distLextreme(annMax)
plotLextreme(dlf, log=TRUE)
plotLextreme(dlf, log="xy")
plotLextreme(dlf)

# Object structure:
str(dlf, max.lev=2)
printL(dlf)

# discharge levels for default return periods:
dlf$returnlev

# Estimate discharge that could occur every 80 years (at least empirically):
Q80 <- distLextreme(dlf=dlf, RPs=80)$returnlev
round(sort(Q80[1:17,1]),1)
# 99 to 143 m^3/s can make a relevant difference in engineering!
# That's why the rows weighted by GOF are helpful. Weights are given as in
plotLweights(dlf) # See also section weighted mean below
# For confidence intervals see ?distLexBoot

# Return period of a given discharge value, say 120 m^3/s:
round0(sort(1/(1-sapply(dlf$parameter, plmomco, x=120) )  ),1)
# exponential:                   every 29 years
# gev (general extreme value dist):  59,
# Weibull:                       every 73 years only
```

```
# BM vs POT --------------------------------------------------------------
# Return levels by Block Maxima approach vs Peak Over Threshold approach:
# BM distribution theoretically converges to GEV, POT to GPD

data(rain, package="ismev")
days <- seq(as.Date("1914-01-01"), as.Date("1961-12-30"), by="days")
BM <- tapply(rain, format(days,"%Y"), max)  ;  rm(days)
dlfBM <- plotLextreme(distLextreme(BM, emp=FALSE), ylim=lim0(100), log=TRUE, nbest=10)
plotLexBoot(distLexBoot(dlfBM, quiet=TRUE), ylim=lim0(100))
plotLextreme(dlfBM, log=TRUE, ylim=lim0(100))

dlfPOT99 <- distLextreme(rain, npy=365.24, trunc=0.99, emp=FALSE)
dlfPOT99 <- plotLextreme(dlfPOT99, ylim=lim0(100), log=TRUE, nbest=10, main="POT 99")
printL(dlfPOT99)

# using only nonzero values (normally yields better fits, but not here)
rainnz <- rain[rain>0]
dlfPOT99nz <- distLextreme(rainnz, npy=length(rainnz)/48, trunc=0.99, emp=FALSE)
dlfPOT99nz <- plotLextreme(dlfPOT99nz, ylim=lim0(100), log=TRUE, nbest=10,
                           main=paste("POT 99 x>0, npy =", round(dlfPOT99nz$npy,2)))

## Not run:  ## Excluded from CRAN R CMD check because of computing time

dlfPOT99boot <- distLexBoot(dlfPOT99, prop=0.4)
printL(dlfPOT99boot)
plotLexBoot(dlfPOT99boot)


dlfPOT90 <- distLextreme(rain, npy=365.24, trunc=0.90, emp=FALSE)
dlfPOT90 <- plotLextreme(dlfPOT90, ylim=lim0(100), log=TRUE, nbest=10, main="POT 90")

dlfPOT50 <- distLextreme(rain, npy=365.24, trunc=0.50, emp=FALSE)
dlfPOT50 <- plotLextreme(dlfPOT50, ylim=lim0(100), log=TRUE, nbest=10, main="POT 50")

## End(Not run)

ig99 <- ismev::gpd.fit(rain, dlfPOT99$threshold)
ismev::gpd.diag(ig99); title(main=paste(99, ig99$threshold))
## Not run:
ig90 <- ismev::gpd.fit(rain, dlfPOT90$threshold)
ismev::gpd.diag(ig90); title(main=paste(90, ig90$threshold))
ig50 <- ismev::gpd.fit(rain, dlfPOT50$threshold)
ismev::gpd.diag(ig50); title(main=paste(50, ig50$threshold))

## End(Not run)


# Plotting options -------------------------------------------------------
plotLextreme(dlf=dlf)
# Line colors / select distributions to be plotted:
plotLextreme(dlf, nbest=17, distcols=heat.colors(17), lty=1:5) # lty is recycled
plotLextreme(dlf, selection=c("gev", "gam", "gum"), distcols=4:6, PPcol=3, lty=3:2)
plotLextreme(dlf, selection=c("gpa","glo","wei","exp"), pch=c(NA,NA,6,8),
```

```
                    order=TRUE, cex=c(1,0.6, 1,1), log=TRUE, PPpch=c(16,NA), n_pch=20)
# use n_pch to say how many points are drawn per line (important for linear axis)

plotLextreme(dlf, legarg=list(cex=0.5, x="bottom", box.col="red", col=3))
# col in legarg list is (correctly) ignored
## Not run:
## Excluded from package R CMD check because it's time consuming

plotLextreme(dlf, PPpch=c(1,NA)) # only Weibull plotting positions
# add different dataset to existing plot:
distLextreme(Nile/15, add=TRUE, PPpch=NA, distcols=1, selection="wak", legend=FALSE)

# Logarithmic axis
plotLextreme(distLextreme(Nile), log=TRUE, nbest=8)




# weighted mean based on Goodness of fit (GOF) --------------------------------
# Add discharge weighted average estimate continuously:
plotLextreme(dlf, nbest=17, legend=FALSE)
abline(h=115.6, v=50)
RP <- seq(1, 70, len=100)
DischargeEstimate <- distLextreme(dlf=dlf, RPs=RP, plot=FALSE)$returnlev
lines(RP, DischargeEstimate["weighted2",], lwd=3, col="orange")

# Or, on log scale:
plotLextreme(dlf, nbest=17, legend=FALSE, log=TRUE)
abline(h=115.9, v=50)
RP <- unique(round(logSpaced(min=1, max=70, n=200, plot=FALSE),2))
DischargeEstimate <- distLextreme(dlf=dlf, RPs=RP)$returnlev
lines(RP, DischargeEstimate["weighted2",], lwd=5)


# Minima ----------------------------------------------------------------------

browseURL("https://nrfa.ceh.ac.uk/data/station/meanflow/39072")
qfile <- system.file("extdata/discharge39072.csv", package="berryFunctions")
Q <- read.table(qfile, skip=19, header=TRUE, sep=",", fill=TRUE)[,1:2]
rm(qfile)
colnames(Q) <- c("date","discharge")
Q$date <- as.Date(Q$date)
plot(Q, type="l")
Qmax <- tapply(Q$discharge, format(Q$date,"%Y"), max)
plotLextreme(distLextreme(Qmax, quiet=TRUE))
Qmin <- tapply(Q$discharge, format(Q$date,"%Y"), min)
dlf <- distLextreme(-Qmin, quiet=TRUE, RPs=c(2,5,10,20,50,100,200,500))
plotLextreme(dlf, ylim=c(0,-31), yaxs="i", yaxt="n", ylab="Q annual minimum", nbest=14)
axis(2, -(0:3*10), 0:3*10, las=1)
-dlf$returnlev[c(1:14,21), ]
# Some distribution functions are an obvious bad choice for this, so I use
# weighted 3: Values weighted by GOF of dist only for the best half.
# For the Thames in Windsor, we will likely always have > 9 m^3/s streamflow
```

```
# compare extremeStat with other packages: --------------------------------------
library(extRemes)
plot(fevd(annMax))
par(mfrow=c(1,1))
return.level(fevd(annMax, type="GEV")) # "GP", "PP", "Gumbel", "Exponential"
distLextreme(dlf=dlf, RPs=c(2,20,100))$returnlev["gev",]
# differences are small, but noticeable...
# if you have time for a more thorough control, please pass me the results!


# yet another dataset for testing purposes:
Dresden_AnnualMax <- c(403, 468, 497, 539, 542, 634, 662, 765, 834, 847, 851, 873,
885, 983, 996, 1020, 1028, 1090, 1096, 1110, 1173, 1180, 1180,
1220, 1270, 1285, 1329, 1360, 1360, 1387, 1401, 1410, 1410, 1456,
1556, 1580, 1610, 1630, 1680, 1734, 1740, 1748, 1780, 1800, 1820,
1896, 1962, 2000, 2010, 2238, 2270, 2860, 4500)
plotLextreme(distLextreme(Dresden_AnnualMax))

## End(Not run) # end dontrun
```

---

distLfit	*Fit distributions via L-moments*

Description

Fit several distributions via L-moments with lmomco::`lmom2par` and compute goodness of fit measures.

Usage

```
distLfit(
  dat,
  datname = deparse(substitute(dat)),
  selection = NULL,
  speed = TRUE,
  ks = FALSE,
  truncate = 0,
  threshold = berryFunctions::quantileMean(dat, truncate),
  progbars = length(dat) > 200,
  time = TRUE,
  quiet = FALSE,
  ssquiet = quiet,
  ...
)
```

Arguments

dat	Vector with values
datname	Character string for main, xlab etc. DEFAULT: deparse(substitute(dat))
selection	Selection of distributions. Character vector with types as in `lmom2par`. Overrides speed. DEFAULT: NULL
speed	If TRUE, several distributions are omitted, for the reasons shown in lmomco::`dist.list`(). DEFAULT: TRUE
ks	Include ks.test results and CDF R^2 in dlf$gof? Computing is much faster when FALSE. DEFAULT: FALSE
truncate	Number between 0 and 1. POT Censored `distLquantile`: fit to highest values only (truncate lower proportion of x). Probabilities are adjusted accordingly. DEFAULT: 0
threshold	POT cutoff value. If you want correct percentiles, set this only via truncate, see Details of `q_gpd`. DEFAULT: `quantileMean`(x, truncate)
progbars	Show progress bars for each loop? DEFAULT: TRUE if n > 200
time	`message` execution time? DEFAULT: TRUE
quiet	Suppress notes? DEFAULT: FALSE
ssquiet	Suppress sample size notes? DEFAULT: quiet
...	Further arguments passed to `distLweights` like weightc, order=FALSE

Value

invisible dlf object, see `printL`.

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Sept 2014, July 2015, Dec 2016

See Also

`plotLfit`, `distLweights`, `plotLweights`, extRemes::`fevd`, MASS::`fitdistr`.
More complex estimates of quality of fits: Fard, M.N.P. and Holmquist, B. (2013, Chilean Journal of
Statistics): Powerful goodness-of-fit tests for the extreme value distribution. https://chjs.mat.utfsm.cl/volumes/04/01/Fard_Hc

Examples

```
data(annMax)
# basic usage on real data (annual discharge maxima in Austria)
dlf <- distLfit(annMax)
str(dlf, max.lev=2)
printL(dlf)
plotLfit(dlf)

# arguments that can be passed to plotting function:
plotLfit(dlf, lty=2, col=3, nbest=17, legargs=list(lwd=3), main="booh!")
```

```
set.seed(42)
dlf_b <- distLfit(rbeta(100, 5, 2))
plotLfit(dlf_b, nbest=10, legargs=c(x="left"))
plotLfit(dlf_b, selection=c("gpa", "glo", "gev", "wak"))
plotLfit(dlf_b, selection=c("gpa", "glo", "gev", "wak"), order=TRUE)
plotLfit(dlf_b, distcols=c("orange",3:6), lty=1:3) # lty is recycled
plotLfit(dlf_b, cdf=TRUE)
plotLfit(dlf_b, cdf=TRUE, histargs=list(do.points=FALSE), sel="nor")


# logarithmic axes:
set.seed(1)
y <- 10^rnorm(300, mean=2, sd=0.3) # if you use 1e4, distLfit will be much slower
hist(y, breaks=20)
berryFunctions::logHist(y, col=8)
dlf <- distLfit(log10(y))
plotLfit(dlf, breaks=50)
plotLfit(dlf, breaks=50, log=TRUE)


# Goodness of fit: how well do the distributions fit the original data?
# measured by RMSE of cumulated distribution function and ?ecdf
# RMSE: root of average of ( errors squared )  ,   errors = line distances
dlf <- distLfit(annMax, ks=TRUE)
plotLfit(dlf, cdf=TRUE, sel=c("wak", "revgum"))
x <- sort(annMax)
segments(x0=x, y0=lmomco::plmomco(x, dlf$parameter$revgum), y1=ecdf(annMax)(x), col=2)
segments(x0=x, y0=lmomco::plmomco(x, dlf$parameter$wak), y1=ecdf(annMax)(x), col=4, lwd=2)
# weights by three different weighting schemes, see distLweights:
plotLweights(dlf)
plotLfit(distLfit(annMax               ), cdf=TRUE, nbest=17)$gof
plotLfit(distLfit(annMax, truncate=0.7), cdf=TRUE, nbest=17)$gof
pairs(dlf$gof[,-(2:5)]) # measures of goodness of fit are correlated quite well here.
dlf$gof

# Kolmogorov-Smirnov Tests for normal distribution return slightly different values:
library(lmomco)
ks.test(annMax, "pnorm", mean(annMax), sd(annMax) )$p.value
ks.test(annMax, "cdfnor", parnor(lmoms(annMax)))$p.value


# Fit all available distributions (30):
## Not run: # this takes a while...
d_all <- distLfit(annMax, speed=FALSE, progbars=TRUE) # 20 sec
printL(d_all)
plotLfit(d_all, nbest=30, distcols=grey(1:22/29), xlim=c(20,140))
plotLfit(d_all, nbest=30, ylim=c(0,0.04), xlim=c(20,140))
plotLweights(d_all)
d_all$gof

## End(Not run)
```

---

distLquantile   *distribution quantiles*

Description

Parametric quantiles of distributions fitted to a sample.

Usage

```
distLquantile(
  x = NULL,
  probs = c(0.8, 0.9, 0.99),
  truncate = 0,
  threshold = quantileMean(dlf$dat_full[is.finite(dlf$dat_full)], truncate),
  sanerange = NA,
  sanevals = NA,
  selection = NULL,
  order = TRUE,
  dlf = NULL,
  datname = deparse(substitute(x)),
  list = FALSE,
  empirical = TRUE,
  qemp.type = 8,
  weighted = empirical,
  gpd = empirical,
  speed = TRUE,
  quiet = FALSE,
  ssquiet = quiet,
  ttquiet = quiet,
  gpquiet = missing(quiet) | quiet,
  ...
)
```

Arguments

x	Sample for which parametric quantiles are to be calculated. If it is NULL (the default), dat from dlf is used. DEFAULT: NULL
probs	Numeric vector of probabilities with values in [0,1]. DEFAULT: c(0.8,0.9,0.99)
truncate	Number between 0 and 1 (proportion of sample discarded). Censored quantile: fit to highest values only (truncate lower proportion of x). Probabilities are adjusted accordingly. DEFAULT: 0
threshold	POT cutoff value. If you want correct percentiles, set this only via truncate, see Details of q_gpd. DEFAULT: quantileMean(x, truncate)
sanerange	Range outside of which results should be changed to sanevals. This can capture numerical errors in small samples (notably GPD_MLE_extRemes). If NA, this is ignored. Attention: the RMSE column is also checked and changed. DEFAULT: NA

sanevals	Values to be used below [1] and above [2] sanerange. DEFAULT: NA
selection	Distribution type, eg. "gev" or "wak", see lmomco::dist.list. Can be a vector. If NULL (the default), all types present in dlf$distnames are used. DEFAULT: NULL
order	Logical: sort by RMSE, even if selection is given? See distLweights. DEFAULT: TRUE
dlf	dlf object described in extremeStat. Use this to save computing time for large datasets where you already have dlf. DEFAULT: NULL
datname	Character string: data name, important if list=TRUE. DEFAULT: deparse(substitute(x))
list	Return full dlflist with output attached as element quant? If FALSE (the default), just the matrix with quantile estimates is returned. DEFAULT: FALSE
empirical	Add rows "empirical" and "quantileMean" in the output matrix? Uses quantile with qemp.type (ignoring truncation) and quantileMean. DEFAULT: TRUE
qemp.type	Method passed to quantile for row "empirical". Only used if empirical=TRUE. DEFAULT: 8 (NOT the stats::quantile default)
weighted	Include weighted averages across distribution functions to the output? DEFAULT: empirical, so additional options can all be excluded with emp=F.
gpd	Include GPD quantile estimation via q_gpd? Note that the 'GPD_LMO_lmomco' result differs slightly from 'gpa', especially if truncate=0. This comes from using x>threshold (all 'GPD_*' distributions) or x>=threshold ('gpa' and all other distributions in extremeStat). DEFAULT: empirical
speed	Compute q_gpd only for fast methods? Currently, only the Bayesian method is excluded. DEFAULT: TRUE
quiet	Suppress notes? If it is actually set to FALSE (not missing), gpquiet is set to FALSE to print all the warnings including stacks. DEFAULT: FALSE
ssquiet	Suppress sample size notes? DEFAULT: quiet
ttquiet	Suppress truncation!=threshold note? DEFAULT: quiet
gpquiet	Suppress warnings in q_gpd? DEFAULT: TRUE if quiet is not specified, else quiet
...	Arguments passed to distLfit and distLweights like weightc, ks=TRUE

Details

Very high quantiles (99% and higher) need large sample sizes for quantile to yield a robust estimate. Theoretically, at least 1/(1-probs) values must be present, e.g. 10'000 for Q99.99%. With smaller sample sizes (eg n=35), they underestimate the actual (but unknown) quantile. Parametric quantiles need only small sample sizes. They don't have a systematical underestimation bias, but have higher variability.

Value

if list=FALSE (default): invisible matrix with distribution quantile values . if list=TRUE: invisible dlf object, see printL

Note

NAs are always removed from x in `distLfit`

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, March + July 2015, Feb 2016

References

On GPD: https://stats.stackexchange.com/questions/69438

See Also

`q_gpd`, `distLfit`, require("truncdist") Xian Zhou, Liuquan Sun and Haobo Ren (2000): Quantile estimation for left truncated and right censored data, Statistica Sinica 10 https://www3.stat.sinica.edu.tw/statistica/oldpdf/A10n411.pdf

Examples

```
data(annMax) # Annual Discharge Maxima (streamflow)

distLquantile(annMax, emp=FALSE)[,] # several distribution functions in lmomco

## Not run:
## Taken out from CRAN package check because it's slow
distLquantile(annMax, truncate=0.8, probs=0.95)[,] # POT (annMax already block maxima)
dlf <- distLquantile(annMax, probs=0.95, list=TRUE)
plotLquantile(dlf, linargs=list(lwd=3), nbest=5, breaks=10)
dlf$quant
# Parametric 95% quantile estimates range from 92 to 111!
# But the best fitting distributions all lie aroud 103.

# compare General Pareto Fitting methods
# Theoretically, the tails of distributions converge to GPD (General Pareto)
# q_gpd compares several R packages for fitting and quantile estimation:
dlq <- distLquantile(annMax, weighted=FALSE, quiet=TRUE, probs=0.97, list=TRUE)
dlq$quant
plotLquantile(dlq) # per default best fitting distribution functions
plotLquantile(dlq, row=c("wak","GPD*"), nbest=14)
#pdf("dummy.pdf", width=9)
plotLquantile(dlq, row="GPD*", nbest=13, xlim=c(102,110),
        linargs=list(lwd=3), heights=seq(0.02, 0.005, len=14))
#dev.off()


# Sanity checks: important for very small samples:
x1 <- c(2.6, 2.5, 2.9, 3, 5, 2.7, 2.7, 5.7, 2.8, 3.1, 3.6, 2.6, 5.8, 5.6, 5.7, 5.3)
q1 <- distLquantile(x1, sanerange=c(0,500), sanevals=c(NA,500))
x2 <- c(6.1, 2.4, 4.1, 2.4, 6, 6.3, 2.9, 6.8, 3.5)
```

```
q2 <- distLquantile(x2, sanerange=c(0,500), sanevals=c(NA,500), quiet=FALSE)
x3 <- c(4.4, 3, 1.8, 7.3, 2.1, 2.1, 1.8, 1.8)
q3 <- distLquantile(x3, sanerange=c(0,500), sanevals=c(NA,500))

# weighted distribution quantiles are calculated by different weighting schemes:
plotLweights(dlf)

# If speed is important and parameters are already available, pass them via dlf:
distLquantile(dlf=dlf, probs=0:5/5, selection=c("wak","gev","kap"))
distLquantile(dlf=dlf, truncate=0.3, list=TRUE)$truncate

# censored (truncated, trimmed) quantile, Peak Over Treshold (POT) method:
qwak <- distLquantile(annMax, sel="wak", prob=0.95, emp=FALSE, list=TRUE)
plotLquantile(qwak, ylim=c(0,0.06) ); qwak$quant
qwak2 <-distLquantile(annMax, sel="wak", prob=0.95, emp=FALSE, list=TRUE, truncate=0.6)
plotLquantile(qwak2, add=TRUE, distcols="blue")


# Simulation of truncation effect
library(lmomco)
#set.seed(42)
rnum <- rlmomco(n=1e3, para=dlf$parameter$gev)
myprobs <- c(0.9, 0.95, 0.99, 0.999)
mytrunc <- seq(0, 0.9, length.out=20)
trunceffect <- sapply(mytrunc, function(mt) distLquantile(rnum, selection="gev",
                                 probs=myprobs, truncate=mt, quiet=TRUE,
                                 pempirical=FALSE)["gev",])
# If more values are truncated, the function runs faster

op <- par(mfrow=c(2,1), mar=c(2,4.5,2,0.5), cex.main=1)
dlf1 <- distLquantile(rnum, sel="gev", probs=myprobs, emp=FALSE, list=TRUE)
dlf2 <- distLquantile(rnum, sel="gev", probs=myprobs, emp=FALSE, list=TRUE, truncate=0.3)
plotLquantile(dlf1, ylab="", xlab="")
plotLquantile(dlf2, add=TRUE, distcols=4)
legend("right", c("fitted GEV", "fitted with truncate=0.3"), lty=1, col=c(2,4), bg="white")
par(mar=c(3,4.5,3,0.5))
plot(mytrunc, trunceffect[1,], ylim=range(trunceffect), las=1, type="l",
     main=c("High quantiles of 1000 random numbers from gev distribution",
            "Estimation based on proportion of lower values truncated"),
     xlab="", ylab="parametric quantile")
title(xlab="Proportion censored", mgp=c(1.8,1,0))
for(i in 2:4) lines(mytrunc, trunceffect[i,])
library("berryFunctions")
textField(rep(0.5,4), trunceffect[,11], paste0("Q",myprobs*100,"%") )
par(op)

trunc <- seq(0,0.1,len=200)
dd <- pbsapply(trunc, function(t) distLquantile(annMax,
         selection="gpa", weight=FALSE, truncate=t, prob=0.99, quiet=T)[c(1,3),])
 plot(trunc, dd[1,], type="o", las=1)
lines(trunc, dd[2,], type="o", col=2)
```

```
set.seed(3); rnum <- rlmomco(n=1e3, para=dlf$parameter$gpa)
qd99 <- evir::quant(rnum, p=0.99, start=15, end=1000, ci=0.5, models=30)
axis(3, at=seq(-1000,0, length=6), labels=0:5/5, pos=par("usr")[3])
title(xlab="Proportion truncated", line=-3)
mytrunc <- seq(0, 0.9, length.out=30)
trunceffect <- sapply(mytrunc, function(mt) distLquantile(rnum, selection="gpa",
                      probs=0.99, truncate=mt, plot=FALSE, quiet=TRUE,
                      empirical=FALSE, gpd=TRUE))
lines(-1000*(1-mytrunc), trunceffect[1,], col=4)
lines(-1000*(1-mytrunc), trunceffect[2,], col=3) # interesting...
for(i in 3:13) lines(-1000*(1-mytrunc), trunceffect[i,], col=3) # interesting...

# If you want the estimates only for one single truncation, use
q_gpd(rnum, probs=myprobs, truncate=0.5)



## End(Not run) # end dontrun
```

---

distLweights   *Compute distribution weights from GOF*

Description

Determine distribution function weights from RMSE for weighted averages. The weights are inverse to RMSE: weight1 for all dists, weight2 places zero weight on the worst fitting function, weight3 on the worst half of functions.

Usage

```
distLweights(
  RMSE,
  order = TRUE,
  onlydn = TRUE,
  weightc = NA,
  quiet = FALSE,
  ...
)
```

Arguments

RMSE   Numeric: Named vector with goodness of fit values (RMSE). Can also be a
              data.frame, in which case the column rmse or RMSE is used.

order   Logical: should result be ordered by RMSE? If order=FALSE, the order of ap-
              pearance in RMSE is kept (alphabetic or selection in [distLfit](#)). DEFAULT:
              TRUE

onlydn   Logical: weight only distributions from lmomco::[dist.list](#)? DEFAULT: TRUE
              (all other RMSEs are set to 0)

weightc	Optional: a named vector with custom weights for each distribution. Are internally normalized to sum=1 after removing nonfitted dists. Names match the parameter names from RMSE. DEFAULT: NA
quiet	Logical: Suppress messages. DEFAULT: FALSE
...	Ignored arguments (so a set of arguments can be passed to distLfit and distLquantile and arguments used only in the latter will not throw errors)

Value

data.frame

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Dec 2016

See Also

[distLfit](), [distLquantile]()

Examples

```
# weights from RMSE vector:
RMSE <- c(gum=0.20, wak=0.17, gam=0.21, gev=0.15)
distLweights(RMSE)
distLweights(RMSE, order=FALSE)

# weights from RMSE in data.frame:
df <- data.frame("99.9%"=2:5, RMSE=sample(3:6))
rownames(df) <- letters[1:4]
df ;  distLweights(df, onlydn=FALSE)

# custom weights:
set.seed(42); x <- data.frame(A=1:5, RMSE=runif(5)) ; x
distLweights(x) # two warnings
distLweights(x, weightc=c("1"=3, "3"=5), onlydn=FALSE)
distLweights(x, weightc=c("1"=3, "3"=5), order=FALSE, onlydn=FALSE)

# real life example:
data(annMax)
cw <- c("gpa"=7, "gev"=3, "wak"=6, "wei"=4, "kap"=3.5, "gum"=3, "ray"=2.1,
        "ln3"=2, "pe3"=2.5, "gno"=4, "gam"=5)
dlf <- distLfit(annMax, weightc=cw, quiet=TRUE, order=FALSE)
plotLweights(dlf)


# GOF judgement by RMSE, not R2 --------
# Both RMSE and R2 are computed with ECDF and TCDF
# R2 may be very good (see below), but fit needs to be close to 1:1 line,
# which is better measured by RMSE

dlf <- distLfit(annMax, ks=TRUE)
```

```
op <- par(mfrow=c(1,2), mar=c(3,4,0.5,0.5), mgp=c(1.9,0.7,0))
yy <- nrow(dlf$gof):1 # depends on length of lmomco::dist.list()
plot(dlf$gof$RMSE, yy, yaxt="n", ylab="", type="o"); axis(2, yy, rownames(dlf$gof), las=1)
plot(dlf$gof$R2,   yy, yaxt="n", ylab="", type="o"); axis(2, yy, rownames(dlf$gof), las=1)
par(op)
sel <- c("wak","lap","nor","revgum")
plotLfit(dlf, selection=sel, cdf=TRUE)
dlf$gof[sel,-(2:7)]

x <- sort(annMax, decreasing=TRUE)
ECDF <- ecdf(x)(x)
TCDF <- sapply(sel, function(d) lmomco::plmomco(x,dlf$parameter[[d]]))

plot(TCDF[,"lap"],    ECDF, col="cyan", asp=1, las=1)
points(TCDF[,"nor"],    ECDF, col="green")
#points(TCDF[,"wak"],    ECDF, col="blue")
#points(TCDF[,"revgum"], ECDF, col="red")
abline(a=0, b=1, lwd=3, lty=3)
legend("bottomright", c("lap good RMSE bad R2", "nor bad RMSE good R2"),
       col=c("cyan","green"), lwd=2)
berryFunctions::linReg(TCDF[,"lap"], ECDF, add=TRUE, digits=3, col="cyan", pos1="topleft")
berryFunctions::linReg(TCDF[,"nor"], ECDF, add=TRUE, digits=3, col="green", pos1="left")


# more distinct example (but with fake data)
set.seed(42); x <- runif(30)
y1 <-     x+rnorm(30,sd=0.09)
y2 <- 1.5*x+rnorm(30,sd=0.01)-0.3
plot(x,x, asp=1, las=1, main="High cor (R2) does not necessarily mean good fit!")
berryFunctions::linReg(x, y2, add=TRUE, digits=4, pos1="topleft")
points(x,y2, col="red", pch=3)
points(x,y1, col="blue")
berryFunctions::linReg(x, y1, add=TRUE, digits=4, col="blue", pos1="left")
abline(a=0, b=1, lwd=3, lty=3)
```

---

extremeStat	*Extreme value statistics on a linear scale*

Description

Fit (via L moments), plot (on a linear scale) and compare (by goodness of fit) several (extreme value) distributions. Compute high quantiles even in small samples and estimate extrema at given return periods.

Open the Vignette for an introduction to the package: vignette("extremeStat")

This package heavily relies on and thankfully acknowledges the package lmomco by WH Asquith.

Package overview

The main functions in the extremeStat package are:

```
                  distLweights                        -> plotLweights
                  distLfit                             -> plotLfit
                  q_gpd + q_weighted -> distLquantile  -> plotLquantile
                  distLextreme                         -> plotLextreme
                  distLexBoot
```

They create and modify a list object printed by (and documented in) `printL`.

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, 2014-2016

See Also

If you are looking for more detailed (uncertainty) analysis, eg confidence intervals, check out the package extRemes, especially the function `fevd`. `https://cran.r-project.org/package=extRemes`
Intro slides: `https://sites.lsa.umich.edu/eva2015/wp-content/uploads/sites/44/2015/06/Intro2EVT.pdf`
Parameter fitting and distribution functions: `https://cran.r-project.org/package=lmomco`
Distributions: `https://web.archive.org/web/20110807225801/https://www.rmetrics.org/files/Meielisalp2009/Presentations/Scott.pdf` and: `https://cran.r-project.org/view=Distributions`
R in Hydrology: `https://abouthydrology.blogspot.de/2012/08/r-resources-for-hydrologists.html`

Examples

```
data(annMax) # annual discharge maxima from a stream in Austria
plot(annMax, type="l")
dle <- distLextreme(annMax)
dle$returnlev
```

---

plotLexBoot   *Bootstrapping uncertainty intervals for return periods*

Description

plot bootstrap uncertainty intervals for `plotLextreme`.

Usage

```
plotLexBoot(dlf, selection = NULL, add = FALSE, log = TRUE, ...)
```

Arguments

dlf	dlf object, as returned by distLexBoot
selection	Character vector with distribution function names to be used. Suggested to keep this low. DEFAULT: NULL
add	Add to existing plot? DEFAULT: FALSE
log	Plot on a logarithmic axis. DEFAULT: TRUE
...	Further arguments passed to plotLextreme. If add=TRUE, they are instead passed to berryFunctions::ciBand

Value

invisible dlf object, see printL

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Dec 2016

See Also

distLexBoot

Examples

```
# see distLexBoot
```

---

plotLextreme   *Plot extreme value statistics*

Description

Plots distributions fitted by L-moments and adds plotting positions by Weibull and Gringorton. This is an auxiliary graphing function to distLextreme

Usage

```
plotLextreme(
  dlf,
  selection = NULL,
  order = FALSE,
  add = FALSE,
  nbest = 5,
  log = "",
  xlim = NULL,
  ylim = NULL,
  las = 1,
```

```
    main = dlf$datname,
    xlab = "Return Period RP  [a]",
    ylab = "Discharge HQ  [m\U00B3/s]",
    PPcol = "black",
    PPpch = c(16, 3),
    PPcex = 1,
    distcols = berryFunctions::rainbow2(nbest),
    lty = 1,
    lwd = 1,
    pch = NA,
    cex = 1,
    n_pch = 15,
    legend = TRUE,
    rmse = 4,
    legargs = NULL,
    quiet = FALSE,
    logargs = NULL,
    ...
)
```

Arguments

dlf	List as returned by [distLextreme](#) or [distLexBoot](#)
selection	Selection of distributions. Character vector with type as in [lmom2par](#). DE-FAULT: NULL
order	If selection is given, should legend and colors be ordered by gof anyways? DE-FAULT: FALSE
add	If TRUE, plot is not called before adding lines. This lets you add lines to an existing plot. DEFAULT: FALSE
nbest	Number of distributions plotted, in order of goodness of fit. Overwritten internally if selection is given. DEFAULT: 5
log	Charstring ("x", "y", "xy") for logarithmic axes. See logargs. DEFAULT: ""
xlim	X-axis limits. DEFAULT: xlim of plotting positions
ylim	Y-lim. DEFAULT: from min to extended max
las	LabelAxisStyle to orient labels, see [par](#). DEFAULT: 1
main	Title of plot. DEFAULT: dlf$datname
xlab	X axis label. DEFAULT: "Return Period RP [a]"
ylab	Y axis label. Please note that the ubuntu pdf viewer might be unable to display unicode superscript. DEFAULT: "Discharge HQ [m3/s]"
PPcol	Plotting Position point colors, vector of length two for Weibull and Gringorton, recycled. PP are not used for fitting distributions, but for plotting only. DE-FAULT: "black"
PPpch	point characters for plotting positions after Weibull and Gringorton, respectively. NA to suppress in plot and legend. DEFAULT: c(16,3)
PPcex	Character EXpansion of plotting points. DEFAULT: 1

distcols	Color for each distribution added with [lines](). Recycled, if necessary. DEFAULT: [rainbow2]()
lty	Line TYpe for plotted distributions. Is recycled to from a vector of length nbest, i.e. a value for each dist. DEFAULT: 1
lwd	Line WiDth of distribution lines. Recycled vector of length nbest. DEFAULT: 1
pch	Point CHaracter of points added at regular intervals. This makes lines more distinguishable from each other. NA to suppress. Recycled vector of length nbest. DEFAULT: NA
cex	if pch != NA, size of points. Recycled vector of length nbest. DEFAULT: 1
n_pch	Number of points spread evenly along the line. Recycled vector of length nbest. DEFAULT: 15
legend	Logical. Add a legend? DEFAULT: TRUE
rmse	Integer. If rmse > 0, RMSE values are added to legend. They are rounded to rmse digits. DEFAULT: 4
legargs	list of arguments passed to [legend]() except for legend, col, pch, lwd, lty. DEFAULT: NULL
quiet	Suppress notes? DEFAULT: FALSE
logargs	list of arguments passed to berryFunctions::[logAxis]().
...	Further arguments passed to [plot]() like yaxt="n", ...

Value

invisible dlf object, see [printL]()

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, March 2015, updated heavily Aug 2015

See Also

[distLextreme](), [plotLfit]()

Examples

```
#see
?distLextreme
```

---

plotLfit | *Plot distributions fitted with L-moments*

Description

Plot histogram and distribution densities *or* ecdf with cumulated probability

Usage

```
plotLfit(
  dlf,
  nbest = 5,
  selection = NULL,
  order = TRUE,
  rmse = 4,
  cdf = FALSE,
  log = FALSE,
  supportends = TRUE,
  breaks = 20,
  xlim = extendrange(dlf$dat, f = 0.15),
  ylim = NULL,
  col = "grey",
  main = paste(if (cdf) "Cumulated", "density distributions of", dlf$datname),
  xlab = dlf$datname,
  ylab = if (cdf) "(Empirical) Cumulated Density (CDF)" else
    "Probability Density Function (PDF)",
  las = 1,
  distcols = berryFunctions::rainbow2(nbest),
  lty = 1,
  add = FALSE,
  logargs = NULL,
  legend = TRUE,
  legargs = NULL,
  histargs = NULL,
  ...
)
```

Arguments

dlf	List as returned by [distLfit](), containing the elements dat, parameter, gof, datname
nbest	Number of distributions plotted, in order of goodness of fit. DEFAULT: 5
selection	Names of distributions in dlf$parameter that will be drawn. Overrides nbest. DEFAULT: NULL
order	Logical: order legend and colors by RMSE, even if dlf$gof is unordered or selection is given? DEFAULT: TRUE

rmse	Integers. If rmse != 0, RMSE values are added to legend. They are rounded to rmse digits. DEFAULT: 4
cdf	If TRUE, plot cumulated DF instead of probability density. DEFAULT: FALSE
log	If TRUE, logAxis is called. Only makes sense if dlf$dat is already logarithmic and ranges eg. from -2 to 3. DEFAULT: FALSE
supportends	If TRUE, dots are placed at the support bounds. DEFAULT: TRUE
breaks	hist breaks. DEFAULT: 20
xlim, ylim	hist or ecdf axis limits.
col	hist bar color or ecdf point color. DEFAULT: "grey"
main, xlab, ylab	
	hist or ecdf main, xlab, ylab. DEFAULT: abstractions from dlf$datname
las	Label Axis Style for orientation of numbers along axes. DEFAULT: 1
distcols	Color for each distribution added with lines. DEFAULT: rainbow2
lty	Line TYpe for plotted distributions. Recycled vector of length nbest. DEFAULT: 1
add	If TRUE, hist/ecdf is not called before adding lines. This lets you add lines highly customized one by one. DEFAULT: FALSE
logargs	List of arguments passed to logAxis if log=TRUE. DEFAULT: NULL
legend	Should legend be called? DEFAULT: TRUE
legargs	List of arguments passed to legend except for legend and col. DEFAULT: NULL
histargs	List of arguments passed to hist or ecdf except for x, freq. DEFAULT: NULL
...	Further arguments passed to lines, like type, pch, ...

Details

By default, this plots density instead of CDF, because the distributions are easier to discern and tail behavior is easier to judge visually.

Value

invisible dlf object, see printL

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Sept 2014

See Also

distLfit, plotLquantile

Examples

```
# See distLfit
```

---

plotLquantile   *Plot quantiles of distributions fitted with L-moments*

Description

Plot quantiles of distributions fitted with L-moments

Usage

```
plotLquantile(
  dlf,
  nbest = 5,
  selection = NULL,
  order = FALSE,
  rows = NULL,
  heights = stats::quantile(par("usr")[3:4], 0.2),
  distcols = dlfplot$distcols,
  linargs = NULL,
  ...
)
```

Arguments

dlf	List as returned by [distLquantile](#), containing the elements dat, parameter, gof, datname, quant
nbest, selection, order	
	Distributions to be plotted, see [plotLfit](#)
rows	Rowname(s) of dlf$quant that should be drawn instead of the selection / nbest highest ranking distribution functions. 'GPD*' will select all the gpd fits. heights and distcols must then accordingly have at least 13 elements (or will be recycled). DEFAULT: NULL
heights	Coordinates of quantile line ends, recycled if necessary. DEFAULT: 20% of plot height.
distcols	Color for each distribution added with [lines](#). DEFAULT: dlfplot$distcols
linargs	Arguments passed to [lines](#). DEFAULT: NULL
...	Further arguments passed to [plotLfit](#)

Value

invisible dlf object, see [printL](#)

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Dec 2016

See Also

distLquantile, plotLfit

Examples

```
# See distLquantile
```

---

plotLweights   *Distribution rank comparison*

Description

Plot rank comparison of fitted distributions calculated by distLfit.

Usage

```
plotLweights(
  dlf,
  type = "o",
  col = RColorBrewer::brewer.pal(5, "Set2"),
  pch = c(1:4, NA),
  lty = 1,
  lwd = 1,
  legargs = NULL,
  main = "Distribution function GOF and weights",
  xlab = "Weight / RMSE",
  ylab = "",
  xlim = range(gof[, grep("weight", colnames(gof))], na.rm = TRUE),
  ...
)
```

Arguments

dlf	List as returned by distLfit, containing the element gof
type, col, pch, lty, lwd	
	Vectors with 5 values for line customization. Recycled if necessary.
legargs	List of arguments passed to legend, like cex, bg, etc.
main, xlab, ylab	
	plot title and axis labels
xlim	Range of x axis. DEFAULT: range(gof$weight*)
...	Further arguments passed to plot.

Value

None.

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Sept 2014

See Also

[distLweights](#), [distLfit](#)

Examples

```
# see distLweights and distLfit
```

---

printL   *print dlf objects*

Description

print list objects created in this package

Usage

```
printL(dlf, digits = 1)
```

Arguments

dlf	List as explained in section Details
digits	number of digits [round](#)ed to. DEFAULT: 1

Details

The common object to share between functions (see overview in [extremeStat](#)) is a list with the following elements:

dat	numeric vector with (extreme) values, with all NAs and values below threshold removed
dat_full	original input data complete with NAs
datname	character string for main, xlab etc
parameter	list (usually of length 17 if speed=TRUE in [distLfit](#)) with parameters of each distribution
gof	dataframe with 'Goodness of Fit' measures, sorted by RMSE of theoretical and empirical cumulated
distnames	character vector with selected distribution names
distfailed	Names of nonfitted distributions or ""
distcols	colors for distnames (for plotting). If not given manually, determined by berryFunctions::[rainbo](#)
distselector	character string with function name creating the selection
truncate, threshold	Truncation percentage and threshold value, relevant for [distLquantile](#)

optionally, it can also contain:

returnlev, npy	dataframe with values of distributions for given return periods (RPs), num
RPweibull, RPgringorton	Return periods according to plotting positions, added in [plotLextreme](#)
quant	Quantile estimates from [distLquantile](#)
exBootRPs, qexBootSim, exBootCI, exBootCL	objects from [distLexBoot](#)

Value

none, prints via [message](#).

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Sept 2014, March + July 2015, Dec 2016

See Also

[extremeStat](#)

Examples

```
# see
?distLextreme
```

---

quantGPD   *Fast GPD quantile estimate*

Description

Fast GPD quantile estimate through L-moments

Usage

```
quantGPD(
  x,
  probs = c(0.8, 0.9, 0.99),
  truncate = 0,
  threshold = berryFunctions::quantileMean(x, truncate),
  addn = TRUE,
  quiet = FALSE,
  ...
)
```

Arguments

x	Vector with numeric values. NAs are silently ignored.
probs	Probabilities. DEFAULT: c(0.8,0.9,0.99)
truncate, threshold	
	Truncation proportion or threshold. DEFAULT: 0, computed See q_gpd.
addn	Logical: add element with sample size (after truncation). DEFAULT: TRUE
quiet	Should messages from this function be suppressed? DEFAULT: FALSE
...	Further arguments passed to lmomco::pargpa

Value

Vector with quantiles

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Jun 2017

See Also

q_gpd for a comparison across R packages and methods, distLquantile to compare distributions

Examples

```
data(annMax)
quantile(annMax, 0.99)
quantGPD(annMax, 0.99)

## Not run:  # Excluded from CRAN checks to reduce checking time
data(rain, package="ismev") ;  rain <- rain[rain>0]
hist(rain, breaks=50, col=7)
tr <- seq(0,0.999, len=50)
qu <- pbapply::pbsapply(tr, quantGPD, x=rain, probs=c(0.9,0.99,0.999) ) # 30 s
plot(tr, qu[3,], ylim=range(rain), las=1, type="l")
lines(tr, qu[2,], col=2); lines(tr, qu[1,], col=4)

tr <- seq(0.88,0.999, len=50)
qu <- pbapply::pbsapply(tr, quantGPD, x=rain, probs=c(0.9,0.99,0.999) ) # 5 s
plot(tr, qu[3,], ylim=range(rain), las=1, type="l")
lines(tr, qu[2,], col=2); lines(tr, qu[1,], col=4);
tail(qu["n",])

library(microbenchmark)
data(rain, package="ismev"); rain <- rain[rain>0]
mb <- microbenchmark(quantGPD(rain[1:200], truncate=0.8, probs=0.99, addn=F),
distLquantile(rain[1:200], sel="gpa", emp=F, truncate=0.8, quiet=T, probs=0.99)[1,1]
)
boxplot(mb)
# since computing the lmoments takes most of the computational time,
# there's not much to optimize in large samples like n=2000
```

```
## End(Not run)
```

---

q_gpd   *GPD quantile of sample*

Description

Compute quantile of General Pareto Distribution fitted to sample by peak over threshold (POT) method using threshold from truncation proportion, comparing several R packages doing this

Usage

```
q_gpd(
  x,
  probs = c(0.8, 0.9, 0.99),
  truncate = 0,
  threshold = berryFunctions::quantileMean(x, truncate),
  package = "extRemes",
  method = NULL,
  list = FALSE,
  undertruncNA = TRUE,
  quiet = FALSE,
  ttquiet = quiet,
  efquiet = quiet,
  ...
)
```

Arguments

x	Vector with numeric values. NAs are silently ignored.
probs	Probabilities of truncated (Peak over threshold) quantile. DEFAULT: c(0.8,0.9,0.99)
truncate	Truncation percentage (proportion of sample discarded). DEFAULT: 0
threshold	POT cutoff value. If you want correct percentiles, set this only via truncate, see Details. DEFAULT: [quantileMean](x, truncate)
package	Character string naming package to be used. One of c("lmomco","evir","evd","extRemes","fExtremes","is DEFAULT: "extRemes"
method	method passed to the fitting function, if applicable. Defaults are internally specified (See Details), depending on package, if left to the DEFAULT: NULL.
list	Return result from the fitting function with the quantiles added to the list as element quant and some information in elements starting with q_gpd_. DEFAULT: FALSE
undertruncNA	Return NAs for probs below truncate? Highly recommended to leave this at the DEFAULT: TRUE

quiet	Should messages from this function be suppressed? DEFAULT: FALSE
ttquiet	Should truncation!=threshold messages from this function be suppressed? DE-FAULT: quiet
efquiet	Should warnings in function calls to the external packages be suppressed via [options](warn=-1)? The usual type of warning is: NAs produced in log(...). DEFAULT: quiet
...	Further arguments passed to the fitting function listed in section Details.

Details

Depending on the value of "package", this fits the GPD using
lmomco::[pargpa](#)
evir::[gpd](#)
evd::[fpot](#)
extRemes::[fevd](#)
fExtremes::[gpdFit](#)
ismev::[gpd.fit](#)
Renext::[Renouv](#) or Renext::[fGPD](#)

The method defaults (and other possibilities) are
lmomco: none, only L-moments
evir: "pwm" (probability-weighted moments), or "ml" (maximum likelihood)
evd: none, only Maximum-likelihood fitting implemented
extRemes: "MLE", or "GMLE", "Bayesian", "Lmoments"
fExtremes: "pwm", or "mle"
ismev: none, only Maximum-likelihood fitting implemented
Renext: "r" for [Renouv](#) (since distname.y = "gpd", evd::fpot is used), or 'f' for [fGPD](#) (with minimum POTs added)

The Quantiles are always given with probs in regard to the full (uncensored) sample. If e.g. truncate is 0.90, the distribution function is fitted to the top 10% of the sample. The 95th percentile of the full sample is equivalent to the 50% quantile of the subsample actually used for fitting. For computation, the probabilities are internally updated with p2=(p-t)/(1-t) but labeled with the original p. If you truncate 90% of the sample, you cannot compute the 70th percentile anymore, thus undertruncNA should be left to TRUE.
If not exported by the packages, the quantile functions are extracted from their source code (Nov 2016).

Value

Named vector of quantile estimates for each value of probs,
or if(list): list with element q_gpd_quant and info-elements added. q_gpd_n_geq is number of values greater than or equal to q_gpd_threshold. gt is only greater than.

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Feb 2016

References

<https://stackoverflow.com/q/27524131>, <https://stats.stackexchange.com/q/129438>

See Also

distLquantile which compares results for all packages
Other related packages (not implemented):
<https://cran.r-project.org/package=gPdtest>
<https://cran.r-project.org/package=actuar>
<https://cran.r-project.org/package=fitdistrplus>
<https://cran.r-project.org/package=lmom>

Examples

```
data(annMax)
q_gpd(annMax)
q_gpd(annMax, truncate=0.6)
q_gpd(annMax, truncate=0.85)
q_gpd(annMax, truncate=0.91)

q_gpd(annMax, package="evir")
q_gpd(annMax, package="evir", method="ml")
q_gpd(annMax, package="evd")
q_gpd(annMax, package="extRemes")
q_gpd(annMax, package="extRemes", method="GMLE")
#q_gpd(annMax, package="extRemes", method="Bayesian") # computes a while
q_gpd(annMax, package="extRemes", method="Lmoments")
q_gpd(annMax, package="extRemes", method="nonsense") # NAs
q_gpd(annMax, package="fExtremes")                        # log warnings
q_gpd(annMax, package="fExtremes", efquiet=TRUE)     # silenced warnings
q_gpd(annMax, package="fExtremes", method= "mle")
q_gpd(annMax, package="ismev")
q_gpd(annMax, package="Renext")
q_gpd(annMax, package="Renext", method="f")
berryFunctions::is.error(q_gpd(annMax, package="nonsense"), force=TRUE)

# compare all at once with
d <- distLquantile(annMax); d
# d <- distLquantile(annMax, speed=FALSE); d # for Bayesian also

q_gpd(annMax, truncate=0.85, package="evd")            # Note about quantiles
q_gpd(annMax, truncate=0.85, package="evir")
q_gpd(annMax, truncate=0.85, package="evir", quiet=TRUE) # No note
q_gpd(annMax, truncate=0.85, package="evir", undertruncNA=FALSE)

q_gpd(annMax, truncate=0.85, package="evir", list=TRUE)
```

```
str(  q_gpd(annMax, truncate=0.85, probs=0.6, package="evir", list=TRUE) )# NAs
str(  q_gpd(annMax, package="evir",     list=TRUE)   )
str(  q_gpd(annMax, package="evd",      list=TRUE)   )
str(  q_gpd(annMax, package="extRemes", list=TRUE)   )
str(  q_gpd(annMax, package="fExtremes", list=TRUE)  )
str(  q_gpd(annMax, package="ismev",    list=TRUE)   )
str(  q_gpd(annMax, package="Renext",   list=TRUE)   )

q_gpd(annMax, package="evir", truncate=0.9, method="ml") # NAs (MLE fails often)

trunc <- seq(0,0.9,len=500)
library("pbapply")
quant <- pbsapply(trunc, function(tr) q_gpd(annMax, pack="evir", method = "pwm",
                                            truncate=tr, quiet=TRUE))
quant <- pbsapply(trunc, function(tr) q_gpd(annMax, pack="lmomco", truncate=tr, quiet=TRUE))
plot(trunc, quant["99%",], type="l", ylim=c(80,130), las=1)
lines(trunc, quant["90%",])
lines(trunc, quant["80%",])
plot(trunc, quant["RMSE",], type="l", las=1)

## Not run:
## Not run in checks because simulation takes too long

trunc <- seq(0,0.9,len=200)
dlfs <- pblapply(trunc, function(tr) distLfit(annMax, truncate=tr, quiet=TRUE, order=FALSE))
rmses <- sapply(dlfs, function(x) x$gof$RMSE)
plot(trunc, trunc, type="n", ylim=range(rmses,na.rm=TRUE), las=1, ylab="rmse")
cols <- rainbow2(17)[rank(rmses[,1])]
for(i in 1:17) lines(trunc, rmses[i,], col=cols[i])

dlfs2 <- lapply(0:8/10, function(tr) distLfit(annMax, truncate=tr, quiet=TRUE))
pdf("dummy.pdf")
dummy <- sapply(dlfs2, function(x)
{plotLfit(x, cdf=TRUE, main=x$truncate, ylim=0:1, xlim=c(20,135), nbest=1)
title(sub=round(x$gof$RMSE[1],4))
})
dev.off()

# truncation effect
mytruncs <- seq(0, 0.9, len=150)
oo <- options(show.error.messages=FALSE, warn=-1)
myquants <- sapply(mytruncs, function(t) q_gpd(annMax, truncate=t, quiet=TRUE))
options(oo)
plot(1, type="n", ylim=range(myquants, na.rm=TRUE), xlim=c(0,0.9), las=1,
     xlab="truncated proportion", ylab="estimated quantiles")
abline(h=quantileMean(annMax, probs=c(0.8,0.9,0.99)))
for(i in 1:3) lines(mytruncs, myquants[i,], col=i)
text(0.3, c(87,97,116), rownames(myquants), col=1:3)


# Underestimation in small samples
# create known population:
dat <- extRemes::revd(1e5, scale=50, shape=-0.02, threshold=30, type="GP")
```

```
op <- par(mfrow=c(1,2), mar=c(2,2,1,1))
hist(dat, breaks=50, col="tan")
berryFunctions::logHist(dat, breaks=50, col="tan")
par(op)

# function to estimate empirical and GPD quantiles from subsamples
samsizeeffect <- function(n, nrep=30, probs=0.999, trunc=0.5, Q=c(0.4,0.5,0.6))
{
res <- replicate(nrep, {
subsample <- sample(dat, n)
qGPD <- q_gpd(subsample, probs=probs, truncate=trunc)
qEMP <- berryFunctions::quantileMean(subsample, probs=probs, truncate=trunc)
c(qGPD=qGPD, qEMP=qEMP)})
apply(res, MARGIN=1, berryFunctions::quantileMean, probs=Q)
}

# Run and plot simulations
samplesize <- c(seq(20, 150, 10), seq(200,800, 100))
results <- pbapply::pblapply(samplesize, samsizeeffect)
res <- function(row, col) sapply(results, function(x) x[row,col])
berryFunctions::ciBand(yu=res(3,1),yl=res(1,1),ym=res(2,1),x=samplesize,
  main="99.9% Quantile underestimation", xlab="subsample size", ylim=c(200,400), colm=4)
berryFunctions::ciBand(yu=res(3,2),yl=res(1,2),ym=res(2,2),x=samplesize, add=TRUE)
abline(h=berryFunctions::quantileMean(dat, probs=0.999))
text(300, 360, "empirical quantile of full sample")
text(300, 340, "GPD parametric estimate", col=4)
text(300, 300, "empirical quantile estimate", col="green3")


## End(Not run) # end of dontrun
```

---

q_weighted   *Compute weighted averages of quantile estimates*

Description

Compute weighted averages of quantile estimates

Usage

```
q_weighted(quant, weights = distLweights(quant, ...), onlyc = FALSE, ...)
```

Arguments

quant	Data.frame as in [distLquantile](#) output.
weights	Data.frame as in [distLweights](#) output.
onlyc	Logical: only return custom weighted quantile estimates as a vector? Useful to add those to existing results. See examples. DEFAULT: FALSE

...	Arguments passed to `distLweights` like weightc, onlydn=FALSE. order will be ignored, as q_weighted only adds/changes the rows weighted*.

Value

data.frame with rows "weighted*" added.

Author(s)

Berry Boessenkool, <berry-b@gmx.de>, Dec 2016

See Also

`distLquantile`

Examples

```
x <- data.frame(A=1:5, RMSE=runif(5))
distLweights(x, onlydn=FALSE)

q_weighted(x,  onlydn=FALSE)
q_weighted(x,  distLweights(x, weightc=c("1"=3, "3"=5), order=FALSE, onlydn=FALSE)  )

## Not run:  # time consuming
x <- rexp(190)
d <- distLquantile(x)
d2 <- q_weighted(d)
stopifnot(all(d==d2, na.rm=TRUE))

# fast option for adding custom weighted estimates:
cw <- runif(17)
names(cw) <- c("exp", "gam", "gev", "glo", "gno", "gpa", "gum", "kap", "lap",
               "ln3", "nor", "pe3", "ray", "revgum", "rice", "wak", "wei")
dw <- distLweights(d, weightc=cw)
qw1 <- q_weighted(d, weightc=cw); qw1
qw2 <- q_weighted(d, weights=dw); qw2
stopifnot(all(qw1==qw2, na.rm=TRUE))
q_weighted(d, weights=dw, onlyc=TRUE)
q_weighted(d, weights=data.frame(weightc=cw), onlyc=TRUE)

system.time(pbreplicate(5000, q_weighted(d, weightc=cw)))             # 8.5 secs
system.time(pbreplicate(5000, q_weighted(d, weights=dw, onlyc=TRUE))) # 0.8 secs

## End(Not run)
```

---

weightp   *distribution weights*

Description

Weights for weighted average as in the submission of revisions for the paper [https://nhess.](https://nhess.copernicus.org/articles/17/1623/2017/nhess-17-1623-2017-discussion.html)
[copernicus.org/articles/17/1623/2017/nhess-17-1623-2017-discussion.html](https://nhess.copernicus.org/articles/17/1623/2017/nhess-17-1623-2017-discussion.html)

Format

named num [1:17]

Source

See paper revisions (not yet online at moment of extremeStat update) (<berry-b@gmx.de>)

Examples

```
data(weightp)
data.frame(weightp)
barplot(weightp, horiz=TRUE, las=1)
stopifnot(   all.equal(sum(weightp), 1)   )

data(annMax) ; data(weightp)
dlf <- distLfit(annMax, weightc=weightp)
dlf$gof
quant <- distLquantile(annMax, weightc=weightp)
quant
```

Index