

Package ‘famish’

December 8, 2025

Type Package

Title Flexibly Tune Families of Probability Distributions

Version 0.2.0

Description Fits probability distributions to data and plugs into the ‘probaverse’ suite of R packages so distribution objects are ready for further manipulation and evaluation. Supports methods such as maximum likelihood and L-moments, and provides diagnostics including empirical ranking and quantile score.

License MIT + file LICENSE

Encoding UTF-8

Imports checkmate, distionary, fitdistrplus, ismev, lmom, rlang, stats, vctrs

RoxygenNote 7.3.3

Suggests knitr, rmarkdown, testthat (>= 3.0.0)

Config/testthat/edition 3

URL <https://famish.probaverse.com/>,
<https://github.com/probaverse/famish/>

VignetteBuilder knitr

NeedsCompilation no

Author Vincenzo Coia [aut, cre, cph]

Maintainer Vincenzo Coia <vincenzo.coia@gmail.com>

Repository CRAN

Date/Publication 2025-12-08 08:20:11 UTC

Contents

fit_dst	2
fit_dst_family_wrappers	4
quantile_score	6
uscore	7

Index

10

fit_dst*Fit a distribution***Description**

Estimation of probability distributions available in the 'distortion' package. Wraps the 'lmom' package when fitting by L-moments, the 'ismev' package when fitting the GP/GEV/Gumbel by MLE, and the 'fitdistrplus' package for other combinations.

Usage

```
fit_dst(
  family,
  x,
  method = c("mle", "mge", "mme", "lmom", "lmom-log"),
  na_action = c("null", "drop", "fail"),
  on_unres = c("null", "fail")
)
```

Arguments

<code>family</code>	Name of the target distribution family, such as "norm", "gev", "pois". See details. Character vector of length 1.
<code>x</code>	Numeric vector containing the observations to fit.
<code>method</code>	Estimation method to use. Valid choices include "mle", "mge", "mme", "lmom", and "lmom-log". The default is "mle", although beware that not all families support the "mle" method yet (pearson3 and lp3).
<code>na_action</code>	Strategy for dealing with NA values in <code>x</code> . "null" returns a Null distribution (<code>distortion::dst_null()</code>); "drop" silently removes missing observations before fitting; and "fail" aborts with an error.
<code>on_unres</code>	Behaviour when no distribution can be resolved for the supplied inputs. "null" (default) yields a Null distribution (<code>distortion::dst_null()</code>) distribution with a warning, whereas "fail" propagates an error.

Details

The `fit_dst()` function is currently a lightweight fitting wrapper, with pre-specified behaviour for certain `family` / `method` combinations. A full list of families and their compatible estimation methods is available via specific family wrappers, such as `fit_dst_norm()`, `fit_dst_pois()`, etc.

Here is how fitting is implemented.

- For `method = "lmom"` and distribution families 'gamma', 'gev', 'gp', 'gumbel', 'lnorm', 'norm', 'pearson3', and 'weibull', the 'lmom' package is wrapped by first calling `lmom::sam1mu()` on the data vector `x` to calculate the L-moments, then the relevant `lmom::pel*`() function is called to estimate parameters.

- For `method = "lmom"` and distribution families `'exp'`, `'pois'`, `'bern'`, `'geom'`, `'chisq'`, and `'unif'`, the method of L-moments is manually implemented. All of these families except `'unif'` have a single parameter for which only the mean is needed (and thus is equivalent to the `'mme'` method). The `'unif'` family has minimum and maximum parameter values calculated as $11 - 3 * 12$ and $11 + 3 * 12$, where 11 and 12 are the first and second L-moments (see Hosking, 1990, Table 1).
- For `method = "lmom-log"`, only the `'lnorm'` and `'lp3'` families are supported, otherwise no distribution will be resolved. The method fits the distributions via the `'lmom'` method on the log scale. That is, `'norm'` and `'pearson3'` distributions are fit on the log of the data, for which the respective `'lnorm'` or `'lp3'` distribution is obtained.
- For `method = "mle"` and distribution families `'gev'`, `'gp'`, or `'gumbel'`, the `'ismev'` package is used to fit the distribution by maximum likelihood estimation. This is done by invoking the functions `ismev::gev.fit()`, `ismev::gpd.fit()` (with `threshold = 0`), and `ismev::gum.fit()` (respectively).
- For `method = "mle"` and distribution family `'bern'` and `'degenerate'`, the MLE is calculated manually. For `'bern'`, the parameter is estimated as the mean of the 0-1 data; for `'degenerate'`, the unique data value.
- For `method = "mme"` and `"lmom"`, the `'cauchy'` family fails to fit because Cauchy distributions don't have finite moments (Feller, 1971).
- For families `'empirical'` and `'finite'`, the empirical distribution is fit to the supplied data.
- For the `'null'` family, a Null distribution is returned.
- For any other combination of `family` and `method`, the `fitdistrplus::fitdist()` function is called by inserting the data `x`, the `family` name, and the `method`. Some distributions require starting values for the parameters. For the families `'t'`, `'f'`, and `'chisq'`, this is done by moment matching (`'mme'`). For `'gev'`, `'gp'`, and `'gumbel'`, the MLE is used as starting values (through `method = "mle"`).

To understand what the distribution families are, see the documentation in the `'dictionary'` package through the `dst_*`() functions. For example, the `'lp3'` family can be found at `?dictionary::dst_lp3()`. Note that the Gumbel distribution is not available yet in `'dictionary'`, but is simply the `'gev'` family with `shape = 0`.

To understand the estimation methods, see the `'lmom'` package for the `"lmom"` method. For the `"lmom-log"` method, it is the same as `"lmom"`, but via the log of the data and the corresponding log-transformed distributions. For all others, see the `'fitdistrplus'` package documentation.

Handling of missing or unresolvable data:

When `na_action = "drop"`, the function operates on the subset of `x` without missing values (via `x <- x[!is.na(x)]`). This takes priority over behaviour indicated in `on_unres`.

If fitting fails, a Null distribution is output if `on_unres = "null"` (the default), or an error is thrown if `on_unres = "fail"`. Fitting can fail due to not having enough data, not being able to isolate a single distribution, or various other reasons that would typically otherwise result in an error or NA parameters in the wrapped fitting method.

Value

A distribution object of class `"dst"` encapsulating the fitted distribution.

References

- Hosking, J. R. M. (1990). L-moments: Analysis and estimation of distributions using linear combinations of order statistics. *Journal of the Royal Statistical Society: Series B (Methodological)*, 52(1), 105–124.
- Feller, W. (1971). *An Introduction to Probability Theory and Its Applications* (Vol. 2, 2nd ed.). Wiley.

See Also

`fit_dst_*`() helpers such as [fit_dst_norm\(\)](#).

Examples

```
fit_dst("norm", x = 1:10, method = "mle")
fit_dst("gev", x = c(1, 4, 3, NA, 5), method = "lmom", na_action = "drop")
fit_dst("pois", x = c(1, 4, 3, NA, 5), na_action = "null")

# "lnorm" with "lmom-log" shares parameters with "norm" fit by "lmom".
fit_dst("lnorm", x = 1:10, method = "lmom-log")
fit_dst("norm", x = log(1:10), method = "lmom")
```

fit_dst_family_wrappers

Fit Distributions by Family

Description

Convenience wrappers around `fit_dst()` for each distribution family supported by the package.

Usage

```
fit_dst_bern(x, method = c("mle", "lmom"), ...)
fit_dst_beta(x, method = c("mle", "mge", "mme"), ...)
fit_dst_cauchy(x, method = c("mle", "mge"), ...)
fit_dst_chisq(x, method = c("mle", "mge", "lmom"), ...)
fit_dst_degenerate(x, method = "mle", ...)
fit_dst_empirical(x, ...)
fit_dst_exp(x, method = c("mle", "mge", "mme", "lmom"), ...)
fit_dst_f(x, method = c("mle", "mge"), ...)
```

```

fit_dst_finite(x, ...)

fit_dst_gamma(x, method = c("mle", "mge", "mme", "lmom"), ...)

fit_dst_geom(x, method = c("mle", "mme", "lmom"), ...)

fit_dst_gev(x, method = c("mle", "lmom"), ...)

fit_dst_gp(x, method = c("mle", "lmom"), ...)

fit_dst_gumbel(x, method = c("mle", "lmom"), ...)

fit_dst_lnorm(x, method = c("mle", "mge", "mme", "lmom", "lmom-log"), ...)

fit_dst_lp3(x, method = "lmom-log", ...)

fit_dst_nbinom(x, method = c("mle", "mme"), ...)

fit_dst_norm(x, method = c("mle", "mge", "mme", "lmom"), ...)

fit_dst_null(x, ...)

fit_dst_pearson3(x, method = "lmom", ...)

fit_dst_pois(x, method = c("mle", "mme", "lmom"), ...)

fit_dst_t(x, method = c("mle", "mge"), ...)

fit_dst_unif(x, method = c("mle", "mge", "mme", "lmom"), ...)

fit_dst_weibull(x, method = c("mle", "mge", "lmom"), ...)

```

Arguments

<code>x</code>	Numeric vector of observations to fit.
<code>method</code>	Estimation method to use. Available options depend on the distribution family and are enforced via <code>rlang::arg_match()</code> .
<code>...</code>	Additional arguments passed on to <code>fit_dst()</code> .

Details

Each helper simply forwards to `fit_dst()` with the associated `family` value indicated by the function suffix.

Some families do not have a unique fitting method where it is not applicable. These are the 'finite' ones (including 'degenerate' and 'empirical'), and the 'Null' distribution.

Missing Combinations:

Some combinations of families and methods are not supported, when one might think that they should be. These combinations are:

- 'gp' fit by 'mge'
- 'lp3' fit by 'mle' and 'mge'
- 'pearson3' fit by 'mle' and 'mge'
- 'gev' fit by 'mge'

They are not included because of how the `fitdistrplus::fitdist()` function looks for those distributional representations (i.e., `plp3()`, `dlp3()`, etc.): since we cannot guarantee that it will find the correct ones, these combinations are omitted for safety.

To elaborate, the current version of the wrapped `fitdistrplus::fitdist()` function looks for these representations starting from the `namespace:fitdistrplus` environment. `famish` cannot control what's found in that search path until the global environment, but that's too late because that's where behaviour becomes conditional on the user's actions. This is not a problem for other distributions because these are found in the `namespace:base` environment before the global environment is reached.

Value

A distribution object made by the 'distionary' package.

See Also

[fit_dst\(\)](#)

Examples

```
# Calls can be quite simple.
fit_dst_norm(1:10)
fit_dst_gumbel(2:6)

# Still have access to the functionality available through `fit_dst()`
x <- c(1, 4, 3, NA, 5)
fit_dst_lnorm(x, method = "lmom", na_action = "null")
fit_dst_lnorm(x, method = "lmom", na_action = "drop")

# Fitting by l-moments on the log-scale not the same as original scale.
fit_dst_lnorm(x, method = "lmom-log", na_action = "drop")
```

Description

Computes the asymmetric absolute loss commonly used to assess quantile forecasts. Lower scores indicate a better match between the estimated quantile and the observed value at level tau.

Usage

```
quantile_score(x, xhat, tau)
```

Arguments

- x Numeric vector of observed values.
- xhat Numeric vector of estimated quantiles.
- tau Numeric vector of quantile levels in (0, 1).

Details

The score minimises to zero when the observation equals the estimated quantile, so that smaller scores indicate a better fitting model. Positive residuals are penalised by a factor of tau, and negative residuals by tau - 1. This loss is also known as the stick function, check loss, asymmetric absolute deviation, or pinball loss.

For observation x, estimate x_hat, and level tau, the score (c.f. Gneiting, 2011) is

$$S_\tau(x, \hat{x}) = \begin{cases} \tau|x - \hat{x}|, & x \geq \hat{x}, \\ (1 - \tau)|x - \hat{x}|, & x < \hat{x}. \end{cases}$$

Vector recycling of all three arguments follows the rules in `vctrs::vec_recycle_common()`.

Value

Numeric vector of quantile scores corresponding to each element of the recycled inputs.

References

- Gneiting, T. (2011). Making and evaluating point forecasts. *Journal of the American Statistical Association*, 106(494), 746–762.

Examples

```
quantile_score(c(5, 15, 10), xhat = 7, tau = 0.8)
quantile_score(c(5, 15, 10), xhat = c(6, 19, 12), tau = c(0.2, 0.9, 0.5))
```

Description

Converts a numeric vector to its rank-based scores. For uscore() (uniform scores), values become roughly equally spaced between 0 and 1, keeping their order. nscore() calculates normal scores by spacing the uniform scores along a standard normal distribution; rpscore() calculates empirical return periods by spacing the uniform scores u by 1 / (1 - u).

Usage

```
uscore(x, pos = "Hazen", na.rm = FALSE)

nscore(x, pos = "Hazen", na.rm = FALSE)

rpscore(x, pos = "Hazen", na.rm = FALSE)
```

Arguments

<code>x</code>	Numeric vector.
<code>pos</code>	Positional adjustment for uniform scores. See Details. Can be a single numeric, or could be named after one of the proponents behind a choice of the numeric: "Weibull", "Beard", "Gringorten", or "Hazen".
<code>na.rm</code>	Logical indicating whether NA and NaN values should be removed from the output.

Details

Uniform scores are calculated by $(\text{rank}+a)/(n+1+2*a)$, where `rank` is the ranked `x` values, and `a` is the positional adjustment `pos`. Alternatively, could be named after an individual associated with a choice of `a`:

- Weibull (1939) proposed $a = 0$.
- Beard (1943) proposed $a = -0.31$.
- Gringorten (1963) proposed $a = -0.44$.
- Hazen (1914) proposed $a = -0.5$.

Value

Vector of uniform scores corresponding to values in `x`.

Author(s)

Thanks to Dr. Harry Joe for providing a starting framework for the `uscore()` function.

References

- Beard, L. R. (1943). Statistical analysis in hydrology. *Transactions of the American Society of Civil Engineers*, 108, 1110–1160.
- Gringorten, I. I. (1963). A plotting rule for extreme probability paper. *Journal of Geophysical Research*, 68(3), 813–814.
- Hazen, A. (1914). Storage to be provided in impounding reservoirs for municipal water supply. *Transactions of the American Society of Civil Engineers*, 77, 1539–1640.
- Weibull, W. (1939). A statistical theory of the strength of materials. *IVB-Handl.*, 151.

Examples

```
x <- c(0.3, 0.56, NA, 0.1, -12)
uscore(x)
uscore(x, pos = "Gringorten")
nscore(x, pos = -0.4)
rpscore(x)
```

Index

fit_dst, 2
fit_dst(), 6
fit_dst_bern (fit_dst_family_wrappers), 4
fit_dst_beta (fit_dst_family_wrappers), 4
fit_dst_cauchy (fit_dst_family_wrappers), 4
fit_dst_chisq (fit_dst_family_wrappers), 4
fit_dst_degenerate (fit_dst_family_wrappers), 4
fit_dst_empirical (fit_dst_family_wrappers), 4
fit_dst_exp (fit_dst_family_wrappers), 4
fit_dst_f (fit_dst_family_wrappers), 4
fit_dst_family_wrappers, 4
fit_dst_finite (fit_dst_family_wrappers), 4
fit_dst_gamma (fit_dst_family_wrappers), 4
fit_dst_geom (fit_dst_family_wrappers), 4
fit_dst_gev (fit_dst_family_wrappers), 4
fit_dst_gp (fit_dst_family_wrappers), 4
fit_dst_gumbel (fit_dst_family_wrappers), 4
fit_dst_lnorm (fit_dst_family_wrappers), 4
fit_dst_lp3 (fit_dst_family_wrappers), 4
fit_dst_nbinom (fit_dst_family_wrappers), 4
fit_dst_norm (fit_dst_family_wrappers), 4
fit_dst_norm(), 4
fit_dst_null (fit_dst_family_wrappers), 4
fit_dst_pearson3 (fit_dst_family_wrappers), 4
fit_dst_pois (fit_dst_family_wrappers), 4
fit_dst_t (fit_dst_family_wrappers), 4
fit_dst_unif (fit_dst_family_wrappers), 4
fit_dst_weibull (fit_dst_family_wrappers), 4
nscore (uscore), 7
quantile_score, 6
rpscore (uscore), 7
uscore, 7