Package 'GROAN'

November 28, 2022

Type Package

Title Genomic Regression Workbench

Version 1.3.1

Date 2022-11-28

Author Nelson Nazzicari & Filippo Biscarini

Maintainer Nelson Nazzicari <nelson.nazzicari@gmail.com>

Description Workbench for testing genomic regression accuracy on
(optionally noisy) phenotypes.

License GPL-3 | file LICENSE

Depends R (>= 2.10)

Imports plyr, rrBLUP

Suggests BGLR, e1071, ggplot2, knitr, randomForest, rmarkdown

VignetteBuilder knitr

Encoding UTF-8

LazyData TRUE

RoxygenNote 7.1.1

NeedsCompilation no

Repository CRAN

Date/Publication 2022-11-28 12:30:02 UTC

R topics documented:

---

addRegressor   *Add an extra regressor to a Workbench*

Description

This function adds a regressor to an existing GROAN.Workbench object.

Usage

```
addRegressor(wb, regressor, regressor.name = regressor, ...)
```

Arguments

wb	the GROAN.Workbench instance to be updated
regressor	regressor function
regressor.name	string that will be used in reports. Keep in mind that when deciding names.
...	extra parameters are passed to the regressor function

Value

an updated instance of the original GROAN.Workbench

See Also

createWorkbench GROAN.run

Examples

```
#creating a Workbench with all default arguments
wb = createWorkbench()
#adding a second regressor
wb = addRegressor(wb, regressor = phenoRegressor.dummy, regressor.name = 'dummy')

## Not run:
#trying to add again a regressor with the same name would result in a naming conflict error
wb = addRegressor(wb, regressor = phenoRegressor.dummy, regressor.name = 'dummy')
## End(Not run)
```

---

are.compatible	*Check two GROAN.NoisyDataSet for dimension compatibility*

Description

This function verifies that the two passed GROAN.NoisyDataSet objects have the same dimensions and can thus be used in the same experiment (typically training models on one and testing on the other). The function returns a TRUE/FALSE. In verbose mode the function also prints messages detailing the comparisons.

Usage

```
are.compatible(nds1, nds2, verbose = FALSE)
```

Arguments

nds1   the first GROAN.NoisyDataSet to be tested

nds2   the second GROAN.NoisyDataSet to be tested

verbose   boolean, if TRUE the function prints messages detailing the comparison.

Value

TRUE if the passed GROAN.NoisyDataSet are dimensionally compatible, FALSE otherwise

---

createNoisyDataset *Noisy Data Set Constructor*

Description

This function creates a GROAN.NoisyDataset object (or fails trying). The class will contain all noisy data set components: genotypes and/or covariance matrix, phenotypes, strata (optional), a noise injector function and its parameters.

You can have a general description of the created object using the overridden print.GROAN.NoisyDataset function.

Usage

```
createNoisyDataset(
  name,
  genotypes = NULL,
  covariance = NULL,
  phenotypes,
  strata = NULL,
  extraCovariates = NULL,
  ploidy = 2,
  allowFractionalGenotypes = FALSE,
  noiseInjector = noiseInjector.dummy,
  ...
)
```

Arguments

name	A string defining the dataset name, used later do identify this particular instance in reports and result files. It is advisable for it to be it somewhat meaningful (to you, GROAN simply reports it as it is)
genotypes	Matrix or dataframe containing SNP genotypes, one row per sample (N), one column per marker (M), 0/1/2 format (for diploids) or 0/1/2.../ploidy in case of polyploids
covariance	matrix of covariances between samples of this dataset. It is usually a square (NxN) matrix, but rectangular matrices (NxW) are accepted to incapsulate covariances between samples in this set and samples of other sets. Please note that some regression models expect the covariance to be square and will fail on rectangular ones
phenotypes	numeric array, N slots
strata	array of M slots, describing the strata each data point belongs to. This is used for stratified crossvalidation (see createWorkbench)
extraCovariates	
	dataframe of optional extra covariates (N lines, one column per extra covariate). Numeric ones will be normalized, string and categorical ones will be transformed in stub TRUE/FALSE variables (one per possible value, see model.matrix).

ploidy	number of haploid sets in the cell. Defaults to 2 (diploid).
allowFractionalGenotypes	
	if TRUE non-integer values for genotypes can be allowed. Defaults to FALSE
noiseInjector	name of a noise injector function, defaults to noiseInjector.dummy
...	further arguments are passed along to noiseInjector

Value

a GROAN.NoisyDataset object.

See Also

GROAN.run createWorkbench

Examples

```
#For more complete examples see the package vignette
#creating a noisy dataset with normal noise
nds = createNoisyDataset(
  name = 'PEA, normal noise',
  genotypes = GROAN.KI$SNPs,
  phenotypes = GROAN.KI$yield,
  noiseInjector = noiseInjector.norm,
  mean = 0,
  sd = sd(GROAN.KI$yield) * 0.5
)
```

---

createRunId	*Generate a random run id*

Description

This function returns a partially random alphanumeric string that can be used to identify a single run.

Usage

```
createRunId()
```

Value

a partially random alphanumeric string

---

createWorkbench *Workbench constructor*

Description

This function creates a GROAN.Workbench instance (or fails trying). The created object contains:
a) one regressor with its own specific configuration
b) the experiment parameters (number of repetitions, number of folds in case of crossvalidation, stratification...)
You can have a general description of the created object using the overridden print.GROAN.Workbench function.
It is possible to add other regressors to the created GROAN.Workbench object using addRegressor.
Once the GROAN.Workbench is created it must be passed to GROAN.run to start the experiment.

Usage

```
createWorkbench(
  folds = 10,
  reps = 5,
  stratified = FALSE,
  outfolder = NULL,
  outfile.name = "accuracy.csv",
  saveHyperParms = FALSE,
  saveExtraData = FALSE,
  regressor = phenoRegressor.rrBLUP,
  regressor.name = "default regressor",
  ...
)
```

Arguments

folds   number of folds for crossvalidation, defaults to 10. If NULL no crossvalidation happens and all training data will be used. In this case a second dataset, for test, is needed (see GROAN.run for details)

reps   number of times the whole test must be repeated, defaults to 5

stratified boolean indicating whether GROAN should take into account data strata. This have two effects. First, the crossvalidation becomes stratified, meaning that folds will be split so that training and test sets will contain the same proportions of each data stratum. Second, prediction accuracy will be assessed (also) by strata. If no strata are present in the GROAN.NoisyDataSet object and stratified==TRUE all samples will be considered belonging to the same strata ("dummyStrata"). If stratified is FALSE (the default) GROAN will simply ignore the strata, even if present in the GROAN.NoisyDataSet.

outfolder folder where to save the data. If NULL (the default) nothing will be saved. Filenames are standardized. If existing, accuracy and hyperparameter files will be

updated, otherwise are created. ExtraData cannot be updated, so unique file-
names will be generated using runId (see GROAN.run)

`outfile.name`	file name to be used to save the accuracies in a text file. Defaults to "accu-racy.csv". Ignored if `outfolder` is NULL
`saveHyperParms`	boolean indicating if the hyperparameters from regressor training should be saved in `outfolder`. Defaults to FALSE.
`saveExtraData`	boolean indicating if extradata from regressor training should be saved in `outfolder` as R objects (using the save function). Defaults to FALSE.
`regressor`	regressor function. Defaults to phenoRegressor.rrBLUP
`regressor.name`	string that will be used in reports. Keep that in mind when deciding names. Defaults to "default regressor"
`...`	extra parameter are passed to regressor function

Value

An instance of GROAN.Workbench

See Also

addRegressor GROAN.run createNoisyDataset

Examples

```
#creating a Workbench with all default arguments
wb1 = createWorkbench()
#another Workbench, with different crossvalidation
wb2 = createWorkbench(folds=5, reps=20)
#a third one, with a different regressor and extra parameters passed to regressor function
wb3 = createWorkbench(regressor=phenoRegressor.BGLR, regressor.name='Bayesian Lasso', type='BL')
```

---

getNoisyPhenotype   *Generate an instance of noisy phenotypes*

Description

Given a `Noisy Dataset` object, this function applies the noise injector to the data and returns a noisy version of it. It is useful for inspecting the noisy injector effects.

Usage

```
getNoisyPhenotype(nds)
```

Arguments

`nds`	a `Noisy Dataset` object

Value

the phenotypes contained in `nds` with added noise.

GROAN.AI   *Example data for pea AI lines*

Description

This list contains all data required to run GROAN examples. It refers to a pea experiment with 105 lines coming from a biparental Attika x Isard cross.

Usage

```
GROAN.AI
```

Format

A list with the following fields:

- *"GROAN.AI$yield"*: named array with 105 slots, containing data on grain yield [t/ha]
- *"GROAN.AI$SNPs"*: data frame with 105 rows and 647 variables. Each row is a pea AI line, each column a SNP marker. Values can either be 0, 1, or 2, representing the three possible genotypes (AA, Aa, and aa, respectively).
- *"GROAN.AI$kinship"*: square dataframe containing the realized kinships between all pairs of each of the 105 pea AI lines. Values were computed following the Astle & Balding metric. Higher values represent a higher degree of genetic similarity between lines. This metric mainly accounts for additive genetic contributions (as an alternative to dominant contributions).

Source

Annicchiarico et al., *GBS-Based Genomic Selection for Pea Grain Yield under Severe Terminal Drought*, The Plant Genome, Volume 10. doi: 10.3835/plantgenome2016.07.0072

GROAN.KI   *Example data for pea KI lines*

Description

This list contains all data required to run GROAN examples. It refers to a pea experiment with 103 lines coming from a biparental Kaspa x Isard cross.

Usage

```
GROAN.KI
```

Format

A list with the following fields:

- *"GROAN.KI$yield"*: named array with 103 slots, containing data on grain yield [t/ha]

- *"GROAN.KI$SNPs"*: data frame with 103 rows and 647 variables. Each row is a pea KI line, each column a SNP marker. Values can either be 0, 1, or 2, representing the three possible genotypes (AA, Aa, and aa, respectively).

- *"GROAN.KI$kinship"*: square dataframe containing the realized kinships between all pairs of each of the 103 pea KI lines. Values were computed following the Astle & Balding metric. Higher values represent a higher degree of genetic similarity between lines. This metric mainly accounts for additive genetic contributions (as an alternative to dominant contributions).

Source

Annicchiarico et al., *GBS-Based Genomic Selection for Pea Grain Yield under Severe Terminal Drought*, The Plant Genome, Volume 10. doi: 10.3835/plantgenome2016.07.0072

GROAN.pea.kinship   *[DEPRECATED]*

Description

This piece of data is deprecated and will be dismissed in next release. Please use GROAN.KI instead.

Usage

```
GROAN.pea.kinship
```

Format

A data frame with 103 rows and 103 variables. Row and column names are pea KI lines.

Source

Annicchiarico et al., *GBS-Based Genomic Selection for Pea Grain Yield under Severe Terminal Drought*, The Plant Genome, Volume 10. doi: 10.3835/plantgenome2016.07.0072

GROAN.pea.SNPs   *[DEPRECATED]*

Description

This piece of data is deprecated and will be dismissed in next release. Please use GROAN.KI
instead.

Usage

```
GROAN.pea.SNPs
```

Format

A data frame with 103 rows and 647 variables. Each row represent a pea KI line, each column a
SNP marker

Source

Annicchiarico et al., *GBS-Based Genomic Selection for Pea Grain Yield under Severe Terminal
Drought*, The Plant Genome, Volume 10. doi: 10.3835/plantgenome2016.07.0072

GROAN.pea.yield   *[DEPRECATED]*

Description

This piece of data is deprecated and will be dismissed in next release. Please use GROAN.KI
instead.

Usage

```
GROAN.pea.yield
```

Format

A named array with 103 slots.

Source

Annicchiarico et al., *GBS-Based Genomic Selection for Pea Grain Yield under Severe Terminal
Drought*, The Plant Genome, Volume 10. doi: 10.3835/plantgenome2016.07.0072

GROAN.run   *Compare Genomic Regressors on a Noisy Dataset*

Description

This function runs the experiment described in a [GROAN.Workbench](#) object, training regressor(s) on the data contained in a [GROAN.NoisyDataSet](#) object via parameter nds. The prediction accuracy is estimated either through crossvalidation or on separate test dataset supplied via parameter nds.test. It returns a GROAN.Result object, which have a [summary](#) function for quick inspection and can be fed to [plotResult](#) for visual comparisons. In case of crossvalidation the test dataset in the result object will report the [CV] suffix.

The experiment statistics are computed via [measurePredictionPerformance](#).

Each time this function is invoked it will refer to a runId - an alphanumeric string identifying each specific run. The runId is usually generated internally, but it is possible to pass it if the intention is to join results from different runs for analysis purposes.

Usage

```
GROAN.run(nds, wb, nds.test = NULL, run.id = createRunId())
```

Arguments

nds	a GROAN.NoisyDataSet object, containing the data (genotypes, phenotypes and so forth) plus a noiseInjector function
wb	a GROAN.Workbench object, containing the regressors to be tested together with the description of the experiment
nds.test	either a GROAN.NoisyDataSet or a list of GROAN.NoisyDataSet. The regression model(s) trained on nds will be tested on nds.test
run.id	an alphanumeric string identifying this specific run. If not passed it is generated using [createRunId](#)

Value

a GROAN.Result object

See Also

[measurePredictionPerformance](#)

Examples

```
## Not run:
#Complete examples are found in the vignette
vignette('GROAN.vignette', package='GROAN')

#Minimal example
#1) creating a noisy dataset with normal noise
```

```
nds = createNoisyDataset(
  name = 'PEA KI, normal noise',
  genotypes = GROAN.KI$SNPs,
  phenotypes = GROAN.KI$yield,
  noiseInjector = noiseInjector.norm,
  mean = 0,
  sd = sd(GROAN.KI$yield) * 0.5
)

#2) creating a GROAN.WorkBench using default regressor and crossvalidation preset
wb = createWorkbench()

#3) running the experiment
res = GROAN.run(nds, wb)

#4) examining results
summary(res)
plotResult(res)

## End(Not run)
```

---

measurePredictionPerformance
                            Measure Performance of a Prediction

---

Description

This method returns several performance metrics for the passed predictions.

Usage

```
measurePredictionPerformance(truevals, predvals)
```

Arguments

truevals	true values
predvals	predicted values

Value

A named array with the following fields:

pearson Pearson's correlation

spearman Spearmans' correlation (order based)

rmse Root Mean Square Error

mae Mean Absolute Error

coeff_det Coefficient of determination

ndcg10, ndcg20, ndcg50, ndcg100 mean Normalized Discounted Cumulative Gain with k equal to 0.1, 0.2, 0.5 and 1

---

ndcg   *Function to calculate mean Normalized Discounted Cumulative Gain (NDCG)*

Description

This function calculates NDCG from the vectors of observed and predicted values and the chosen proportion k of top observations (rank).

Usage

```
ndcg(y, y_hat, k = 0.2)
```

Arguments

y	true values
y_hat	predicted values
k	relevant proportion of rank (top)

Value

a real value in [0,1]

---

noiseInjector.dummy   *Noise Injector dummy function*

Description

This noise injector does not add any noise. Passed `phenotypes` are simply returned. This function is useful when comparing different regressors on the same dataset without the effect of extra injected noise.

Usage

```
noiseInjector.dummy(phenotypes)
```

Arguments

phenotypes	input phenotypes. This object will be returned without checks.

Value

the same passed `phenotypes`

See Also

Other noiseInjectors: noiseInjector.norm(), noiseInjector.swapper(), noiseInjector.unif()

Examples

```
phenos = rnorm(10)
all(phenos == noiseInjector.dummy(phenos)) #TRUE
```

---

noiseInjector.norm   *Inject Normal Noise*

Description

This function adds to the passed `phenotypes` array noise sampled from a normal distribution with the specified mean and standard deviation.
The function can interest the totality of the passed phenotype array or a random subset of it (commanded by `subset` parameter).

Usage

```
noiseInjector.norm(phenotypes, mean = 0, sd = 1, subset = 1)
```

Arguments

phenotypes	an array of numbers.
mean	mean of the normal distribution.
sd	standard deviation of the normal distribution.
subset	integer in [0,1], the proportion of original dataset to be injected

Value

An array, of the same size as phenotypes, where normal noise has been added to the original phenotype values.

See Also

Other noiseInjectors: noiseInjector.dummy(), noiseInjector.swapper(), noiseInjector.unif()

Examples

```
#a sinusoid signal
phenos = sin(seq(0,5, 0.1))
plot(phenos, type='p', pch=16, main='Original (black) vs. Injected (red), 100% affected')

#adding normal noise to all samples
phenos.noise = noiseInjector.norm(phenos, sd = 0.2)
points(phenos.noise, type='p', col='red')
```

```
#adding noise only to 30% of the samples
plot(phenos, type='p', pch=16, main='Original (black) vs. Injected (red), 30% affected')
phenos.noise.subset = noiseInjector.norm(phenos, sd = 0.2, subset = 0.3)
points(phenos.noise.subset, type='p', col='red')
```

---

noiseInjector.swapper   *Swap phenotypes between samples*

Description

This function introduces swap noise, i.e. a number of couples of samples will have their phenotypes swapped.
The number of couples is computed so that the total fraction of interested phenotypes approximates subset.

Usage

```
noiseInjector.swapper(phenotypes, subset = 0.1)
```

Arguments

phenotypes   an array of numbers

subset   fraction of phenotypes to be interested by noise.

Value

the same passed phenotypes, but with some elements swapped

See Also

Other noiseInjectors: noiseInjector.dummy(), noiseInjector.norm(), noiseInjector.unif()

Examples

```
#a set of phenotypes
phenos = 1:10
#swapping two elements
phenos.sw2 = noiseInjector.swapper(phenos, 0.2)
#swapping four elements
phenos.sw4 = noiseInjector.swapper(phenos, 0.4)
#swapping four elements again, since 30% of 10 elements
#is rounded to 4 (two couples)
phenos.sw4.again = noiseInjector.swapper(phenos, 0.3)
```

---

noiseInjector.unif   *Inject Uniform Noise*

Description

This function adds to the passed `phenotypes` array noise sampled from a uniform distribution with the specified range.
The function can interest the totality of the passed phenotype array or a random subset of it (commanded by `subset` parameter).

Usage

```
noiseInjector.unif(phenotypes, min = 0, max = 1, subset = 1)
```

Arguments

phenotypes	an array of numbers.
min, max	lower and upper limits of the distribution. Must be finite.
subset	integer in [0,1], the proportion of original dataset to be injected

Value

An array, of the same size as phenotypes, where uniform noise has been added to the original phenotype values.

See Also

Other noiseInjectors: [noiseInjector.dummy](), [noiseInjector.norm](), [noiseInjector.swapper]()

Examples

```
#a sinusoid signal
phenos = sin(seq(0,5, 0.1))
plot(phenos, type='p', pch = 16, main='Original (black) vs. Injected (red), 100% affected')

#adding normal noise to all samples
phenos.noise = noiseInjector.unif(phenos, min=0.1, max=0.3)
points(phenos.noise, type='p', col='red')

#adding noise only to 30% of the samples
plot(phenos, type='p', pch = 16, main='Original (black) vs. Injected (red), 30% affected')
phenos.noise.subset = noiseInjector.unif(phenos, min=0.1, max=0.3, subset = 0.3)
points(phenos.noise.subset, type='p', col='red')
```

phenoRegressor.BGLR   *Regression using BGLR package*

Description

This is a wrapper around [BGLR](#). As such, it won't work if BGLR package is not installed.
Genotypes are modeled using the specified `type`. If `type` is 'RKHS' (and only in this case) the
covariance/kinship matrix `covariances` is required, and it will be modeled as matrix K in BGLR
terms. In all other cases genotypes and covariances are put in the model as X matrices.
Extra covariates, if present, are modeled as FIXED effects.

Usage

```
phenoRegressor.BGLR(
  phenotypes,
  genotypes,
  covariances,
  extraCovariates,
  type = c("FIXED", "BRR", "BL", "BayesA", "BayesB", "BayesC", "RKHS"),
  ...
)
```

Arguments

phenotypes	phenotypes, a numeric array (n x 1), missing values are predicted
genotypes	SNP genotypes, one row per phenotype (n), one column per marker (m), values in 0/1/2 for diploids or 0/1/2/...ploidy for polyploids. Can be NULL if `covariances` is present.
covariances	square matrix (n x n) of covariances. Can be NULL if `genotypes` is present.
extraCovariates	
	extra covariates set, one row per phenotype (n), one column per covariate (w). If NULL no extra covariates are considered.
type	character literal, one of the following: FIXED (Flat prior), BRR (Gaussian prior), BL (Double-Exponential prior), BayesA (scaled-t prior), BayesB (two component mixture prior with a point of mass at zero and a scaled-t slab), BayesC (two component mixture prior with a point of mass at zero and a Gaussian slab)
...	extra parameters are passed to [BGLR](#)

Value

The function returns a list with the following fields:

- `predictions` : an array of (n) predicted phenotypes, with NAs filled and all other positions repredicted (useful for calculating residuals)
- `hyperparams` : empty, returned for compatibility
- `extradata` : list with information on trained model, coming from [BGLR](#)

See Also

BGLR

Other phenoRegressors: phenoRegressor.RFR(), phenoRegressor.SVR(), phenoRegressor.dummy(),
phenoRegressor.rrBLUP(), phenoregressor.BGLR.multikinships()

Examples

```
## Not run:
#using the GROAN.KI dataset, we regress on the dataset and predict the first ten phenotypes
phenos = GROAN.KI$yield
phenos[1:10]  = NA

#calling the regressor with Bayesian Lasso
results = phenoRegressor.BGLR(
  phenotypes = phenos,
  genotypes = GROAN.KI$SNPs,
  covariances = NULL,
  extraCovariates = NULL,
  type = 'BL', nIter = 2000 #BGLR-specific parameters
)

#examining the predictions
plot(GROAN.KI$yield, results$predictions,
     main = 'Train set (black) and test set (red) regressions',
     xlab = 'Original phenotypes', ylab = 'Predicted phenotypes')
points(GROAN.KI$yield[1:10], results$predictions[1:10], pch=16, col='red')

#printing correlations
test.set.correlation  = cor(GROAN.KI$yield[1:10], results$predictions[1:10])
train.set.correlation = cor(GROAN.KI$yield[-(1:10)], results$predictions[-(1:10)])
writeLines(paste(
  'test-set correlation :', test.set.correlation,
  '\ntrain-set correlation:', train.set.correlation
))

## End(Not run)
```

---

phenoregressor.BGLR.multikinships
Multi-matrix GBLUP using BGLR

Description

This regressor implements Genomic BLUP using Bayesian methods from BGLR package, but allows to use more than one covariance matrix.

Usage

```
phenoregressor.BGLR.multikinships(
  phenotypes,
  genotypes = NULL,
  covariances,
  extraCovariates,
  type = "RKHS",
  ...
)
```

Arguments

phenotypes	phenotypes, a numeric array (n x 1), missing values are predicted
genotypes	added for compatibility with the other GROAN regressors, must be NULL
covariances	square matrix (n x n) of covariances.
extraCovariates	the extra covariance matrices to be added in the GBLUP model, collated in a single matrix-like structure, with optionally first column as an ignored intercept (supported for compatibility). See details, below.
type	character literal, one of the following: FIXED (Flat prior), BRR (Gaussian prior), BL (Double-Exponential prior), BayesA (scaled-t prior), BayesB (two component mixture prior with a point of mass at zero and a scaled-t slab), BayesC (two component mixture prior with a point of mass at zero and a Gaussian slab), RKHS (Gaussian processes, default)
...	extra parameters are passed to [BGLR](#)

Details

In its simplest form, GBLUP is defined as:

$$y = 1\mu + Zu + e$$

with

$$var(y) = K\sigma_u^2 + I\sigma_e^2$$

Where μ is the overall mean, K is the incidence matrix relating individual weights u to y, and e is a vector of residuals with zero mean and covariance matrix $I\sigma_e^2$

It is possible to extend the above model to include different types of kinship matrices, each capturing different links between genotypes and phenotypes:

$$y = 1\mu + Z1u1 + Z2u2 + \cdots + e$$

with

$$var(y) = K1\sigma_u1^2 + K2\sigma_u2^2 + \cdots + I\sigma_e^2$$

This function receives the first kinship matrix $K1$ via the `covariances` argument and an arbitrary number of extra matrices via the `extraCovariates` built as follow:

```
#given the following defined variables
y = <some values, Nx1 array>
K1 = <NxN kinship matrix>
K2 = <another NxN kinship matrix>
K3 = <a third NxN kinship matrix>

#invoking the multi kinship GBLUP
y_hat = phenoregressor.BGLR.multikinships(
  phenotypes = y,
  covariances = K1,
  extraCovariates = cbind(K2, K3)
)
```

Value

The function returns a list with the following fields:

- predictions : an array of (n) predicted phenotypes, with NAs filled and all other positions repredicted (useful for calculating residuals)
- hyperparams : empty, returned for compatibility
- extradata : list with information on trained model, coming from BGLR

See Also

BGLR

Other phenoRegressors: phenoRegressor.BGLR(), phenoRegressor.RFR(), phenoRegressor.SVR(), phenoRegressor.dummy(), phenoRegressor.rrBLUP()

---

phenoRegressor.dummy   *Regression dummy function*

Description

This function is for development purposes. It returns, as "predictions", an array of random numbers. It accept the standard inputs and produces a formally correct output. It is, obviously, quite fast.

Usage

```
phenoRegressor.dummy(phenotypes, genotypes, covariances, extraCovariates)
```

Arguments

phenotypes   phenotypes, numeric array (n x 1), missing values are predicted

genotypes   SNP genotypes, one row per phenotype (n), one column per marker (m), values in 0/1/2 for diploids or 0/1/2/...ploidy for polyploids. Can be NULL if covariances is present.

covariances   square matrix (n x n) of covariances. Can be NULL if genotypes is present.

extraCovariates

                 extra covariates set, one row per phenotype (n), one column per covariate (w).
                 If NULL no extra covariates are considered.

Value

The function should return a list with the following fields:

- `predictions` : an array of (k) predicted phenotypes

- `hyperparams` : named array of hyperparameters selected during training

- `extradata` : any extra information

See Also

Other phenoRegressors: `phenoRegressor.BGLR()`, `phenoRegressor.RFR()`, `phenoRegressor.SVR()`, `phenoRegressor.rrBLUP()`, `phenoregressor.BGLR.multikinships()`

Examples

```
#genotypes are not really investigated. Only
#number of test phenotypes is used.
phenoRegressor.dummy(
  phenotypes = c(1:10, NA, NA, NA),
  genotypes = matrix(nrow = 13, ncol=30)
)
```

---

phenoRegressor.RFR   *Random Forest Regression using package randomForest*

Description

This is a wrapper around randomForest and related functions. As such, this function will not work if randomForest package is not installed. There is no distinction between regular covariates (genotypes) and extra covariates (fixed effects) in random forest. If extra covariates are passed, they are put together with genotypes, side by side. Same thing happens with covariances matrix. This can bring to the scientifically questionable but technically correct situation of regressing on a big matrix made of SNP genotypes, covariances and other covariates, all collated side by side. The function makes no distinction, and it's up to the user understand what is correct in each specific experiment.

WARNING: this function can be *very* slow, especially when called on thousands of SNPs.

Usage

```
phenoRegressor.RFR(
  phenotypes,
  genotypes,
  covariances,
  extraCovariates,
  ntree = ceiling(length(phenotypes)/5),
  ...
)
```

Arguments

phenotypes	phenotypes, a numeric array (n x 1), missing values are predicted
genotypes	SNP genotypes, one row per phenotype (n), one column per marker (m), values in 0/1/2 for diploids or 0/1/2/...ploidy for polyploids. Can be NULL if `covariances` is present.
covariances	square matrix (n x n) of covariances. Can be NULL if genotypes is present.
extraCovariates	extra covariates set, one row per phenotype (n), one column per covariate (w). If NULL no extra covariates are considered.
ntree	number of trees to grow, defaults to a fifth of the number of samples (rounded up). As per `randomForest` documentation, it should not be set to too small a number, to ensure that every input row gets predicted at least a few times
...	any extra parameter is passed to randomForest::randomForest()

Value

The function returns a list with the following fields:

- predictions : an array of (k) predicted phenotypes
- hyperparams : named vector with the following keys: ntree (number of grown trees) and mtry (number of variables randomly sampled as candidates at each split)
- extradata : the object returned by randomForest::randomForest(), containing the full trained forest and the used parameters

See Also

[randomForest](#)

Other phenoRegressors: [phenoRegressor.BGLR()](#), [phenoRegressor.SVR()](#), [phenoRegressor.dummy()](#), [phenoRegressor.rrBLUP()](#), [phenoregressor.BGLR.multikinships()](#)

Examples

```
## Not run:
#using the GROAN.KI dataset, we regress on the dataset and predict the first ten phenotypes
phenos = GROAN.KI$yield
phenos[1:10]  = NA
```

```
#calling the regressor with random forest
results = phenoRegressor.RFR(
  phenotypes = phenos,
  genotypes = GROAN.KI$SNPs,
  covariances = NULL,
  extraCovariates = NULL,
  ntree = 20,
  mtry = 200 #randomForest-specific parameters
)

#examining the predictions
plot(GROAN.KI$yield, results$predictions,
     main = 'Train set (black) and test set (red) regressions',
     xlab = 'Original phenotypes', ylab = 'Predicted phenotypes')
points(GROAN.KI$yield[1:10], results$predictions[1:10], pch=16, col='red')

#printing correlations
test.set.correlation  = cor(GROAN.KI$yield[1:10], results$predictions[1:10])
train.set.correlation = cor(GROAN.KI$yield[-(1:10)], results$predictions[-(1:10)])
writeLines(paste(
  'test-set correlation :', test.set.correlation,
  '\ntrain-set correlation:', train.set.correlation
))

## End(Not run)
```

---

phenoRegressor.rrBLUP   *SNP-BLUP or G-BLUP using rrBLUP package*

Description

This is a wrapper around rrBLUP function `mixed.solve`. It can either work with genotypes (in form of a SNP matrix) or with kinships (in form of a covariance matrix). In the first case the function will implement a SNP-BLUP, in the second a G-BLUP. An error is returned if both SNPs and covariance matrix are passed.

In rrBLUP terms, genotypes are modeled as random effects (matrix Z), covariances as matrix K, and extra covariates, if present, as fixed effects (matrix X).

Please note that this function won't work if rrBLUP package is not installed.

Usage

```
phenoRegressor.rrBLUP(
  phenotypes,
  genotypes = NULL,
  covariances = NULL,
  extraCovariates = NULL,
  ...
)
```

Arguments

phenotypes	phenotypes, a numeric array (n x 1), missing values are predicted
genotypes	SNP genotypes, one row per phenotype (n), one column per marker (m), values in 0/1/2 for diploids or 0/1/2/...ploidy for polyploids. Can be NULL if `covariances` is present.
covariances	square matrix (n x n) of covariances.
extraCovariates	
	optional extra covariates set, one row per phenotype (n), one column per covariate (w). If NULL no extra covariates are considered.
...	extra parameters are passed to rrBLUP::mixed.solve

Value

The function returns a list with the following fields:

- `predictions` : an array of (k) predicted phenotypes
- `hyperparams` : named vector with the following keys: Vu, Ve, beta, LL
- `extradata` : list with information on trained model, coming from mixed.solve

See Also

mixed.solve

Other phenoRegressors: `phenoRegressor.BGLR()`, `phenoRegressor.RFR()`, `phenoRegressor.SVR()`, `phenoRegressor.dummy()`, `phenoregressor.BGLR.multikinships()`

Examples

```
## Not run:
#using the GROAN.KI dataset, we regress on the dataset and predict the first ten phenotypes
phenos = GROAN.KI$yield
phenos[1:10]  = NA

#calling the regressor with ridge regression BLUP on SNPs and kinship
results.SNP.BLUP = phenoRegressor.rrBLUP(
  phenotypes = phenos,
  genotypes = GROAN.KI$SNPs,
  SE = TRUE, return.Hinv = TRUE #rrBLUP-specific parameters
)
results.G.BLUP = phenoRegressor.rrBLUP(
  phenotypes = phenos,
  covariances = GROAN.KI$kinship,
  SE = TRUE, return.Hinv = TRUE #rrBLUP-specific parameters
)

#examining the predictions
plot(GROAN.KI$yield, results.SNP.BLUP$predictions,
     main = '[SNP-BLUP] Train set (black) and test set (red) regressions',
     xlab = 'Original phenotypes', ylab = 'Predicted phenotypes')
abline(a=0, b=1)
```

```
points(GROAN.KI$yield[1:10], results.SNP.BLUP$predictions[1:10], pch=16, col='red')

plot(GROAN.KI$yield, results.G.BLUP$predictions,
     main = '[G-BLUP] Train set (black) and test set (red) regressions',
     xlab = 'Original phenotypes', ylab = 'Predicted phenotypes')
abline(a=0, b=1)
points(GROAN.KI$yield[1:10], results.G.BLUP$predictions[1:10], pch=16, col='red')

#printing correlations
correlations = data.frame(
  model = 'SNP-BLUP',
  test_set_correlations = cor(GROAN.KI$yield[1:10], results.SNP.BLUP$predictions[1:10]),
  train_set_correlations = cor(GROAN.KI$yield[-(1:10)], results.SNP.BLUP$predictions[-(1:10)])
)
correlations = rbind(correlations, data.frame(
  model = 'G-BLUP',
  test_set_correlations = cor(GROAN.KI$yield[1:10], results.G.BLUP$predictions[1:10]),
  train_set_correlations = cor(GROAN.KI$yield[-(1:10)], results.G.BLUP$predictions[-(1:10)])
))
print(correlations)

## End(Not run)
```

---

phenoRegressor.SVR   *Support Vector Regression using package e1071*

Description

This is a wrapper around several functions from e1071 package (as such, it won't work if e1071 package is not installed). This function implements Support Vector Regressions, meaning that the data points are projected in a transformed higher dimensional space where linear regression is possible.

phenoRegressor.SVR can operate in three modes: run, train and tune.
In **run** mode you need to pass the function an already tuned/trained SVR model, typically obtained either directly from e1071 functions (e.g. from svm, best.svm and so forth) or from a previous run of phenoRegressor.SVR in a different mode. The passed model is applied to the passed dataset and predictions are returned.
In **train** mode a SVR model will be trained on the passed dataset using the passed hyper parameters. The trained model will then be used for predictions.
In **tune** mode you need to pass one or more sets of hyperparameters. The best combination of hyperparameters will be selected through crossvalidation. The best performing SVR model will be used for final predictions. This mode can be very slow.

There is no distinction between regular covariates (genotypes) and extra covariates (fixed effects) in Support Vector Regression. If extra covariates are passed, they are put together with genotypes, side by side. Same thing happens with covariances matrix. This can bring to the scientifically questionable but technically correct situation of regressing on a big matrix made of SNP genotypes, covariances and other covariates, all collated side by side. The function makes no distinction, and

it's up to the user understand what is correct in each specific experiment.

Usage

```
phenoRegressor.SVR(
  phenotypes,
  genotypes,
  covariances,
  extraCovariates,
  mode = c("tune", "train", "run"),
  tuned.model = NULL,
  scale.pheno = TRUE,
  scale.geno = FALSE,
  ...
)
```

Arguments

phenotypes	phenotypes, a numeric array (n x 1), missing values are predicted
genotypes	SNP genotypes, one row per phenotype (n), one column per marker (m), values in 0/1/2 for diploids or 0/1/2/...ploidy for polyploids. Can be NULL if `covariances` is present.
covariances	square matrix (n x n) of covariances. Can be NULL if genotypes is present.
extraCovariates	
	extra covariates set, one row per phenotype (n), one column per covariate (w). If NULL no extra covariates are considered.
mode	this parameter decides what will happen with the passed dataset • mode = "tune" : hyperparameters will be tuned on a grid (you may want to specify its values using extra params) with a call to `e1071::tune.svm`. Use this option if you have no idea about the optimal choice of hyperparameters. This mode can be very slow. • mode = "train" : an SVR will be trained on the train dataset using the passed hyperparameters (if you know them). This more invokes `e1071::train` • mode = "run" : you already have a tuned and trained SVR (put it into `tuned.model`) and want to use it. The fastest mode.
tuned.model	a tuned and trained SVR to be used for prediction. This object is only used if `mode` is equal to "run".
scale.pheno	if TRUE (default) the phenotypes will be scaled and centered (before tuning or before applying the passed tuned model).
scale.geno	if TRUE the genotypes will be scaled and centered (before tuning or before applying the passed tuned model. It is usually not a good idea, since it leads to worse results. Defaults to FALSE.
...	all extra parameters are passed to `e1071::svm` or `e1071::tune.svm`

Value

The function returns a list with the following fields:

- `predictions` : an array of (n) predicted phenotypes

- `hyperparams` : named vector with the following keys: gamma, cost, coef0, nu, epsilon. Some of the values may not make sense given the selected model, and will contain default values from e1071 library.

- `extradata` : depending on `mode` parameter, `extradata` will contain one of the following: 1) a SVM object returned by e1071::tune.svm, containing both the best performing model and the description of the training process 2) a newly trained SVR model 3) the same object passed as `tuned.model`

See Also

svm, tune.svm, best.svm from e1071 package

Other phenoRegressors: `phenoRegressor.BGLR()`, `phenoRegressor.RFR()`, `phenoRegressor.dummy()`, `phenoRegressor.rrBLUP()`, `phenoregressor.BGLR.multikinships()`

Examples

```
## Not run:
### WARNING ###
#The 'tuning' part of the example can take quite some time to run,
#depending on the computational power.

#using the GROAN.KI dataset, we regress on the dataset and predict the first ten phenotypes
phenos = GROAN.KI$yield
phenos[1:10] = NA

#--------- TUNE ---------
#tuning the SVR on a grid of hyperparameters
results.tune = phenoRegressor.SVR(
  phenotypes = phenos,
  genotypes = GROAN.KI$SNPs,
  covariances = NULL,
  extraCovariates = NULL,
  mode = 'tune',
  kernel = 'linear', cost = 10^(-3:+3) #SVR-specific parameters
)

#examining the predictions
plot(GROAN.KI$yield, results.tune$predictions,
     main = 'Mode = TUNING\nTrain set (black) and test set (red) regressions',
     xlab = 'Original phenotypes', ylab = 'Predicted phenotypes')
points(GROAN.KI$yield[1:10], results.tune$predictions[1:10], pch=16, col='red')

#printing correlations
test.set.correlation  = cor(GROAN.KI$yield[1:10], results.tune$predictions[1:10])
train.set.correlation = cor(GROAN.KI$yield[-(1:10)], results.tune$predictions[-(1:10)])
writeLines(paste(
```

```
  'test-set correlation :', test.set.correlation,
  '\ntrain-set correlation:', train.set.correlation
))

#--------- TRAIN ---------
#training the SVR, hyperparameters are given
results.train = phenoRegressor.SVR(
  phenotypes = phenos,
  genotypes = GROAN.KI$SNPs,
  covariances = NULL,
  extraCovariates = NULL,
  mode = 'train',
  kernel = 'linear', cost = 0.01 #SVR-specific parameters
)

#examining the predictions
plot(GROAN.KI$yield, results.train$predictions,
     main = 'Mode = TRAIN\nTrain set (black) and test set (red) regressions',
     xlab = 'Original phenotypes', ylab = 'Predicted phenotypes')
points(GROAN.KI$yield[1:10], results.train$predictions[1:10], pch=16, col='red')

#printing correlations
test.set.correlation  = cor(GROAN.KI$yield[1:10], results.train$predictions[1:10])
train.set.correlation = cor(GROAN.KI$yield[-(1:10)], results.train$predictions[-(1:10)])
writeLines(paste(
  'test-set correlation :', test.set.correlation,
  '\ntrain-set correlation:', train.set.correlation
))

#--------- RUN ---------
#we recover the trained model from previous run, predictions will be exactly the same
results.run = phenoRegressor.SVR(
  phenotypes = phenos,
  genotypes = GROAN.KI$SNPs,
  covariances = NULL,
  extraCovariates = NULL,
  mode = 'run',
  tuned.model = results.train$extradata
)

#examining the predictions
plot(GROAN.KI$yield, results.run$predictions,
     main = 'Mode = RUN\nTrain set (black) and test set (red) regressions',
     xlab = 'Original phenotypes', ylab = 'Predicted phenotypes')
points(GROAN.KI$yield[1:10], results.run$predictions[1:10], pch=16, col='red')

#printing correlations
test.set.correlation  = cor(GROAN.KI$yield[1:10], results.run$predictions[1:10])
train.set.correlation = cor(GROAN.KI$yield[-(1:10)], results.run$predictions[-(1:10)])
writeLines(paste(
  'test-set correlation :', test.set.correlation,
  '\ntrain-set correlation:', train.set.correlation
))
```

```
## End(Not run)
```

---

plotResult   *Plot results of a run*

Description

This function uses ggplot2 package (which must be installed) to graphically render the result of a run. The function receive as input the output of GROAN.run and returns a ggplot2 object (that can be further customized). Currently implemented types of plot are:

- box : boxplot, showing the distribution of repetitions. See geom_boxplot
- bar : barplot, showing the average over repetitions. See stat_summary
- bar_conf95 : same as 'bar', but with 95% confidence intervals

Usage

```
plotResult(
  res,
 variable = c("pearson", "spearman", "rmse", "time_per_fold", "coeff_det", "mae"),
  x.label = c("both", "train_only", "test_only"),
  plot.type = c("box", "bar", "bar_conf95"),
  strata = c("no_strata", "avg_strata", "single")
)
```

Arguments

res	a result data frame containing the output of GROAN.run
variable	name of the variable to be used as y values
x.label	select what to put on x-axis between both train and test dataset (default), train dataset only or test dataset only
plot.type	a string indicating the type of plot to be obtained
strata	string determining behaviour toward strata. If 'no_strata' will plot accuracies not considering strata. If 'avg_strata' will average single strata accuracies. If 'single' each strata will be represented separately.

Value

a ggplot2 object

---

print.GROAN.NoisyDataset
                    Print a GROAN Noisy Dataset object

Description

Short description for class GROAN.NoisyDataset, created with createNoisyDataset.

Usage

```
## S3 method for class 'GROAN.NoisyDataset'
print(x, ...)
```

Arguments

x   object of class GROAN.NoisyDataset.

...   ignored, put here to match S3 function signature

Value

This function returns the original GROAN.NoisyDataset object invisibly (via invisible(x))

---

print.GROAN.Workbench   *Print a GROAN Workbench object*

Description

Short description for class GROAN.Workbench, created with createWorkbench.

Usage

```
## S3 method for class 'GROAN.Workbench'
print(x, ...)
```

Arguments

x   object of class GROAN.Workbench.

...   ignored, put here to match S3 function signature

Value

This function returns the original GROAN.Workbench object invisibly (via invisible(x))

summary.GROAN.NoisyDataset

Summary for GROAN Noisy Dataset object

Description

Returns a dataframe with some description of an object created with createNoisyDataset.

Usage

```
## S3 method for class 'GROAN.NoisyDataset'
summary(object, ...)
```

Arguments

object   instance of class GROAN.NoisyDataset.

...   additional arguments ignored, added for compatibility to generic summary func-
                tion

Value

a data frame with GROAN.NoisyDataset stats.

summary.GROAN.Result   *Summary of GROAN.Result*

Description

Performance metrics are averaged over repetitions, so that a data.frame is produced with one row
per dataset/regressor/extra_covariates/strata/samples/markers/folds combination.

Usage

```
## S3 method for class 'GROAN.Result'
summary(object, ...)
```

Arguments

object   an object returned from GROAN.run

...   additional arguments ignored, added for compatibility to generic summary func-
                tion

Value

a data.frame with averaged statistics

Index