

# Package ‘keyed’

February 6, 2026

**Title** Explicit Key Assumptions for Flat-File Data

**Version** 0.1.3

**Language** en-US

**Description** Helps make implicit data assumptions explicit by attaching keys to flat-file data that error when those assumptions are violated. Designed for CSV-first workflows without database infrastructure or version control. Provides key definition, assumption checks, join diagnostics, and optional drift detection against reference snapshots.

**License** MIT + file LICENSE

**URL** <https://github.com/gcol33/keyed>, <https://gillescolling.com/keyed/>

**BugReports** <https://github.com/gcol33/keyed/issues>

**Encoding** UTF-8

**RoxygenNote** 7.3.3

**Depends** R (>= 4.1.0)

**Imports** cli, dplyr (>= 1.0.0), digest, pillar, rlang (>= 1.0.0), tibble, vctrs (>= 0.5.0)

**Suggests** joinspy, knitr, rmarkdown, testthat (>= 3.0.0), uuid

**VignetteBuilder** knitr

**Config/testthat/edition** 3

**NeedsCompilation** no

**Author** Gilles Colling [aut, cre, cph] (ORCID: <<https://orcid.org/0000-0003-3070-6066>>)

**Maintainer** Gilles Colling <gilles.colling051@gmail.com>

**Repository** CRAN

**Date/Publication** 2026-02-06 19:50:08 UTC

## Contents

add_id . . . . .	2
bind_id . . . . .	3
bind_keyed . . . . .	4
check_drift . . . . .	5
check_id . . . . .	5
check_id_disjoint . . . . .	6
clear_all_snapshots . . . . .	7
clear_snapshot . . . . .	7
commit_keyed . . . . .	8
compare_ids . . . . .	9
compare_keys . . . . .	9
compare_structure . . . . .	10
diagnose_join . . . . .	11
extend_id . . . . .	12
find_duplicates . . . . .	12
get_id . . . . .	13
get_key_cols . . . . .	14
has_id . . . . .	14
has_key . . . . .	15
key . . . . .	15
key_is_valid . . . . .	16
key_status . . . . .	16
list_snapshots . . . . .	17
lock_complete . . . . .	17
lock_coverage . . . . .	18
lock_no_na . . . . .	18
lock_nrow . . . . .	19
lock_unique . . . . .	20
make_id . . . . .	20
remove_id . . . . .	21
summary.keyed_df . . . . .	22
unkey . . . . .	22

## Index

**add\_id**

*Add identity column*

### Description

Adds a stable UUID column to each row. This is an opt-in feature for tracking row lineage through transformations.

### Usage

```
add_id(.data, .id = ".id", .overwrite = FALSE)
```

## Arguments

- |            |  |
|------------|--|
| .data      | A data frame.  |
| .id        | Column name for the ID (default: ".id").   |
| .overwrite | If TRUE, overwrite existing ID column. If FALSE (default), error if column exists. |

## Details

IDs are generated using a hash of row content plus a random salt, making them stable for identical rows within a session but unique across different data frames.

If the `uuid` package is available, it will be used for true UUIDs. Otherwise, a hash-based ID is generated.

## Value

Data frame with ID column added.

## Examples

```
df <- data.frame(x = 1:3, y = c("a", "b", "c"))
df <- add_id(df)
df
```

---

bind\_id

*Bind data frames with ID handling*

---

## Description

Binds data frames while properly handling ID columns. Checks for overlapping IDs, combines the data, and fills in missing IDs.

## Usage

```
bind_id(..., .id = ".id")
```

## Arguments

- |     |                                       |
|-----|---------------------------------------|
| ... | Data frames to bind.                  |
| .id | Column name for IDs (default: ".id"). |

## Details

This function:

1. Checks if IDs overlap between datasets (warns if so)
2. Binds rows using `dplyr::bind_rows()`
3. Fills missing IDs using `extend_id()`

Use this instead of `dplyr::bind_rows()` when working with ID columns.

## Value

Combined data frame with valid IDs for all rows.

## Examples

```
df1 <- add_id(data.frame(x = 1:3))
df2 <- data.frame(x = 4:6)
combined <- bind_id(df1, df2)
```

`bind_keyed`

*Bind rows of keyed data frames*

## Description

Bind keyed data frames

## Usage

```
bind_keyed(..., .id = NULL)
```

## Arguments

- ... Data frames to bind.
- .id Optional column name to identify source.

## Details

Wrapper for `dplyr::bind_rows()` that attempts to preserve key metadata.

## Value

A keyed data frame if key is preserved and unique, otherwise tibble.

---

check_drift	<i>Check for drift from committed snapshot</i>
-------------	--

---

## Description

Compares current data against its committed reference snapshot. Returns diagnostic information about changes.

## Usage

```
check_drift(.data, reference = NULL)
```

## Arguments

- |           |  |
|-----------|--|
| .data     | A data frame with a snapshot reference.  |
| reference | Optional content hash to compare against. If NULL, uses the attached snapshot reference. |

## Value

A drift report (class `keyed_drift_report`), or NULL if no snapshot found.

## Examples

```
df <- key(data.frame(id = 1:3, x = c("a", "b", "c")), id)
df <- commit_keyed(df)

# Modify the data
df$x[1] <- "modified"
check_drift(df)
```

---

check_id	<i>Check ID integrity</i>
----------	---------------------------

---

## Description

Validates ID column for common issues: missing values, duplicates, and suspicious formats.

## Usage

```
check_id(.data, .id = ".id")
```

## Arguments

- |       |                               |
|-------|-------------------------------|
| .data | A data frame with ID column.  |
| .id   | Column name (default: ".id"). |

**Value**

Invisibly returns a list with:

- `valid`: TRUE if no issues found
- `n_na`: count of NA values
- `n_duplicates`: count of duplicate IDs
- `format_ok`: TRUE if IDs look like proper UUIDs/hashes

**Examples**

```
df <- add_id(data.frame(x = 1:3))
check_id(df)
```

<code>check_id_disjoint</code>	<i>Check IDs are disjoint across datasets</i>
--------------------------------	---

**Description**

Verifies that ID columns don't overlap between datasets. Useful before binding datasets to ensure no ID collisions.

**Usage**

```
check_id_disjoint(..., .id = ".id")
```

**Arguments**

- |                  |                                       |
|------------------|---------------------------------------|
| <code>...</code> | Data frames to check.                 |
| <code>.id</code> | Column name for IDs (default: ".id"). |

**Value**

Invisibly returns a list with:

- `disjoint`: TRUE if no overlaps found
- `overlaps`: character vector of overlapping IDs (if any)

**Examples**

```
df1 <- add_id(data.frame(x = 1:3))
df2 <- add_id(data.frame(x = 4:6))
check_id_disjoint(df1, df2)
```

---

`clear_all_snapshots`    *Clear all snapshots from cache*

---

### Description

Clear all snapshots from cache

### Usage

```
clear_all_snapshots(confirm = TRUE)
```

### Arguments

`confirm`    If TRUE, require confirmation.

### Value

No return value, called for side effects.

---

`clear_snapshot`    *Clear snapshot for a data frame*

---

### Description

Removes the snapshot reference from a data frame.

### Usage

```
clear_snapshot(.data, purge = FALSE)
```

### Arguments

`.data`    A data frame.

`purge`    If TRUE, also remove the snapshot from cache.

### Value

Data frame without snapshot reference.

---

<code>commit_keyed</code>	<i>Commit a keyed data frame as reference</i>
---------------------------	---

---

## Description

Stores a hash-based snapshot of the current data state. Only one active reference per data frame (identified by its content hash).

## Usage

```
commit_keyed(.data, name = NULL)
```

## Arguments

- |                    |   |
|--------------------|---|
| <code>.data</code> | A data frame (preferably keyed).                            |
| <code>name</code>  | Optional name for the snapshot. If NULL, derived from data. |

## Details

The snapshot stores:

- Row count
- Column names and types
- Hash of key columns (if keyed)
- Hash of full content

Snapshots are stored in memory for the session. They are keyed by content hash, so identical data shares the same snapshot.

## Value

Invisibly returns `.data` with snapshot metadata attached.

## Examples

```
df <- key(data.frame(id = 1:3, x = c("a", "b", "c")), id)
df <- commit_keyed(df)

# Later, check for drift
df2 <- df
df2$x[1] <- "z"
check_drift(df2)
```

---

`compare_ids`*Compare IDs between data frames*

---

## Description

Compares IDs between two data frames to detect lost rows.

## Usage

```
compare_ids(before, after, .id = ".id")
```

## Arguments

<code>before</code>	Data frame before transformation.
<code>after</code>	Data frame after transformation.
<code>.id</code>	Column name for IDs (default: ".id").

## Value

A list with:

- `lost`: IDs present in `before` but not `after`
- `gained`: IDs present in `after` but not `before`
- `preserved`: IDs present in both

## Examples

```
df1 <- add_id(data.frame(x = 1:5))
df2 <- df1[1:3, ]
compare_ids(df1, df2)
```

---

`compare_keys`*Compare key values between two data frames*

---

## Description

Identifies keys that are new, removed, or common between two keyed data frames. Does not compare values, only key membership.

## Usage

```
compare_keys(x, y, by = NULL)
```

**Arguments**

- x First keyed data frame.
- y Second keyed data frame.
- by Column(s) to compare. If NULL, uses the key from x.

**Value**

A key comparison object.

**Examples**

```
df1 <- key(data.frame(id = 1:3, x = 1:3), id)
df2 <- key(data.frame(id = 2:4, x = 2:4), id)
compare_keys(df1, df2)
```

<code>compare_structure</code>	<i>Compare structure of two data frames</i>
--------------------------------	---

**Description**

Compares the structural properties of two data frames without comparing actual values. Useful for detecting schema drift.

**Usage**

```
compare_structure(x, y)
```

**Arguments**

- x First data frame.
- y Second data frame.

**Value**

A structure comparison object.

**Examples**

```
df1 <- data.frame(id = 1:3, x = c("a", "b", "c"))
df2 <- data.frame(id = 1:5, x = c("a", "b", "c", "d", "e"), y = 1:5)
compare_structure(df1, df2)
```

---

diagnose_join	<i>Diagnose a join before executing</i>
---------------	---

---

## Description

Analyzes join cardinality without performing the full join. Useful for detecting many-to-many joins that would explode row count.

## Usage

```
diagnose_join(x, y, by = NULL, use_joinspy = TRUE)
```

## Arguments

x	Left data frame.
y	Right data frame.
by	Join specification.
use_joinspy	If TRUE (default), use joinspy for enhanced diagnostics when available. Set to FALSE to use built-in diagnostics only.

## Details

If the joinspy package is installed, this function delegates to `joinspy::join_spy()` for enhanced diagnostics including whitespace detection, encoding issues, and detailed match analysis.

## Value

A list with cardinality information, or a `JoinReport` object if `joinspy` is used.

## See Also

`joinspy::join_spy()` for enhanced diagnostics (if installed)

## Examples

```
x <- data.frame(id = c(1, 1, 2), a = 1:3)
y <- data.frame(id = c(1, 1, 2), b = 4:6)
diagnose_join(x, y, by = "id", use_joinspy = FALSE)
```

extend_id	<i>Extend IDs to new rows</i>
-----------	-------------------------------

### Description

Adds IDs to rows where the ID column is NA, preserving existing IDs. Useful after binding new data to an existing dataset with IDs.

### Usage

```
extend_id(.data, .id = ".id")
```

### Arguments

.data	A data frame with an ID column (possibly with NAs).
.id	Column name for IDs (default: ".id").

### Value

Data frame with IDs filled in for NA rows.

### Examples

```
# Original data with IDs
old <- add_id(data.frame(x = 1:3))

# New data without IDs
new <- data.frame(.id = NA_character_, x = 4:5)

# Combine and extend
combined <- dplyr::bind_rows(old, new)
extend_id(combined)
```

find_duplicates	<i>Find duplicate keys</i>
-----------------	----------------------------

### Description

Identifies rows with duplicate key values.

### Usage

```
find_duplicates(.data, ...)
```

**Arguments**

- .data A data frame.
- ... Column names to check. If empty, uses the key columns.

**Value**

Data frame containing only the rows with duplicate keys, with a .n column showing the count.

**Examples**

```
df <- data.frame(id = c(1, 1, 2, 3, 3, 3), x = letters[1:6])
find_duplicates(df, id)
```

---

get_id	<i>Get ID column</i>
--------	----------------------

---

**Description**

Get ID column

**Usage**

```
get_id(.data, .id = ".id")
```

**Arguments**

- .data A data frame.
- .id Column name (default: ".id").

**Value**

Character vector of IDs, or NULL if not present.

---

get\_key\_cols            *Get key column names*

---

**Description**

Get key column names

**Usage**

```
get_key_cols(.data)
```

**Arguments**

.data            A keyed data frame.

**Value**

Character vector of column names, or NULL if no key.

---

has\_id            *Check if data frame has IDs*

---

**Description**

Check if data frame has IDs

**Usage**

```
has_id(.data, .id = ".id")
```

**Arguments**

.data            A data frame.  
.id            Column name to check (default: ".id").

**Value**

Logical.

---

has_key	<i>Check if data frame has a key</i>
---------	--------------------------------------

---

## Description

Check if data frame has a key

## Usage

```
has_key(.data)
```

## Arguments

.data            A data frame.

## Value

Logical.

---

key	<i>Define a key for a data frame</i>
-----	--------------------------------------

---

## Description

Attaches key metadata to a data frame, marking which column(s) form the unique identifier for rows. Keys are validated for uniqueness at creation.

## Usage

```
key(.data, ..., .validate = TRUE, .strict = FALSE)
```

```
key(.data) <- value
```

## Arguments

.data            A data frame or tibble.  
. . .            Column names (unquoted) that form the key. Can be a single column or multiple columns for a composite key.  
.validate       If TRUE (default), check that the key is unique.  
.strict          If TRUE, error on non-unique keys. If FALSE (default), warn but still attach the key.  
value            Character vector of column names to use as key.

**Value**

A keyed data frame (class keyed\_df).

**Examples**

```
df <- data.frame(id = 1:3, x = c("a", "b", "c"))
key(df, id)

# Composite key
df2 <- data.frame(country = c("US", "US", "UK"), year = c(2020, 2021, 2020), val = 1:3)
key(df2, country, year)
```

---

**key\_is\_valid**

*Check if the key is still valid*

---

**Description**

Checks whether the key columns still exist and are still unique.

**Usage**

```
key_is_valid(.data)
```

**Arguments**

.data            A keyed data frame.

**Value**

Logical. Returns FALSE with a warning if key is invalid.

---

**key\_status**

*Get key status summary*

---

**Description**

Returns diagnostic information about a keyed data frame.

**Usage**

```
key_status(.data)
```

**Arguments**

.data            A data frame.

**Value**

A key status object with diagnostic information.

**Examples**

```
df <- key(data.frame(id = 1:3, x = c("a", "b", "c")), id)
key_status(df)
```

---

list_snapshots	<i>List all snapshots in cache</i>
----------------	------------------------------------

---

**Description**

List all snapshots in cache

**Usage**

```
list_snapshots()
```

**Value**

Data frame with snapshot information.

---

lock_complete	<i>Assert that data is complete (no missing values anywhere)</i>
---------------	--

---

**Description**

Checks that all columns have no NA values.

**Usage**

```
lock_complete(.data, .strict = FALSE)
```

**Arguments**

.data	A data frame.
.strict	If TRUE, error on failure. If FALSE (default), warn.

**Value**

Invisibly returns .data (for piping).

`lock_coverage`      *Assert minimum coverage of values*

### Description

Checks that the fraction of non-NA values meets a threshold. Useful after joins to verify expected coverage.

### Usage

```
lock_coverage(.data, threshold, ..., .strict = FALSE)
```

### Arguments

<code>.data</code>	A data frame.
<code>threshold</code>	Minimum fraction of non-NA values (0 to 1).
<code>...</code>	Column names (unquoted) to check. If empty, checks all columns.
<code>.strict</code>	If TRUE, error on failure. If FALSE (default), warn.

### Value

Invisibly returns `.data` (for piping).

### Examples

```
df <- data.frame(id = 1:10, x = c(1:8, NA, NA))
lock_coverage(df, 0.8, x)
lock_coverage(df, 0.9, x) # warns (only 80% coverage)
```

`lock_no_na`      *Assert that columns have no missing values*

### Description

Checks that specified columns contain no NA values.

### Usage

```
lock_no_na(.data, ..., .strict = FALSE)
```

### Arguments

<code>.data</code>	A data frame.
<code>...</code>	Column names (unquoted) to check. If empty, checks all columns.
<code>.strict</code>	If TRUE, error on failure. If FALSE (default), warn.

**Value**

Invisibly returns .data (for piping).

**Examples**

```
df <- data.frame(id = 1:3, x = c("a", NA, "c"))
lock_no_na(df, id)
lock_no_na(df, x) # warns
```

---

**lock\_nrow***Assert row count within expected range*

---

**Description**

Checks that the number of rows is within an expected range. Useful for sanity checks after filtering or joins.

**Usage**

```
lock_nrow(.data, min = 1, max = Inf, expected = NULL, .strict = FALSE)
```

**Arguments**

.data	A data frame.
min	Minimum expected rows (inclusive). Default 1.
max	Maximum expected rows (inclusive). Default Inf.
expected	Exact expected row count. If provided, overrides min/max.
.strict	If TRUE, error on failure. If FALSE (default), warn.

**Value**

Invisibly returns .data (for piping).

**Examples**

```
df <- data.frame(id = 1:100)
lock_nrow(df, min = 50, max = 200)
lock_nrow(df, expected = 100)
```

<code>lock_unique</code>	<i>Assert that columns are unique</i>
--------------------------	---------------------------------------

**Description**

Checks that the combination of specified columns has unique values. This is a point-in-time assertion that either passes silently or fails.

**Usage**

```
lock_unique(.data, ..., .strict = FALSE)
```

**Arguments**

- .data            A data frame.
- ...              Column names (unquoted) to check for uniqueness.
- .strict          If TRUE, error on failure. If FALSE (default), warn.

**Value**

Invisibly returns .data (for piping).

**Examples**

```
df <- data.frame(id = 1:3, x = c("a", "b", "c"))
lock_unique(df, id)

# Fails with warning
df2 <- data.frame(id = c(1, 1, 2), x = c("a", "b", "c"))
lock_unique(df2, id)
```

<code>make_id</code>	<i>Create ID from columns</i>
----------------------	-------------------------------

**Description**

Creates an ID column by combining values from one or more columns. Unlike [add\\_id\(\)](#), this produces deterministic IDs based on column values.

**Usage**

```
make_id(.data, ..., .id = ".id", .sep = "|")
```

**Arguments**

- .data A data frame.
- ... Columns to combine into the ID.
- .id Column name for the ID (default: ".id").
- .sep Separator between column values (default: "|").

**Value**

Data frame with ID column added.

**Examples**

```
df <- data.frame(country = c("US", "UK", "US"), year = c(2020, 2020, 2021))
make_id(df, country, year)
#>   .id      country year
#> 1 US|2020  US      2020
#> 2 UK|2020  UK      2020
#> 3 US|2021  US      2021
```

---

remove\_id

*Remove ID column*

---

**Description**

Remove ID column

**Usage**

```
remove_id(.data, .id = ".id")
```

**Arguments**

- .data A data frame.
- .id Column name to remove (default: ".id").

**Value**

Data frame without the ID column.

---

summary.keyed\_df      *Summary method for keyed data frames*

---

**Description**

Summary method for keyed data frames

**Usage**

```
## S3 method for class 'keyed_df'  
summary(object, ...)
```

**Arguments**

object	A keyed data frame.
...	Additional arguments (ignored).

**Value**

Invisibly returns a summary list.

---

unkey      *Remove key from a data frame*

---

**Description**

Remove key from a data frame

**Usage**

```
unkey(.data)
```

**Arguments**

.data	A keyed data frame.
-------	---------------------

**Value**

A tibble without key metadata.

# Index

add\_id, 2  
add\_id(), 20  
bind\_id, 3  
bind\_keyed, 4  
check\_drift, 5  
check\_id, 5  
check\_id\_disjoint, 6  
clear\_all\_snapshots, 7  
clear\_snapshot, 7  
commit\_keyed, 8  
compare\_ids, 9  
compare\_keys, 9  
compare\_structure, 10  
diagnose\_join, 11  
dplyr::bind\_rows(), 4  
extend\_id, 12  
extend\_id(), 4  
find\_duplicates, 12  
get\_id, 13  
get\_key\_cols, 14  
has\_id, 14  
has\_key, 15  
key, 15  
key<- (key), 15  
key\_is\_valid, 16  
key\_status, 16  
list\_snapshots, 17  
lock\_complete, 17  
lock\_coverage, 18  
lock\_no\_na, 18  
lock\_nrow, 19  
lock\_unique, 20  
make\_id, 20  
remove\_id, 21  
summary.keyed\_df, 22  
unkey, 22