

# Package ‘kindling’

February 4, 2026

**Type** Package

**Title** Higher-Level Interface of ‘torch’ Package to Auto-Train Neural Networks

**Version** 0.2.0

**Description** Provides a higher-level interface to the ‘torch’ package for defining, training, and fine-tuning neural networks, including its depth, powered by code generation. This package currently supports few to several architectures, namely feedforward (multi-layer perceptron) and recurrent neural networks (Recurrent Neural Networks (RNN), Long Short-Term Memory (LSTM), Gated Recurrent Unit (GRU)), while also reduces boilerplate ‘torch’ code while enabling seamless integration with ‘torch’. The model methods to train neural networks from this package also bridges to titanic ML frameworks in R, namely ‘tidymodels’ ecosystem, which enables the ‘parsnip’ model specifications, workflows, recipes, and tuning tools.

**License** MIT + file LICENSE

**Encoding** UTF-8

**Imports** purrr, torch, rlang, cli, glue, vctrs, parsnip (>= 1.0.0), tibble, tidyr, dplyr, stats, NeuralNetTools, vip, ggplot2, tune, dials, hardhat

**Suggests** testthat (>= 3.0.0), magrittr, box, recipes, workflows, rsample, yardstick, mlbench, modeldata, knitr, rmarkdown, DiceDesign, lhs, sfd

**Config/testthat.edition** 3

**RoxygenNote** 7.3.3

**Depends** R (>= 4.1.0)

**URL** <https://kindling.joshuamarie.com>,  
<https://github.com/joshuamarie/kindling>

**BugReports** <https://github.com/joshuamarie/kindling/issues>

**VignetteBuilder** knitr

**NeedsCompilation** no

**Author** Joshua Marie [aut, cre]  
**Maintainer** Joshua Marie <joshua.marie.k@gmail.com>  
**Repository** CRAN  
**Date/Publication** 2026-02-04 16:00:08 UTC

## Contents

act_funcs . . . . .	2
args . . . . .	3
ffnn . . . . .	3
ffnn_generator . . . . .	6
ffnn_wrapper . . . . .	10
grid_depth . . . . .	11
kindling-varimp . . . . .	15
mlp_kindling . . . . .	17
ordinal_gen . . . . .	20
rnn_kindling . . . . .	21
table_summary . . . . .	23

<b>Index</b>	<b>26</b>
--------------	-----------

---

act_funcs	<i>Activation Functions Specification Helper</i>
-----------	--

---

### Description

This function is a DSL function, kind of like `ggplot2::aes()`, that helps to specify activation functions for neural network layers. It validates that activation functions exist in `torch` and that any parameters match the function's formal arguments.

### Usage

```
act_funcs(...)
```

### Arguments

...	Activation function specifications. Can be:
	<ul style="list-style-type: none"> <li>• Bare symbols: <code>relu</code>, <code>tanh</code></li> <li>• Character strings (simple): <code>"relu"</code>, <code>"tanh"</code></li> <li>• Character strings (with params): <code>"softshrink(lambda = 0.1)"</code>, <code>"rrelu(lower = 1/5, upper = 1/4)"</code></li> <li>• Named with parameters: <code>softmax = args(dim = 2L)</code></li> </ul>

### Value

A `vctrs` vector with class "activation\_spec" containing validated activation specifications.

---

args

*Activation Function Arguments Helper*

---

## Description

Type-safe helper to specify parameters for activation functions. All parameters must be named and match the formal arguments of the corresponding torch activation function.

## Usage

args(...)

## Arguments

...                    Named arguments for the activation function.

## Value

A list with class "activation\_args" containing the parameters.

---

ffnn

*Base models for Neural Network Training in kindling*

---

## Description

Base models for Neural Network Training in kindling

## Usage

```
ffnn(  
  formula = NULL,  
  data = NULL,  
  hidden_neurons,  
  activations = NULL,  
  output_activation = NULL,  
  bias = TRUE,  
  epochs = 100,  
  batch_size = 32,  
  penalty = 0,  
  mixture = 0,  
  learn_rate = 0.001,  
  optimizer = "adam",  
  optimizer_args = list(),  
  loss = "mse",  
  validation_split = 0,  
  device = NULL,
```

```

verbose = FALSE,
cache_weights = FALSE,
...,
x = NULL,
y = NULL
)

rnn(
  formula = NULL,
  data = NULL,
  hidden_neurons,
  rnn_type = "lstm",
  activations = NULL,
  output_activation = NULL,
  bias = TRUE,
  bidirectional = TRUE,
  dropout = 0,
  epochs = 100,
  batch_size = 32,
  penalty = 0,
  mixture = 0,
  learn_rate = 0.001,
  optimizer = "adam",
  optimizer_args = list(),
  loss = "mse",
  validation_split = 0,
  device = NULL,
  verbose = FALSE,
  cache_weights = FALSE,
  ...,
  x = NULL,
  y = NULL
)

```

## Arguments

<code>formula</code>	Formula. Model formula (e.g., $y \sim x1 + x2$ ).
<code>data</code>	Data frame. Training data.
<code>hidden_neurons</code>	Integer vector. Number of neurons in each hidden layer.
<code>activations</code>	Activation function specifications. See <code>act_funs()</code> .
<code>output_activation</code>	Optional. Activation for output layer.
<code>bias</code>	Logical. Use bias weights. Default TRUE.
<code>epochs</code>	Integer. Number of training epochs. Default 100.
<code>batch_size</code>	Integer. Batch size for training. Default 32.
<code>penalty</code>	Numeric. Regularization penalty (lambda). Default 0 (no regularization).

mixture	Numeric. Elastic net mixing parameter (0-1). Default 0.
learn_rate	Numeric. Learning rate for optimizer. Default 0.001.
optimizer	Character. Optimizer type ("adam", "sgd", "rmsprop"). Default "adam".
optimizer_args	Named list. Additional arguments passed to the optimizer. Default list().
loss	Character. Loss function ("mse", "mae", "cross_entropy", "bce"). Default "mse".
validation_split	Numeric. Proportion of data for validation (0-1). Default 0.
device	Character. Device to use ("cpu", "cuda", "mps"). Default NULL (auto-detect).
verbose	Logical. Print training progress. Default FALSE.
cache_weights	Logical. Cache weight matrices for faster variable importance. Default FALSE.
...	Additional arguments. Can be used to pass x and y for direct interface.
x	When not using formula: predictor data (data.frame or matrix).
y	When not using formula: outcome data (vector, factor, or matrix).
rnn_type	Character. Type of RNN ("rnn", "lstm", "gru"). Default "lstm".
bidirectional	Logical. Use bidirectional RNN. Default TRUE.
dropout	Numeric. Dropout rate between layers. Default 0.

## Value

An object of class "ffnn\_fit" containing the trained model and metadata.

## FFNN

Train a feed-forward neural network using the torch package.

## RNN

Train a recurrent neural network using the torch package.

## Examples

```
if (torch::torch_is_installed()) {
  # Formula interface (original)
  model_reg = ffnn(
    Sepal.Length ~ .,
    data = iris[, 1:4],
    hidden_neurons = c(64, 32),
    activations = "relu",
    epochs = 50
  )

  # XY interface (new)
  model_xy = ffnn(
    hidden_neurons = c(64, 32),
    activations = "relu",
    epochs = 50,
```

```

x = iris[, 2:4],
y = iris$Sepal.Length
)
}

if (torch::torch_is_installed()) {
  # Formula interface (original)
  model_rnn = rnn(
    Sepal.Length ~ .,
    data = iris[, 1:4],
    hidden_neurons = c(64, 32),
    rnn_type = "lstm",
    activations = "relu",
    epochs = 50
  )

  # XY interface (new)
  model_xy = rnn(
    hidden_neurons = c(64, 32),
    rnn_type = "gru",
    epochs = 50,
    x = iris[, 2:4],
    y = iris$Sepal.Length
  )
}

```

**ffnn\_generator***Functions to generate nn\_module (language) expression***Description**

Functions to generate nn\_module (language) expression

**Usage**

```

ffnn_generator(
  nn_name = "DeepFFN",
  hd_neurons,
  no_x,
  no_y,
  activations = NULL,
  output_activation = NULL,
  bias = TRUE
)

```

```
rnn_generator(
  nn_name = "DeepRNN",
  hd_neurons,
  no_x,
  no_y,
  rnn_type = "lstm",
  bias = TRUE,
  activations = NULL,
  output_activation = NULL,
  bidirectional = TRUE,
  dropout = 0,
  ...
)
```

## Arguments

nn_name	Character. Name of the generated RNN module class. Default is "DeepRNN".
hd_neurons	Integer vector. Number of neurons in each hidden RNN layer.
no_x	Integer. Number of input features.
no_y	Integer. Number of output features.
activations	Activation function specifications for each hidden layer. Can be: <ul style="list-style-type: none"> <li>• NULL: No activation functions.</li> <li>• Character vector: e.g., c("relu", "sigmoid").</li> <li>• List: e.g., act_funs(relu, elu, softshrink = args(lambda = 0.5)).</li> <li>• activation_spec object from act_funs().</li> </ul> If the length of activations is 1L, this will be the activation throughout the architecture.
output_activation	Optional. Activation function for the output layer. Same format as activations but should be a single activation.
bias	Logical. Whether to use bias weights. Default is TRUE
rnn_type	Character. Type of RNN to use. Must be one of "rnn", "lstm", or "gru". Default is "lstm".
bidirectional	Logical. Whether to use bidirectional RNN layers. Default is TRUE.
dropout	Numeric. Dropout rate between RNN layers. Default is 0.
...	Additional arguments (currently unused).

## Details

The generated FFNN module will have the specified number of hidden layers, with each layer containing the specified number of neurons. Activation functions can be applied after each hidden layer as specified. This can be used for both classification and regression tasks.

The generated module properly namespaces all torch functions to avoid polluting the global namespace.

The generated RNN module will have the specified number of recurrent layers, with each layer containing the specified number of hidden units. Activation functions can be applied after each RNN layer as specified. The final output is taken from the last time step and passed through a linear layer.

The generated module properly namespaces all torch functions to avoid polluting the global namespace.

### **Value**

A torch module expression representing the FFNN.

A torch module expression representing the RNN.

### **Feed-Forward Neural Network Module Generator**

The `ffnn_generator()` function generates a feed-forward neural network (FFNN) module expression from the `torch` package in R. It allows customization of the FFNN architecture, including the number of hidden layers, neurons, and activation functions.

### **Recurrent Neural Network Module Generator**

The `rnn_generator()` function generates a recurrent neural network (RNN) module expression from the `torch` package in R. It allows customization of the RNN architecture, including the number of hidden layers, neurons, RNN type, activation functions, and other parameters.

### **Examples**

```
# FFNN
if (torch::torch_is_installed()) {
  # Generate an MLP module with 3 hidden layers
  ffnn_mod = ffnn_generator(
    nn_name = "MyFFNN",
    hd_neurons = c(64, 32, 16),
    no_x = 10,
    no_y = 1,
    activations = 'relu'
  )

  # Evaluate and instantiate
  model = eval(ffnn_mod)()

  # More complex: With different activations
  ffnn_mod2 = ffnn_generator(
    nn_name = "MyFFNN2",
    hd_neurons = c(128, 64, 32),
    no_x = 20,
    no_y = 5,
    activations = act_funcs(
      relu,
      selu,
      sigmoid
    )
  )
}
```

```
)  
  
# Even more complex: Different activations and customized argument  
# for the specific activation function  
ffnn_mod2 = ffnn_generator(  
    nn_name = "MyFFNN2",  
    hd_neurons = c(128, 64, 32),  
    no_x = 20,  
    no_y = 5,  
    activations = act_funcs(  
        relu,  
        selu,  
        softshrink = args(lambd = 0.5)  
    )  
)  
  
# Customize output activation (softmax is useful for classification tasks)  
ffnn_mod3 = ffnn_generator(  
    hd_neurons = c(64, 32),  
    no_x = 10,  
    no_y = 3,  
    activations = 'relu',  
    output_activation = act_funcs(softmax = args(dim = 2L))  
)  
}  
} else {  
    message("Torch not fully installed - skipping example")  
}  
  
## RNN  
if (torch::torch_is_installed()) {  
    # Basic LSTM with 2 layers  
    rnn_mod = rnn_generator(  
        nn_name = "MyLSTM",  
        hd_neurons = c(64, 32),  
        no_x = 10,  
        no_y = 1,  
        rnn_type = "lstm",  
        activations = 'relu'  
    )  
  
    # Evaluate and instantiate  
    model = eval(rnn_mod)()  
  
    # GRU with different activations  
    rnn_mod2 = rnn_generator(  
        nn_name = "MyGRU",  
        hd_neurons = c(128, 64, 32),  
        no_x = 20,  
        no_y = 5,  
        rnn_type = "gru",  
        activations = act_funcs(relu, elu, relu),
```

```

        bidirectional = FALSE
    )

} else {
    message("Torch not fully installed - skipping example")
}

## Not run:
## Parameterized activation and dropout
# (Will throw an error due to `nnf_tanh()` not being available in `{torch}`)
# rnn_mod3 = rnn_generator(
#     hd_neurons = c(100, 50, 25),
#     no_x = 15,
#     no_y = 3,
#     rnn_type = "lstm",
#     activations = act_funcs(
#         relu,
#         leaky_relu = args(negative_slope = 0.01),
#         tanh
#     ),
#     bidirectional = TRUE,
#     dropout = 0.3
# )

## End(Not run)

```

## Description

Basemodels-tidymodels wrappers

## Usage

```

ffnn_wrapper(formula, data, ...)
rnn_wrapper(formula, data, ...)

```

## Arguments

<code>formula</code>	A formula specifying the model (e.g., $y \sim x1 + x2$ )
<code>data</code>	A data frame containing the training data
<code>...</code>	Additional arguments passed to the underlying training function

## Details

These wrapper functions are designed to interface with the `{tidymodels}` ecosystem, particularly for use with `tune::tune_grid()` and workflows. They handle the conversion of tuning parameters (especially list-column parameters from `grid_depth()`) into the format expected by `ffnn()` and `rnn()`.

## Value

- `ffnn_wrapper()` returns an object of class `"ffnn_fit"` containing the trained feedforward neural network model and metadata. See `ffnn()` for details.
- `rnn_wrapper()` returns an object of class `"rnn_fit"` containing the trained recurrent neural network model and metadata. See `rnn()` for details.

## FFNN (MLP) Wrapper for `{tidymodels}` interface

This is a function to interface into `{tidymodels}` (do not use this, use `kindling::ffnn()` instead).

## RNN Wrapper for `{tidymodels}` interface

This is a function to interface into `{tidymodels}` (do not use this, use `kindling::rnn()` instead).

`grid_depth`

*Depth-Aware Grid Generation for Neural Networks*

## Description

`grid_depth()` extends standard grid generation to support multi-layer neural network architectures. It creates heterogeneous layer configurations by generating list columns for `hidden_neurons` and `activations`.

## Usage

```
grid_depth(
  x,
  ...,
  n_hlayer = 2L,
  size = 5L,
  type = c("regular", "random", "latin_hypercube", "max_entropy", "audze_eglais"),
  original = TRUE,
  levels = 3L,
  variogram_range = 0.5,
  iter = 1000L
)
## S3 method for class 'parameters'
grid_depth(
  x,
```

```

...,
n_hlayer = 2L,
size = 5L,
type = c("regular", "random", "latin_hypercube", "max_entropy", "audze_eglais"),
original = TRUE,
levels = 3L,
variogram_range = 0.5,
iter = 1000L
)

## S3 method for class 'list'
grid_depth(
  x,
  ...,
  n_hlayer = 2L,
  size = 5L,
  type = c("regular", "random", "latin_hypercube", "max_entropy", "audze_eglais"),
  original = TRUE,
  levels = 3L,
  variogram_range = 0.5,
  iter = 1000L
)

## S3 method for class 'workflow'
grid_depth(
  x,
  ...,
  n_hlayer = 2L,
  size = 5L,
  type = c("regular", "random", "latin_hypercube", "max_entropy", "audze_eglais"),
  original = TRUE,
  levels = 3L,
  variogram_range = 0.5,
  iter = 1000L
)

## S3 method for class 'model_spec'
grid_depth(
  x,
  ...,
  n_hlayer = 2L,
  size = 5L,
  type = c("regular", "random", "latin_hypercube", "max_entropy", "audze_eglais"),
  original = TRUE,
  levels = 3L,
  variogram_range = 0.5,
  iter = 1000L
)

```

```

## S3 method for class 'param'
grid_depth(
  x,
  ...,
  n_hlayer = 2L,
  size = 5L,
  type = c("regular", "random", "latin_hypercube", "max_entropy", "audze_eglais"),
  original = TRUE,
  levels = 3L,
  variogram_range = 0.5,
  iter = 1000L
)

## Default S3 method:
grid_depth(
  x,
  ...,
  n_hlayer = 2L,
  size = 5L,
  type = c("regular", "random", "latin_hypercube", "max_entropy", "audze_eglais"),
  original = TRUE,
  levels = 3L,
  variogram_range = 0.5,
  iter = 1000L
)

```

## Arguments

<code>x</code>	A <code>parameters</code> object, list, workflow, or model spec. Can also be a single <code>param</code> object if <code>...</code> contains additional <code>param</code> objects.
<code>...</code>	One or more <code>param</code> objects (e.g., <code>hidden_neurons()</code> , <code>epochs()</code> ). If <code>x</code> is a <code>parameters</code> object, <code>...</code> is ignored. None of the objects can have <code>unknown()</code> values.
<code>n_hlayer</code>	Integer vector specifying number of hidden layers to generate (e.g., <code>2:4</code> for 2, 3, or 4 layers), or a <code>param</code> object created with <code>n_hlayers()</code> . Default is 2.
<code>size</code>	Integer. Number of parameter combinations to generate.
<code>type</code>	Character. Type of grid: "regular", "random", "latin_hypercube", "max_entropy", or "audze_eglais".
<code>original</code>	Logical. Should original parameter ranges be used?
<code>levels</code>	Integer. Levels per parameter for regular grids.
<code>variogram_range</code>	Numeric. Range for audze_eglais design.
<code>iter</code>	Integer. Iterations for max_entropy optimization.

## Details

This function is specifically for `{kindling}` models. The `n_hlayer` parameter determines network depth and creates list columns for `hidden_neurons` and `activations`, where each element is a vector of length matching the sampled depth.

When `n_hlayer` is a parameter object (created with `n_hlayers()`), it will be treated as a tunable parameter and sampled according to its defined range.

## Value

A tibble with list columns for `hidden_neurons` and `activations`, where each element is a vector of length `n_hlayer`.

## Examples

```
## Not run:
library(dials)
library(workflows)
library(tune)

# Method 1: Fixed depth
grid = grid_depth(
  hidden_neurons(c(32L, 128L)),
  activations(c("relu", "elu")),
  epochs(c(50L, 200L)),
  n_hlayer = 2:3,
  type = "random",
  size = 20
)

# Method 2: Tunable depth using parameter object
grid = grid_depth(
  hidden_neurons(c(32L, 128L)),
  activations(c("relu", "elu")),
  epochs(c(50L, 200L)),
  n_hlayer = n_hlayers(range = c(2L, 4L)),
  type = "random",
  size = 20
)

# Method 3: From workflow
wf = workflow() |>
  add_model(mlp_kindling(hidden_neurons = tune(), activations = tune())) |>
  add_formula(y ~ .)
grid = grid_depth(wf, n_hlayer = 2:4, type = "latin_hypercube", size = 15)

## End(Not run)
```

## Description

This file implements methods for variable importance generics from NeuralNetTools and vip packages.

## Usage

```
## S3 method for class 'ffnn_fit'
garson(mod_in, bar_plot = FALSE, ...)

## S3 method for class 'ffnn_fit'
olden(mod_in, bar_plot = TRUE, ...)

## S3 method for class 'ffnn_fit'
vi_model(object, type = c("olden", "garson"), ...)
```

## Arguments

mod_in	A fitted model object of class "ffnn_fit".
bar_plot	Logical. Whether to plot variable importance (default TRUE).
...	Additional arguments passed to methods.
object	A fitted model object of class "ffnn_fit".
type	Type of algorithm to extract the variable importance. This must be one of the strings: <ul style="list-style-type: none"><li>• 'olden'</li><li>• 'garson'</li></ul>

## Value

A data frame for both "garson" and "olden" classes with columns:

x_names	Character vector of predictor variable names
y_names	Character string of response variable name
rel_imp	Numeric vector of relative importance scores (percentage)

The data frame is sorted by importance in descending order.

A tibble with columns "Variable" and "Importance" (via `vip::vi()` / `vip::vi_model()` only).

## Garson's Algorithm for FFNN Models

{kindling} inherits NeuralNetTools::garson to extract the variable importance from the fitted `ffnn()` model.

### Olden's Algorithm for FFNN Models

{kindling} inherits NeuralNetTools::olden to extract the variable importance from the fitted ffnn() model.

### Variable Importance via {vip} Package

You can directly use vip::vi() and vip::vi\_model() to extract the variable importance from the fitted ffnn() model.

## References

- Beck, M.W. 2018. NeuralNetTools: Visualization and Analysis Tools for Neural Networks. Journal of Statistical Software. 85(11):1-20.
- Garson, G.D. 1991. Interpreting neural network connection weights. Artificial Intelligence Expert. 6(4):46-51.
- Goh, A.T.C. 1995. Back-propagation neural networks for modeling complex systems. Artificial Intelligence in Engineering. 9(3):143-151.
- Olden, J.D., Jackson, D.A. 2002. Illuminating the 'black-box': a randomization approach for understanding variable contributions in artificial neural networks. Ecological Modelling. 154:135-150.
- Olden, J.D., Joy, M.K., Death, R.G. 2004. An accurate comparison of methods for quantifying variable importance in artificial neural networks using simulated data. Ecological Modelling. 178:389-397.

## Examples

```
if (torch::torch_is_installed()) {
  model_mlp = ffnn(
    Species ~ .,
    data = iris,
    hidden_neurons = c(64, 32),
    activations = "relu",
    epochs = 100,
    verbose = FALSE,
    cache_weights = TRUE
  )

  # Directly use `NeuralNetTools::garson`
  model_mlp |>
    garson()

  # Directly use `NeuralNetTools::olden`
  model_mlp |>
    olden()
} else {
  message("Torch not fully installed - skipping example")
}
```

```

# kindling also supports `vip::vi()` / `vip::vi_model()`
if (torch::torch_is_installed()) {
  model_mlp = ffnn(
    Species ~ .,
    data = iris,
    hidden_neurons = c(64, 32),
    activations = "relu",
    epochs = 100,
    verbose = FALSE,
    cache_weights = TRUE
  )

  model_mlp |>
    vip::vi(type = 'garson') |>
    vip::vip()
} else {
  message("Torch not fully installed - skipping example")
}

```

mlp\_kindling

*Multi-Layer Perceptron (Feedforward Neural Network) via kindling*

## Description

`mlp_kindling()` defines a feedforward neural network model that can be used for classification or regression. It integrates with the `tidymodels` ecosystem and uses the `torch` backend via `kindling`.

## Usage

```

mlp_kindling(
  mode = "unknown",
  engine = "kindling",
  hidden_neurons = NULL,
  activations = NULL,
  output_activation = NULL,
  bias = NULL,
  epochs = NULL,
  batch_size = NULL,
  penalty = NULL,
  mixture = NULL,
  learn_rate = NULL,
  optimizer = NULL,
  optimizer_args = NULL,
  loss = NULL,
  validation_split = NULL,
  device = NULL,
  verbose = NULL
)

```

## Arguments

<code>mode</code>	A single character string for the type of model. Possible values are "unknown", "regression", or "classification".
<code>engine</code>	A single character string specifying what computational engine to use for fitting. Currently only "kindling" is supported.
<code>hidden_neurons</code>	An integer vector for the number of units in each hidden layer. Can be tuned.
<code>activations</code>	A character vector of activation function names for each hidden layer (e.g., "relu", "tanh", "sigmoid"). Can be tuned.
<code>output_activation</code>	A character string for the output activation function. Can be tuned.
<code>bias</code>	Logical for whether to include bias terms. Can be tuned.
<code>epochs</code>	An integer for the number of training iterations. Can be tuned.
<code>batch_size</code>	An integer for the batch size during training. Can be tuned.
<code>penalty</code>	A number for the regularization penalty (lambda). Default 0 (no regularization). Higher values increase regularization strength. Can be tuned.
<code>mixture</code>	A number between 0 and 1 for the elastic net mixing parameter. Default 0 (pure L2/Ridge regularization). <ul style="list-style-type: none"> <li>• 0: Pure L2 regularization (Ridge)</li> <li>• 1: Pure L1 regularization (Lasso)</li> <li>• <math>0 &lt; \text{mixture} &lt; 1</math>: Elastic net (combination of L1 and L2) Only relevant when <code>penalty</code> <math>&gt; 0</math>. Can be tuned.</li> </ul>
<code>learn_rate</code>	A number for the learning rate. Can be tuned.
<code>optimizer</code>	A character string for the optimizer type ("adam", "sgd", "rmsprop"). Can be tuned.
<code>optimizer_args</code>	A named list of additional arguments passed to the optimizer. Cannot be tuned.
<code>loss</code>	A character string for the loss function ("mse", "mae", "cross_entropy", "bce"). Cannot be tuned.
<code>validation_split</code>	A number between 0 and 1 for the proportion of data used for validation. Can be tuned.
<code>device</code>	A character string for the device to use ("cpu", "cuda", "mps"). If NULL, auto-detects available GPU. Cannot be tuned.
<code>verbose</code>	Logical for whether to print training progress. Default FALSE. Cannot be tuned.

## Details

This function creates a model specification for a feedforward neural network that can be used within tidymodels workflows. The model supports:

- Multiple hidden layers with configurable units
- Various activation functions per layer
- GPU acceleration (CUDA, MPS, or CPU)

- Hyperparameter tuning integration
- Both regression and classification tasks

The `hidden_neurons` parameter accepts an integer vector where each element represents the number of neurons in that hidden layer. For example, `hidden_neurons = c(128, 64, 32)` creates a network with three hidden layers.

The `device` parameter controls where computation occurs:

- `NULL` (default): Auto-detect best available device (CUDA > MPS > CPU)
- `"cuda"`: Use NVIDIA GPU
- `"mps"`: Use Apple Silicon GPU
- `"cpu"`: Use CPU only

When tuning, you can use special tune tokens:

- For `hidden_neurons`: use `tune("hidden_neurons")` with a custom range
- For `activation`: use `tune("activation")` with values like `"relu"`, `"tanh"`

## Value

A model specification object with class `mlp_kinding`.

## Examples

```
if (torch::torch_is_installed()) {
  box::use(
    recipes[recipe],
    workflows[workflow, add_recipe, add_model],
    tune[tune],
    parsnip[fit]
  )

  # Model specs
  mlp_spec = mlp_kinding(
    mode = "classification",
    hidden_neurons = c(128, 64, 32),
    activation = c("relu", "relu", "relu"),
    epochs = 100
  )

  # If you want to tune
  mlp_tune_spec = mlp_kinding(
    mode = "classification",
    hidden_neurons = tune(),
    activation = tune(),
    epochs = tune(),
    learn_rate = tune()
  )
  wf = workflow() |>
    add_recipe(recipe(Species ~ ., data = iris)) |>
    add_model(mlp_spec)
```

```

    fit_wf = fit(wf, data = iris)
} else {
  message("Torch not fully installed - skipping example")
}

```

**ordinal\_gen***Ordinal Suffixes Generator***Description**

This function is originally from `numform::f_ordinal()`.

**Usage**

```
ordinal_gen(x)
```

**Arguments**

x	Vector of numbers. Could be a string equivalent
---	---

**Value**

Returns a string vector with ordinal suffixes.

**This is how you use it**

```
kindling:::ordinal_gen(1:10)
```

Note: This is not exported into public namespace. So please, refer to `numform::f_ordinal()` instead.

**References**

Rinker, T. W. (2021). `numform`: A publication style number and plot formatter version 0.7.0.  
<https://github.com/trinker/numform>

---

rnn_kindling	<i>Recurrent Neural Network via kindling</i>
--------------	--

---

## Description

`rnn_kindling()` defines a recurrent neural network model that can be used for classification or regression on sequential data. It integrates with the tidymodels ecosystem and uses the torch backend via kindling.

## Usage

```
rnn_kindling(  
  mode = "unknown",  
  engine = "kindling",  
  hidden_neurons = NULL,  
  rnn_type = NULL,  
  activations = NULL,  
  output_activation = NULL,  
  bias = NULL,  
  bidirectional = NULL,  
  dropout = NULL,  
  epochs = NULL,  
  batch_size = NULL,  
  penalty = NULL,  
  mixture = NULL,  
  learn_rate = NULL,  
  optimizer = NULL,  
  optimizer_args = NULL,  
  loss = NULL,  
  validation_split = NULL,  
  device = NULL,  
  verbose = NULL  
)
```

## Arguments

<code>mode</code>	A single character string for the type of model. Possible values are "unknown", "regression", or "classification".
<code>engine</code>	A single character string specifying what computational engine to use for fitting. Currently only "kindling" is supported.
<code>hidden_neurons</code>	An integer vector for the number of units in each hidden layer. Can be tuned.
<code>rnn_type</code>	A character string for the type of RNN cell ("rnn", "lstm", "gru"). Cannot be tuned.
<code>activations</code>	A character vector of activation function names for each hidden layer (e.g., "relu", "tanh", "sigmoid"). Can be tuned.

<code>output_activation</code>	A character string for the output activation function. Can be tuned.
<code>bias</code>	Logical for whether to include bias terms. Can be tuned.
<code>bidirectional</code>	A logical indicating whether to use bidirectional RNN. Can be tuned.
<code>dropout</code>	A number between 0 and 1 for dropout rate between layers. Can be tuned.
<code>epochs</code>	An integer for the number of training iterations. Can be tuned.
<code>batch_size</code>	An integer for the batch size during training. Can be tuned.
<code>penalty</code>	A number for the regularization penalty (lambda). Default 0 (no regularization). Higher values increase regularization strength. Can be tuned.
<code>mixture</code>	A number between 0 and 1 for the elastic net mixing parameter. Default 0 (pure L2/Ridge regularization). <ul style="list-style-type: none"> <li>• 0: Pure L2 regularization (Ridge)</li> <li>• 1: Pure L1 regularization (Lasso)</li> <li>• <math>0 &lt; \text{mixture} &lt; 1</math>: Elastic net (combination of L1 and L2) Only relevant when <code>penalty</code> <math>&gt; 0</math>. Can be tuned.</li> </ul>
<code>learn_rate</code>	A number for the learning rate. Can be tuned.
<code>optimizer</code>	A character string for the optimizer type ("adam", "sgd", "rmsprop"). Can be tuned.
<code>optimizer_args</code>	A named list of additional arguments passed to the optimizer. Cannot be tuned.
<code>loss</code>	A character string for the loss function ("mse", "mae", "cross_entropy", "bce"). Cannot be tuned.
<code>validation_split</code>	A number between 0 and 1 for the proportion of data used for validation. Can be tuned.
<code>device</code>	A character string for the device to use ("cpu", "cuda", "mps"). If NULL, auto-detects available GPU. Cannot be tuned.
<code>verbose</code>	Logical for whether to print training progress. Default FALSE. Cannot be tuned.

## Details

This function creates a model specification for a recurrent neural network that can be used within `tidymodels` workflows. The model supports:

- Multiple RNN types: basic RNN, LSTM, and GRU
- Bidirectional processing
- Dropout regularization
- GPU acceleration (CUDA, MPS, or CPU)
- Hyperparameter tuning integration
- Both regression and classification tasks

The `device` parameter controls where computation occurs:

- NULL (default): Auto-detect best available device (CUDA > MPS > CPU)
- "cuda": Use NVIDIA GPU
- "mps": Use Apple Silicon GPU
- "cpu": Use CPU only

**Value**

A model specification object with class `rnn_kinding`.

**Examples**

```
if (torch::torch_is_installed()) {
    box::use(
        recipes[recipe],
        workflows[workflow, add_recipe, add_model],
        parsnip[fit]
    )

    # Model specs
    rnn_spec = rnn_kinding(
        mode = "classification",
        hidden_neurons = c(64, 32),
        rnn_type = "lstm",
        activation = c("relu", "elu"),
        epochs = 100,
        bidirectional = TRUE
    )

    wf = workflow() |>
        add_recipe(recipe(Species ~ ., data = iris)) |>
        add_model(rnn_spec)

    fit_wf = fit(wf, data = iris)
    fit_wf
} else {
    message("Torch not fully installed - skipping example")
}
```

**table\_summary**

*Summarize and Display a Two-Column Data Frame as a Formatted Table*

**Description**

This function takes a two-column data frame and formats it into a summary-like table. The table can be optionally split into two parts, centered, and given a title. It is useful for displaying summary information in a clean, tabular format. The function also supports styling with ANSI colors and text formatting through the `{cli}` package and column alignment options.

**Usage**

```
table_summary(
    data,
```

```

title = NULL,
l = NULL,
header = FALSE,
center_table = FALSE,
border_char = "-",
style = list(),
align = NULL,
...
)

```

## Arguments

<code>data</code>	A data frame with exactly two columns. The data to be summarized and displayed.
<code>title</code>	A character string. An optional title to be displayed above the table.
<code>l</code>	An integer. The number of rows to include in the left part of a split table. If <code>NULL</code> , the table is not split.
<code>header</code>	A logical value. If <code>TRUE</code> , the column names of <code>data</code> are displayed as a header.
<code>center_table</code>	A logical value. If <code>TRUE</code> , the table is centered in the terminal.
<code>border_char</code>	Character used for borders. Default is "\u2500".
<code>style</code>	A list controlling the visual styling of table elements using ANSI formatting. Can include the following components: <ul style="list-style-type: none"> <li>• <code>left_col</code>: Styling for the left column values.</li> <li>• <code>right_col</code>: Styling for the right column values.</li> <li>• <code>border_text</code>: Styling for the border.</li> <li>• <code>title</code>: Styling for the title.</li> <li>• <code>sep</code>: Separator character between left and right column.</li> </ul> Each style component can be either a predefined style string (e.g., "blue", "red_italic", "bold") or a function that takes a context list with/without a <code>value</code> element and returns the styled text.
<code>align</code>	Controls the alignment of column values. Can be specified in three ways: <ul style="list-style-type: none"> <li>• A single string: affects only the left column (e.g., "left", "center", "right").</li> <li>• A vector of two strings: affects both columns in order (e.g., <code>c("left", "right")</code>).</li> <li>• A list with named components: explicitly specifies alignment for each column</li> </ul>
...	Additional arguments (currently unused).

## Value

This function does not return a value. It prints the formatted table to the console.

## Examples

```
# Create a sample data frame
df = data.frame(
  Category = c("A", "B", "C", "D", "E"),
  Value = c(10, 20, 30, 40, 50)
)

# Display the table with a title and header
table_summary(df, title = "Sample Table", header = TRUE)

# Split the table after the second row and center it
table_summary(df, l = 2, center_table = TRUE)

# Use styling and alignment
table_summary(
  df, header = TRUE,
  style = list(
    left_col = "blue_bold",
    right_col = "red",
    title = "green",
    border_text = "yellow"
  ),
  align = c("center", "right")
)

# Use custom styling with lambda functions
table_summary(
  df, header = TRUE,
  style = list(
    left_col = \$(ctx) cli::col_red(ctx), # ctx$value is another option
    right_col = \$(ctx) cli::col_blue(ctx)
  ),
  align = list(left_col = "left", right_col = "right")
)
```

# Index

act\_funcs, 2  
args, 3  
  
ffnn, 3  
ffnn(), 11  
ffnn\_generator, 6  
ffnn\_wrapper, 10  
  
garson.ffnn\_fit (kindling-varimp), 15  
grid\_depth, 11  
grid\_depth(), 11  
  
kindling-varimp, 15  
  
mlp\_kindling, 17  
  
olden.ffnn\_fit (kindling-varimp), 15  
ordinal\_gen, 20  
  
rnn (ffnn), 3  
rnn(), 11  
rnn\_generator (ffnn\_generator), 6  
rnn\_kindling, 21  
rnn\_wrapper (ffnn\_wrapper), 10  
  
table\_summary, 23  
tune::tune\_grid(), 11  
  
vi\_model.ffnn\_fit (kindling-varimp), 15