

Package ‘localLLM’

October 15, 2025

Type Package

Title Running Local LLMs with 'llama.cpp' Backend

Version 1.0.1

Date 2025-10-08

Description The 'localLLM' package provides R bindings to the 'llama.cpp' library for running large language models.

The package uses a lightweight architecture where the C++ backend library is downloaded at runtime rather than bundled with the package.

Package features include text generation, reproducible generation, and parallel inference.

License MIT + file LICENSE

Depends R (>= 3.6.0)

Imports Rcpp (>= 1.0.14), tools, utils

Suggests testthat (>= 3.0.0), covr

URL <https://github.com/EddieYang211/localLLM>

BugReports <https://github.com/EddieYang211/localLLM/issues>

SystemRequirements C++17, libcurl (optional, for model downloading)

Encoding UTF-8

LazyData true

RoxygenNote 7.3.2

NeedsCompilation yes

Author Eddie Yang [aut] (ORCID: <<https://orcid.org/0000-0002-3696-3226>>),
Yaosheng Xu [aut, cre] (ORCID: <<https://orcid.org/0009-0006-8138-369X>>)

Maintainer Yaosheng Xu <xu2009@purdue.edu>

Repository CRAN

Date/Publication 2025-10-15 19:10:08 UTC

Contents

localLLM-package	2
ag_news_sample	4
apply_chat_template	5
apply_gemma_chat_template	6
backend_free	7
backend_init	7
context_create	8
detokenize	9
download_model	10
generate	11
generate_parallel	13
get_lib_path	14
get_model_cache_dir	15
install_localLLM	15
lib_is_installed	16
list_cached_models	16
model_load	17
quick_llama	19
quick_llama_reset	21
set_hf_token	21
smart_chat_template	22
tokenize	23
tokenize_test	24

Index

25

localLLM-package *R Interface to llama.cpp with Runtime Library Loading*

Description

Provides R bindings to the llama.cpp library for running large language models locally. This package uses an innovative lightweight architecture where the C++ backend library is downloaded at runtime rather than bundled with the package, enabling zero-configuration AI inference in R with enterprise-grade performance.

Details

The localLLM package brings state-of-the-art language models to R users through a carefully designed four-layer architecture that combines ease of use with high performance.

```
## Quick Start 1. Install the R package: install.packages("localLLM") 2. Download backend library: install_localLLM() 3. Start generating text: quick_llama("Hello, how are you?")
```

Key Features

- **Zero Configuration:** One-line setup with automatic model downloading

- **High Performance:** Native C++ inference engine with GPU support
- **Cross Platform:** Pre-compiled binaries for Windows, macOS, and Linux
- **Memory Efficient:** Smart caching and memory management
- **Production Ready:** Robust error handling and comprehensive documentation

Architecture Overview The package uses a layered design:

- **High-Level API:** `quick_llama` for simple text generation
- **Mid-Level API:** `model_load`, `generate` for detailed control
- **Low-Level API:** Direct access to tokenization and context management
- **C++ Backend:** `llama.cpp` engine with dynamic loading

Main Functions

- `install_localLLM` - Download and install backend library
- `quick_llama` - High-level text generation (recommended for beginners)
- `model_load` - Load GGUF models with smart caching
- `context_create` - Create inference contexts
- `generate` - Generate text with full parameter control
- `tokenize / detokenize` - Text <-> Token conversion
- `apply_chat_template` - Format conversations for chat models

Example Workflows

Basic Text Generation

```
# Simple one-liner
response <- quick_llama("Explain quantum computing")

# With custom parameters
creative_text <- quick_llama("Write a poem about AI",
                             temperature = 0.9,
                             max_tokens = 150)
```

Advanced Usage with Custom Models

```
# Load your own model
model <- model_load("path/to/your/model.gguf")
ctx <- context_create(model, n_ctx = 4096)

# Direct text generation with auto-tokenization
output <- generate(ctx, "The future of AI is", max_tokens = 100)
```

Batch Processing

```
# Process multiple prompts efficiently
prompts <- c("Summarize AI trends", "Explain machine learning", "What is deep learning?")
responses <- quick_llama(prompts)
```

Supported Model Formats The package works with GGUF format models from various sources:

- Hugging Face Hub (automatic download)
- Local .gguf files
- Custom quantized models
- Ollama-compatible models

Performance Tips

- Use n_gpu_layers = -1 to fully utilize GPU acceleration
- Set n_threads to match your CPU cores for optimal performance
- Use larger n_ctx values for longer conversations
- Enable use_mlock for frequently used models to prevent swapping

Author(s)

Eddie Yang and Yaosheng Xu <xu2009@purdue.edu>

References

<https://github.com/EddieYang211/localLLM>

See Also

Useful links:

- <https://github.com/EddieYang211/localLLM>
- Report bugs at <https://github.com/EddieYang211/localLLM/issues>

ag_news_sample

AG News classification sample

Description

A 100-row subset of the AG News Topic Classification dataset consisting of 25 documents from each of the four classes (World, Sports, Business, Sci/Tech). The sample is intended for quick demonstrations and tests without requiring the full external dataset.

Usage

`data(ag_news_sample)`

Format

A data frame with 100 rows and 3 character columns:

class News topic label ("World", "Sports", "Business", or "Sci/Tech").

title Headline of the news article.

description Short description for the article.

Details

The sample was obtained from `textdata::dataset_ag_news()` (Zhang et al., 2015) using a fixed random seed to ensure reproducibility. It is provided solely for illustrative purposes.

Source

Zhang, X., Zhao, J., & LeCun, Y. (2015). "Character-level Convolutional Networks for Text Classification." arXiv:1509.01626. Original data distributed via the AG News Topic Classification dataset.

See Also

`[textdata::dataset_ag_news()]`

`apply_chat_template` *Apply Chat Template to Format Conversations*

Description

Formats conversation messages using the model's built-in chat template or a custom template. This is essential for chat models that expect specific formatting for multi-turn conversations.

Usage

```
apply_chat_template(model, messages, template = NULL, add_assistant = TRUE)
```

Arguments

<code>model</code>	A model object created with model_load
<code>messages</code>	List of chat messages, each with 'role' and 'content' fields. Role should be 'user', 'assistant', or 'system'
<code>template</code>	Optional custom template string (default: NULL, uses model's built-in template)
<code>add_assistant</code>	Whether to add assistant prompt suffix for response generation (default: TRUE)

Value

Formatted prompt string ready for text generation

See Also

[model_load](#), [quick_llama](#), [generate](#)

Examples

```
## Not run:
# Load a chat model
model <- model_load("path/to/chat_model.gguf")

# Format a conversation
messages <- list(
  list(role = "system", content = "You are a helpful assistant."),
  list(role = "user", content = "What is machine learning?"),
  list(role = "assistant", content = "Machine learning is..."),
  list(role = "user", content = "Give me an example.")
)

# Apply chat template
formatted_prompt <- apply_chat_template(model, messages)

# Generate response
response <- quick_llama(formatted_prompt)

## End(Not run)
```

`apply_gemma_chat_template`

Apply Gemma-Compatible Chat Template

Description

Creates a properly formatted chat template for Gemma models, which use `<start_of_turn>` and `<end_of_turn>` markers instead of ChatML format. This function addresses compatibility issues with `apply_chat_template()` when used with Gemma models.

Usage

```
apply_gemma_chat_template(messages, add_assistant = TRUE)
```

Arguments

<code>messages</code>	A list of message objects, each with 'role' and 'content' fields
<code>add_assistant</code>	Whether to add the assistant turn prefix (default: TRUE)

Value

A character string with properly formatted Gemma chat template

Examples

```
## Not run:  
messages <- list(  
  list(role = "system", content = "You are a helpful assistant."),  
  list(role = "user", content = "Hello!")  
)  
formatted <- apply_gemma_chat_template(messages)  
  
## End(Not run)
```

backend_free*Free localLLM backend*

Description

Clean up backend resources. Usually called automatically.

Usage

```
backend_free()
```

Value

No return value, called for side effects (frees backend resources).

backend_init*Initialize localLLM backend*

Description

Initialize the backend library. This should be called once before using other functions.

Usage

```
backend_init()
```

Value

No return value, called for side effects (initializes backend).

context_create *Create Inference Context for Text Generation*

Description

Creates a context object that manages the computational state for text generation. The context maintains the conversation history and manages memory efficiently for processing input tokens and generating responses. Each model can have multiple contexts with different settings.

Usage

```
context_create(  
    model,  
    n_ctx = 2048L,  
    n_threads = 4L,  
    n_seq_max = 1L,  
    verbosity = 1L  
)
```

Arguments

<code>model</code>	A model object returned by model_load
<code>n_ctx</code>	Maximum context length in tokens (default: 2048). This determines how many tokens of conversation history can be maintained. Larger values require more memory but allow for longer conversations. Must not exceed the model's maximum context length
<code>n_threads</code>	Number of CPU threads for inference (default: 4). Set to the number of available CPU cores for optimal performance. Only affects CPU computation
<code>n_seq_max</code>	Maximum number of parallel sequences (default: 1). Used for batch processing multiple conversations simultaneously. Higher values require more memory
<code>verbosity</code>	Control backend logging during context creation (default: 1L). Larger values print more information: 0 emits only errors, 1 includes warnings, 2 adds informational logs, and 3 enables the most verbose debug output.

Value

A context object (external pointer) used for text generation with [generate](#)

See Also

[model_load](#), [generate](#), [tokenize](#)

Examples

```
## Not run:  
# Load model and create basic context  
model <- model_load("path/to/model.gguf")  
ctx <- context_create(model)  
  
# Create context with larger buffer for long conversations  
long_ctx <- context_create(model, n_ctx = 4096)  
  
# High-performance context with more threads  
fast_ctx <- context_create(model, n_ctx = 2048, n_threads = 8)  
  
# Context for batch processing multiple conversations  
batch_ctx <- context_create(model, n_ctx = 2048, n_seq_max = 4)  
  
# Create context with minimal verbosity (quiet mode)  
quiet_ctx <- context_create(model, verbosity = 2L)  
  
## End(Not run)
```

detokenize

Convert Token IDs Back to Text

Description

Converts a sequence of integer token IDs back into human-readable text. This is the inverse operation of tokenization and is typically used to convert model output tokens into text that can be displayed to users.

Usage

```
detokenize(model, tokens)
```

Arguments

model	A model object created with model_load . Must be the same model that was used for tokenization to ensure proper decoding
tokens	Integer vector of token IDs to convert back to text. These are typically generated by tokenize or generate

Value

Character string containing the decoded text corresponding to the input tokens

See Also

[tokenize](#), [generate](#), [model_load](#)

Examples

```
## Not run:
# Load model
model <- model_load("path/to/model.gguf")

# Tokenize then detokenize (round-trip)
original_text <- "Hello, how are you today?"
tokens <- tokenize(model, original_text)
Recovered_text <- detokenize(model, tokens)
print(Recovered_text) # Should match original_text

# Generate and display text
ctx <- context_create(model)
generated_text <- generate(ctx, "The weather is", max_tokens = 10)

# Inspect individual tokens
single_token <- c(123) # Some token ID
token_text <- detokenize(model, single_token)
print(paste("Token", single_token, "represents:", token_text))

## End(Not run)
```

`download_model` *Download a model manually*

Description

Download a model manually

Usage

```
download_model(
  model_url,
  output_path = NULL,
  show_progress = TRUE,
  verify_integrity = TRUE,
  max_retries = 3,
  hf_token = NULL
)
```

Arguments

<code>model_url</code>	URL of the model to download (currently only supports https://)
<code>output_path</code>	Local path where to save the model (optional, will use cache if not provided)
<code>show_progress</code>	Whether to show download progress (default: TRUE)
<code>verify_integrity</code>	Verify file integrity after download (default: TRUE)

max_retries	Maximum number of download retries (default: 3)
hf_token	Optional Hugging Face access token to use for this download. Defaults to the existing ‘HF_TOKEN’ environment variable.

Value

The path where the model was saved

Examples

```
## Not run:
# Download to specific location
download_model(
  "https://example.com/model.gguf",
  file.path(tempdir(), "my_model.gguf")
)

# Download to cache (path will be returned)
cached_path <- download_model("https://example.com/model.gguf")

## End(Not run)
```

Description

Generates text using a loaded language model context with automatic tokenization. Simply provide a text prompt and the model will handle tokenization internally. This function now has a unified API with [generate_parallel](#).

Usage

```
generate(
  context,
  prompt,
  max_tokens = 100L,
  top_k = 40L,
  top_p = 1,
  temperature = 0,
  repeat_last_n = 0L,
  penalty_repeat = 1,
  seed = 1234L,
  clean = FALSE
)
```

Arguments

context	A context object created with context_create
prompt	Character string containing the input text prompt
max_tokens	Maximum number of tokens to generate (default: 100). Higher values produce longer responses
top_k	Top-k sampling parameter (default: 40). Limits vocabulary to k most likely tokens. Use 0 to disable
top_p	Top-p (nucleus) sampling parameter (default: 1.0). Cumulative probability threshold for token selection
temperature	Sampling temperature (default: 0.0). Set to 0 for greedy decoding. Higher values increase creativity
repeat_last_n	Number of recent tokens to consider for repetition penalty (default: 0). Set to 0 to disable
penalty_repeat	Repetition penalty strength (default: 1.0). Values >1 discourage repetition. Set to 1.0 to disable
seed	Random seed for reproducible generation (default: 1234). Use positive integers for deterministic output
clean	If TRUE, strip common chat-template control tokens from the generated text (default: FALSE).

Value

Character string containing the generated text

See Also

[quick_llama](#), [generate_parallel](#), [context_create](#)

Examples

```
## Not run:
# Load model and create context
model <- model_load("path/to/model.gguf")
ctx <- context_create(model, n_ctx = 2048)

response <- generate(ctx, "Hello, how are you?", max_tokens = 50)

# Creative writing with higher temperature
story <- generate(ctx, "Once upon a time", max_tokens = 200, temperature = 0.8)

# Prevent repetition
no_repeat <- generate(ctx, "Tell me about AI",
                      repeat_last_n = 64,
                      penalty_repeat = 1.1)

# Clean output (remove special tokens)
clean_output <- generate(ctx, "Explain quantum physics", clean = TRUE)
```

```
## End(Not run)
```

generate_parallel *Generate Text in Parallel for Multiple Prompts*

Description

Generate Text in Parallel for Multiple Prompts

Usage

```
generate_parallel(
  context,
  prompts,
  max_tokens = 100L,
  top_k = 40L,
  top_p = 1,
  temperature = 0,
  repeat_last_n = 0L,
  penalty_repeat = 1,
  seed = 1234L,
  progress = FALSE,
  clean = FALSE
)
```

Arguments

<code>context</code>	A context object created with context_create
<code>prompts</code>	Character vector of input text prompts
<code>max_tokens</code>	Maximum number of tokens to generate (default: 100)
<code>top_k</code>	Top-k sampling parameter (default: 40). Limits vocabulary to k most likely tokens
<code>top_p</code>	Top-p (nucleus) sampling parameter (default: 1.0). Cumulative probability threshold for token selection
<code>temperature</code>	Sampling temperature (default: 0.0). Set to 0 for greedy decoding. Higher values increase creativity
<code>repeat_last_n</code>	Number of recent tokens to consider for repetition penalty (default: 0). Set to 0 to disable
<code>penalty_repeat</code>	Repetition penalty strength (default: 1.0). Values >1 discourage repetition. Set to 1.0 to disable
<code>seed</code>	Random seed for reproducible generation (default: 1234). Use positive integers for deterministic output
<code>progress</code>	If TRUE, displays a console progress bar indicating batch completion status while generations are running (default: FALSE).
<code>clean</code>	If TRUE, remove common chat-template control tokens from each generated text (default: FALSE).

Value

Character vector of generated texts

<code>get_lib_path</code>	<i>Get Backend Library Path</i>
---------------------------	---------------------------------

Description

Returns the full path to the installed localLLM backend library.

Usage

```
get_lib_path()
```

Details

This function will throw an error if the backend library is not installed. Use `lib_is_installed` to check installation status first.

Value

Character string containing the path to the backend library file.

See Also

`lib_is_installed`, `install_localLLM`

Examples

```
## Not run:  
# Get the library path (only if installed)  
if (lib_is_installed()) {  
  lib_path <- get_lib_path()  
  message("Library is at: ", lib_path)  
}  
  
## End(Not run)
```

get_model_cache_dir *Get the model cache directory*

Description

Get the model cache directory

Usage

`get_model_cache_dir()`

Value

Path to the directory where models are cached

install_localLLM *Install localLLM Backend Library*

Description

This function downloads and installs the pre-compiled C++ backend library required for the local-LLM package to function.

Usage

`install_localLLM()`

Details

This function downloads platform-specific pre-compiled binaries from GitHub releases. The backend library is stored in the user's data directory and loaded at runtime. Internet connection is required for the initial download.

Value

Returns NULL invisibly. Called for side effects.

See Also

[lib_is_installed](#), [get_lib_path](#)

Examples

```
## Not run:  
# Install the backend library  
install_localLLM()  
  
## End(Not run)
```

`lib_is_installed` *Check if Backend Library is Installed*

Description

Checks whether the localLLM backend library has been downloaded and installed.

Usage

```
lib_is_installed()
```

Value

Logical value indicating whether the backend library is installed.

See Also

[install_localLLM](#), [get_lib_path](#)

Examples

```
# Check if backend library is installed
if (lib_is_installed()) {
  message("Backend library is ready")
} else {
  message("Please run install_localLLM() first")
}
```

`list_cached_models` *List cached models on disk*

Description

Enumerates the models that have been downloaded to the local cache. This is useful when you want to reuse a previously downloaded model but no longer remember the original URL. The cache directory can be overridden with the ‘LOCALLLM_CACHE_DIR’ environment variable or via the ‘cache_dir’ argument.

Usage

```
list_cached_models(cache_dir = NULL)
```

Arguments

<code>cache_dir</code>	Optional cache directory to inspect. Defaults to the package cache used by ‘model_load()’.
------------------------	--

Value

A data frame with one row per cached model and the columns ‘name’ (file name), ‘path’ (absolute path), ‘size_bytes’, and ‘modified’. Returns an empty data frame when no models are cached.

model_load*Load Language Model with Automatic Download Support*

Description

Loads a GGUF format language model from local path or URL with intelligent caching and download management. Supports various model sources including Hugging Face, Ollama repositories, and direct HTTPS URLs. Models are automatically cached to avoid repeated downloads.

Usage

```
model_load(  
  model_path,  
  cache_dir = NULL,  
  n_gpu_layers = 0L,  
  use_mmap = TRUE,  
  use_mlock = FALSE,  
  show_progress = TRUE,  
  force_redownload = FALSE,  
  verify_integrity = TRUE,  
  check_memory = TRUE,  
  hf_token = NULL,  
  verbosity = 1L  
)
```

Arguments

model_path	Path to local GGUF model file, URL, or cached model name. Supported URL formats: <ul style="list-style-type: none">• <code>https://</code> - Direct download from web servers If you previously downloaded a model through this package you can supply the cached file name (or a distinctive fragment of it) instead of the full path or URL. The loader will search the local cache and offer any matches.
cache_dir	Custom directory for downloaded models (default: NULL uses system cache directory)
n_gpu_layers	Number of transformer layers to offload to GPU (default: 0 for CPU-only). Set to -1 to offload all layers, or a positive integer for partial offloading
use_mmap	Enable memory mapping for efficient model loading (default: TRUE). Disable only if experiencing memory issues
use_mlock	Lock model in physical memory to prevent swapping (default: FALSE). Enable for better performance but requires sufficient RAM

show_progress Display download progress bar for remote models (default: TRUE)
force_redownload Force re-download even if cached version exists (default: FALSE). Useful for updating to newer model versions
verify_integrity Verify file integrity using checksums when available (default: TRUE)
check_memory Check if sufficient system memory is available before loading (default: TRUE)
hf_token Optional Hugging Face access token to set during model resolution. Defaults to the existing ‘HF_TOKEN’ environment variable.
verbosity Control backend logging during model loading (default: 1L). Larger numbers print more detail: 0 shows only errors, 1 adds warnings, 2 prints informational messages, and 3 enables the most verbose debug output.

Value

A model object (external pointer) that can be used with [context_create](#), [tokenize](#), and other model functions

See Also

[context_create](#), [download_model](#), [get_model_cache_dir](#), [list_cached_models](#)

Examples

```
## Not run:
# Load local GGUF model
model <- model_load("/path/to/my_model.gguf")

# Download from Hugging Face and cache locally
hf_path = "https://huggingface.co/Qwen/Qwen3-0.6B-GGUF/resolve/main/Qwen3-0.6B-Q8_0.gguf"
model <- model_load(hf_path)

# Load with GPU acceleration (offload 10 layers)
model <- model_load("/path/to/model.gguf", n_gpu_layers = 10)

# Download to custom cache directory
model <- model_load(hf_path,
                     cache_dir = file.path(tempdir(), "my_models"))

# Force fresh download (ignore cache)
model <- model_load(hf_path,
                     force_redownload = TRUE)

# High-performance settings for large models
model <- model_load("/path/to/large_model.gguf",
                     n_gpu_layers = -1,      # All layers on GPU
                     use_mlock = TRUE)      # Lock in memory

# Load with minimal verbosity (quiet mode)
model <- model_load("/path/to/model.gguf", verbosity = 2L)
```

```
## End(Not run)
```

quick_llama	<i>Quick LLaMA Inference</i>
-------------	------------------------------

Description

A high-level convenience function that provides one-line LLM inference. Automatically handles model downloading, loading, and text generation with optional chat template formatting and system prompts for instruction-tuned models.

Usage

```
quick_llama(  
  prompt,  
  model = .get_default_model(),  
  n_threads = NULL,  
  n_gpu_layers = "auto",  
  n_ctx = 2048L,  
  verbosity = 1L,  
  max_tokens = 100L,  
  top_k = 40L,  
  top_p = 1,  
  temperature = 0,  
  repeat_last_n = 0L,  
  penalty_repeat = 1,  
  min_p = 0.05,  
  system_prompt = "You are a helpful assistant.",  
  auto_format = TRUE,  
  chat_template = NULL,  
  stream = FALSE,  
  seed = 1234L,  
  progress = interactive(),  
  clean = TRUE,  
  ...  
)
```

Arguments

prompt	Character string or vector of prompts to process
model	Model URL or path (default: Llama 3.2 3B Instruct Q5_K_M)
n_threads	Number of threads (default: auto-detect)
n_gpu_layers	Number of GPU layers (default: auto-detect)
n_ctx	Context size (default: 2048)

verbosity	Backend logging verbosity (default: 1L). Higher values show more detail: 0 prints only errors, 1 adds warnings, 2 includes informational messages, and 3 enables the most verbose debug output.
max_tokens	Maximum tokens to generate (default: 100)
top_k	Top-k sampling (default: 40). Limits vocabulary to k most likely tokens
top_p	Top-p sampling (default: 1.0). Set to 0.9 for nucleus sampling
temperature	Sampling temperature (default: 0.0). Higher values increase creativity
repeat_last_n	Number of recent tokens to consider for repetition penalty (default: 0). Set to 0 to disable
penalty_repeat	Repetition penalty strength (default: 1.0). Set to 1.0 to disable
min_p	Minimum probability threshold (default: 0.05)
system_prompt	System prompt to add to conversation (default: "You are a helpful assistant.")
auto_format	Whether to automatically apply chat template formatting (default: TRUE)
chat_template	Custom chat template to use (default: NULL uses model's built-in template)
stream	Whether to stream output (default: auto-detect based on interactive())
seed	Random seed for reproducibility (default: 1234)
progress	Show a console progress bar when running parallel generation. Default: interactive(). Has no effect for single-prompt runs.
clean	Whether to strip chat-template control tokens from the generated output. Defaults to TRUE.
...	Additional parameters passed to generate() or generate_parallel()

Value

Character string (single prompt) or named list (multiple prompts)

See Also

[model_load](#), [generate](#), [generate_parallel](#), [install_localLLM](#)

Examples

```
## Not run:
# Simple usage with default settings (deterministic)
response <- quick_llama("Hello, how are you?")

# Raw text generation without chat template
raw_response <- quick_llama("Complete this: The capital of France is",
                             auto_format = FALSE)

# Custom system prompt
code_response <- quick_llama("Write a Python hello world program",
                             system_prompt = "You are a Python programming expert.")

# Creative writing with higher temperature
creative_response <- quick_llama("Tell me a story",
```

```

temperature = 0.8,
max_tokens = 200)

# Prevent repetition
no_repeat <- quick_llama("Explain AI",
                           repeat_last_n = 64,
                           penalty_repeat = 1.1)

# Multiple prompts (parallel processing)
responses <- quick_llama(c("Summarize AI", "Explain quantum computing"),
                           max_tokens = 150)

## End(Not run)

```

quick_llama_reset *Reset quick_llama state*

Description

Clears cached model and context objects, forcing fresh initialization on the next call to quick_llama().

Usage

```
quick_llama_reset()
```

Value

No return value, called for side effects (resets cached state).

set_hf_token *Configure Hugging Face access token*

Description

Utility helper to manage the ‘HF_TOKEN’ environment variable used for authenticated downloads from Hugging Face. The token is set for the current R session, and it can optionally be persisted to a ‘.Renvironment’ file for future sessions. The token is not printed back to the console.

Usage

```
set_hf_token(token, persist = FALSE, r environ_path = NULL)
```

Arguments

- token** Character scalar. Your Hugging Face access token, typically starting with ‘hf_’.
persist Logical flag controlling whether to persist the token to a startup file. Defaults to ‘FALSE’.
renviron_path Optional path to the ‘.Renviron’ file to update when ‘persist = TRUE’. Must be supplied explicitly when persisting.

Value

Invisibly returns the currently active token value.

Examples

```
## Not run:
set_hf_token("hf_xxx")
tmp_env <- file.path(tempdir(), ".Renviron_localLLM")
set_hf_token("hf_xxx", persist = TRUE, renviron_path = tmp_env)

## End(Not run)
```

smart_chat_template *Smart Chat Template Application*

Description

Automatically detects the model type and applies the appropriate chat template. For Gemma models, uses the Gemma-specific format. For other models, falls back to the standard apply_chat_template function.

Usage

```
smart_chat_template(model, messages, template = NULL, add_assistant = TRUE)
```

Arguments

- model** A model object created with model_load
messages A list of message objects
template Custom template (passed to apply_chat_template if not Gemma)
add_assistant Whether to add assistant turn prefix

Value

Formatted chat template string

tokenize	<i>Convert Text to Token IDs</i>
----------	----------------------------------

Description

Converts text into a sequence of integer token IDs that the language model can process. This is the first step in text generation, as models work with tokens rather than raw text. Different models may use different tokenization schemes (BPE, SentencePiece, etc.).

Usage

```
tokenize(model, text, add_special = TRUE)
```

Arguments

model	A model object created with model_load
text	Character string or vector to tokenize. Can be a single text or multiple texts
add_special	Whether to add special tokens like BOS (Beginning of Sequence) and EOS (End of Sequence) tokens (default: TRUE). These tokens help models understand text boundaries

Value

Integer vector of token IDs corresponding to the input text. These can be used with [generate](#) for text generation or [detokenize](#) to convert back to text

See Also

[detokenize](#), [generate](#), [model_load](#)

Examples

```
## Not run:  
# Load model  
model <- model_load("path/to/model.gguf")  
  
# Basic tokenization  
tokens <- tokenize(model, "Hello, world!")  
print(tokens) # e.g., c(15339, 11, 1917, 0)  
  
# Tokenize without special tokens (for model inputs)  
raw_tokens <- tokenize(model, "Continue this text", add_special = FALSE)  
  
# Tokenize multiple texts  
batch_tokens <- tokenize(model, c("First text", "Second text"))  
  
# Check tokenization of specific phrases  
question_tokens <- tokenize(model, "What is AI?")
```

```
print(length(question_tokens)) # Number of tokens  
## End(Not run)
```

tokenize_test *Test tokenize function (debugging)*

Description

Test tokenize function (debugging)

Usage

```
tokenize_test(model)
```

Arguments

model A model object

Value

Integer vector of tokens for "H"

Index

- * **datasets**
 - ag_news_sample, [4](#)
- * **package**
 - localLLM-package, [2](#)
- ag_news_sample, [4](#)
- apply_chat_template, [3, 5](#)
- apply_gemma_chat_template, [6](#)
- backend_free, [7](#)
- backend_init, [7](#)
- context_create, [3, 8, 12, 13, 18](#)
- detokenize, [3, 9, 23](#)
- download_model, [10, 18](#)
- generate, [3, 5, 8, 9, 11, 20, 23](#)
- generate_parallel, [11, 12, 13, 20](#)
- get_lib_path, [14, 15, 16](#)
- get_model_cache_dir, [15, 18](#)
- install_localLLM, [3, 14, 15, 16, 20](#)
- lib_is_installed, [14, 15, 16](#)
- list_cached_models, [16, 18](#)
- localLLM(localLLM-package), [2](#)
- localLLM-package, [2](#)
- model_load, [3, 5, 8, 9, 17, 20, 23](#)
- quick_llama, [3, 5, 12, 19](#)
- quick_llama_reset, [21](#)
- set_hf_token, [21](#)
- smart_chat_template, [22](#)
- tokenize, [3, 8, 9, 18, 23](#)
- tokenize_test, [24](#)