Package 'maestro'

January 7, 2025

Type Package

Title Orchestration of Data Pipelines

Version 0.5.0

Maintainer Will Hipson <will.e.hipson@gmail.com>

Description Framework for creating and orchestrating data pipelines. Organize, orchestrate, and monitor multiple pipelines in a single project. Use tags to decorate functions with scheduling parameters and configuration.

License MIT + file LICENSE

URL <https://github.com/whipson/maestro>,

<https://whipson.github.io/maestro/>

BugReports <https://github.com/whipson/maestro/issues>

Encoding UTF-8

LazyData true

Imports cli (>= 3.3.0), dplyr (>= 1.1.0), glue, lifecycle, logger, lubridate (>= 1.9.1), purrr (>= 1.0.0), R.utils, R6, rlang (>= 1.0.0), roxygen2, tictoc, timechange, utils

RoxygenNote 7.3.2

Depends R (>= 4.1.0)

Suggests asciicast, DiagrammeR, furrr, future, knitr, quarto, rmarkdown, rstudioapi, testthat (>= 3.0.0), withr

Config/testthat/edition 3

VignetteBuilder knitr, quarto

NeedsCompilation no

Author Will Hipson [cre, aut, cph] (<<https://orcid.org/0000-0002-3931-2189>>),
Ryan Garnett [aut, ctb, cph]

Repository CRAN

Date/Publication 2025-01-07 13:40:02 UTC

Contents

---

build_schedule   *Build a schedule table*

Description

Builds a schedule data.frame for scheduling pipelines in run_schedule().

Usage

```
build_schedule(pipeline_dir = "./pipelines", quiet = FALSE)
```

Arguments

pipeline_dir	path to directory containing the pipeline scripts
quiet	silence metrics to the console (default = FALSE)

Details

This function parses the maestro tags of functions located in pipeline_dir which is conventionally called 'pipelines'. An orchestrator requires a schedule table to determine which pipelines are to run and when. Each row in a schedule table is a pipeline name and its scheduling parameters such as its frequency.

The schedule table is mostly intended to be used by run_schedule() immediately. In other words, it is not recommended to make changes to it.

Value

MaestroSchedule

Examples

```
# Creating a temporary directory for demo purposes! In practice, just
# create a 'pipelines' directory at the project level.

if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline("my_new_pipeline", pipeline_dir, open = FALSE)
  build_schedule(pipeline_dir = pipeline_dir)
}
```

---

  create_maestro   *Creates a new maestro project*

Description

Creates a new maestro project

Usage

```
create_maestro(path, type = "R", overwrite = FALSE, quiet = FALSE, ...)
```

Arguments

path	file path for the orchestrator script
type	file type for the orchestrator (supports R, Quarto, and RMarkdown)
overwrite	whether to overwrite an existing orchestrator or maestro project
quiet	whether to silence messages in the console (default = FALSE)
...	unused

Value

invisible

Examples

```
# Creates a new maestro project with an R orchestrator
if (interactive()) {
  new_proj_dir <- tempdir()
  create_maestro(new_proj_dir, type = "R", overwrite = TRUE)

  create_maestro(new_proj_dir, type = "Quarto", overwrite = TRUE)
}
```

---

create_orchestrator   *Create a new orchestrator*

Description

Create a new orchestrator

Usage

```
create_orchestrator(
  path,
  type = c("R", "Quarto", "RMarkdown"),
  open = interactive(),
  quiet = FALSE,
  overwrite = FALSE
)
```

Arguments

path	file path for the orchestrator script
type	file type for the orchestrator (supports R, Quarto, and RMarkdown)
open	whether or not to open the script upon creation
quiet	whether to silence messages in the console (default = FALSE)
overwrite	whether to overwrite an existing orchestrator or maestro project

Value

invisible

---

create_pipeline   *Create a new pipeline in a pipelines directory*

Description

Allows the creation of new pipelines (R scripts) and fills in the maestro tags as specified.

Usage

```
create_pipeline(
  pipe_name,
  pipeline_dir = "pipelines",
  frequency = "1 day",
  start_time = Sys.Date(),
  tz = "UTC",
```

```
    log_level = "INFO",
    quiet = FALSE,
    open = interactive(),
    overwrite = FALSE,
    skip = FALSE,
    inputs = NULL,
    outputs = NULL
)
```

Arguments

`pipe_name`	name of the pipeline and function
`pipeline_dir`	directory containing the pipeline scripts
`frequency`	how often the pipeline should run (e.g., 1 day, daily, 3 hours, 4 months). Fills in maestroFrequency tag
`start_time`	start time of the pipeline schedule. Fills in maestroStartTime tag
`tz`	timezone that pipeline will be scheduled in. Fills in maestroTz tag
`log_level`	log level for the pipeline (e.g., INFO, WARN, ERROR). Fills in maestroLogLevel tag
`quiet`	whether to silence messages in the console (default = `FALSE`)
`open`	whether or not to open the script upon creation
`overwrite`	whether or not to overwrite an existing pipeline of the same name and location.
`skip`	whether to skip the pipeline when running in the orchestrator (default = `FALSE`)
`inputs`	vector of names of pipelines that input into this pipeline (default = `NULL` for no inputs)
`outputs`	vector of names of pipelines that receive output from this pipeline (default = `NULL` for no outputs)

Value

invisible

Examples

```
if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline(
    "extract_data",
    pipeline_dir = pipeline_dir,
    frequency = "1 hour",
    open = FALSE,
    quiet = TRUE,
    overwrite = TRUE
  )

  create_pipeline(
    "new_job",
```

```
      pipeline_dir = pipeline_dir,
      frequency = "20 minutes",
      start_time = as.POSIXct("2024-06-21 12:20:00"),
      log_level = "ERROR",
      open = FALSE,
      quiet = TRUE,
      overwrite = TRUE
    )
  }
```

---

get_artifacts   *Get the artifacts (return values) of the pipelines in a MaestroSchedule*
                                 object.

Description

Artifacts are return values from pipelines. They are accessible as a named list where the names
correspond to the names of the pipeline.

Usage

```
get_artifacts(schedule)
```

Arguments

schedule   object of type MaestroSchedule created using `build_schedule()`

Value

named list

Examples

```
if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline("my_new_pipeline", pipeline_dir, open = FALSE)
  schedule <- build_schedule(pipeline_dir = pipeline_dir)

  schedule <- run_schedule(
    schedule,
    orch_frequency = "1 day",
    quiet = TRUE
  )

  get_artifacts(schedule)

  # Alternatively, use the underlying R6 method
  schedule$get_artifacts()
}
```

---

get_schedule *Get the schedule from a MaestroSchedule object*

Description

A schedule is represented as a table where each row is a pipeline and the columns contain scheduling parameters such as the frequency and start time.

Usage

```
get_schedule(schedule)
```

Arguments

schedule   object of type MaestroSchedule created using `build_schedule()`

Details

The schedule table is used internally in a MaestroSchedule object but can be accessed using this function or accessing the R6 method of the MaestroSchedule object.

Value

data.frame

Examples

```
if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline("my_new_pipeline", pipeline_dir, open = FALSE)
  schedule <- build_schedule(pipeline_dir = pipeline_dir)

  get_schedule(schedule)

  # Alternatively, use the underlying R6 method
  schedule$get_schedule()
}
```

---

get_status *Get the statuses of the pipelines in a MaestroSchedule object*

Description

A status data.frame contains the names and locations of the pipelines as well as information around whether they were invoked, the status (error, warning, etc.), and the run time.

Usage

```
get_status(schedule)
```

Arguments

schedule   object of type MaestroSchedule created using `build_schedule()`

Value

data.frame

Examples

```
if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline("my_new_pipeline", pipeline_dir, open = FALSE)
  schedule <- build_schedule(pipeline_dir = pipeline_dir)

  schedule <- run_schedule(
    schedule,
    orch_frequency = "1 day",
    quiet = TRUE
  )

  get_status(schedule)

  # Alternatively, use the underlying R6 method
  schedule$get_status()
}
```

---

invoke   *Manually run a pipeline regardless of schedule*

Description

Instantly run a single pipeline from the schedule. This is useful for testing purposes or if you want to just run something one-off.

Usage

```
invoke(schedule, pipe_name, resources = list(), ...)
```

Arguments

schedule   object of type MaestroSchedule created using `build_schedule()`

pipe_name   name of a single pipe name from the schedule

resources   named list of shared resources made available to pipelines as needed

...   other arguments passed to `run_schedule()`

Details

Scheduling parameters such as the frequency, start time, and specifiers are ignored. The pipeline will be run even if maestroSkip is present.

Value

invisible

Examples

```
if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline("my_new_pipeline", pipeline_dir, open = FALSE)
  schedule <- build_schedule(pipeline_dir = pipeline_dir)

  invoke(schedule, "my_new_pipeline")
}
```

---

last_build_errors *Retrieve latest maestro build errors*

Description

Gets the latest schedule build errors following use of build_schedule(). If the build succeeded or build_schedule() has not been run it will be NULL.

Usage

```
last_build_errors()
```

Value

error messages

Examples

```
last_build_errors()
```

---

last_run_errors   *Retrieve latest maestro pipeline errors*

Description

Gets the latest pipeline errors following use of run_schedule(). If the all runs succeeded or run_schedule() has not been run it will be NULL.

Usage

```
last_run_errors()
```

Value

error messages

Examples

```
last_run_errors()
```

---

last_run_messages   *Retrieve latest maestro pipeline messages*

Description

Gets the latest pipeline messages following use of run_schedule(). If there are no messages or run_schedule() has not been run it will be NULL.

Usage

```
last_run_messages()
```

Details

Note that setting maestroLogLevel to something greater than INFO will ignore messages.

Value

messages

Examples

```
last_run_messages()
```

`last_run_warnings`   *Retrieve latest maestro pipeline warnings*

Description

Gets the latest pipeline warnings following use of `run_schedule()`. If there are no warnings or `run_schedule()` has not been run it will be `NULL`.

Usage

```
last_run_warnings()
```

Details

Note that setting `maestroLogLevel` to something greater than `WARN` will ignore warnings.

Value

warning messages

Examples

```
last_run_warnings()
```

`MaestroSchedule`   *Class for a schedule of pipelines*

Description

Class for a schedule of pipelines

Class for a schedule of pipelines

Public fields

`PipelineList` object of type MaestroPipelineList

Methods

Public methods:

- `MaestroSchedule$new()`
- `MaestroSchedule$print()`
- `MaestroSchedule$run()`
- `MaestroSchedule$get_schedule()`
- `MaestroSchedule$get_status()`
- `MaestroSchedule$get_artifacts()`

- `MaestroSchedule$get_network()`
- `MaestroSchedule$show_network()`
- `MaestroSchedule$clone()`

Method new(): Create a MaestroSchedule object

Usage:
`MaestroSchedule$new(Pipelines = NULL)`

Arguments:
`Pipelines` list of MaestroPipelines

Returns: MaestroSchedule

Method print(): Print the schedule object

Usage:
`MaestroSchedule$print()`

Returns: print

Method run(): Run a MaestroSchedule

Usage:
`MaestroSchedule$run(..., quiet = FALSE, run_all = FALSE, n_show_next = 5)`

Arguments:
`...` arguments passed to MaestroPipelineList$run

`quiet` whether or not to silence console messages

`run_all` run all pipelines regardless of the schedule (default is `FALSE`) - useful for testing.

`n_show_next` show the next n scheduled pipes

Returns: invisible

Method get_schedule(): Get the schedule as a data.frame

Usage:
`MaestroSchedule$get_schedule()`

Returns: data.frame

Method get_status(): Get status of the pipelines as a data.frame

Usage:
`MaestroSchedule$get_status()`

Returns: data.frame

Method get_artifacts(): Get artifacts (return values) from the pipelines

Usage:
`MaestroSchedule$get_artifacts()`

Returns: list

Method get_network(): Get the network structure of the pipelines as an edge list (will be empty if there are no DAG pipelines)

Usage:

```
MaestroSchedule$get_network()
```

Returns: data.frame

Method `show_network()`: Visualize the DAG relationships between pipelines in the schedule

Usage:

```
MaestroSchedule$show_network()
```

Returns: interactive visualization

Method `clone()`: The objects of this class are cloneable with this method.

Usage:

```
MaestroSchedule$clone(deep = FALSE)
```

Arguments:

deep Whether to make a deep clone.

Examples

```
if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline("my_new_pipeline", pipeline_dir, open = FALSE)
  schedule <- build_schedule(pipeline_dir = pipeline_dir)
}
```

---

maestro_tags | *Maestro Tags*

Description

maestro tags are roxygen2 comments for configuring the scheduling and execution of pipelines.

Details

maestro tags follow the format #' @maestroTagName

Tag List:

tagName	description	value	examples (con
maestroFrequency	Time unit for scheduling	string	1 hour, daily, 3
maestroLogLevel	Threshold for logging when logging is requested	string	INFO, WARN
maestroSkip	Skips the pipeline when running (presence of tag indicates to skip)	n/a	
maestroStartTime	Start time of the pipeline; sets the point in time for recurrence	date or timestamp	1970-01-01 00
maestroTz	Timezone of the start time	string	UTC, America
maestroHours	Hours of day to run pipeline	ints	0 12 23
maestroDays	Days of week or days of month to run pipeline	ints or strings	1 14 30, Mon
maestroMonths	Months of year to run pipeline	ints	1 3 9 12

maestroInputs	Pipelines that input into this pipeline	strings	my_upstream_
maestroOutputs	Pipelines that take the output from this pipeline	strings	my_downstrea
maestro	Generic tag for identifying a maestro pipeline with all defaults	n/a.	

---

run_schedule   *Run a schedule*

Description

Given a schedule in a maestro project, runs the pipelines that are scheduled to execute based on the current time.

Usage

```
run_schedule(
  schedule,
  orch_frequency = ”1 day”,
  check_datetime = lubridate::now(tzone = ”UTC”),
  resources = list(),
  run_all = FALSE,
  n_show_next = 5,
  cores = 1,
  logging = lifecycle::deprecated(),
  log_file = lifecycle::deprecated(),
  log_file_max_bytes = 1e+06,
  quiet = FALSE,
  log_to_console = FALSE,
  log_to_file = FALSE
)
```

Arguments

schedule	object of type MaestroSchedule created using build_schedule()
orch_frequency	of the orchestrator, a single string formatted like "1 day", "2 weeks", "hourly", etc.
check_datetime	datetime against which to check the running of pipelines (default is current system time in UTC)
resources	named list of shared resources made available to pipelines as needed
run_all	run all pipelines regardless of the schedule (default is FALSE) - useful for testing. Does not apply to pipes with a maestroSkip tag.
n_show_next	show the next n scheduled pipes
cores	number of cpu cores to run if running in parallel. If > 1, furrr is used and a multisession plan must be executed in the orchestrator (see details)

logging	whether or not to write the logs to a file (deprecated in 0.5.0 - use `log_to_file` and/or `log_to_console` arguments instead)
log_file	path to the log file (ignored if `log_to_file == FALSE`) (deprecated in 0.5.0 - use `log_to_file`)
log_file_max_bytes	
	numeric specifying the maximum number of bytes allowed in the log file before purging the log (within a margin of error)
quiet	silence metrics to the console (default = FALSE). Note this does not affect messages generated from pipelines when `log_to_console = TRUE`.
log_to_console	whether or not to include pipeline messages, warnings, errors to the console (default = FALSE) (see Logging & Console Output section)
log_to_file	either a boolean to indicate whether to create and append to a `maestro.log` or a character path to a specific log file. If `FALSE` or `NULL` it will not log to a file.

Details

Pipeline schedule logic:

The function `run_schedule()` examines each pipeline in the schedule table and determines whether it is scheduled to run at the current time using some simple time arithmetic. We assume `run_schedule(schedule, check_datetime = Sys.time())`, but this need not be the case.

Output:

`run_schedule()` returns the same MaestroSchedule object with modified attributes. Use `get_status()` to examine the status of each pipeline and use `get_artifacts()` to get any return values from the pipelines as a list.

Pipelines with arguments (resources):

If a pipeline takes an argument that doesn't include a default value, these can be supplied in the orchestrator via `run_schedule(resources = list(arg1 = val))`. The name of the argument used by the pipeline must match the name of the argument in the list. Currently, each named resource must refer to a single object. In other words, you can't have two pipes using the same argument but requiring different values.

Running in parallel:

Pipelines can be run in parallel using the `cores` argument. First, you must run `future::plan(future::multisession)` in the orchestrator. Then, supply the desired number of cores to the `cores` argument. Note that console output appears different in multicore mode.

Logging & Console Output:

By default, `maestro` suppresses pipeline messages, warnings, and errors from appearing in the console, but messages coming from `print()` and other console logging packages like `cli` and `logger` are not suppressed and will be interwoven into the output generated from `run_schedule()`. Messages from `cat()` and related functions are always suppressed due to the nature of how those functions operate with standard output.

Users are advised to make use of R's `message()`, `warning()`, and `stop()` functions in their pipelines for managing conditions. Use `log_to_console = TRUE` to print these to the console.

Maestro can generate a log file that is appended to each time the orchestrator is run. Use `log_to_file` = TRUE or `log_to_file = '[path-to-file]'` and maestro will create/append to a file in the project directory. This log file will be appended to until it exceeds the byte size defined in `log_file_max_bytes` argument after which the log file is deleted.

Value

MaestroSchedule object

Examples

```
if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline("my_new_pipeline", pipeline_dir, open = FALSE)
  schedule <- build_schedule(pipeline_dir = pipeline_dir)

  # Runs the schedule every 1 day
  run_schedule(
    schedule,
    orch_frequency = "1 day",
    quiet = TRUE
  )

  # Runs the schedule every 15 minutes
  run_schedule(
    schedule,
    orch_frequency = "15 minutes",
    quiet = TRUE
  )
}
```

---

show_network   *Visualize the schedule as a DAG*

Description

Create an interactive network visualization to show the dependency structure of pipelines in the schedule. This is only useful if there are pipelines in the schedule that take inputs/outputs from other pipelines.

Usage

```
show_network(schedule)
```

Arguments

schedule   object of type MaestroSchedule created using `build_schedule()`

Details

Note that running this function on a schedule with all independent pipelines will produce a network visual with no connections.

This function requires the installation of `DiagrammeR` which is not automatically installed with `maestro`.

Value

DiagrammeR visualization

Examples

```
if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline("my_new_pipeline", pipeline_dir, open = FALSE)
  schedule <- build_schedule(pipeline_dir = pipeline_dir)

  schedule <- run_schedule(
    schedule,
    orch_frequency = "1 day",
    quiet = TRUE
  )

  show_network(schedule)
}
```

---

suggest_orch_frequency

Suggest orchestrator frequency based on a schedule

Description

Suggests a frequency to run the orchestrator based on the frequencies of the pipelines in a schedule.

Usage

```
suggest_orch_frequency(
  schedule,
  check_datetime = lubridate::now(tzone = "UTC")
)
```

Arguments

schedule   MaestroSchedule object created by `build_schedule()`

check_datetime datetime against which to check the running of pipelines (default is current system time in UTC)

Details

This function attempts to find the smallest interval of time between all pipelines. If the smallest interval is less than 15 minutes, it just uses the smallest interval.

Note this function is intended to be used interactively when deciding how often to schedule the orchestrator. Programmatic use is not recommended.

Value

frequency string

Examples

```
if (interactive()) {
  pipeline_dir <- tempdir()
  create_pipeline("my_new_pipeline", pipeline_dir, open = FALSE)
  schedule <- build_schedule(pipeline_dir = pipeline_dir)
  suggest_orch_frequency(schedule)
}
```

Index