Package 'pamr'

July 1, 2024

Title Pam: Prediction Analysis for Microarrays

Version 1.57

Description Some functions for sample classification in microarrays.

Depends R (>= 3.5)

License GPL-2

Date 2024-06-28

RoxygenNote 7.3.1

Encoding UTF-8

Imports cluster, survival

NeedsCompilation no

Author Trevor Hastie [aut],
Rob Tibshirani [aut],
Balasubramanian Narasimhan [aut, cre],
Gilbert Chu [aut]

Maintainer Balasubramanian Narasimhan <naras@stanford.edu>

Repository CRAN

Date/Publication 2024-07-01 06:40:02 UTC

Contents

---

khan   *Khan microarray data*

Description

The khan data frame has 2308 rows and 65 columns. These are one of the datasets data used in the
Tibshirani et al paper in PNAS on nearest shrunken centroids.

Details

The first two columns of gene ids and names and the remaining columns are gene expression values
for 63 samples. An attribute cancer_type contains the cancer type for each sample.

---

pamr.adaptthresh   *A function to adaptive choose threshold scales, for use in pamr.train*

Description

A function to adaptive choose threshold scales, for use in pamr.train

Usage

```
pamr.adaptthresh(object, ntries = 10, reduction.factor = 0.9, full.out = FALSE)
```

Arguments

object   The result of a call to pamr.train

ntries   Number of iterations to use in algorithm

reduction.factor
                  Amount by which a scaling is reduced in one step of the algorithm

full.out   Should full output be returned? Default FALSE

Details

pamr.adaptthresh Adaptively searches for set of good threshold scales. The baseline (default) scale is 1 for each class. The idea is that for easy to classify classes, the threshold scale can be increased without increasing the error rate for that class, and resulting in fewer genes needed for the classification rule. The scalings from pamr.adaptthresh are then used in pamr.train, and pamr.cv. The results may be better than those obtained with the default values of threshold.scale.

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

References

Robert Tibshirani, Trevor Hastie, Balasubramanian Narasimhan, and Gilbert Chu. "Diagnosis of multiple cancer types by shrunken centroids of gene expression" PNAS 2002 99:6567-6572 (May 14).

Robert Tibshirani, Trevor Hastie, Balasubramanian Narasimhan, and Gilbert Chu (2002). Class prediction by nearest shrunken centroids,with applications to DNA microarrays. Stanford tech report.

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y)
mytrain <-   pamr.train(mydata)
new.scales <- pamr.adaptthresh(mytrain)


mytrain2 <- pamr.train(mydata, threshold.scale=new.scales)

myresults2 <- pamr.cv(mytrain2, mydata)
```

---

pamr.batchadjust   *A function to mean-adjust microarray data by batches*

Description

A function to mean-adjust microarray data by batches

Usage

```
pamr.batchadjust(data)
```

Arguments

data   The input data. A list with components: x- an expression genes in the rows,
                    samples in the columns, and y- a vector of the class labels for each sample, and
                    batchlabels- a vector of batch labels for each sample.

Details

pamr.batchadjust does a genewise one-way ANOVA adjustment for expression values. Let
$x(i,j)$ be the expression for gene i in sample j. Suppose sample j in in batch b, and let B be
the set of all samples in batch b. Then pamr.batchadjust adjusts $x(i,j)$ to $x(i,j) - mean[x(i,j)]$
where the mean is taken over all samples j in B.

Value

A data object of the same form as the input data, with x replaced by the adjusted x

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
#generate some data
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
batchlabels <- sample(c(1:5),size=20,replace=TRUE)
mydata <- list(x=x,y=factor(y),batchlabels=factor(batchlabels))

mydata2 <- pamr.batchadjust(mydata)
```

---

pamr.confusion   *A function giving a table of true versus predicted values, from a nearest*
                            shrunken centroid fit.

Description

A function giving a table of true versus predicted values, from a nearest shrunken centroid fit.

Usage

```
pamr.confusion(fit, threshold, extra = TRUE)
```

Arguments

`fit`	The result of a call to pamr.train or pamr.cv
`threshold`	The desired threshold value
`extra`	Should the classwise and overall error rates be returned? Default TRUE

Details

`pamr.confusion` Gives a cross-tabulation of true versus predicted classes for the fit returned by pamr.train or pamr.cv, at the specified threshold.

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y)
mytrain <-   pamr.train(mydata)
mycv <- pamr.cv(mytrain,mydata)
pamr.confusion(mytrain,  threshold=2)
pamr.confusion(mycv,  threshold=2)
```

---

pamr.confusion.survival

Compute confusin matrix from pamr survival fit

Description

computes confusion matrix for (survival.time,censoring) outcome based on fit object "fit" and class predictions "yhat" soft response probabilities for (survival.time,censoring) are first estimated using Kaplan-Meier method applied to training data

Usage

```
pamr.confusion.survival(fit, survival.time, censoring.status, yhat)
```

Arguments

`fit`	The result of a call to pamr.train or pamr.cv
`survival.time`	Survival time
`censoring.status`	
	censoring status
`yhat`	class predictions

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

---

pamr.cv	*A function to cross-validate the nearest shrunken centroid classifier*

Description

A function to cross-validate the nearest shrunken centroid classifier produced by pamr.train

Usage

```
pamr.cv(fit, data, nfold = NULL, folds = NULL, ...)
```

Arguments

fit	The result of a call to pamr.train
data	A list with at least two components: x- an expression genes in the rows, samples in the columns), and y- a vector of the class labels for each sample. Same form as data object used by pamr.train.
nfold	Number of cross-validation folds. Default is the smallest class size
folds	A list with nfold components, each component a vector of indices of the samples in that fold. By default a (random) balanced cross-validation is used
...	Any additional arguments that are to be passed to pamr.train

Details

`pamr.cv` carries out cross-validation for a nearest shrunken centroid classifier.

Value

A list with components

threshold	A vector of the thresholds tried in the shrinkage
errors	The number of cross-validation errors for each threshold value
loglik	The cross-validated multinomial log-likelihood value for each threshold value
size	A vector of the number of genes that survived the thresholding, for each threshold value tried.
.	
yhat	A matrix of size n by nthreshold, containing the cross-validated class predictions for each threshold value, in each column
prob	A matrix of size n by nthreshold, containing the cross-validated class probabilities for each threshold value, in each column
folds	The cross-validation folds used
cv.objects	Train objects (output of pamr.train), from each of the CV folds
call	The calling sequence used

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)

mydata <- list(x=x,y=factor(y), geneid=as.character(1:nrow(x)),
 genenames=paste("g",as.character(1:nrow(x)),sep=""))

mytrain <-   pamr.train(mydata)
mycv <- pamr.cv(mytrain,mydata)
```

---

pamr.decorrelate	*A function to decorrelate (adjust) the feature matrix with respect to some additional predictors*

Description

A function to decorrelate (adjust) the feature matrix with respect to some additional predictors

Usage

```
pamr.decorrelate(
  x,
  adjusting.predictors,
  xtest = NULL,
  adjusting.predictors.test = NULL
)
```

Arguments

x	Matrix of training set feature values, with genes in the rows, samples in the columns
adjusting.predictors	
	List of training set predictors to be used for adjustment
xtest	Optional matrix of test set feature values, to be adjusted in the same way as the training set
adjusting.predictors.test	
	Optional list of test set predictors to be used for adjustment

Details

pamr.decorrelate Does a least squares regression of each row of x on the adjusting predictors, and returns the residuals. If xtest is provided, it also returns the adjusted version of xtest, using the training set least squares regression model for adjustment

Value

A list with components

x.adj	Adjusted x matrix
xtest.adj	Adjusted xtest matrix, if xtest we provided

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

References

Robert Tibshirani, Trevor Hastie, Balasubramanian Narasimhan, and Gilbert Chu Diagnosis of multiple cancer types by shrunken centroids of gene expression PNAS 99: 6567-6572. Available at www.pnas.org

Examples

```
#generate some data
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)

x<-matrix(rnorm(1000*20),ncol=20)
y<-c(rep(1,10),rep(2,10))
adjusting.predictors=list(pred1=rnorm(20), pred2=as.factor(sample(c(1,2),replace
=TRUE,size=20)))
xtest=matrix(rnorm(1000*10),ncol=10)
adjusting.predictors.test=list(pred1=rnorm(10), pred2=as.factor(sample(c(1,2),replace
=TRUE,size=10)))

# decorrelate training x wrt adjusting predictors

x.adj=pamr.decorrelate(x,adjusting.predictors)$x.adj
# train classifier with adjusted x

d=list(x=x.adj,y=y)
a<-pamr.train(d)

# decorrelate training and test x wrt adjusting predictors, then make
#predictions for test set

temp <- pamr.decorrelate(x,adjusting.predictors, xtest=xtest,
                          adjusting.predictors.test=adjusting.predictors.test)

d=list(x=temp$x.adj,y=y)
```

```
a<-pamr.train(d)
aa<-pamr.predict(a,temp$xtest.adj, threshold=.5)
```

---

pamr.fdr	*A function to estimate false discovery rates for the nearest shrunken centroid classifier*

Description

A function to estimate false discovery rates for the nearest shrunken centroid classifier

Usage

```
pamr.fdr(
  trained.obj,
  data,
  nperms = 100,
  xl.mode = c("regular", "firsttime", "onetime", "lasttime"),
  xl.time = NULL,
  xl.prevfit = NULL
)
```

Arguments

trained.obj	The result of a call to pamr.train
data	Data object; same as the one passed to pamr.train
nperms	Number of permutations for estimation of FDRs. Default is 100
xl.mode	Used by Excel interface
xl.time	Used by Excel interface
xl.prevfit	Used by Excel interface

Details

`pamr.fdr` estimates false discovery rates for a nearest shrunken centroid classifier

Value

A list with components:

results	Matrix of estimates FDRs for various various threshold values. Reported are both the median and 90th percentile of the FDR over permutations
pi0	The estimated proportion of genes that are null, i.e. not significantly different

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)

mydata <- list(x=x,y=factor(y), geneid=as.character(1:nrow(x)),
               genenames=paste("g",as.character(1:nrow(x)),sep=""))

mytrain <-   pamr.train(mydata)
myfdr <- pamr.fdr(mytrain, mydata)
```

---

pamr.geneplot	*A function to plot the genes that surive the thresholding from the nearest shrunken centroid classifier*

Description

A function to plot the genes that survive the thresholding, from the nearest shrunken centroid classifier produced by pamr.train

Usage

```
pamr.geneplot(fit, data, threshold)
```

Arguments

fit	The result of a call to pamr.train
data	The input data. In the same format as the input data for pamr.train
threshold	The desired threshold value

Details

pamr.geneplot Plots the raw gene expression for genes that survive the specified threshold. Plot is stratified by class. Plot is set up to display only up to about 20 or 25 genes, otherwise it gets too crowded. Hence threshold should be chosen to yield at most about 20 or 25 genes.

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y)
mytrain <-   pamr.train(mydata)
pamr.geneplot(mytrain, mydata, threshold=1.6)
```

| pamr.indeterminate | *A function that takes estimate class probabilities and produces a class prediction or indeterminate prediction* |

Description

A function that takes estimate class probabilities and produces a class prediction or indeterminate prediction

Usage

```
pamr.indeterminate(prob, mingap = 0)
```

Arguments

prob	Estimated class probabilities, from pamr.predict with type="posterior")
mingap	Minimum difference between highest and second highest probability. If difference is < mingap, prediction is set to indeterminate (NA)

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y)
mytrain <-   pamr.train(mydata)
prob<- pamr.predict(mytrain, mydata$x , threshold=1, type="posterior")
pamr.indeterminate(prob,mingap=.75)
```

| pamr.listgenes | *A function to list the genes that survive the thresholding, from the nearest shrunken centroid classifier* |

Description

A function to list the genes that survive the thresholding, from the nearest shrunken centroid classifier produced by pamr.train

Usage

```
pamr.listgenes(fit, data, threshold, fitcv = NULL, genenames = FALSE)
```

Arguments

fit	The result of a call to pamr.train
data	The input data. In the same format as the input data for pamr.train
threshold	The desired threshold value
fitcv	Optional object, result of a call to pamr.cv
genenames	Include genenames in the list? If yes, they are taken from "data". Default is false (geneid is always included in the list).

Details

pamr.listgenes List the geneids, and standardized centroids for each class, for genes surviving at the given threshold. If fitcv is provided, the function also reports the average rank of the gene in the cross-validation folds, and the proportion of times that the gene is chosen (at the given threshold) in the cross-validation folds.

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
#generate some data
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)

mydata <- list(x=x,y=factor(y), geneid=as.character(1:nrow(x)),
               genenames=paste("g",as.character(1:nrow(x)),sep=""))


#train classifier
```

```
mytrain<-   pamr.train(mydata)

pamr.listgenes(mytrain, mydata, threshold=1.6)
```

---

pamr.makeclasses   *A function to interactively define classes from a clustering tree*

Description

function to interactively define classes from a clustering tree

Usage

```
pamr.makeclasses(data, sort.by.class = FALSE, ...)
```

Arguments

data	The input data. A list with components: x- an expression genes in the rows, samples in the columns, and y- a vector of the class labels for each sample, and batchlabels- a vector of batch labels for each sample. This object if the same form as that produced by pamr.from.excel.
sort.by.class	Optional argument. If true, the clustering tree is forced to put all samples in the same class (as defined by the class labels y in 'data') together in the tree. This is useful if a regrouping of classes is desired. Eg: given classes 1,2,3,4 you want to define new classes (1,3) vs (2,4) or 2 vs (1,3)
...	Any additional arguments to be passed to hclust

Details

pamr.makeclasses Using this function the user interactively defines a new set of classes, to be used in pamr.train, pamr.cv etc. After invoking pamr.makeclasses, a clustering tree is drawn. This callss the R function `hclust`, and any arguments for `hclust` can be passed to it. Using the left button, the user clicks at the junction point defining the subgroup 1. More groups can be added to class 1 by clicking on further junction points. The user ends the definition of class 1 by clicking on the rightmost button (in Windows, an additional menu appears and he chooses Stop). This process is continued for classes 2,3 etc. Note that some sample may be left out of the new classes. Two consecutive clicks of the right button ends the definition for all classes.

At the end, the clustering is redrawn, with the new class labels shown.

Note: this function is "fragile". The user must click close to the junction point, to avoid confusion with other junction points. Classes 1,2,3.. cannot have samples in common (if they do, an Error message will appear). If the function is confused about the desired choices, it will complain and ask the user to rerun pamr.makeclasses. The user should also check that the labels on the final redrawn cluster tree agrees with the desired classes.

Value

A vector of class labels 1,2,3... If a component is NA (missing), then the sample is not assigned
to any class. This vector should be assigned to the newy component of data, for use in pamr.train etc.
Note that pamr.train uses the class labels in the component newy`` if it is present. Otherwise it uses the data label

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
#generate some data
x <- matrix(rnorm(1000*40),ncol=40)
y <- sample(c(1:4),size=40,replace=TRUE)
batchlabels <- sample(c(1:5),size=40,replace=TRUE)

mydata <- list(x=x,y=factor(y),batchlabels=factor(batchlabels),
               geneid=as.character(1:nrow(x)),
               genenames=paste("g",as.character(1:nrow(x)),sep=""))

# mydata$newy <- pamr.makeclasses(mydata) Run this and define some new classes

train <- pamr.train(mydata)
```

---

pamr.menu   *A function that interactively leads the user through a PAM analysis*

Description

A function that interactively leads the user through a PAM analysis

Usage

```
pamr.menu(data)
```

Arguments

data   A list with at least two components: x- an expression genes in the rows, samples
                in the columns), and y- a vector of the class labels for each sample. Same form
                as data object used by pamr.train.

Details

pamr.menu provides a menu for training, cross-validating and plotting a nearest shrunken centroid
analysis.

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y)
#  pamr.menu(mydata)
```

---

pamr.plotcen	*A function to plot the shrunken class centroids, from the nearest shrunken centroid classifier*

Description

A function to plot the shrunken class centroids, from the nearest shrunken centroid classifier produced by pamr.train

Usage

```
pamr.plotcen(fit, data, threshold)
```

Arguments

fit	The result of a call to pamr.train
data	The input data, in the same form as that used by pamr.train
threshold	The desired threshold value

Details

`pamr.plotcen` plots the shrunken class centroids for each class, for genes surviving the threshold for at least once class. If genenames are included in "data", they are added to the plot. Note: for many classes and long gene names, this plot may need some manual prettying.

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y,genenames=as.character(1:1000))
mytrain <-   pamr.train(mydata)
mycv <- pamr.cv(mytrain,mydata)
pamr.plotcen(mytrain, mydata,threshold=1.6)
```

---

pamr.plotcv	*A function to plot the cross-validated error curves from the nearest shrunken centroid classifier*

Description

A function to plot the cross-validated error curves the nearest shrunken centroid classifier

Usage

```
pamr.plotcv(fit)
```

Arguments

fit   The result of a call to pamr.cv

Details

`pamr.plotcv` plots the cross-validated misclassification error curves, from nearest shrunken centroid classifier. An overall plot, and a plot by class, are produced.

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y)
mytrain <-   pamr.train(mydata)
mycv <-  pamr.cv(mytrain, mydata)
pamr.plotcv(mycv)
```

---

pamr.plotcvprob	*A function to plot the cross-validated sample probabilities from the nearest shrunken centroid classifier*

Description

A function to plot the cross-validated sample probabilities from the nearest shrunken centroid classifier

Usage

```
pamr.plotcvprob(fit, data, threshold)
```

Arguments

fit	The result of a call to pamr.cv
data	A list with at least two components: x- an expression genes in the rows, samples in the columns), and y- a vector of the class labels for each sample. Same form as data object used by pamr.train.
threshold	Threshold value to be used

Details

`pamr.plotcvprob` plots the cross-validated sample probabilities the from nearest shrunken centroid classifier, stratified by the true classses.

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y)
mytrain <-   pamr.train(mydata)
mycv <-  pamr.cv(mytrain,mydata)
pamr.plotcvprob(mycv,mydata,threshold=1.6)
```

pamr.plotfdr	*A function to plot the FDR curve from the nearest shrunken centroid classifier*

Description

A function to plot the FDR curve the nearest shrunken centroid classifier

Usage

```
pamr.plotfdr(fdrfit, call.win.metafile = FALSE)
```

Arguments

fdrfit   The result of a call to pamr.fdr

call.win.metafile
                Used by Excel interface

Details

`pamr.plotfdr` plots the FDR curves from nearest shrunken centroid classifier. The median FDR (solid line) and upper 90 percentile (broken line) are shown

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:2),size=20,replace=TRUE)
x[1:50,y==2]=x[1:50,y==2]+3
mydata <- list(x=x,y=y)
mytrain <-   pamr.train(mydata)
myfdr <-  pamr.fdr(mytrain, mydata)
pamr.plotfdr(myfdr)
```

| pamr.plotstrata | *A function to plot the survival curves in each Kaplan Meier stratum* |

Description

A function to plot the survival curves in each Kaplan Meier stratum

Usage

```
pamr.plotstrata(fit, survival.time, censoring.status)
```

Arguments

fit   The result of a call to pamr.train

survival.time   Vector of survival times
censoring.status
                    Vector of censoring status values

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
gendata<-function(n=100, p=2000){
  tim <- 3*abs(rnorm(n))
  u<-runif(n,min(tim),max(tim))
  y<-pmin(tim,u)
   ic<-1*(tim<u)
m <- median(tim)
x<-matrix(rnorm(p*n),ncol=n)
  x[1:100, tim>m] <-  x[1:100, tim>m]+3
  return(list(x=x,y=y,ic=ic))
}

# generate training data; 2000 genes, 100 samples

junk<-gendata(n=100)
y<-junk$y
ic<-junk$ic
x<-junk$x
d <- list(x=x,survival.time=y, censoring.status=ic,
geneid=as.character(1:nrow(x)), genenames=paste("g",as.character(1:nrow(x)),sep=
""))

# train model
a3<- pamr.train(d, ngroup.survival=2)
```

```
pamr.plotstrata(a3, d$survival.time, d$censoring.status)
```

---

pamr.plotsurvival   *A function to plots Kaplan-Meier curves stratified by a group variable*

Description

A function to plots Kaplan-Meier curves stratified by a group variable

Usage

```
pamr.plotsurvival(group, survival.time, censoring.status)
```

Arguments

group   A grouping factor

survival.time   Vector of survival times
censoring.status
                 Vector of censoring status values: 1=died, 0=censored

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
gendata<-function(n=100, p=2000){
  tim <- 3*abs(rnorm(n))
  u<-runif(n,min(tim),max(tim))
  y<-pmin(tim,u)
   ic<-1*(tim<u)
m <- median(tim)
x<-matrix(rnorm(p*n),ncol=n)
  x[1:100, tim>m] <-  x[1:100, tim>m]+3
  return(list(x=x,y=y,ic=ic))
}

# generate training data; 2000 genes, 100 samples

junk<-gendata(n=100)
y<-junk$y
ic<-junk$ic
x<-junk$x
d <- list(x=x,survival.time=y, censoring.status=ic,
geneid=as.character(1:nrow(x)), genenames=paste("g",as.character(1:nrow(x)),sep=
""))
```

```
# train model
a3<- pamr.train(d, ngroup.survival=2)

#make class predictions

yhat <- pamr.predict(a3,d$x, threshold=1.0)

pamr.plotsurvival(yhat, d$survival.time, d$censoring.status)
```

pamr.predict	*A function giving prediction information, from a nearest shrunken centroid fit.*

Description

A function giving prediction information, from a nearest shrunken centroid fit

Usage

```
pamr.predict(
  fit,
  newx,
  threshold,
  type = c("class", "posterior", "centroid", "nonzero"),
  prior = fit$prior,
  threshold.scale = fit$threshold.scale
)
```

Arguments

fit	The result of a call to pamr.train
newx	Matrix of features at which predictions are to be made
threshold	The desired threshold value
type	Type of prediction desired: class predictions, posterior probabilities, (unshrunken) class centroids, vector of genes surviving the threshold
prior	Prior probabilities for each class. Default is that specified in "fit"
threshold.scale	Additional scaling factors to be applied to the thresholds. Vector of length equal to the number of classes. Default is that specified in "fit".

Details

pamr.predict Give a cross-tabulation of true versus predicted classes for the fit returned by pamr.train or pamr.cv, at the specified threshold

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y)
mytrain <-   pamr.train(mydata)
mycv <- pamr.cv(mytrain,mydata)
pamr.predict(mytrain, mydata$x , threshold=1)
```

---

pamr.predictmany	*A function giving prediction information for many threshold values, from a nearest shrunken centroid fit.*

Description

A function giving prediction information for many threshold values, from a nearest shrunken centroid fit

Usage

```
pamr.predictmany(
  fit,
  newx,
  threshold = fit$threshold,
  prior = fit$prior,
  threshold.scale = fit$threshold.scale,
  ...
)
```

Arguments

fit	The result of a call to pamr.train
newx	Matrix of features at which predictions are to be made
threshold	The desired threshold values
prior	Prior probabilities for each class. Default is that specified in "fit"
threshold.scale	
	Additional scaling factors to be applied to the thresholds. Vector of length equal to the number of classes. Default is that specified in "fit".
...	Additional arguments to be passed to pamr.predict

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=y)
mytrain <-   pamr.train(mydata)

pamr.predictmany(mytrain, mydata$x)
```

---

pamr.surv.to.class2 | *A function to assign observations to categories, based on their survival times.*

Description

A function to assign observations to categories, based on their survival times.

Usage

```
pamr.surv.to.class2(
  y,
  icens,
  cutoffs = NULL,
  n.class = NULL,
  class.names = NULL,
  newy = y,
  newic = icens
)
```

Arguments

y	vector of survival times
icens	Vector of censorng status values: 1=died, 0=censored
cutoffs	Survival time cutoffs for categories. Default NULL
n.class	Number of classes to create: if cutoffs is NULL, n.class equal classes are created.
class.names	Character names for classes
newy	New set of survival times, for which probabilities are computed (see below). Default is y
newic	New set of censoring statuses, for which probabilities are computed (see below). Default is icens

Details

pamr.pamr.surv.to.class2 splits observations into categories based on their survival times and the Kaplan-Meier estimates. For example if n.class=2, it makes two categories, one below the median survival, the other above. For each observation (newy, ic), it then computes the probability of that observation falling in each category. For an uncensored observation that probability is just 1 or 0 depending on when the death occurred. For a censored observation, the probabilities are based on the Kaplan Meier and are typically between 0 and 1.

Value

class	The category labels
prob	The estimates class probabilities
cutoffs	The cutoffs used

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
gendata<-function(n=100, p=2000){
  tim <- 3*abs(rnorm(n))
  u<-runif(n,min(tim),max(tim))
  y<-pmin(tim,u)
   ic<-1*(tim<u)
m <- median(tim)
x<-matrix(rnorm(p*n),ncol=n)
  x[1:100, tim>m] <-  x[1:100, tim>m]+3
  return(list(x=x,y=y,ic=ic))
}

# generate training data; 2000 genes, 100 samples

junk<-gendata(n=100)
y<-junk$y
ic<-junk$ic
x<-junk$x
d <- list(x=x,survival.time=y, censoring.status=ic,
geneid=as.character(1:nrow(x)), genenames=paste("g",as.character(1:nrow(x)),sep=
""))

# train model
a3<- pamr.train(d, ngroup.survival=2)

# generate test data
junkk<- gendata(n=500)

dd <- list(x=junkk$x, survival.time=junkk$y, censoring.status=junkk$ic)

# compute soft labels
```

```
proby <-  pamr.surv.to.class2(dd$survival.time, dd$censoring.status,
              n.class=a3$ngroup.survival)$prob
```

---

pamr.test.errors.surv.compute

A function giving a table of true versus predicted values, from a nearest shrunken centroid fit from survival data.

Description

A function giving a table of true versus predicted values, from a nearest shrunken centroid fit from survival data.

Usage

```
pamr.test.errors.surv.compute(proby, yhat)
```

Arguments

proby   Survival class probabilities, from pamr.surv.to.class2

yhat   Estimated class labels, from pamr.predict

Details

`pamr.test.errors.surv.compute` computes the erros between the true 'soft" class labels proby and the estimated ones "yhat"

Author(s)

Trevor Hastie, Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

Examples

```
gendata<-function(n=100, p=2000){
  tim <- 3*abs(rnorm(n))
  u<-runif(n,min(tim),max(tim))
  y<-pmin(tim,u)
   ic<-1*(tim<u)
m <- median(tim)
x<-matrix(rnorm(p*n),ncol=n)
  x[1:100, tim>m] <-  x[1:100, tim>m]+3
  return(list(x=x,y=y,ic=ic))
}

# generate training data; 2000 genes, 100 samples
```

```
junk<-gendata(n=100)
y<-junk$y
ic<-junk$ic
x<-junk$x
d <- list(x=x,survival.time=y, censoring.status=ic,
geneid=as.character(1:nrow(x)), genenames=paste("g",as.character(1:nrow(x)),sep=
""))

# train model
a3<- pamr.train(d, ngroup.survival=2)

# generate test data
junkk<- gendata(n=500)

dd <- list(x=junkk$x, survival.time=junkk$y, censoring.status=junkk$ic)

# compute soft labels
proby <-  pamr.surv.to.class2(dd$survival.time, dd$censoring.status,
            n.class=a3$ngroup.survival)$prob


# make class predictions for test data
yhat <- pamr.predict(a3,dd$x, threshold=1.0)

# compute test errors

pamr.test.errors.surv.compute(proby, yhat)
```

---

pamr.train   *A function to train a nearest shrunken centroid classifier*

Description

A function that computes a nearest shrunken centroid for gene expression (microarray) data

Usage

```
pamr.train(
  data,
  gene.subset = NULL,
  sample.subset = NULL,
  threshold = NULL,
  n.threshold = 30,
  scale.sd = TRUE,
  threshold.scale = NULL,
  se.scale = NULL,
  offset.percent = 50,
```

```
    hetero = NULL,
    prior = NULL,
    remove.zeros = TRUE,
    sign.contrast = "both",
    ngroup.survival = 2
)
```

Arguments

data	The input data. A list with components: x- an expression genes in the rows, samples in the columns), and y- a vector of the class labels for each sample. Optional components- genenames, a vector of gene names, and geneid- a vector of gene identifiers.
gene.subset	Subset of genes to be used. Can be either a logical vector of length total number of genes, or a list of integers of the row numbers of the genes to be used
sample.subset	Subset of samples to be used. Can be either a logical vector of length total number of samples, or a list of integers of the column numbers of the samples to be used.
threshold	A vector of threshold values for the centroid shrinkage.Default is a set of 30 values chosen by the software
n.threshold	Number of threshold values desired (default 30)
scale.sd	Scale each threshold by the wthin class standard deviations? Default: true
threshold.scale	Additional scaling factors to be applied to the thresholds. Vector of length equal to the number of classes. Default- a vectors of ones.
se.scale	Vector of scaling factors for the within class standard errors. Default is sqrt(1/n.class-1/n), where n is the overall sample size and n.class is the sample sizes in each class. This default adjusts for different class sizes.
offset.percent	Fudge factor added to the denominator of each t-statistic, expressed as a percentile of the gene standard deviation values. This is a small positive quantity to penalize genes with expression values near zero, which can result in very large ratios. This factor is expecially impotant for Affy data. Default is the median of the standard deviations of each gene.
hetero	Should a heterogeneity transformation be done? If yes, hetero must be set to one of the class labels (see Details below). Default is no (hetero=NULL)
prior	Vector of length the number of classes, representing prior probabilities for each of the classes. The prior is used in Bayes rule for making class prediction. Default is NULL, and prior is then taken to be n.class/n, where n is the overall sample size and n.class is the sample sizes in each class.
remove.zeros	Remove threshold values yielding zero genes? Default TRUE
sign.contrast	Directions of allowed deviations of class-wise average gene expression from the overall average gene expression. Default is both'' (positive or negative). Can also be set to positive" or "negative".
ngroup.survival	Number of groups formed for survival data. Default 2

Details

pamr.train fits a nearest shrunken centroid classifier to gene expression data. Details may be found in the PNAS paper referenced below. One feature not described there is "heterogeneity analysis". Suppose there are two classes labelled "A" and "B". CLass "A" is considered a normal class, and "B" an abnormal class. Setting hetero="A" transforms expression values $x[i,j]$ to $|x[i,j] - mean(x[i,j])|$ where the mean is taken only over samples in class "A". The transformed feature values are then used in Pam. This is useful when the abnormal class "B" is heterogeneous, i.e. a given gene might have higher expresion than normal for some class "B" samples, and lower for others. With more than 2 classes, each class is centered on the class specified by hetero.

Value

A list with components

y	The outcome classes.
yhat	A matrix of predicted classes, each column representing the results from one threshold.
.	
prob	A array of predicted class probabilities. of dimension n by nclass by n.threshold. n is the number samples, nclass is the number of classes, n.threshold is the number of thresholds tried
centroids	A matrix of (unshrunken) class centroids, n by nclass
hetero	Value of hetero used in call to pamr.train
norm.cent	Centroid of "normal" group, if hetero was specified
centroid.overall	
	A vector containing the (unshrunken) overall centroid (all classes together)
sd	A vector of the standard deviations for each gene
threshold	A vector of the threshold tried in the shrinkage
nonzero	A vector of the number of genes that survived the thresholding, for each threshold value tried
threshold.scale	
	A vector of threshold scale factors that were used
se.scale	A vector of standard error scale factors that were used
call	The calling sequence used
prior	The prior probabilities used
errors	The number of trainin errors for each threshold value

Author(s)

Trevor Hastie,Robert Tibshirani, Balasubramanian Narasimhan, and Gilbert Chu

References

Robert Tibshirani, Trevor Hastie, Balasubramanian Narasimhan, and Gilbert Chu Diagnosis of multiple cancer types by shrunken centroids of gene expression PNAS 99: 6567-6572. Available at www.pnas.org

Examples

```
#generate some data
suppressWarnings(RNGversion("3.5.0"))
set.seed(120)
x <- matrix(rnorm(1000*20),ncol=20)
y <- sample(c(1:4),size=20,replace=TRUE)
mydata <- list(x=x,y=factor(y))

#train classifier
results<-   pamr.train(mydata)

# train classifier on all  data except class 4
results2 <- pamr.train(mydata,sample.subset=(mydata$y!=4))

# train classifier on  only the first 500 genes
results3 <- pamr.train(mydata,gene.subset=1:500)
```

Index