

Package ‘RPIV’

October 19, 2025

Title Residual Prediction Test for Well-Specification of Instrumental Variable Models

Version 1.0.0

Description A test for the well-specification of the linear instrumental variable model. The test is based on trying to predict the residuals of a two-stage least-squares regression using a random forest. Details can be found in Scheidegger, Londschen and Bühlmann (2025) ``A residual prediction test for the well-specification of linear instrumental variable models''
[<doi:10.48550/arXiv.2506.12771>](https://doi.org/10.48550/arXiv.2506.12771).

URL <https://github.com/cyrillsch/RPIV>

License GPL (>= 3)

Encoding UTF-8

RoxygenNote 7.3.3

Imports ranger, stats

Suggests testthat (>= 3.0.0)

Config/testthat/edition 3

NeedsCompilation no

Author Cyril Scheidegger [aut, cre, cph] (ORCID:
[<https://orcid.org/0009-0005-2851-1384>](https://orcid.org/0009-0005-2851-1384))

Maintainer Cyril Scheidegger <cyrill.scheidegger@stat.math.ethz.ch>

Repository CRAN

Date/Publication 2025-10-19 13:10:07 UTC

Contents

RPIV_test	2
Index	5

RPIV_test*Residual Prediction Test for Linear Instrumental Variable Models*

Description

Performs a hypothesis test for the well-specification of linear instrumental variable (IV) model. More specifically, it tests the null-hypothesis $H_0 : \exists \beta \in \mathbb{R}^p$ s.t. $\mathbb{E}[Y - X^T \beta | Z] = 0$. It uses sample splitting and a random forest to try to predict the two-stage least-squares residuals from the instruments Z .

Usage

```
RPIV_test(
  Y,
  X,
  C = NULL,
  Z,
  frac_A = NULL,
  gamma = 0.05,
  variance_estimator = "heteroskedastic",
  clustering = NULL,
  upper_clip_quantile = 0.8,
  regr_par = list(),
  fit_intercept = TRUE
)
```

Arguments

Y	A numeric vector. The outcome variable.
X	A numeric matrix or vector. The endogenous explanatory variables.
C	A numeric matrix, vector or NULL. The additional exogenous explanatory variables (optional).
Z	A numeric matrix or vector. The instruments.
frac_A	A numeric scalar between 0 and 1 or NULL. The fraction of the sample used for training (sample splitting). Default is $\min(0.5, \exp(1)/\log(n))$, where n is the sample size.
gamma	A non-negative scalar. If the variance estimator is less than gamma times the noise level (as estimated as by the mean of the squared residuals), gamma times the noise level is used as variance estimator.
variance_estimator	Character string or vector. One or more of "homoskedastic", "heteroskedastic", "cluster". Specifies the types of variance estimation used.
clustering	A vector of cluster identifiers or NULL. Observations with the same value of clustering belong to the same cluster. Required if variance_estimator includes "cluster".

upper_clip_quantile	A scalar between 0 and 1. The estimated weight-function will be clipped at the corresponding quantile of the random forest predictions on the auxiliary sample. Use 0 to use the sign of the predictions. Default is 0.8.
regr_par	A list of parameters passed to the random forest regression model. Supports num.trees, num.mtry (number of different mtry values to try out) or a vector mtry, a vector max.depth, num.min.node.size (number of different min.node.size values to try out) or a vector min.node.size.
fit_intercept	Logical. Should an intercept be included in the model? Default is TRUE.

Details

The RPIV test splits the sample into an auxiliary and a main sample. On the auxiliary sample, a random forest is used to predict the two-stage least squares residuals from the instruments. The test statistic is the scalar product of the two-stage least-squares residuals with a clipped and rescaled version of the learned function evaluated on the main sample divided by an estimator of its standard deviation.

If clustering is supplied, sample splitting is done at cluster level (also for variance_estimator "homoskedastic" or "heteroskedastic").

Value

If a single variance estimator is used, returns a list with:

- p_value** p-value of the residual prediction test.
- test_statistic** The value of the test statistic.
- var_fraction** The estimated variance fraction, i.e., variance estimator divided by noise level estimate.
- T_null** The value of the initial test statistic. If var_fraction $\geq \gamma$, it is equal to test_statistic, otherwise, it has larger absolute value.
- variance_estimator** The variance estimator used.

If multiple estimators are supplied, returns a named list of such results for each estimator.

References

Cyrill Scheidegger, Malte Londschen and Peter Bühlmann. A residual prediction test for the well-specification of linear instrumental variable models. Preprint, [doi:10.48550/arXiv.2506.12771](https://doi.org/10.48550/arXiv.2506.12771), 2025.

Examples

```
set.seed(1)
n <- 100
Z <- rnorm(n)
H <- rnorm(n)
C <- rnorm(n)
X <- Z + rnorm(n) + H
Y1 <- X - C - H + rnorm(n)
```

```
Y2 <- X - C - H + rnorm(n) + Z^2  
RPIV_test(Y1, X, C, Z)  
RPIV_test(Y2, X, C, Z)
```

Index

RPIV_test, [2](#)