Package 'RRI'

October 12, 2022

Type Package

Title Residual Randomization Inference for Regression Models

Version 1.1

Author Panos Toulis

Maintainer Panos Toulis <panos.toulis@chicagobooth.edu>

Description Testing and inference for regression models using residual randomization methods. The basis of inference is an invariance assumption on the regression errors, e.g., clustered errors, or doubly-clustered errors.

License GPL-2

Encoding UTF-8

LazyData true

Imports Rcpp (>= 1.0.1)

LinkingTo Rcpp, RcppArmadillo

RoxygenNote 7.0.2

NeedsCompilation yes

Repository CRAN

Date/Publication 2019-12-19 06:40:03 UTC

R topics documented:

Index **16**

---

check_model *Checks whether the input* model *is valid.*

Description

Checks whether the input model is valid.

Usage

```
check_model(model)
```

Arguments

model   A model object. See example_model for details.

---

example_clustering *An example* clustering *object. A clustering is a* List *that splits indexes 1..#num_datapoints to clusters. Each* List *element corresponds to one cluster. The clustering is not necessarily a partition but it usually is.*

Description

An example clustering object. A clustering is a List that splits indexes 1..#num_datapoints to clusters. Each List element corresponds to one cluster. The clustering is not necessarily a partition but it usually is.

Usage

```
example_clustering()
```

Value

A List for the clustering of indexes 1..#num_datapoints.

---

example_model   *Example regression model and H0.*

Description

Example regression model and H0.

Usage

```
example_model(n = 100)
```

Arguments

n   Number of datapoints.

Value

List of (y, X, lam, lam0) that corresponds to regression model and null hypothesis:

- y = n-length vector of outcomes
- X = n x p covariate matrix;
- lam = p-vector of coefficients
- lam0 = real number.

The null we are testing through this specification is

H0: lam' beta = lam[1] * beta[1] + ... + lam[p] * beta[p] = lam0,

where beta are the model parameters in the regression, y = X beta + e. By default this example sets p = 2-dim model, lam = (0, 1) and lam0 = 0. In this specification, H0: beta[2] = 0.

Examples

```
model = example_model()
lm(model$y ~ model$X + 0)
```

---

fastLm   *Fast least squares*

Description

This functions fits the regression y ~ X using Armadillo `solve` function.

Usage

```
fastLm(y, X)
```

Arguments

y	Vector of outcomes.
X	Matrix of covariates (first column should be 1's)

Value

List of regression output with elements `coef`, `stderr`.

---

get_clustered_eps   *Calculate residuals restricted under H0*

Description

Given regression `model` and `clustering`, this function calculates the OLS residuals under the linear null hypothesis, and assigns them to the specified clusters.

Usage

```
get_clustered_eps(model, clustering)
```

Arguments

model	A regression `model`. See example_model for details.
clustering	A `List` that specifies a clustering of indexes 1...n (#datapoints). See example_clustering for details.

Value

A `List` of the restricted residuals clustered according to `clustering`.

Examples

```
m = example_model(n=100)
cl = list(1:50, 51:100)
er = get_clustered_eps(m, cl)
stopifnot(length(er) == length(cl))
stopifnot(length(er[[1]]) == 50)
```

OLS_c	*Fast least squares*

Description

Fast OLS as in fastLm but returns only the fitted coefficients.

Usage

```
OLS_c(y, X)
```

Arguments

y	Vector of outcomes.
X	Matrix of covariates (first column should be 1's)

Value

Vector of coefficients.

---

one_sided_test	*One-sided testing*

Description

Decides to reject or not based on observed test statistic value tobs and randomization values tvals.

Usage

```
one_sided_test(tobs, tvals, alpha, tol = 1e-14)
```

Arguments

tobs	The observed value of the test statistic (scalar).
tvals	Vector of randomization values of the test statistic (to compare with tobs).
alpha	Desired level of the test (between 0 to 1).
tol	Used to check whether tobs is equal to the 1-alpha quantile of tvals.

Details

The test may randomize to achieve the specified level alpha when there are very few randomization values.

Value

Test decision (binary).

See Also

Testing Statistical Hypotheses (Ch. 15, Lehman and Romano, 2006)

---

out_pval	*Calculates p-value or test decision*

Description

Depending on `ret_pval` this function returns either a p-value for the test or the binary decision.

Usage

```
out_pval(rtest_out, ret_pval, alpha)
```

Arguments

rtest_out	A List with elements tobs, tvals (see [one_sided_test](#) for details.)
ret_pval	A Boolean indicating whether to return a p-value (TRUE) or not.
alpha	Desired test level (from 0 to 1).

Details

Returns 1 if the test rejects, 0 otherwise.

Value

Binary decision if `ret_pval` is TRUE, or the p-value otherwise.

---

restricted_OLS_c	*Fast least squares with linear constraint*

Description

This functions fits the regression y ~ X under a linear constraint on the model parameters. The constraint is Q * beta = c where beta are the regression model parameters, and Q, c are inputs.

Usage

```
restricted_OLS_c(y, X, bhat, Q, c)
```

Arguments

y	Vector of outcomes.
X	Matrix of covariates (first column should be 1's)
bhat	Unconstrained OLS-fitted coefficients.
Q	Matrix of linear constraints (k x p).
c	Vector of constraint values (k x 1).

Value

Vector of fitted OLS coefficients under linear constraint.

See Also

Advanced Econometrics (Section 1.4, Takeshi Amemiya, 1985)

---

rrinf	*Generic residual randomization confidence intervals*

Description

This function is a wrapper over rrtest and gives confidence intervals for all parameters.

Usage

```
rrinf(
  y,
  X,
  g_invar,
  cover = 0.95,
  num_R = 999,
  control = list(num_se = 6, num_breaks = 60)
)
```

Arguments

y	Vector of outcomes (length n)
X	Covariate matrix (n x p). First column should be ones to include intercept.
g_invar	Function that transforms residuals. Accepts n-vector and returns n-vector.
cover	Number from [0, 1] that denotes the confidence interval coverage (e.g., 0.95 denotes 95%)
num_R	Number of test statistic values to calculate in the randomization test (similar to no. of bootstrap samples).
control	A List that constrols the scope of the test inversion.

Details

This function has similar funtionality as standard confint. It generates confidence intervals by testing plausible values for each parameter. The plausible values are generated as follows. For some parameter beta_i we test successively

H0: beta_i = hat_beta_i - num_se * se_i

...up to...

H0: beta_i = hat_beta_i + num_se * se_i

broken in `num_breaks` intervals. Here, hat_beta_i is the OLS estimate of beta_i and se_i is the standard error. We then report the minimum and maximum values in this search space which we cannot reject at level `alpha`. This forms the desired confidence interval.

Value

Matrix that includes the confidence interval endpoints, and the interval midpoint estimate.

Note

If the confidence interval appears to be a point or is empty, then this means that the nulls we consider are implausible. We can try to improve the search through `control.tinv`. For example, we can both increase `num_se` to increase the width of search, and increase `num_breaks` to make the search space finer.

See Also

Life after bootstrap: residual randomization inference in regression models (Toulis, 2019)

https://www.ptoulis.com/residual-randomization

Examples

```
set.seed(123)
X = cbind(rep(1, 100), runif(100))
beta = c(-1, 1)
y = X %*% beta + rnorm(100)
g_invar = function(e) sample(e)  # Assume exchangeable errors.
M = rrinf(y, X, g_invar, control=list(num_se=4, num_breaks=20))
M  # Intervals cover true values
```

---

rrinfBase	*Generic residual randomization inference This function provides the basis for all other rrinf* functions.*

Description

Generic residual randomization inference This function provides the basis for all other rrinf* functions.

Usage

```
rrinfBase(y, X, g_or_clust, cover, num_R, control.tinv)
```

Arguments

y	Vector of outcomes (length n)
X	Covariate matrix (n x p). First column should be ones to include intercept.
g_or_clust	Either `clustering` or an invariance function that transforms residuals.
cover	Number from [0, 1] that denotes the confidence interval coverage (e.g., 0.95 denotes 95%)
num_R	Number of test statistic values to calculate in the randomization test (similar to no. of bootstrap samples).
control.tinv	A `List` that determines the test inversion.

Details

This function has similar funtionality as standard [confint.](#) It does so by testing plausible values for each parameter. The plausible values can be controlled as follows. For some parameter beta_i we will test successively

H0: beta_i = hat_beta_i - num_se * se_i

...up to...

H0: beta_i = hat_beta_i + num_se * se_i

broken in `num_breaks` intervals. Here, hat_beta_i is the OLS estimate of beta_i and se_i is the standard error.

The `g_or_clust` object should either be (i) a g-invariance function $R^n \to R^n$; or (ii) a list(type, cl) where type=c("perm", "sign", "double") and cl=clustering (see [example_clustering](#) for details).

https://www.ptoulis.com/residual-randomization

Value

Matrix that includes the confidence interval endpoints, and the interval midpoint estimate.

---

rrinf_clust	*Residual randomization inference based on cluster invariances*

Description

This function is a wrapper over [rrtest_clust](#) and gives confidence intervals for all parameters assuming a particular cluster invariance on the errors.

Usage

```
rrinf_clust(
  y,
  X,
  type,
  clustering = NULL,
  cover = 0.95,
  num_R = 999,
  control = list(num_se = 6, num_breaks = 60)
)
```

Arguments

y	Vector of outcomes (length n)
X	Covariate matrix (n x p). First column should be ones to include intercept.
type	A string, either "perm", "sign" or "double".
clustering	A List that specifies a clustering of datapoint indexes 1, ..., n. See [example_clustering](#) for details.
cover	Number from [0, 1] that denotes the confidence interval coverage (e.g., 0.95 denotes 95%)
num_R	Number of test statistic values to calculate in the randomization test (similar to no. of bootstrap samples).
control	A List that controls the scope of the test inversion.

Details

This function has similar funtionality as standard [confint](#). It generates confidence intervals by testing plausible values for each parameter. The plausible values are generated as follows. For some parameter beta_i we test successively

H0: beta_i = hat_beta_i - num_se * se_i

...up to...

H0: beta_i = hat_beta_i + num_se * se_i

broken in num_breaks intervals. Here, hat_beta_i is the OLS estimate of beta_i and se_i is the standard error. We then report the minimum and maximum values in this search space which we cannot reject at level alpha. This forms the desired confidence interval.

Value

Matrix that includes the OLS estimate, and confidence interval endpoints.

Note

If the confidence interval appears to be a point or is empty, then this means that the nulls we consider are implausible. We can try to improve the search through control.tinv. For example, we can both increase num_se to increase the width of search, and increase num_breaks to make the search space finer.

See [rrtest_clust](#) for a description of type and clustering.

See Also

Life after bootstrap: residual randomization inference in regression models (Toulis, 2019)

https://www.ptoulis.com/residual-randomization

Examples

```
# Heterogeneous example
set.seed(123)
n = 200
X = cbind(rep(1, n), 1:n/n)
beta = c(-1, 0.2)
ind = c(rep(0, 0.9*n), rep(1, .1*n))  # cluster indicator
y = X %*% beta + rnorm(n, sd= (1-ind) * 0.1 + ind * 5) # heteroskedastic
confint(lm(y ~ X + 0))  # normal OLS CI is imprecise

cl = list(which(ind==0), which(ind==1))  # define the clustering
rrinf_clust(y, X, "perm", cl)  # improved CI through clustered errors
```

---

rrtest	*Generic residual randomization test*

Description

This function tests the specified linear hypothesis in `model` assuming the errors are distributionally invariant with respect to stochastic function `g_invar`.

Usage

```
rrtest(model, g_invar, num_R = 999, alpha = 0.05, val_type = "decision")
```

Arguments

model	Regression model and hypothesis. See example_model for details.
g_invar	Stochastic function that transforms residuals. Accepts n-vector and returns n-vector.
num_R	Number of test statistic values to calculate in the randomization test.
alpha	Nominal test level (between 0 to 1).
val_type	The type of return value.

Details

For the regression y = X * beta + e, this function is testing the following linear null hypothesis:

H0: lam' beta = lam[1] * beta[1] + ... + lam[p] * beta[p] = lam0,

where y, X, lam, lam0 are specified in `model`. The assumption is that the errors, e, have some form of cluster invariance. Specifically:

(e_1, e_2, ..., e_n) ~ g_invar(e_1, e_2, ..., e_n),

where ~ denotes equality in distribution, and `g_invar` is the supplied invariance function.

Value

If `val_type = "decision"` (default) we get the test binary decision (1=REJECT H0).

If `val_type = "pval"` we get the test p-value.

If `val_type = "full"` we get the full test output, i.e., a `List` with elements `tobs`, `tvals`, the observed and randomization values of the test statistic, respectively.

Note

There is no guarantee that an arbitrary `g_invar` will produce valid tests. The rrtest_clust function has such guarantees under mild assumptions.

See Also

Life after bootstrap: residual randomization inference in regression models (Toulis, 2019)

https://www.ptoulis.com/residual-randomization

Examples

```
model = example_model(n = 100)  # test H0: beta2 = 0 (here, H0 is true)
g_invar = function(e) sample(e)   # Assume errors are exchangeable.
rrtest(model, g_invar) # same as rrtest_clust(model, "perm")
```

---

rrtest_clust	*Residual randomization test under cluster invariances*

Description

This function tests the specified linear hypothesis in `model` assuming that the errors have some form of cluster invariance determined by `type` within the clusters determined by `clustering`.

Usage

```
rrtest_clust(
  model,
  type,
  clustering = NULL,
  num_R = 999,
  alpha = 0.05,
  val_type = "decision"
)
```

Arguments

`model`	Regression model and hypothesis. See [example_model](#) for details.
`type`	A `character`, either "perm", "sign" or "double".
`clustering`	A `List` that specifies a clustering of datapoint indexes 1, ..., n. See [example_clustering](#). If NULL it takes default value according to `type` (see Note)
`num_R`	Number of test statistic values to calculate in the test.
`alpha`	Nominal test level (between 0 to 1).
`val_type`	The type of return value.

Details

For the regression y = X * beta + e, this function is testing the following linear null hypothesis:

H0: lam' beta = lam[1] * beta[1] + ... + lam[p] * beta[p] = lam0,

where y, X, lam, lam0 are specified in `model`. The assumption is that the errors, e, have some form of cluster invariance. Specifically:

- If `type` = "perm" then the errors are assumed exchangeable within the specified clusters:
  (e_1, e_2, ..., e_n) ~ cluster_perm(e_1, e_2, ..., e_n),
  where ~ denotes equality in distribution, and cluster_perm is any random permutation within the clusters defined by `clustering`. Internally, the test repeatedly calculates a test statistic by randomly permuting the residuals within clusters.
- If `type` = "sign" then the errors are assumed sign-symmetric within the specified clusters:
  (e_1, e_2, ..., e_n) ~ cluster_signs(e_1, e_2, ..., e_n),
  where cluster_signs is a random signs flip of residuals on the cluster level. Internally, the test repeatedly calculates a test statistic by randomly flipping the signs of cluster residuals.
- If `type` = "double" then the errors are assumed both exchangeable and sign symmetric within the specified clusters:
  (e_1, e_2, ..., e_n) ~ cluster_signs(cluster_perm(e_1, e_2, ..., e_n)),
  Internally, the test repeatedly calculates a test statistic by permuting and randomly flipping the signs of residuals on the cluster level.

Value

If `val_type` = "decision" (default) we get the test binary decision (1=REJECT H0).

If `val_type` = "pval" we get the test p-value.

If `val_type` = "full" we get the full test output, i.e., a `List` with elements `tobs`, `tvals`, the observed and randomization values of the test statistic, respectively.

Note

If `clustering` is NULL then it will be assigned a default value:

- `list(1:n)` if `type` = "perm", where n is the number of datapoints;
- `as.list(1:n)` if `type` = "sign" or "double".

As in bootstrap `num_R` is usually between 1000-5000.

See Also

Life after bootstrap: residual randomization inference in regression models (Toulis, 2019)

https://www.ptoulis.com/residual-randomization

Examples

```
# 1. Validity example
set.seed(123)
n = 50
X = cbind(rep(1, n), 1:n/n)
beta = c(0, 0)
rej = replicate(200, {
  y = X %*% beta  + rt(n, df=5)
  model = list(y=y, X=X, lam=c(0, 1), lam0=0)  # H0: beta2 = 0
  rrtest_clust(model, "perm")
})
mean(rej)  # Should be ~ 5% since H0 is true.

# 2. Heteroskedastic example
set.seed(123)
n = 200
X = cbind(rep(1, n), 1:n/n)
beta = c(-1, 0.2)
ind = c(rep(0, 0.9*n), rep(1, .1*n))  # cluster indicator
y = X %*% beta + rnorm(n, sd= (1-ind) * 0.1 + ind * 5) # heteroskedastic
confint(lm(y ~ X + 0))  # normal OLS does not reject H0: beta2 = 0
cl = list(which(ind==0), which(ind==1))
model = list(y=y, X=X, lam=c(0, 1), lam0=0)

rrtest_clust(model, "sign")  # errors are sign symmetric regardless of cluster.
# Cluster sign test does not reject because of noise.

rrtest_clust(model, "perm", cl)  # errors are exchangeable within clusters
# Cluster permutation test rejects because inference is sharper.
```

---

r_test_c   *Residual randomization test*

Description

Implements the residual randomization test. The hypothesis tested is

Usage

```
r_test_c(y, X, lam, lam0, cluster_eps_r, use_perm, use_sign, num_R)
```

Arguments

y	Vector of outcomes (n x 1).
X	Matrix of covariates (n x p). First column should be 1's.
lam	Vector of coefficients in linear H0 (p x 1).
lam0	Scalar value for linear H0.
cluster_eps_r	A `List` with restricted residuals. See [get_clustered_eps](#).
use_perm	Boolean flag whether to use permutations within clusters.
use_sign	Boolean flag whether to use sign flips across clusters.
num_R	Integer of how many randomization values to calculate.

Details

H0: lam' beta = lam[1] * beta[1] + ... + lam[p] * beta[p] = lam0.

Value

A `List` with the observed test statistic value (`tobs`), and the randomization values (`tvals`)

---

two_sided_test	*Two-sided testing*

Description

Decides to reject or not based on observed test statistic value `tobs` and randomization values `tvals`. The test may randomize to achieve the specified level `alpha` when there are very few randomization values.

Usage

```
two_sided_test(tobs, tvals, alpha)
```

Arguments

tobs	The observed value of the test statistic (scalar).
tvals	Vector of randomization values of the test statistic (to compare with `tobs`).
alpha	Desired level of the test (between 0 to 1).

Value

Test decision (binary).

See Also

Testing Statistical Hypotheses (Ch. 15, Lehman and Romano, 2006)

Index