Package 'sparklyr'

April 29, 2024

Type Package

Title R Interface to Apache Spark

Version 1.8.6

Maintainer Edgar Ruiz <edgar@rstudio.com>

Description R interface to Apache Spark, a fast and general
engine for big data processing, see <https://spark.apache.org/>. This
package supports connecting to local and remote Apache Spark clusters,
provides a 'dplyr' compatible back-end, and provides an interface to
Spark's built-in machine learning algorithms.

License Apache License 2.0 | file LICENSE

URL https://spark.posit.co/

BugReports https://github.com/sparklyr/sparklyr/issues

Depends R (>= 3.2)

Imports config (>= 0.2), DBI (>= 1.0.0), dbplyr (>= 2.5.0), dplyr (>=
1.0.9), generics, globals, glue, httr (>= 1.2.1), jsonlite (>=
1.4), methods, openssl (>= 0.8), purrr, rlang (>= 0.1.4),
rstudioapi (>= 0.10), tidyr (>= 1.2.0), tidyselect, uuid,
vctrs, withr, xml2

Suggests arrow (>= 0.17.0), broom, diffobj, foreach, ggplot2,
iterators, janeaustenr, Lahman, mlbench, nnet, nycflights13,
R6, r2d3, RCurl, reshape2, shiny (>= 1.0.1), parsnip, testthat,
rprojroot

Encoding UTF-8

RoxygenNote 7.3.1

SystemRequirements Spark: 2.x, or 3.x

Collate 'spark_data_build_types.R' 'arrow_data.R' 'spark_invoke.R'
'browse_url.R' 'spark_connection.R' 'avro_utils.R'
'config_settings.R' 'config_spark.R' 'connection_instances.R'
'connection_progress.R' 'connection_shinyapp.R'
'spark_version.R' 'connection_spark.R' 'core_arrow.R'
'core_config.R' 'core_connection.R' 'core_deserialize.R'

'core_gateway.R' 'core_invoke.R' 'core_jobj.R'
'core_serialize.R' 'core_utils.R' 'core_worker_config.R'
'utils.R' 'sql_utils.R' 'data_copy.R' 'data_csv.R'
'spark_schema_from_rdd.R' 'spark_apply_bundle.R'
'spark_apply.R' 'tables_spark.R' 'tbl_spark.R' 'spark_sql.R'
'spark_dataframe.R' 'dplyr_spark.R' 'sdf_interface.R'
'data_interface.R' 'databricks_connection.R'
'dbi_spark_connection.R' 'dbi_spark_result.R'
'dbi_spark_table.R' 'do_spark.R' 'dplyr_do.R' 'dplyr_hof.R'
'dplyr_join.R' 'dplyr_spark_data.R' 'dplyr_spark_table.R'
'stratified_sample.R' 'sdf_sql.R' 'dplyr_sql.R'
'dplyr_sql_translation.R' 'dplyr_verbs.R' 'imports.R'
'install_spark.R' 'install_spark_versions.R'
'install_spark_windows.R' 'install_tools.R' 'java.R'
'jobs_api.R' 'kubernetes_config.R' 'shell_connection.R'
'livy_connection.R' 'livy_install.R' 'livy_invoke.R'
'livy_service.R' 'ml_clustering.R'
'ml_classification_decision_tree_classifier.R'
'ml_classification_gbt_classifier.R'
'ml_classification_linear_svc.R'
'ml_classification_logistic_regression.R'
'ml_classification_multilayer_perceptron_classifier.R'
'ml_classification_naive_bayes.R'
'ml_classification_one_vs_rest.R'
'ml_classification_random_forest_classifier.R'
'ml_model_helpers.R' 'ml_clustering_bisecting_kmeans.R'
'ml_clustering_gaussian_mixture.R' 'ml_clustering_kmeans.R'
'ml_clustering_lda.R' 'ml_clustering_power_iteration.R'
'ml_constructor_utils.R' 'ml_evaluate.R'
'ml_evaluation_clustering.R' 'ml_evaluation_prediction.R'
'ml_evaluator.R' 'ml_feature_binarizer.R'
'ml_feature_bucketed_random_projection_lsh.R'
'ml_feature_bucketizer.R' 'ml_feature_chisq_selector.R'
'ml_feature_count_vectorizer.R' 'ml_feature_dct.R'
'ml_feature_sql_transformer.R' 'ml_feature_dplyr_transformer.R'
'ml_feature_elementwise_product.R'
'ml_feature_feature_hasher.R' 'ml_feature_hashing_tf.R'
'ml_feature_idf.R' 'ml_feature_imputer.R'
'ml_feature_index_to_string.R' 'ml_feature_interaction.R'
'ml_feature_lsh_utils.R' 'ml_feature_max_abs_scaler.R'
'ml_feature_min_max_scaler.R' 'ml_feature_minhash_lsh.R'
'ml_feature_ngram.R' 'ml_feature_normalizer.R'
'ml_feature_one_hot_encoder.R'
'ml_feature_one_hot_encoder_estimator.R' 'ml_feature_pca.R'
'ml_feature_polynomial_expansion.R'
'ml_feature_quantile_discretizer.R' 'ml_feature_r_formula.R'
'ml_feature_regex_tokenizer.R' 'ml_feature_robust_scaler.R'
'ml_feature_standard_scaler.R'

'ml_feature_stop_words_remover.R' 'ml_feature_string_indexer.R'
'ml_feature_string_indexer_model.R' 'ml_feature_tokenizer.R'
'ml_feature_vector_assembler.R' 'ml_feature_vector_indexer.R'
'ml_feature_vector_slicer.R' 'ml_feature_word2vec.R'
'ml_fpm_fpgrowth.R' 'ml_fpm_prefixspan.R' 'ml_helpers.R'
'ml_mapping_tables.R' 'ml_metrics.R' 'ml_model_als.R'
'ml_model_bisecting_kmeans.R' 'ml_model_constructors.R'
'ml_model_decision_tree.R' 'ml_model_gaussian_mixture.R'
'ml_model_generalized_linear_regression.R'
'ml_model_gradient_boosted_trees.R'
'ml_model_isotonic_regression.R' 'ml_model_kmeans.R'
'ml_model_lda.R' 'ml_model_linear_regression.R'
'ml_model_linear_svc.R' 'ml_model_logistic_regression.R'
'ml_model_naive_bayes.R' 'ml_model_one_vs_rest.R'
'ml_model_random_forest.R' 'ml_model_utils.R'
'ml_param_utils.R' 'ml_persistence.R' 'ml_pipeline.R'
'ml_pipeline_utils.R' 'ml_print_utils.R'
'ml_recommendation_als.R'
'ml_regression_aft_survival_regression.R'
'ml_regression_decision_tree_regressor.R'
'ml_regression_gbt_regressor.R'
'ml_regression_generalized_linear_regression.R'
'ml_regression_isotonic_regression.R'
'ml_regression_linear_regression.R'
'ml_regression_random_forest_regressor.R' 'ml_stat.R'
'ml_summary.R' 'ml_transformation_methods.R'
'ml_transformer_and_estimator.R' 'ml_tuning.R'
'ml_tuning_cross_validator.R'
'ml_tuning_train_validation_split.R' 'ml_utils.R'
'ml_validator_utils.R' 'mutation.R' 'na_actions.R'
'new_model_multilayer_perceptron.R' 'params_validator.R'
'precondition.R' 'project_template.R' 'qubole_connection.R'
'reexports.R' 'sdf_dim.R' 'sdf_distinct.R' 'sdf_ml.R'
'sdf_saveload.R' 'sdf_sequence.R' 'sdf_stat.R'
'sdf_streaming.R' 'tidyr_utils.R' 'sdf_unnest_longer.R'
'sdf_wrapper.R' 'sdf_unnest_wider.R' 'sdf_utils.R'
'spark_compile.R' 'spark_context_config.R' 'spark_extensions.R'
'spark_gateway.R' 'spark_gen_embedded_sources.R'
'spark_globals.R' 'spark_hive.R' 'spark_home.R' 'spark_ide.R'
'spark_submit.R' 'spark_update_embedded_sources.R'
'spark_utils.R' 'spark_verify_embedded_sources.R'
'stream_data.R' 'stream_job.R' 'stream_operations.R'
'stream_shiny.R' 'stream_view.R' 'synapse_connection.R'
'test_connection.R' 'tidiers_ml_aft_survival_regression.R'
'tidiers_ml_als.R' 'tidiers_ml_isotonic_regression.R'
'tidiers_ml_lda.R' 'tidiers_ml_linear_models.R'
'tidiers_ml_logistic_regression.R'
'tidiers_ml_multilayer_perceptron.R' 'tidiers_ml_naive_bayes.R'

'tidiers_ml_svc_models.R' 'tidiers_ml_tree_models.R'
'tidiers_ml_unsupervised_models.R' 'tidiers_pca.R'
'tidiers_utils.R' 'tidyr_fill.R' 'tidyr_nest.R'
'tidyr_pivot_utils.R' 'tidyr_pivot_longer.R'
'tidyr_pivot_wider.R' 'tidyr_separate.R' 'tidyr_unite.R'
'tidyr_unnest.R' 'worker_apply.R' 'worker_connect.R'
'worker_connection.R' 'worker_invoke.R' 'worker_log.R'
'worker_main.R' 'yarn_cluster.R' 'yarn_config.R' 'yarn_ui.R'
'zzz.R'

NeedsCompilation no

Author Javier Luraschi [aut],
Kevin Kuo [aut] (<https://orcid.org/0000-0001-7803-7901>),
Kevin Ushey [aut],
JJ Allaire [aut],
Samuel Macedo [ctb],
Hossein Falaki [aut],
Lu Wang [aut],
Andy Zhang [aut],
Yitao Li [aut] (<https://orcid.org/0000-0002-1261-905X>),
Jozef Hajnala [ctb],
Maciej Szymkiewicz [ctb] (<https://orcid.org/0000-0003-1469-9396>),
Wil Davis [ctb],
Edgar Ruiz [aut, cre],
RStudio [cph],
The Apache Software Foundation [aut, cph]

R **topics documented:**

checkpoint_directory   *Set/Get Spark checkpoint directory*

Description

Set/Get Spark checkpoint directory

Usage

```
spark_set_checkpoint_dir(sc, dir)

spark_get_checkpoint_dir(sc)
```

Arguments

sc	A spark_connection.
dir	checkpoint directory, must be HDFS path of running on cluster

---

collect_from_rds   *Collect Spark data serialized in RDS format into R*

Description

Deserialize Spark data that is serialized using 'spark_write_rds()' into a R dataframe.

Usage

```
collect_from_rds(path)
```

Arguments

path	Path to a local RDS file that is produced by 'spark_write_rds()' (RDS files stored in HDFS will need to be downloaded to local filesystem first (e.g., by running 'hadoop fs -copyToLocal ...' or similar)

See Also

Other Spark serialization routines: spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

compile_package_jars   *Compile Scala sources into a Java Archive (jar)*

Description

Compile the scala source files contained within an R package into a Java Archive (jar) file that can be loaded and used within a Spark environment.

Usage

```
compile_package_jars(..., spec = NULL)
```

Arguments

`...`	Optional compilation specifications, as generated by `spark_compilation_spec`. When no arguments are passed, `spark_default_compilation_spec` is used instead.
`spec`	An optional list of compilation specifications. When set, this option takes precedence over arguments passed to `...`.

`connection_config`	*Read configuration values for a connection*

Description

Read configuration values for a connection

Usage

```
connection_config(sc, prefix, not_prefix = list())
```

Arguments

`sc`	`spark_connection`
`prefix`	Prefix to read parameters for (e.g. `spark.context.`, `spark.sql.`, etc.)
`not_prefix`	Prefix to not include.

Value

Named list of config parameters (note that if a prefix was specified then the names will not include the prefix)

`copy_to.spark_connection`	
	Copy an R Data Frame to Spark

Description

Copy an R `data.frame` to Spark, and return a reference to the generated Spark DataFrame as a `tbl_spark`. The returned object will act as a `dplyr`-compatible interface to the underlying Spark table.

Usage

```
## S3 method for class 'spark_connection'
copy_to(
  dest,
  df,
  name = spark_table_name(substitute(df)),
  overwrite = FALSE,
  memory = TRUE,
  repartition = 0L,
  ...
)
```

Arguments

dest	A `spark_connection`.
df	An R `data.frame`.
name	The name to assign to the copied table in Spark.
overwrite	Boolean; overwrite a pre-existing table with the name `name` if one already exists?
memory	Boolean; should the table be cached into memory?
repartition	The number of partitions to use when distributing the table across the Spark cluster. The default (0) can be used to avoid partitioning.
...	Optional arguments; currently unused.

Value

A `tbl_spark`, representing a `dplyr`-compatible interface to a Spark DataFrame.

---

distinct	*Distinct*

Description

See [distinct](#) for more details.

---

download_scalac   *Downloads default Scala Compilers*

Description

`compile_package_jars` requires several versions of the scala compiler to work, this is to match Spark scala versions. To help setup your environment, this function will download the required compilers under the default search path.

Usage

```
download_scalac(dest_path = NULL)
```

Arguments

dest_path   The destination path where scalac will be downloaded to.

Details

See `find_scalac` for a list of paths searched and used by this function to install the required compilers.

---

dplyr_hof   *dplyr wrappers for Apache Spark higher order functions*

Description

These methods implement dplyr grammars for Apache Spark higher order functions

---

ensure   *Enforce Specific Structure for R Objects*

Description

These routines are useful when preparing to pass objects to a Spark routine, as it is often necessary to ensure certain parameters are scalar integers, or scalar doubles, and so on.

Arguments

object	An R object.
allow.na	Are NA values permitted for this object?
allow.null	Are NULL values permitted for this object?
default	If object is NULL, what value should be used in its place? If default is specified, allow.null is ignored (and assumed to be TRUE).

---

fill   *Fill*

Description

See [fill](#) for more details.

---

filter   *Filter*

Description

See [filter](#) for more details.

---

find_scalac   *Discover the Scala Compiler*

Description

Find the `scalac` compiler for a particular version of `scala`, by scanning some common directories containing `scala` installations.

Usage

```
find_scalac(version, locations = NULL)
```

Arguments

version	The `scala` version to search for. Versions of the form `major.minor` will be matched against the `scalac` installation with version `major.minor.patch`; if multiple compilers are discovered the most recent one will be used.
locations	Additional locations to scan. By default, the directories `/opt/scala` and `/usr/local/scala` will be scanned.

ft_binarizer   *Feature Transformation – Binarizer (Transformer)*

Description

Apply thresholding to a column, such that values less than or equal to the threshold are assigned the value 0.0, and values greater than the threshold are assigned the value 1.0. Column output is numeric for compatibility with other modeling functions.

Usage

```
ft_binarizer(
  x,
  input_col,
  output_col,
  threshold = 0,
  uid = random_string("binarizer_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
threshold	Threshold used to binarize continuous features.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

Other feature transformers: ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(), ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(), ft_pca(), ft_polynomial_expansion(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

Examples

```
## Not run:
library(dplyr)

sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

iris_tbl %>%
  ft_binarizer(
    input_col = "Sepal_Length",
    output_col = "Sepal_Length_bin",
    threshold = 5
  ) %>%
  select(Sepal_Length, Sepal_Length_bin, Species)

## End(Not run)
```

---

ft_bucketizer   *Feature Transformation – Bucketizer (Transformer)*

Description

Similar to R's [cut](cut) function, this transforms a numeric column into a discretized column, with
breaks specified through the splits parameter.

Usage

```
ft_bucketizer(
  x,
  input_col = NULL,
  output_col = NULL,
  splits = NULL,
  input_cols = NULL,
  output_cols = NULL,
  splits_array = NULL,
  handle_invalid = "error",
  uid = random_string("bucketizer_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
splits	A numeric vector of cutpoints, indicating the bucket boundaries.

input_cols   Names of input columns.

output_cols   Names of output columns.

splits_array   Parameter for specifying multiple splits parameters. Each element in this array
                can be used to map continuous features into buckets.

handle_invalid (Spark 2.1.0+) Param for how to handle invalid entries. Options are 'skip' (filter
                out rows with invalid values), 'error' (throw an error), or 'keep' (keep invalid
                values in a special additional bucket). Default: "error"

uid   A character string used to uniquely identify the feature transformer.

...   Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns
a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with
the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the
transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_chisq_selector(), ft_count_vectorizer(),
ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(),
ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(),
ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(),
ft_pca(), ft_polynomial_expansion(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(),
ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(),
ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(),
ft_word2vec()

Examples

```
## Not run:
library(dplyr)

sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

iris_tbl %>%
  ft_bucketizer(
    input_col = "Sepal_Length",
    output_col = "Sepal_Length_bucket",
    splits = c(0, 4.5, 5, 8)
  ) %>%
  select(Sepal_Length, Sepal_Length_bucket, Species)

## End(Not run)
```

---

ft_chisq_selector   *Feature Transformation – ChiSqSelector (Estimator)*

Description

Chi-Squared feature selection, which selects categorical features to use for predicting a categorical label

Usage

```
ft_chisq_selector(
  x,
  features_col = "features",
  output_col = NULL,
  label_col = "label",
  selector_type = "numTopFeatures",
  fdr = 0.05,
  fpr = 0.05,
  fwe = 0.05,
  num_top_features = 50,
  percentile = 0.1,
  uid = random_string("chisq_selector_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by [ft_r_formula]().
output_col	The name of the output column.
label_col	Label column name. The column should be a numeric column. Usually this column is output by [ft_r_formula]().
selector_type	(Spark 2.1.0+) The selector type of the ChisqSelector. Supported options: "numTopFeatures" (default), "percentile", "fpr", "fdr", "fwe".
fdr	(Spark 2.2.0+) The upper bound of the expected false discovery rate. Only applicable when selector_type = "fdr". Default value is 0.05.
fpr	(Spark 2.1.0+) The highest p-value for features to be kept. Only applicable when selector_type= "fpr". Default value is 0.05.
fwe	(Spark 2.2.0+) The upper bound of the expected family-wise error rate. Only applicable when selector_type = "fwe". Default value is 0.05.

num_top_features

                Number of features that selector will select, ordered by ascending p-value. If
                the number of features is less than `num_top_features`, then this will select all
                features. Only applicable when selector_type = "numTopFeatures". The default
                value of `num_top_features` is 50.

percentile   (Spark 2.1.0+) Percentile of features that selector will select, ordered by statis-
                tics value descending. Only applicable when selector_type = "percentile". De-
                fault value is 0.1.

uid   A character string used to uniquely identify the feature transformer.

...   Optional arguments; currently unused.

Details

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a
`tbl_spark`.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns
a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with
the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the
transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_count_vectorizer()`, `ft_dct()`,
`ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`,
`ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`,
`ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`,
`ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`,
`ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`,
`ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`,
`ft_word2vec()`

---

ft_count_vectorizer   *Feature Transformation – CountVectorizer (Estimator)*

Description

Extracts a vocabulary from document collections.

Usage

```
ft_count_vectorizer(
  x,
  input_col = NULL,
  output_col = NULL,
  binary = FALSE,
  min_df = 1,
  min_tf = 1,
  vocab_size = 2^18,
  uid = random_string("count_vectorizer_"),
  ...
)
```

```
ml_vocabulary(model)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
binary	Binary toggle to control the output vector values. If `TRUE`, all nonzero counts (after `min_tf` filter applied) are set to 1. This is useful for discrete probabilistic models that model binary events rather than integer counts. Default: `FALSE`
min_df	Specifies the minimum number of different documents a term must appear in to be included in the vocabulary. If this is an integer greater than or equal to 1, this specifies the number of documents the term must appear in; if this is a double in [0,1), then this specifies the fraction of documents. Default: 1.
min_tf	Filter to ignore rare words in a document. For each document, terms with frequency/count less than the given threshold are ignored. If this is an integer greater than or equal to 1, then this specifies a count (of times the term must appear in the document); if this is a double in [0,1), then this specifies a fraction (out of the document's token count). Default: 1.
vocab_size	Build a vocabulary that only considers the top `vocab_size` terms ordered by term frequency across the corpus. Default: `2^18`.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.
model	A `ml_count_vectorizer_model`.

Details

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a `tbl_spark`.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

`ml_vocabulary()` returns a vector of vocabulary built.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_dct   *Feature Transformation – Discrete Cosine Transform (DCT) (Trans-former)*

Description

A feature transformer that takes the 1D discrete cosine transform of a real vector. No zero padding is performed on the input vector. It returns a real vector of the same length representing the DCT. The return vector is scaled such that the transform matrix is unitary (aka scaled DCT-II).

Usage

```
ft_dct(
  x,
  input_col = NULL,
  output_col = NULL,
  inverse = FALSE,
  uid = random_string("dct_"),
  ...
)

ft_discrete_cosine_transform(
  x,
  input_col,
  output_col,
  inverse = FALSE,
  uid = random_string("dct_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
inverse	Indicates whether to perform the inverse DCT (TRUE) or forward DCT (FALSE).
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

`ft_discrete_cosine_transform()` is an alias for `ft_dct` for backwards compatibility.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_elementwise_product

Feature Transformation – ElementwiseProduct (Transformer)

Description

Outputs the Hadamard product (i.e., the element-wise product) of each input vector with a provided "weight" vector. In other words, it scales each column of the dataset by a scalar multiplier.

Usage

```
ft_elementwise_product(
  x,
  input_col = NULL,
  output_col = NULL,
  scaling_vec = NULL,
  uid = random_string("elementwise_product_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
scaling_vec	the vector to multiply with input vectors
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(), ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(), ft_pca(), ft_polynomial_expansion(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

---

ft_feature_hasher   *Feature Transformation – FeatureHasher (Transformer)*

Description

Feature Transformation – FeatureHasher (Transformer)

Usage

```
ft_feature_hasher(
  x,
  input_cols = NULL,
  output_col = NULL,
  num_features = 2^18,
  categorical_cols = NULL,
  uid = random_string("feature_hasher_"),
  ...
)
```

Arguments

`x`	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
`input_cols`	Names of input columns.
`output_col`	Name of output column.
`num_features`	Number of features. Defaults to $2^1 8$.
`categorical_cols`	
	Numeric columns to treat as categorical features. By default only string and boolean columns are treated as categorical, so this param can be used to explicitly specify the numerical columns to treat as categorical.
`uid`	A character string used to uniquely identify the feature transformer.
`...`	Optional arguments; currently unused.

Details

Feature hashing projects a set of categorical or numerical features into a feature vector of specified dimension (typically substantially smaller than that of the original feature space). This is done using the hashing trick https://en.wikipedia.org/wiki/Feature_hashing to map features to indices in the feature vector.

The FeatureHasher transformer operates on multiple columns. Each column may contain either numeric or categorical features. Behavior and handling of column data types is as follows: -Numeric columns: For numeric features, the hash value of the column name is used to map the feature value to its index in the feature vector. By default, numeric features are not treated as categorical (even when they are integers). To treat them as categorical, specify the relevant columns in categoricalCols. -String columns: For categorical features, the hash value of the string "column_name=value" is used to map to the vector index, with an indicator value of 1.0. Thus, categorical features are "one-hot" encoded (similarly to using OneHotEncoder with drop_last=FALSE). -Boolean columns: Boolean values are treated in the same way as string columns. That is, boolean features are represented as "column_name=true" or "column_name=false", with an indicator value of 1.0.

Null (missing) values are ignored (implicitly zero in the resulting feature vector).

The hash function used here is also the MurmurHash 3 used in HashingTF. Since a simple modulo on the hashed value is used to determine the vector index, it is advisable to use a power of two as the num_features parameter; otherwise the features will not be mapped evenly to the vector indices.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_hashing_tf   *Feature Transformation – HashingTF (Transformer)*

Description

Maps a sequence of terms to their term frequencies using the hashing trick.

Usage

```
ft_hashing_tf(
  x,
  input_col = NULL,
  output_col = NULL,
  binary = FALSE,
  num_features = 2^18,
  uid = random_string("hashing_tf_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
binary	Binary toggle to control term frequency counts. If true, all non-zero counts are set to 1. This is useful for discrete probabilistic models that model binary events rather than integer counts. (default = FALSE)
num_features	Number of features. Should be greater than 0. (default = 2^18)
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_idf   *Feature Transformation – IDF (Estimator)*

Description

Compute the Inverse Document Frequency (IDF) given a collection of documents.

Usage

```
ft_idf(
  x,
  input_col = NULL,
  output_col = NULL,
  min_doc_freq = 0,
  uid = random_string("idf_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
min_doc_freq	The minimum number of documents in which a term should appear. Default: 0
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a `tbl_spark`.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns
a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with
the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the
transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`,
`ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_imputer()`,
`ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`,
`ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`,
`ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`,
`ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`,
`ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`,
`ft_word2vec()`

---

ft_imputer   *Feature Transformation – Imputer (Estimator)*

Description

Imputation estimator for completing missing values, either using the mean or the median of the
columns in which the missing values are located. The input columns should be of numeric type.
This function requires Spark 2.2.0+.

Usage

```
ft_imputer(
  x,
  input_cols = NULL,
  output_cols = NULL,
  missing_value = NULL,
  strategy = "mean",
  uid = random_string("imputer_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_cols	The names of the input columns
output_cols	The names of the output columns.
missing_value	The placeholder for the missing values. All occurrences of `missing_value` will be imputed. Note that null values are always treated as missing.

strategy	The imputation strategy. Currently only "mean" and "median" are supported. If "mean", then replace missing values using the mean value of the feature. If "median", then replace missing values using the approximate median value of the feature. Default: mean
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a `tbl_spark`.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_index_to_string   *Feature Transformation – IndexToString (Transformer)*

Description

A Transformer that maps a column of indices back to a new column of corresponding string values. The index-string mapping is either from the ML attributes of the input column, or from user-supplied labels (which take precedence over ML attributes). This function is the inverse of `ft_string_indexer`.

Usage

```
ft_index_to_string(
  x,
  input_col = NULL,
  output_col = NULL,
  labels = NULL,
```

```
    uid = random_string("index_to_string_"),
    ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
labels	Optional param for array of labels specifying index-string mapping.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

ft_string_indexer

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_interaction(), ft_lsh, ft_max_abs_scaler(), ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(), ft_pca(), ft_polynomial_expansion ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

---

ft_interaction   *Feature Transformation – Interaction (Transformer)*

Description

Implements the feature interaction transform. This transformer takes in Double and Vector type columns and outputs a flattened vector of their feature interactions. To handle interaction, we first one-hot encode any nominal features. Then, a vector of the feature cross-products is produced.

Usage

```
ft_interaction(
  x,
  input_cols = NULL,
  output_col = NULL,
  uid = random_string("interaction_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_cols	The names of the input columns
output_col	The name of the output column.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_index_to_string(), ft_lsh, ft_max_abs_scaler(), ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(), ft_pca(), ft_polynomial_expansion(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

---

ft_lsh *Feature Transformation – LSH (Estimator)*

Description

Locality Sensitive Hashing functions for Euclidean distance (Bucketed Random Projection) and Jaccard distance (MinHash).

Usage

```
ft_bucketed_random_projection_lsh(
  x,
  input_col = NULL,
  output_col = NULL,
  bucket_length = NULL,
  num_hash_tables = 1,
  seed = NULL,
  uid = random_string("bucketed_random_projection_lsh_"),
  ...
)

ft_minhash_lsh(
  x,
  input_col = NULL,
  output_col = NULL,
  num_hash_tables = 1L,
  seed = NULL,
  uid = random_string("minhash_lsh_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
bucket_length	The length of each hash bucket, a larger bucket lowers the false negative rate. The number of buckets will be (max L2 norm of input vectors) / bucketLength.
num_hash_tables	
	Number of hash tables used in LSH OR-amplification. LSH OR-amplification can be used to reduce the false negative rate. Higher values for this param lead to a reduced false negative rate, at the expense of added computational complexity.
seed	A random seed. Set this value if you need your results to be reproducible across repeated calls.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

In the case where x is a tbl_spark, the estimator fits against x to obtain a transformer, returning a tbl_spark.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with

the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

ft_lsh_utils

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

`ft_lsh_utils`	*Utility functions for LSH models*

Description

Utility functions for LSH models

Usage

```
ml_approx_nearest_neighbors(
  model,
  dataset,
  key,
  num_nearest_neighbors,
  dist_col = "distCol"
)

ml_approx_similarity_join(
  model,
  dataset_a,
  dataset_b,
  threshold,
  dist_col = "distCol"
)
```

Arguments

`model`	A fitted LSH model, returned by either `ft_minhash_lsh()` or `ft_bucketed_random_projection_lsh()`
`dataset`	The dataset to search for nearest neighbors of the key.
`key`	Feature vector representing the item to search for.
`num_nearest_neighbors`	
	The maximum number of nearest neighbors.

dist_col	Output column for storing the distance between each result row and the key.
dataset_a	One of the datasets to join.
dataset_b	Another dataset to join.
threshold	The threshold for the distance of row pairs.

---

ft_max_abs_scaler   *Feature Transformation – MaxAbsScaler (Estimator)*

Description

Rescale each feature individually to range [-1, 1] by dividing through the largest maximum absolute
value in each feature. It does not shift/center the data, and thus does not destroy any sparsity.

Usage

```
ft_max_abs_scaler(
  x,
  input_col = NULL,
  output_col = NULL,
  uid = random_string("max_abs_scaler_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

In the case where x is a tbl_spark, the estimator fits against x to obtain a transformer, returning a
tbl_spark.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns
a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with
the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the
transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(), ft_pca(), ft_polynomial_expansion(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

features <- c("Sepal_Length", "Sepal_Width", "Petal_Length", "Petal_Width")

iris_tbl %>%
  ft_vector_assembler(
    input_col = features,
    output_col = "features_temp"
  ) %>%
  ft_max_abs_scaler(
    input_col = "features_temp",
    output_col = "features"
  )

## End(Not run)
```

---

ft_min_max_scaler   *Feature Transformation – MinMaxScaler (Estimator)*

Description

Rescale each feature individually to a common range [min, max] linearly using column summary statistics, which is also known as min-max normalization or Rescaling

Usage

```
ft_min_max_scaler(
  x,
  input_col = NULL,
  output_col = NULL,
  min = 0,
  max = 1,
  uid = random_string("min_max_scaler_"),
```

```
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
min	Lower bound after transformation, shared by all features Default: 0.0
max	Upper bound after transformation, shared by all features Default: 1.0
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a `tbl_spark`.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

features <- c("Sepal_Length", "Sepal_Width", "Petal_Length", "Petal_Width")

iris_tbl %>%
  ft_vector_assembler(
    input_col = features,
    output_col = "features_temp"
  ) %>%
```

```
  ft_min_max_scaler(
    input_col = "features_temp",
    output_col = "features"
  )

## End(Not run)
```

---

ft_ngram   *Feature Transformation – NGram (Transformer)*

Description

A feature transformer that converts the input array of strings into an array of n-grams. Null values in the input array are ignored. It returns an array of n-grams where each n-gram is represented by a space-separated string of words.

Usage

```
ft_ngram(
  x,
  input_col = NULL,
  output_col = NULL,
  n = 2,
  uid = random_string("ngram_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
n	Minimum n-gram length, greater than or equal to 1. Default: 2, bigram features
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

When the input is empty, an empty array is returned. When the input array length is less than n (number of elements per n-gram), no n-grams are returned.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_normalizer   *Feature Transformation – Normalizer (Transformer)*

Description

Normalize a vector to have unit norm using the given p-norm.

Usage

```
ft_normalizer(
  x,
  input_col = NULL,
  output_col = NULL,
  p = 2,
  uid = random_string("normalizer_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
p	Normalization in L^p space. Must be >= 1. Defaults to 2.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_one_hot_encoder   *Feature Transformation – OneHotEncoder (Transformer)*

Description

One-hot encoding maps a column of label indices to a column of binary vectors, with at most a single one-value. This encoding allows algorithms which expect continuous features, such as Logistic Regression, to use categorical features. Typically, used with `ft_string_indexer()` to index a column first.

Usage

```
ft_one_hot_encoder(
  x,
  input_cols = NULL,
  output_cols = NULL,
  handle_invalid = NULL,
  drop_last = TRUE,
  uid = random_string("one_hot_encoder_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_cols	The name of the input columns.
output_cols	The name of the output columns.
handle_invalid	(Spark 2.1.0+) Param for how to handle invalid entries. Options are 'skip' (filter out rows with invalid values), 'error' (throw an error), or 'keep' (keep invalid values in a special additional bucket). Default: "error"
drop_last	Whether to drop the last category. Defaults to `TRUE`.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_one_hot_encoder_estimator

Feature Transformation – OneHotEncoderEstimator (Estimator)

Description

A one-hot encoder that maps a column of category indices to a column of binary vectors, with at most a single one-value per row that indicates the input category index. For example with 5 categories, an input value of 2.0 would map to an output vector of [0.0, 0.0, 1.0, 0.0]. The last category is not included by default (configurable via dropLast), because it makes the vector entries sum up to one, and hence linearly dependent. So an input value of 4.0 maps to [0.0, 0.0, 0.0, 0.0].

Usage

```
ft_one_hot_encoder_estimator(
  x,
  input_cols = NULL,
  output_cols = NULL,
  handle_invalid = "error",
  drop_last = TRUE,
  uid = random_string("one_hot_encoder_estimator_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_cols	Names of input columns.
output_cols	Names of output columns.

handle_invalid	(Spark 2.1.0+) Param for how to handle invalid entries. Options are 'skip' (filter out rows with invalid values), 'error' (throw an error), or 'keep' (keep invalid values in a special additional bucket). Default: "error"
drop_last	Whether to drop the last category. Defaults to TRUE.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

In the case where x is a tbl_spark, the estimator fits against x to obtain a transformer, returning a tbl_spark.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(), ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_pca(), ft_polynomial_expansion ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

---

ft_pca   *Feature Transformation – PCA (Estimator)*

Description

PCA trains a model to project vectors to a lower dimensional space of the top k principal components.

Usage

```
ft_pca(
  x,
  input_col = NULL,
  output_col = NULL,
  k = NULL,
  uid = random_string("pca_"),
  ...
)

ml_pca(x, features = tbl_vars(x), k = length(features), pc_prefix = "PC", ...)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
k	The number of principal components
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.
features	The columns to use in the principal components analysis. Defaults to all columns in x.
pc_prefix	Length-one character vector used to prepend names of components.

Details

In the case where x is a tbl_spark, the estimator fits against x to obtain a transformer, returning a tbl_spark.

ml_pca() is a wrapper around ft_pca() that returns a ml_model.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(), ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(), ft_polynomial_expansion(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

Examples

```
## Not run:
library(dplyr)

sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

iris_tbl %>%
  select(-Species) %>%
  ml_pca(k = 2)
```

```
## End(Not run)
```

---

ft_polynomial_expansion

Feature Transformation – PolynomialExpansion (Transformer)

Description

Perform feature expansion in a polynomial space. E.g. take a 2-variable feature vector as an example: (x, y), if we want to expand it with degree 2, then we get (x, x * x, y, x * y, y * y).

Usage

```
ft_polynomial_expansion(
  x,
  input_col = NULL,
  output_col = NULL,
  degree = 2,
  uid = random_string("polynomial_expansion_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
degree	The polynomial degree to expand, which should be greater than equal to 1. A value of 1 means no expansion. Default: 2
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(), ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(), ft_pca(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

---

ft_quantile_discretizer

Feature Transformation – QuantileDiscretizer (Estimator)

Description

ft_quantile_discretizer takes a column with continuous features and outputs a column with binned categorical features. The number of bins can be set using the num_buckets parameter. It is possible that the number of buckets used will be smaller than this value, for example, if there are too few distinct values of the input to create enough distinct quantiles.

Usage

```
ft_quantile_discretizer(
  x,
  input_col = NULL,
  output_col = NULL,
  num_buckets = 2,
  input_cols = NULL,
  output_cols = NULL,
  num_buckets_array = NULL,
  handle_invalid = "error",
  relative_error = 0.001,
  uid = random_string("quantile_discretizer_"),
  weight_column = NULL,
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
num_buckets	Number of buckets (quantiles, or categories) into which data points are grouped. Must be greater than or equal to 2.
input_cols	Names of input columns.

output_cols	Names of output columns.
num_buckets_array	
	Array of number of buckets (quantiles, or categories) into which data points are grouped. Each value must be greater than or equal to 2.
handle_invalid	(Spark 2.1.0+) Param for how to handle invalid entries. Options are 'skip' (filter out rows with invalid values), 'error' (throw an error), or 'keep' (keep invalid values in a special additional bucket). Default: "error"
relative_error	(Spark 2.0.0+) Relative error (see documentation for org.apache.spark.sql.DataFrameStatFunctions.appro here for description). Must be in the range [0, 1]. default: 0.001
uid	A character string used to uniquely identify the feature transformer.
weight_column	If not NULL, then a generalized version of the Greenwald-Khanna algorithm will be run to compute weighted percentiles, with each input having a relative weight specified by the corresponding value in 'weight_column'. The weights can be considered as relative frequencies of sample inputs.
...	Optional arguments; currently unused.

Details

NaN handling: null and NaN values will be ignored from the column during `QuantileDiscretizer` fitting. This will produce a `Bucketizer` model for making predictions. During the transformation, `Bucketizer` will raise an error when it finds NaN values in the dataset, but the user can also choose to either keep or remove NaN values within the dataset by setting `handle_invalid` If the user chooses to keep NaN values, they will be handled specially and placed into their own bucket, for example, if 4 buckets are used, then non-NaN data will be put into buckets[0-3], but NaNs will be counted in a special bucket[4].

Algorithm: The bin ranges are chosen using an approximate algorithm (see the documentation for org.apache.spark.sql.DataFrameStatFunctions.approxQuantile here for a detailed description). The precision of the approximation can be controlled with the `relative_error` parameter. The lower and upper bin bounds will be -Infinity and +Infinity, covering all real values.

Note that the result may be different every time you run it, since the sample strategy behind it is non-deterministic.

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a `tbl_spark`.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

[ft_bucketizer](#)

Other feature transformers: [ft_binarizer()](#), [ft_bucketizer()](#), [ft_chisq_selector()](#), [ft_count_vectorizer()](#), [ft_dct()](#), [ft_elementwise_product()](#), [ft_feature_hasher()](#), [ft_hashing_tf()](#), [ft_idf()](#),

ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(),
ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(),
ft_pca(), ft_polynomial_expansion(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(),
ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(),
ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

---

ft_regex_tokenizer   *Feature Transformation – RegexTokenizer (Transformer)*

Description

A regex based tokenizer that extracts tokens either by using the provided regex pattern to split the
text (default) or repeatedly matching the regex (if gaps is false). Optional parameters also allow
filtering tokens using a minimal length. It returns an array of strings that can be empty.

Usage

```
ft_regex_tokenizer(
  x,
  input_col = NULL,
  output_col = NULL,
  gaps = TRUE,
  min_token_length = 1,
  pattern = "\\s+",
  to_lower_case = TRUE,
  uid = random_string("regex_tokenizer_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
gaps	Indicates whether regex splits on gaps (TRUE) or matches tokens (FALSE).
min_token_length	
	Minimum token length, greater than or equal to 0.
pattern	The regular expression pattern to be used.
to_lower_case	Indicates whether to convert all characters to lowercase before tokenizing.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_robust_scaler   *Feature Transformation – RobustScaler (Estimator)*

Description

RobustScaler removes the median and scales the data according to the quantile range. The quantile range is by default IQR (Interquartile Range, quantile range between the 1st quartile = 25th quantile and the 3rd quartile = 75th quantile) but can be configured. Centering and scaling happen independently on each feature by computing the relevant statistics on the samples in the training set. Median and quantile range are then stored to be used on later data using the transform method. Note that missing values are ignored in the computation of medians and ranges.

Usage

```
ft_robust_scaler(
  x,
  input_col = NULL,
  output_col = NULL,
  lower = 0.25,
  upper = 0.75,
  with_centering = TRUE,
  with_scaling = TRUE,
  relative_error = 0.001,
  uid = random_string("ft_robust_scaler_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
lower	Lower quantile to calculate quantile range.
upper	Upper quantile to calculate quantile range.
with_centering	Whether to center data with median.
with_scaling	Whether to scale the data to quantile range.
relative_error	The target relative error for quantile computation.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a `tbl_spark`.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

`ft_r_formula`	*Feature Transformation – RFormula (Estimator)*

Description

Implements the transforms required for fitting a dataset against an R model formula. Currently we support a limited subset of the R operators, including ~, ., :, +, and -.

Usage

```
ft_r_formula(
  x,
  formula = NULL,
  features_col = "features",
  label_col = "label",
  force_index_label = FALSE,
  uid = random_string("r_formula_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
formula	R formula as a character string or a formula. Formula objects are converted to character strings directly and the environment is not captured.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by `ft_r_formula`.
label_col	Label column name. The column should be a numeric column. Usually this column is output by `ft_r_formula`.
force_index_label	(Spark 2.1.0+) Force to index label whether it is numeric or string type. Usually we index label only when it is string type. If the formula was used by classification algorithms, we can force to index label even it is numeric type by setting this param with true. Default: `FALSE`.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

The basic operators in the formula are:

- ~ separate target and terms
- + concat terms, "+ 0" means removing intercept
- - remove a term, "- 1" means removing intercept
- : interaction (multiplication for numeric values, or binarized categorical values)
- . all columns except target

Suppose a and b are double columns, we use the following simple examples to illustrate the effect of RFormula:

- y ~ a + b means model y ~ w0 + w1 * a + w2 * b where w0 is the intercept and w1, w2 are coefficients.
- y ~ a + b + a:b - 1 means model y ~ w1 * a + w2 * b + w3 * a * b where w1, w2, w3 are coefficients.

RFormula produces a vector column of features and a double or string column of label. Like when formulas are used in R for linear regression, string input columns will be one-hot encoded, and numeric columns will be cast to doubles. If the label column is of type string, it will be first transformed to double with StringIndexer. If the label column does not exist in the DataFrame, the output label column will be created from the specified response variable in the formula.

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a `tbl_spark`.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_sql_transformer   *Feature Transformation – SQLTransformer*

Description

Implements the transformations which are defined by SQL statement. Currently we only support SQL syntax like 'SELECT ... FROM __THIS__ ...' where '__THIS__' represents the underlying table of the input dataset. The select clause specifies the fields, constants, and expressions to display in the output, it can be any select clause that Spark SQL supports. Users can also use Spark SQL built-in function and UDFs to operate on these selected columns.

Usage

```
ft_sql_transformer(
  x,
  statement = NULL,
  uid = random_string("sql_transformer_"),
  ...
)

ft_dplyr_transformer(x, tbl, uid = random_string("dplyr_transformer_"), ...)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
statement	A SQL statement.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.
tbl	A `tbl_spark` generated using `dplyr` transformations.

Details

`ft_dplyr_transformer()` is mostly a wrapper around `ft_sql_transformer()` that takes a `tbl_spark` instead of a SQL statement. Internally, the `ft_dplyr_transformer()` extracts the `dplyr` transformations used to generate `tbl` as a SQL statement or a sampling operation. Note that only single-table `dplyr` verbs are supported and that the `sdf_` family of functions are not.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_standard_scaler	*Feature Transformation – StandardScaler (Estimator)*

Description

Standardizes features by removing the mean and scaling to unit variance using column summary statistics on the samples in the training set. The "unit std" is computed using the corrected sample standard deviation, which is computed as the square root of the unbiased sample variance.

Usage

```
ft_standard_scaler(
  x,
  input_col = NULL,
  output_col = NULL,
  with_mean = FALSE,
  with_std = TRUE,
  uid = random_string("standard_scaler_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
with_mean	Whether to center the data with mean before scaling. It will build a dense output, so take care when applying to sparse input. Default: FALSE
with_std	Whether to scale the data to unit standard deviation. Default: TRUE
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

In the case where x is a tbl_spark, the estimator fits against x to obtain a transformer, returning a tbl_spark.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(), ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(), ft_pca(), ft_polynomial_expansion(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

features <- c("Sepal_Length", "Sepal_Width", "Petal_Length", "Petal_Width")

iris_tbl %>%
  ft_vector_assembler(
    input_col = features,
    output_col = "features_temp"
  ) %>%
  ft_standard_scaler(
    input_col = "features_temp",
    output_col = "features",
    with_mean = TRUE
  )

## End(Not run)
```

---

ft_stop_words_remover   *Feature Transformation – StopWordsRemover (Transformer)*

Description

A feature transformer that filters out stop words from input.

Usage

```
ft_stop_words_remover(
  x,
  input_col = NULL,
  output_col = NULL,
  case_sensitive = FALSE,
  stop_words = ml_default_stop_words(spark_connection(x), "english"),
  uid = random_string("stop_words_remover_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
case_sensitive	Whether to do a case sensitive comparison over the stop words.
stop_words	The words to be filtered out.

uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

ml_default_stop_words

Other feature transformers: [ft_binarizer()](ft_binarizer), [ft_bucketizer()](ft_bucketizer), [ft_chisq_selector()](ft_chisq_selector), [ft_count_vectorizer()](ft_count_vectorizer),
[ft_dct()](ft_dct), [ft_elementwise_product()](ft_elementwise_product), [ft_feature_hasher()](ft_feature_hasher), [ft_hashing_tf()](ft_hashing_tf), [ft_idf()](ft_idf),
[ft_imputer()](ft_imputer), [ft_index_to_string()](ft_index_to_string), [ft_interaction()](ft_interaction), [ft_lsh](ft_lsh), [ft_max_abs_scaler()](ft_max_abs_scaler),
[ft_min_max_scaler()](ft_min_max_scaler), [ft_ngram()](ft_ngram), [ft_normalizer()](ft_normalizer), [ft_one_hot_encoder()](ft_one_hot_encoder), [ft_one_hot_encoder_estimator()](ft_one_hot_encoder_estimator),
[ft_pca()](ft_pca), [ft_polynomial_expansion()](ft_polynomial_expansion), [ft_quantile_discretizer()](ft_quantile_discretizer), [ft_r_formula()](ft_r_formula), [ft_regex_tokenizer()](ft_regex_tokenizer),
[ft_robust_scaler()](ft_robust_scaler), [ft_sql_transformer()](ft_sql_transformer), [ft_standard_scaler()](ft_standard_scaler), [ft_string_indexer()](ft_string_indexer),
[ft_tokenizer()](ft_tokenizer), [ft_vector_assembler()](ft_vector_assembler), [ft_vector_indexer()](ft_vector_indexer), [ft_vector_slicer()](ft_vector_slicer), [ft_word2vec()](ft_word2vec)

---

ft_string_indexer	*Feature Transformation – StringIndexer (Estimator)*

Description

A label indexer that maps a string column of labels to an ML column of label indices. If the input column is numeric, we cast it to string and index the string values. The indices are in `[0, numLabels)`, ordered by label frequencies. So the most frequent label gets index 0. This function is the inverse of ft_index_to_string.

Usage

```
ft_string_indexer(
  x,
  input_col = NULL,
  output_col = NULL,
  handle_invalid = "error",
  string_order_type = "frequencyDesc",
  uid = random_string("string_indexer_"),
  ...
)

ml_labels(model)

ft_string_indexer_model(
  x,
```

```
    input_col = NULL,
    output_col = NULL,
    labels,
    handle_invalid = "error",
    uid = random_string("string_indexer_model_"),
    ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
handle_invalid	(Spark 2.1.0+) Param for how to handle invalid entries. Options are 'skip' (filter out rows with invalid values), 'error' (throw an error), or 'keep' (keep invalid values in a special additional bucket). Default: "error"
string_order_type	
	(Spark 2.3+)How to order labels of string column. The first label after ordering is assigned an index of 0. Options are `"frequencyDesc"`, `"frequencyAsc"`, `"alphabetDesc"`, and `"alphabetAsc"`. Defaults to `"frequencyDesc"`.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.
model	A fitted StringIndexer model returned by `ft_string_indexer()`
labels	Vector of labels, corresponding to indices to be assigned.

Details

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a `tbl_spark`.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

`ml_labels()` returns a vector of labels, corresponding to indices to be assigned.

See Also

[ft_index_to_string](#)

Other feature transformers: [ft_binarizer](#)(), [ft_bucketizer](#)(), [ft_chisq_selector](#)(), [ft_count_vectorizer](#)(), [ft_dct](#)(), [ft_elementwise_product](#)(), [ft_feature_hasher](#)(), [ft_hashing_tf](#)(), [ft_idf](#)(), [ft_imputer](#)(), [ft_index_to_string](#)(), [ft_interaction](#)(), [ft_lsh](#), [ft_max_abs_scaler](#)(), [ft_min_max_scaler](#)(), [ft_ngram](#)(), [ft_normalizer](#)(), [ft_one_hot_encoder](#)(), [ft_one_hot_encoder_estimator](#)(), [ft_pca](#)(), [ft_polynomial_expansion](#)(), [ft_quantile_discretizer](#)(), [ft_r_formula](#)(), [ft_regex_tokenizer](#)(),

ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(),
ft_tokenizer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

---

ft_tokenizer   *Feature Transformation – Tokenizer (Transformer)*

Description

A tokenizer that converts the input string to lowercase and then splits it by white spaces.

Usage

```
ft_tokenizer(
  x,
  input_col = NULL,
  output_col = NULL,
  uid = random_string("tokenizer_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns
a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with
the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the
transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(),
ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(),
ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(),
ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(),
ft_pca(), ft_polynomial_expansion(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(),
ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(),
ft_string_indexer(), ft_vector_assembler(), ft_vector_indexer(), ft_vector_slicer(),
ft_word2vec()

---

ft_vector_assembler   *Feature Transformation – VectorAssembler (Transformer)*

Description

Combine multiple vectors into a single row-vector; that is, where each row element of the newly generated column is a vector formed by concatenating each row element from the specified input columns.

Usage

```
ft_vector_assembler(
  x,
  input_cols = NULL,
  output_col = NULL,
  uid = random_string("vector_assembler_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_cols	The names of the input columns
output_col	The name of the output column.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

Other feature transformers: ft_binarizer(), ft_bucketizer(), ft_chisq_selector(), ft_count_vectorizer(), ft_dct(), ft_elementwise_product(), ft_feature_hasher(), ft_hashing_tf(), ft_idf(), ft_imputer(), ft_index_to_string(), ft_interaction(), ft_lsh, ft_max_abs_scaler(), ft_min_max_scaler(), ft_ngram(), ft_normalizer(), ft_one_hot_encoder(), ft_one_hot_encoder_estimator(), ft_pca(), ft_polynomial_expansion(), ft_quantile_discretizer(), ft_r_formula(), ft_regex_tokenizer(), ft_robust_scaler(), ft_sql_transformer(), ft_standard_scaler(), ft_stop_words_remover(), ft_string_indexer(), ft_tokenizer(), ft_vector_indexer(), ft_vector_slicer(), ft_word2vec()

---

ft_vector_indexer   *Feature Transformation – VectorIndexer (Estimator)*

Description

Indexing categorical feature columns in a dataset of Vector.

Usage

```
ft_vector_indexer(
  x,
  input_col = NULL,
  output_col = NULL,
  handle_invalid = "error",
  max_categories = 20,
  uid = random_string("vector_indexer_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
input_col	The name of the input column.
output_col	The name of the output column.
handle_invalid	(Spark 2.1.0+) Param for how to handle invalid entries. Options are 'skip' (filter out rows with invalid values), 'error' (throw an error), or 'keep' (keep invalid values in a special additional bucket). Default: "error"
max_categories	Threshold for the number of values a categorical feature can take. If a feature is found to have > max_categories values, then it is declared continuous. Must be greater than or equal to 2. Defaults to 20.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Details

In the case where x is a tbl_spark, the estimator fits against x to obtain a transformer, returning a tbl_spark.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator or a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the transformer or estimator appended to it. If a tbl_spark, it will return a tbl_spark with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_slicer()`, `ft_word2vec()`

---

ft_vector_slicer	*Feature Transformation – VectorSlicer (Transformer)*

Description

Takes a feature vector and outputs a new feature vector with a subarray of the original features.

Usage

```
ft_vector_slicer(
  x,
  input_col = NULL,
  output_col = NULL,
  indices = NULL,
  uid = random_string("vector_slicer_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
input_col	The name of the input column.
output_col	The name of the output column.
indices	An vector of indices to select features from a vector column. Note that the indices are 0-based.
uid	A character string used to uniquely identify the feature transformer.
...	Optional arguments; currently unused.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_word2vec()`

---

ft_word2vec	*Feature Transformation – Word2Vec (Estimator)*

Description

Word2Vec transforms a word into a code for further natural language processing or machine learning process.

Usage

```
ft_word2vec(
  x,
  input_col = NULL,
  output_col = NULL,
  vector_size = 100,
  min_count = 5,
  max_sentence_length = 1000,
  num_partitions = 1,
  step_size = 0.025,
  max_iter = 1,
  seed = NULL,
  uid = random_string("word2vec_"),
  ...
)

ml_find_synonyms(model, word, num)
```

Arguments

`x`	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
`input_col`	The name of the input column.
`output_col`	The name of the output column.
`vector_size`	The dimension of the code that you want to transform from words. Default: 100
`min_count`	The minimum number of times a token must appear to be included in the word2vec model's vocabulary. Default: 5

max_sentence_length

> (Spark 2.0.0+) Sets the maximum length (in words) of each sentence in the input data. Any sentence longer than this threshold will be divided into chunks of up to `max_sentence_length` size. Default: 1000

num_partitions   Number of partitions for sentences of words. Default: 1

step_size   Param for Step size to be used for each iteration of optimization (> 0).

max_iter   The maximum number of iterations to use.

seed   A random seed. Set this value if you need your results to be reproducible across repeated calls.

uid   A character string used to uniquely identify the feature transformer.

...   Optional arguments; currently unused.

model   A fitted `Word2Vec` model, returned by `ft_word2vec()`.

word   A word, as a length-one character vector.

num   Number of words closest in similarity to the given word to find.

Details

In the case where x is a `tbl_spark`, the estimator fits against x to obtain a transformer, returning a `tbl_spark`.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` or a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the transformer or estimator appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the transformation applied to it.

`ml_find_synonyms()` returns a DataFrame of synonyms and cosine similarities

See Also

Other feature transformers: `ft_binarizer()`, `ft_bucketizer()`, `ft_chisq_selector()`, `ft_count_vectorizer()`, `ft_dct()`, `ft_elementwise_product()`, `ft_feature_hasher()`, `ft_hashing_tf()`, `ft_idf()`, `ft_imputer()`, `ft_index_to_string()`, `ft_interaction()`, `ft_lsh`, `ft_max_abs_scaler()`, `ft_min_max_scaler()`, `ft_ngram()`, `ft_normalizer()`, `ft_one_hot_encoder()`, `ft_one_hot_encoder_estimator()`, `ft_pca()`, `ft_polynomial_expansion()`, `ft_quantile_discretizer()`, `ft_r_formula()`, `ft_regex_tokenizer()`, `ft_robust_scaler()`, `ft_sql_transformer()`, `ft_standard_scaler()`, `ft_stop_words_remover()`, `ft_string_indexer()`, `ft_tokenizer()`, `ft_vector_assembler()`, `ft_vector_indexer()`, `ft_vector_slicer()`

---

full_join   *Full join*

Description

See `full_join` for more details.

`generic_call_interface`
Generic Call Interface

Description

Generic Call Interface

Arguments

`sc`	spark_connection
`static`	Is this a static method call (including a constructor). If so then the `object` parameter should be the name of a class (otherwise it should be a spark_jobj instance).
`object`	Object instance or name of class (for `static`)
`method`	Name of method
`...`	Call parameters

`get_spark_sql_catalog_implementation`
Retrieve the Spark connection's SQL catalog implementation property

---

Description

Retrieve the Spark connection's SQL catalog implementation property

Usage

```
get_spark_sql_catalog_implementation(sc)
```

Arguments

`sc`	spark_connection

Value

spark.sql.catalogImplementation property from the connection's runtime configuration

---

hive_context_config   *Runtime configuration interface for Hive*

Description

Retrieves the runtime configuration interface for Hive.

Usage

```
hive_context_config(sc)
```

Arguments

sc	A `spark_connection`.

---

hof_aggregate   *Apply Aggregate Function to Array Column*

Description

Apply an element-wise aggregation function to an array column (this is essentially a dplyr wrapper for the `aggregate(array<T>, A, function<A, T, A>[, function<A, R>])`: R built-in Spark SQL functions)

Usage

```
hof_aggregate(
  x,
  start,
  merge,
  finish = NULL,
  expr = NULL,
  dest_col = NULL,
  ...
)
```

Arguments

x	The Spark data frame to run aggregation on
start	The starting value of the aggregation
merge	The aggregation function
finish	Optional param specifying a transformation to apply on the final value of the aggregation
expr	The array being aggregated, could be any SQL expression evaluating to an array (default: the last column of the Spark data frame)
dest_col	Column to store the aggregated result (default: expr)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local")
# concatenates all numbers of each array in `array_column` and add parentheses
# around the resulting string
copy_to(sc, dplyr::tibble(array_column = list(1:5, 21:25))) %>%
  hof_aggregate(
    start = "",
    merge = ~ CONCAT(.y, .x),
    finish = ~ CONCAT("(", .x, ")")
  )

## End(Not run)
```

---

hof_array_sort   *Sorts array using a custom comparator*

Description

Applies a custom comparator function to sort an array (this is essentially a dplyr wrapper to the 'array_sort(expr, func)' higher- order function, which is supported since Spark 3.0)

Usage

```
hof_array_sort(x, func, expr = NULL, dest_col = NULL, ...)
```

Arguments

x	The Spark data frame to be processed
func	The comparator function to apply (it should take 2 array elements as arguments and return an integer, with a return value of -1 indicating the first element is less than the second, 0 indicating equality, or 1 indicating the first element is greater than the second)
expr	The array being sorted, could be any SQL expression evaluating to an array (default: the last column of the Spark data frame)
dest_col	Column to store the sorted result (default: expr)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local", version = "3.0.0")
```

```
copy_to(
  sc,
  dplyr::tibble(
    # x contains 2 arrays each having elements in ascending order
    x = list(1:5, 6:10)
  )
) %>%
  # now each array from x gets sorted in descending order
  hof_array_sort(~ as.integer(sign(.y - .x)))

## End(Not run)
```

hof_exists	*Determine Whether Some Element Exists in an Array Column*

Description

Determines whether an element satisfying the given predicate exists in each array from an array column (this is essentially a dplyr wrapper for the exists(array<T>, function<T, Boolean>): Boolean built-in Spark SQL function)

Usage

```
hof_exists(x, pred, expr = NULL, dest_col = NULL, ...)
```

Arguments

x	The Spark data frame to search
pred	A boolean predicate
expr	The array being searched (could be any SQL expression evaluating to an array)
dest_col	Column to store the search result
...	Additional params to dplyr::mutate

hof_filter	*Filter Array Column*

Description

Apply an element-wise filtering function to an array column (this is essentially a dplyr wrapper for the filter(array<T>, function<T, Boolean>): array<T> built-in Spark SQL functions)

Usage

```
hof_filter(x, func, expr = NULL, dest_col = NULL, ...)
```

Arguments

x	The Spark data frame to filter
func	The filtering function
expr	The array being filtered, could be any SQL expression evaluating to an array (default: the last column of the Spark data frame)
dest_col	Column to store the filtered result (default: expr)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local")
# only keep odd elements in each array in `array_column`
copy_to(sc, dplyr::tibble(array_column = list(1:5, 21:25))) %>%
  hof_filter(~ .x %% 2 == 1)

## End(Not run)
```

---

hof_forall   *Checks whether all elements in an array satisfy a predicate*

Description

Checks whether the predicate specified holds for all elements in an array (this is essentially a dplyr wrapper to the 'forall(expr, pred)' higher- order function, which is supported since Spark 3.0)

Usage

```
hof_forall(x, pred, expr = NULL, dest_col = NULL, ...)
```

Arguments

x	The Spark data frame to be processed
pred	The predicate to test (it should take an array element as argument and return a boolean value)
expr	The array being tested, could be any SQL expression evaluating to an array (default: the last column of the Spark data frame)
dest_col	Column to store the boolean result (default: expr)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

sc <- spark_connect(master = "local", version = "3.0.0")
df <- dplyr::tibble(
  x = list(c(1, 2, 3, 4, 5), c(6, 7, 8, 9, 10)),
  y = list(c(1, 4, 2, 8, 5), c(7, 1, 4, 2, 8)),
)
sdf <- sdf_copy_to(sc, df, overwrite = TRUE)

all_positive_tbl <- sdf %>%
  hof_forall(pred = ~ .x > 0, expr = y, dest_col = all_positive) %>%
  dplyr::select(all_positive)

## End(Not run)
```

---

hof_map_filter   *Filters a map*

Description

Filters entries in a map using the function specified (this is essentially a dplyr wrapper to the 'map_filter(expr, func)' higher- order function, which is supported since Spark 3.0)

Usage

```
hof_map_filter(x, func, expr = NULL, dest_col = NULL, ...)
```

Arguments

x	The Spark data frame to be processed
func	The filter function to apply (it should take (key, value) as arguments and return a boolean value, with FALSE indicating the key-value pair should be discarded and TRUE otherwise)
expr	The map being filtered, could be any SQL expression evaluating to a map (default: the last column of the Spark data frame)
dest_col	Column to store the filtered result (default: expr)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local", version = "3.0.0")
sdf <- sdf_len(sc, 1) %>% dplyr::mutate(m = map(1, 0, 2, 2, 3, -1))
```

```
filtered_sdf <- sdf %>% hof_map_filter(~ .x > .y)

## End(Not run)
```

---

hof_map_zip_with   *Merges two maps into one*

Description

Merges two maps into a single map by applying the function specified to pairs of values with the same key (this is essentially a dplyr wrapper to the 'map_zip_with(map1, map2, func)' higher-order function, which is supported since Spark 3.0)

Usage

```
hof_map_zip_with(x, func, dest_col = NULL, map1 = NULL, map2 = NULL, ...)
```

Arguments

x	The Spark data frame to be processed
func	The function to apply (it should take (key, value1, value2) as arguments, where (key, value1) is a key-value pair present in map1, (key, value2) is a key-value pair present in map2, and return a transformed value associated with key in the resulting map
dest_col	Column to store the query result (default: the last column of the Spark data frame)
map1	The first map being merged, could be any SQL expression evaluating to a map (default: the first column of the Spark data frame)
map2	The second map being merged, could be any SQL expression evaluating to a map (default: the second column of the Spark data frame)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local", version = "3.0.0")

# create a Spark dataframe with 2 columns of type MAP<STRING, INT>
two_maps_tbl <- sdf_copy_to(
  sc,
  dplyr::tibble(
    m1 = c("{\"1\":2,\"3\":4,\"5\":6}", "{\"2\":1,\"4\":3,\"6\":5}"),
    m2 = c("{\"1\":1,\"3\":3,\"5\":5}", "{\"2\":2,\"4\":4,\"6\":6}")
  ),
```

```
      overwrite = TRUE
  ) %>%
    dplyr::mutate(m1 = from_json(m1, "MAP<STRING, INT>"),
                  m2 = from_json(m2, "MAP<STRING, INT>"))

  # create a 3rd column containing MAP<STRING, INT> values derived from the
  # first 2 columns

  transformed_two_maps_tbl <- two_maps_tbl %>%
    hof_map_zip_with(
      func = .(k, v1, v2) %->% (CONCAT(k, "_", v1, "_", v2)),
      dest_col = m3
    )

## End(Not run)
```

---

hof_transform   *Transform Array Column*

Description

Apply an element-wise transformation function to an array column (this is essentially a dplyr wrapper for the `transform(array<T>, function<T, U>): array<U>` and the `transform(array<T>, function<T, Int, U>): array<U>` built-in Spark SQL functions)

Usage

```
hof_transform(x, func, expr = NULL, dest_col = NULL, ...)
```

Arguments

x	The Spark data frame to transform
func	The transformation to apply
expr	The array being transformed, could be any SQL expression evaluating to an array (default: the last column of the Spark data frame)
dest_col	Column to store the transformed result (default: expr)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local")
# applies the (x -> x * x) transformation to elements of all arrays
copy_to(sc, dplyr::tibble(arr = list(1:5, 21:25))) %>%
  hof_transform(~ .x * .x)
```

```
## End(Not run)
```

---

hof_transform_keys   *Transforms keys of a map*

Description

Applies the transformation function specified to all keys of a map (this is essentially a dplyr wrapper
to the 'transform_keys(expr, func)' higher- order function, which is supported since Spark 3.0)

Usage

```
hof_transform_keys(x, func, expr = NULL, dest_col = NULL, ...)
```

Arguments

x	The Spark data frame to be processed
func	The transformation function to apply (it should take (key, value) as arguments and return a transformed key)
expr	The map being transformed, could be any SQL expression evaluating to a map (default: the last column of the Spark data frame)
dest_col	Column to store the transformed result (default: expr)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local", version = "3.0.0")
sdf <- sdf_len(sc, 1) %>% dplyr::mutate(m = map("a", 0L, "b", 2L, "c", -1L))
transformed_sdf <- sdf %>% hof_transform_keys(~ CONCAT(.x, " == ", .y))

## End(Not run)
```

---

hof_transform_values   *Transforms values of a map*

Description

Applies the transformation function specified to all values of a map (this is essentially a dplyr wrapper to the 'transform_values(expr, func)' higher- order function, which is supported since Spark 3.0)

Usage

```
hof_transform_values(x, func, expr = NULL, dest_col = NULL, ...)
```

Arguments

x	The Spark data frame to be processed
func	The transformation function to apply (it should take (key, value) as arguments and return a transformed value)
expr	The map being transformed, could be any SQL expression evaluating to a map (default: the last column of the Spark data frame)
dest_col	Column to store the transformed result (default: expr)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local", version = "3.0.0")
sdf <- sdf_len(sc, 1) %>% dplyr::mutate(m = map("a", 0L, "b", 2L, "c", -1L))
transformed_sdf <- sdf %>% hof_transform_values(~ CONCAT(.x, " == ", .y))

## End(Not run)
```

---

hof_zip_with   *Combines 2 Array Columns*

Description

Applies an element-wise function to combine elements from 2 array columns (this is essentially a dplyr wrapper for the zip_with(array<T>, array<U>, function<T, U, R>): array<R> built-in function in Spark SQL)

Usage

```
hof_zip_with(x, func, dest_col = NULL, left = NULL, right = NULL, ...)
```

Arguments

x	The Spark data frame to process
func	Element-wise combining function to be applied
dest_col	Column to store the query result (default: the last column of the Spark data frame)
left	Any expression evaluating to an array (default: the first column of the Spark data frame)
right	Any expression evaluating to an array (default: the second column of the Spark data frame)
...	Additional params to dplyr::mutate

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local")
# compute element-wise products of 2 arrays from each row of `left` and `right`
# and store the resuling array in `res`
copy_to(
  sc,
  dplyr::tibble(
    left = list(1:5, 21:25),
    right = list(6:10, 16:20),
    res = c(0, 0)
  )
) %>%
  hof_zip_with(~ .x * .y)

## End(Not run)
```

---

inner_join   *Inner join*

Description

See [inner_join](#) for more details.

---

invoke   *Invoke a Method on a JVM Object*

Description

Invoke methods on Java object references. These functions provide a mechanism for invoking various Java object methods directly from R.

Usage

```
invoke(jobj, method, ...)

invoke_static(sc, class, method, ...)

invoke_new(sc, class, ...)
```

Arguments

jobj	An R object acting as a Java object reference (typically, a spark_jobj).
method	The name of the method to be invoked.
...	Optional arguments, currently unused.
sc	A spark_connection.
class	The name of the Java class whose methods should be invoked.

Details

Use each of these functions in the following scenarios:

invoke	Execute a method on a Java object reference (typically, a spark_jobj).
invoke_static	Execute a static method associated with a Java class.
invoke_new	Invoke a constructor associated with a Java class.

Examples

```
sc <- spark_connect(master = "spark://HOST:PORT")
spark_context(sc) %>%
  invoke("textFile", "file.csv", 1L) %>%
  invoke("count")
```

---

jarray   *Instantiate a Java array with a specific element type.*

Description

Given a list of Java object references, instantiate an Array[T] containing the same list of references, where T is a non-primitive type that is more specific than java.lang.Object.

Usage

```
jarray(sc, x, element_type)
```

Arguments

sc	A spark_connection.
x	A list of Java object references.
element_type	A valid Java class name representing the generic type parameter of the Java array to be instantiated. Each element of x must refer to a Java object that is assignable to element_type.

Examples

```
sc <- spark_connect(master = "spark://HOST:PORT")

string_arr <- jarray(sc, letters, element_type = "java.lang.String")
# string_arr is now a reference to an array of type String[]
```

---

jfloat   *Instantiate a Java float type.*

Description

Instantiate a java.lang.Float object with the value specified. NOTE: this method is useful when one has to invoke a Java/Scala method requiring a float (instead of double) type for at least one of its parameters.

Usage

```
jfloat(sc, x)
```

Arguments

sc	A spark_connection.
x	A numeric value in R.

Examples

```
sc <- spark_connect(master = "spark://HOST:PORT")

jflt <- jfloat(sc, 1.23e-8)
# jflt is now a reference to a java.lang.Float object
```

---

jfloat_array   *Instantiate an Array[Float].*

Description

Instantiate an `Array[Float]` object with the value specified. NOTE: this method is useful when one has to invoke a Java/Scala method requiring an `Array[Float]` as one of its parameters.

Usage

```
jfloat_array(sc, x)
```

Arguments

sc	A `spark_connection`.
x	A numeric vector in R.

Examples

```
sc <- spark_connect(master = "spark://HOST:PORT")

jflt_arr <- jfloat_array(sc, c(-1.23e-8, 0, -1.23e-8))
# jflt_arr is now a reference an array of java.lang.Float
```

---

join.tbl_spark   *Join Spark tbls.*

Description

These functions are wrappers around their 'dplyr' equivalents that set Spark SQL-compliant values for the 'suffix' argument by replacing dots ('.') with underscores ('_'). See [join] for a description of the general purpose of the functions.

Usage

```
## S3 method for class 'tbl_spark'
inner_join(
  x,
  y,
  by = NULL,
  copy = FALSE,
  suffix = c("_x", "_y"),
  auto_index = FALSE,
  ...,
  sql_on = NULL
```

```
)

## S3 method for class 'tbl_spark'
left_join(
  x,
  y,
  by = NULL,
  copy = FALSE,
  suffix = c("_x", "_y"),
  auto_index = FALSE,
  ...,
  sql_on = NULL
)

## S3 method for class 'tbl_spark'
right_join(
  x,
  y,
  by = NULL,
  copy = FALSE,
  suffix = c("_x", "_y"),
  auto_index = FALSE,
  ...,
  sql_on = NULL
)

## S3 method for class 'tbl_spark'
full_join(
  x,
  y,
  by = NULL,
  copy = FALSE,
  suffix = c("_x", "_y"),
  auto_index = FALSE,
  ...,
  sql_on = NULL
)
```

Arguments

x, y	A pair of lazy data frames backed by database queries.
by	A join specification created with [join_by()](), or a character vector of variables to join by.
	If NULL, the default, *_join() will perform a natural join, using all variables in common across x and y. A message lists the variables so that you can check they're correct; suppress the message by supplying by explicitly.
	To join on different variables between x and y, use a [join_by()]() specification. For example, join_by(a == b) will match x$a to y$b.

To join by multiple variables, use a [join_by()](#) specification with multiple expressions. For example, join_by(a == b, c == d) will match x$a to y$b and x$c to y$d. If the column names are the same between x and y, you can shorten this by listing only the variable names, like join_by(a, c).

[join_by()](#) can also be used to perform inequality, rolling, and overlap joins. See the documentation at [?join_by](#) for details on these types of joins.

For simple equality joins, you can alternatively specify a character vector of variable names to join by. For example, by = c("a", "b") joins x$a to y$a and x$b to y$b. If variable names differ between x and y, use a named character vector like by = c("x_a" = "y_a", "x_b" = "y_b").

To perform a cross-join, generating all combinations of x and y, see [cross_join()](#).

copy   If x and y are not from the same data source, and copy is TRUE, then y will be copied into a temporary table in same database as x. *_join() will automatically run ANALYZE on the created table in the hope that this will make you queries as efficient as possible by giving more data to the query planner.

This allows you to join tables across srcs, but it's potentially expensive operation so you must opt into it.

suffix   If there are non-joined duplicate variables in x and y, these suffixes will be added to the output to disambiguate them. Should be a character vector of length 2.

auto_index   if copy is TRUE, automatically create indices for the variables in by. This may speed up the join if there are matching indexes in x.

...   Other parameters passed onto methods.

sql_on   A custom join predicate as an SQL expression. Usually joins use column equality, but you can perform more complex queries by supply sql_on which should be a SQL expression that uses LHS and RHS aliases to refer to the left-hand side or right-hand side of the join respectively.

---

j_invoke   *Invoke a Java function.*

Description

Invoke a Java function and force return value of the call to be retrieved as a Java object reference.

Usage

```
j_invoke(jobj, method, ...)

j_invoke_static(sc, class, method, ...)

j_invoke_new(sc, class, ...)
```

Arguments

`jobj`	An R object acting as a Java object reference (typically, a `spark_jobj`).
`method`	The name of the method to be invoked.
`...`	Optional arguments, currently unused.
`sc`	A `spark_connection`.
`class`	The name of the Java class whose methods should be invoked.

`left_join` *Left join*

Description

See [`left_join`](#) for more details.

`list_sparklyr_jars` *list all sparklyr-*.jar files that have been built*

Description

list all sparklyr-*.jar files that have been built

Usage

```
list_sparklyr_jars()
```

`livy_config` *Create a Spark Configuration for Livy*

Description

Create a Spark Configuration for Livy

Usage

```
livy_config(
  config = spark_config(),
  username = NULL,
  password = NULL,
  negotiate = FALSE,
  custom_headers = list(`X-Requested-By` = "sparklyr"),
  proxy = NULL,
  curl_opts = NULL,
  ...
)
```

Arguments

config	Optional base configuration
username	The username to use in the Authorization header
password	The password to use in the Authorization header
negotiate	Whether to use gssnegotiate method or not
custom_headers	List of custom headers to append to http requests. Defaults to `list("X-Requested-By" = "sparklyr")`.
proxy	Either NULL or a proxy specified by httr::use_proxy(). Defaults to NULL.
curl_opts	List of CURL options (e.g., verbose, connecttimeout, dns_cache_timeout, etc, see httr::httr_options() for a list of valid options) – NOTE: these configurations are for libcurl only and separate from HTTP headers or Livy session parameters.
...	additional Livy session parameters

Details

Extends a Spark `spark_config()` configuration with settings for Livy. For instance, `username` and `password` define the basic authentication settings for a Livy session.

The default value of `"custom_headers"` is set to `list("X-Requested-By" = "sparklyr")` in order to facilitate connection to Livy servers with CSRF protection enabled.

Additional parameters for Livy sessions are:

`proxy_user` User to impersonate when starting the session

`jars` jars to be used in this session

`py_files` Python files to be used in this session

`files` files to be used in this session

`driver_memory` Amount of memory to use for the driver process

`driver_cores` Number of cores to use for the driver process

`executor_memory` Amount of memory to use per executor process

`executor_cores` Number of cores to use for each executor

`num_executors` Number of executors to launch for this session

`archives` Archives to be used in this session

`queue` The name of the YARN queue to which submitted

`name` The name of this session

`heartbeat_timeout` Timeout in seconds to which session be orphaned

`conf` Spark configuration properties (Map of key=value)

Note that queue is supported only by version 0.4.0 of Livy or newer. If you are using the older one, specify queue via config (e.g. `config = spark_config(spark.yarn.queue = "my_queue")`).

Value

Named list with configuration data

---

livy_service_start   *Start Livy*

Description

Starts the livy service.

Stops the running instances of the livy service.

Usage

```
livy_service_start(
  version = NULL,
  spark_version = NULL,
  stdout = "",
  stderr = "",
  ...
)

livy_service_stop()
```

Arguments

version	The version of 'livy' to use.
spark_version	The version of 'spark' to connect to.
stdout, stderr	where output to 'stdout' or 'stderr' should be sent. Same options as system2.
...	Optional arguments; currently unused.

---

ml-params   *Spark ML – ML Params*

Description

Helper methods for working with parameters for ML objects.

Usage

```
ml_is_set(x, param, ...)

ml_param_map(x, ...)

ml_param(x, param, allow_null = FALSE, ...)

ml_params(x, params = NULL, allow_null = FALSE, ...)
```

Arguments

x	A Spark ML object, either a pipeline stage or an evaluator.
param	The parameter to extract or set.
...	Optional arguments; currently unused.
allow_null	Whether to allow NULL results when extracting parameters. If FALSE, an error will be thrown if the specified parameter is not found. Defaults to FALSE.
params	A vector of parameters to extract.

---

ml-persistence	*Spark ML – Model Persistence*

Description

Save/load Spark ML objects

Usage

```
ml_save(x, path, overwrite = FALSE, ...)

## S3 method for class 'ml_model'
ml_save(
  x,
  path,
  overwrite = FALSE,
  type = c("pipeline_model", "pipeline"),
  ...
)

ml_load(sc, path)
```

Arguments

x	A ML object, which could be a ml_pipeline_stage or a ml_model
path	The path where the object is to be serialized/deserialized.
overwrite	Whether to overwrite the existing path, defaults to FALSE.
...	Optional arguments; currently unused.
type	Whether to save the pipeline model or the pipeline.
sc	A Spark connection.

Value

ml_save() serializes a Spark object into a format that can be read back into sparklyr or by the Scala or PySpark APIs. When called on ml_model objects, i.e. those that were created via the tbl_spark ‑ formula signature, the associated pipeline model is serialized. In other words, the saved model contains both the data processing (RFormulaModel) stage and the machine learning stage.

ml_load() reads a saved Spark object into sparklyr. It calls the correct Scala load method based on parsing the saved metadata. Note that a PipelineModel object saved from a sparklyr ml_model via ml_save() will be read back in as an ml_pipeline_model, rather than the ml_model object.

---

ml-transform-methods   *Spark ML – Transform, fit, and predict methods (ml_ interface)*

Description

Methods for transformation, fit, and prediction. These are mirrors of the corresponding [sdf-transform-methods](#).

Usage

```
is_ml_transformer(x)

is_ml_estimator(x)

ml_fit(x, dataset, ...)

## Default S3 method:
ml_fit(x, dataset, ...)

ml_transform(x, dataset, ...)

ml_fit_and_transform(x, dataset, ...)

ml_predict(x, dataset, ...)

## S3 method for class 'ml_model_classification'
ml_predict(x, dataset, probability_prefix = "probability_", ...)
```

Arguments

x	A ml_estimator, ml_transformer (or a list thereof), or ml_model object.
dataset	A tbl_spark.
...	Optional arguments; currently unused.
probability_prefix	
	String used to prepend the class probability output columns.

Details

These methods are

Value

When x is an estimator, `ml_fit()` returns a transformer whereas `ml_fit_and_transform()` returns a transformed dataset. When x is a transformer, `ml_transform()` and `ml_predict()` return a transformed dataset. When `ml_predict()` is called on a `ml_model` object, additional columns (e.g. probabilities in case of classification models) are appended to the transformed output for the user's convenience.

---

ml-tuning   *Spark ML – Tuning*

Description

Perform hyper-parameter tuning using either K-fold cross validation or train-validation split.

Usage

```
ml_sub_models(model)

ml_validation_metrics(model)

ml_cross_validator(
  x,
  estimator = NULL,
  estimator_param_maps = NULL,
  evaluator = NULL,
  num_folds = 3,
  collect_sub_models = FALSE,
  parallelism = 1,
  seed = NULL,
  uid = random_string("cross_validator_"),
  ...
)

ml_train_validation_split(
  x,
  estimator = NULL,
  estimator_param_maps = NULL,
  evaluator = NULL,
  train_ratio = 0.75,
  collect_sub_models = FALSE,
  parallelism = 1,
  seed = NULL,
  uid = random_string("train_validation_split_"),
```

```
    ...
)
```

Arguments

model	A cross validation or train-validation-split model.
x	A spark_connection, ml_pipeline, or a tbl_spark.
estimator	A ml_estimator object.
estimator_param_maps	
	A named list of stages and hyper-parameter sets to tune. See details.
evaluator	A ml_evaluator object, see [ml_evaluator]().
num_folds	Number of folds for cross validation. Must be >= 2. Default: 3
collect_sub_models	
	Whether to collect a list of sub-models trained during tuning. If set to FALSE, then only the single best sub-model will be available after fitting. If set to true, then all sub-models will be available. Warning: For large models, collecting all sub-models can cause OOMs on the Spark driver.
parallelism	The number of threads to use when running parallel algorithms. Default is 1 for serial execution.
seed	A random seed. Set this value if you need your results to be reproducible across repeated calls.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; currently unused.
train_ratio	Ratio between train and validation data. Must be between 0 and 1. Default: 0.75

Details

ml_cross_validator() performs k-fold cross validation while ml_train_validation_split() performs tuning on one pair of train and validation datasets.

Value

The object returned depends on the class of x.

- spark_connection: When x is a spark_connection, the function returns an instance of a ml_cross_validator or ml_traing_validation_split object.

- ml_pipeline: When x is a ml_pipeline, the function returns a ml_pipeline with the tuning estimator appended to the pipeline.

- tbl_spark: When x is a tbl_spark, a tuning estimator is constructed then immediately fit with the input tbl_spark, returning a ml_cross_validation_model or a ml_train_validation_split_model object.

For cross validation, ml_sub_models() returns a nested list of models, where the first layer represents fold indices and the second layer represents param maps. For train-validation split, ml_sub_models() returns a list of models, corresponding to the order of the estimator param maps.

ml_validation_metrics() returns a data frame of performance metrics and hyperparameter combinations.

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

# Create a pipeline
pipeline <- ml_pipeline(sc) %>%
  ft_r_formula(Species ~ .) %>%
  ml_random_forest_classifier()

# Specify hyperparameter grid
grid <- list(
  random_forest = list(
    num_trees = c(5, 10),
    max_depth = c(5, 10),
    impurity = c("entropy", "gini")
  )
)

# Create the cross validator object
cv <- ml_cross_validator(
  sc,
  estimator = pipeline, estimator_param_maps = grid,
  evaluator = ml_multiclass_classification_evaluator(sc),
  num_folds = 3,
  parallelism = 4
)

# Train the models
cv_model <- ml_fit(cv, iris_tbl)

# Print the metrics
ml_validation_metrics(cv_model)

## End(Not run)
```

```
ml_aft_survival_regression
```
Spark ML – Survival Regression

Description

Fit a parametric survival regression model named accelerated failure time (AFT) model (see Accelerated failure time model (Wikipedia)) based on the Weibull distribution of the survival time.

Usage

```
ml_aft_survival_regression(
```

```
  x,
  formula = NULL,
  censor_col = "censor",
  quantile_probabilities = c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 0.9, 0.95, 0.99),
  fit_intercept = TRUE,
  max_iter = 100L,
  tol = 1e-06,
  aggregation_depth = 2,
  quantiles_col = NULL,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  uid = random_string("aft_survival_regression_"),
  ...
)

ml_survival_regression(
  x,
  formula = NULL,
  censor_col = "censor",
  quantile_probabilities = c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 0.9, 0.95, 0.99),
  fit_intercept = TRUE,
  max_iter = 100L,
  tol = 1e-06,
  aggregation_depth = 2,
  quantiles_col = NULL,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  uid = random_string("aft_survival_regression_"),
  response = NULL,
  features = NULL,
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
formula	Used when x is a tbl_spark. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
censor_col	Censor column name. The value of this column could be 0 or 1. If the value is 1, it means the event has occurred i.e. uncensored; otherwise censored.
quantile_probabilities	
	Quantile probabilities array. Values of the quantile probabilities array should be in the range (0, 1) and the array should be non-empty.
fit_intercept	Boolean; should the model be fit with an intercept term?
max_iter	The maximum number of iterations to use.

tol	Param for the convergence tolerance for iterative algorithms.
aggregation_depth	
	(Spark 2.1.0+) Suggested depth for treeAggregate (>= 2).
quantiles_col	Quantiles column name. This column will output quantiles of corresponding quantileProbabilities if it is set.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by `ft_r_formula`.
label_col	Label column name. The column should be a numeric column. Usually this column is output by `ft_r_formula`.
prediction_col	Prediction column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; see Details.
response	(Deprecated) The name of the response column (as a length-one character vector.)
features	(Deprecated) The name of features (terms) to use for the model fit.

Details

`ml_survival_regression()` is an alias for `ml_aft_survival_regression()` for backwards compatibility.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

See Also

Other ml algorithms: `ml_decision_tree_classifier()`, `ml_gbt_classifier()`, `ml_generalized_linear_regression`
`ml_isotonic_regression()`, `ml_linear_regression()`, `ml_linear_svc()`, `ml_logistic_regression()`,
`ml_multilayer_perceptron_classifier()`, `ml_naive_bayes()`, `ml_one_vs_rest()`, `ml_random_forest_classifier(`

Examples

```
## Not run:

library(survival)
library(sparklyr)

sc <- spark_connect(master = "local")
ovarian_tbl <- sdf_copy_to(sc, ovarian, name = "ovarian_tbl", overwrite = TRUE)

partitions <- ovarian_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

ovarian_training <- partitions$training
```

```
ovarian_test <- partitions$test

sur_reg <- ovarian_training %>%
  ml_aft_survival_regression(futime ~ ecog_ps + rx + age + resid_ds, censor_col = "fustat")

pred <- ml_predict(sur_reg, ovarian_test)
pred

## End(Not run)
```

---

ml_als   *Spark ML – ALS*

Description

Perform recommendation using Alternating Least Squares (ALS) matrix factorization.

Usage

```
ml_als(
  x,
  formula = NULL,
  rating_col = "rating",
  user_col = "user",
  item_col = "item",
  rank = 10,
  reg_param = 0.1,
  implicit_prefs = FALSE,
  alpha = 1,
  nonnegative = FALSE,
  max_iter = 10,
  num_user_blocks = 10,
  num_item_blocks = 10,
  checkpoint_interval = 10,
  cold_start_strategy = "nan",
  intermediate_storage_level = "MEMORY_AND_DISK",
  final_storage_level = "MEMORY_AND_DISK",
  uid = random_string("als_"),
  ...
)

ml_recommend(model, type = c("items", "users"), n = 1)
```

Arguments

x   A spark_connection, ml_pipeline, or a tbl_spark.

formula	Used when x is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see [ft_r_formula](#) for details. The ALS model requires a specific formula format, please use `rating_col ~ user_col + item_col`.
rating_col	Column name for ratings. Default: "rating"
user_col	Column name for user ids. Ids must be integers. Other numeric types are supported for this column, but will be cast to integers as long as they fall within the integer value range. Default: "user"
item_col	Column name for item ids. Ids must be integers. Other numeric types are supported for this column, but will be cast to integers as long as they fall within the integer value range. Default: "item"
rank	Rank of the matrix factorization (positive). Default: 10
reg_param	Regularization parameter.
implicit_prefs	Whether to use implicit preference. Default: FALSE.
alpha	Alpha parameter in the implicit preference formulation (nonnegative).
nonnegative	Whether to apply nonnegativity constraints. Default: FALSE.
max_iter	Maximum number of iterations.
num_user_blocks	
	Number of user blocks (positive). Default: 10
num_item_blocks	
	Number of item blocks (positive). Default: 10
checkpoint_interval	
	Set checkpoint interval (>= 1) or disable checkpoint (-1). E.g. 10 means that the cache will get checkpointed every 10 iterations, defaults to 10.
cold_start_strategy	
	(Spark 2.2.0+) Strategy for dealing with unknown or new users/items at prediction time. This may be useful in cross-validation or production scenarios, for handling user/item ids the model has not seen in the training data. Supported values: - "nan": predicted value for unknown ids will be NaN. - "drop": rows in the input DataFrame containing unknown ids will be dropped from the output DataFrame containing predictions. Default: "nan".
intermediate_storage_level	
	(Spark 2.0.0+) StorageLevel for intermediate datasets. Pass in a string representation of `StorageLevel`. Cannot be "NONE". Default: "MEMORY_AND_DISK".
final_storage_level	
	(Spark 2.0.0+) StorageLevel for ALS model factors. Pass in a string representation of `StorageLevel`. Default: "MEMORY_AND_DISK".
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; currently unused.
model	An ALS model object
type	What to recommend, one of `items` or `users`
n	Maximum number of recommendations to return

Details

ml_recommend() returns the top n users/items recommended for each item/user, for all items/users. The output has been transformed (exploded and separated) from the default Spark outputs to be more user friendly.

Value

ALS attempts to estimate the ratings matrix R as the product of two lower-rank matrices, X and Y, i.e. X * Yt = R. Typically these approximations are called 'factor' matrices. The general approach is iterative. During each iteration, one of the factor matrices is held constant, while the other is solved for using least squares. The newly-solved factor matrix is then held constant while solving for the other factor matrix.

This is a blocked implementation of the ALS factorization algorithm that groups the two sets of factors (referred to as "users" and "products") into blocks and reduces communication by only sending one copy of each user vector to each product block on each iteration, and only for the product blocks that need that user's feature vector. This is achieved by pre-computing some information about the ratings matrix to determine the "out-links" of each user (which blocks of products it will contribute to) and "in-link" information for each product (which of the feature vectors it receives from each user block it will depend on). This allows us to send only an array of feature vectors between each user block and product block, and have the product block find the users' ratings and update the products based on these messages.

For implicit preference data, the algorithm used is based on "Collaborative Filtering for Implicit Feedback Datasets", available at doi:10.1109/ICDM.2008.22, adapted for the blocked approach used here.

Essentially instead of finding the low-rank approximations to the rating matrix R, this finds the approximations for a preference matrix P where the elements of P are 1 if r is greater than 0 and 0 if r is less than or equal to 0. The ratings then act as 'confidence' values related to strength of indicated user preferences rather than explicit ratings given to items.

The object returned depends on the class of x.

- spark_connection: When x is a spark_connection, the function returns an instance of a ml_als recommender object, which is an Estimator.

- ml_pipeline: When x is a ml_pipeline, the function returns a ml_pipeline with the recommender appended to the pipeline.

- tbl_spark: When x is a tbl_spark, a recommender estimator is constructed then immediately fit with the input tbl_spark, returning a recommendation model, i.e. ml_als_model.

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local")

movies <- data.frame(
  user  = c(1, 2, 0, 1, 2, 0),
  item  = c(1, 1, 1, 2, 2, 0),
```

```
   rating = c(3, 1, 2, 4, 5, 4)
)
movies_tbl <- sdf_copy_to(sc, movies)

model <- ml_als(movies_tbl, rating ~ user + item)

ml_predict(model, movies_tbl)

ml_recommend(model, type = "item", 1)

## End(Not run)
```

---

ml_als_tidiers   *Tidying methods for Spark ML ALS*

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_als'
tidy(x, ...)

## S3 method for class 'ml_model_als'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_als'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.

---

ml_bisecting_kmeans   *Spark ML – Bisecting K-Means Clustering*

Description

A bisecting k-means algorithm based on the paper "A comparison of document clustering techniques" by Steinbach, Karypis, and Kumar, with modification to fit Spark. The algorithm starts from a single cluster that contains all points. Iteratively it finds divisible clusters on the bottom level and bisects each of them using k-means, until there are k leaf clusters in total or no leaf clusters are divisible. The bisecting steps of clusters on the same level are grouped together to increase parallelism. If bisecting all divisible clusters on the bottom level would result more than k leaf clusters, larger clusters get higher priority.

Usage

```
ml_bisecting_kmeans(
  x,
  formula = NULL,
  k = 4,
  max_iter = 20,
  seed = NULL,
  min_divisible_cluster_size = 1,
  features_col = "features",
  prediction_col = "prediction",
  uid = random_string("bisecting_bisecting_kmeans_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
formula	Used when x is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
k	The number of clusters to create
max_iter	The maximum number of iterations to use.
seed	A random seed. Set this value if you need your results to be reproducible across repeated calls.
min_divisible_cluster_size	
	The minimum number of points (if greater than or equal to 1.0) or the minimum proportion of points (if less than 1.0) of a divisible cluster (default: 1.0).
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by `ft_r_formula`.
prediction_col	Prediction column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments, see Details. #' @return The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

Examples

```
## Not run:
library(dplyr)

sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)
```

```
iris_tbl %>%
  select(-Species) %>%
  ml_bisecting_kmeans(k = 4, Species ~ .)

## End(Not run)
```

---

ml_chisquare_test   *Chi-square hypothesis testing for categorical data.*

Description

Conduct Pearson's independence test for every feature against the label. For each feature, the (feature, label) pairs are converted into a contingency matrix for which the Chi-squared statistic is computed. All label and feature values must be categorical.

Usage

```
ml_chisquare_test(x, features, label)
```

Arguments

x	A `tbl_spark`.
features	The name(s) of the feature columns. This can also be the name of a single vector column created using `ft_vector_assembler()`.
label	The name of the label column.

Value

A data frame with one row for each (feature, label) pair with p-values, degrees of freedom, and test statistics.

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

features <- c("Petal_Width", "Petal_Length", "Sepal_Length", "Sepal_Width")

ml_chisquare_test(iris_tbl, features = features, label = "Species")

## End(Not run)
```

ml_clustering_evaluator

Spark ML - Clustering Evaluator

Description

Evaluator for clustering results. The metric computes the Silhouette measure using the squared Euclidean distance. The Silhouette is a measure for the validation of the consistency within clusters. It ranges between 1 and -1, where a value close to 1 means that the points in a cluster are close to the other points in the same cluster and far from the points of the other clusters.

Usage

```
ml_clustering_evaluator(
  x,
  features_col = "features",
  prediction_col = "prediction",
  metric_name = "silhouette",
  uid = random_string("clustering_evaluator_"),
  ...
)
```

Arguments

x	A `spark_connection` object or a `tbl_spark` containing label and prediction columns. The latter should be the output of [sdf_predict]().
features_col	Name of features column.
prediction_col	Name of the prediction column.
metric_name	The performance metric. Currently supports "silhouette".
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; currently unused.

Value

The calculated performance metric

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

partitions <- iris_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

iris_training <- partitions$training
```

```
iris_test <- partitions$test

formula <- Species ~ .

# Train the models
kmeans_model <- ml_kmeans(iris_training, formula = formula)
b_kmeans_model <- ml_bisecting_kmeans(iris_training, formula = formula)
gmm_model <- ml_gaussian_mixture(iris_training, formula = formula)

# Predict
pred_kmeans <- ml_predict(kmeans_model, iris_test)
pred_b_kmeans <- ml_predict(b_kmeans_model, iris_test)
pred_gmm <- ml_predict(gmm_model, iris_test)

# Evaluate
ml_clustering_evaluator(pred_kmeans)
ml_clustering_evaluator(pred_b_kmeans)
ml_clustering_evaluator(pred_gmm)

## End(Not run)
```

---

ml_corr *Compute correlation matrix*

Description

Compute correlation matrix

Usage

```
ml_corr(x, columns = NULL, method = c("pearson", "spearman"))
```

Arguments

x	A tbl_spark.
columns	The names of the columns to calculate correlations of. If only one column is specified, it must be a vector column (for example, assembled using ft_vector_assember()).
method	The method to use, either "pearson" or "spearman".

Value

A correlation matrix organized as a data frame.

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

features <- c("Petal_Width", "Petal_Length", "Sepal_Length", "Sepal_Width")

ml_corr(iris_tbl, columns = features, method = "pearson")

## End(Not run)
```

---

ml_decision_tree_classifier

Spark ML – Decision Trees

Description

Perform classification and regression using decision trees.

Usage

```
ml_decision_tree_classifier(
  x,
  formula = NULL,
  max_depth = 5,
  max_bins = 32,
  min_instances_per_node = 1,
  min_info_gain = 0,
  impurity = "gini",
  seed = NULL,
  thresholds = NULL,
  cache_node_ids = FALSE,
  checkpoint_interval = 10,
  max_memory_in_mb = 256,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  probability_col = "probability",
  raw_prediction_col = "rawPrediction",
  uid = random_string("decision_tree_classifier_"),
  ...
)

ml_decision_tree(
  x,
  formula = NULL,
  type = c("auto", "regression", "classification"),
```

```
    features_col = "features",
    label_col = "label",
    prediction_col = "prediction",
    variance_col = NULL,
    probability_col = "probability",
    raw_prediction_col = "rawPrediction",
    checkpoint_interval = 10L,
    impurity = "auto",
    max_bins = 32L,
    max_depth = 5L,
    min_info_gain = 0,
    min_instances_per_node = 1L,
    seed = NULL,
    thresholds = NULL,
    cache_node_ids = FALSE,
    max_memory_in_mb = 256L,
    uid = random_string("decision_tree_"),
    response = NULL,
    features = NULL,
    ...
  )

ml_decision_tree_regressor(
    x,
    formula = NULL,
    max_depth = 5,
    max_bins = 32,
    min_instances_per_node = 1,
    min_info_gain = 0,
    impurity = "variance",
    seed = NULL,
    cache_node_ids = FALSE,
    checkpoint_interval = 10,
    max_memory_in_mb = 256,
    variance_col = NULL,
    features_col = "features",
    label_col = "label",
    prediction_col = "prediction",
    uid = random_string("decision_tree_regressor_"),
    ...
  )
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
formula	Used when x is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.

max_depth	Maximum depth of the tree (>= 0); that is, the maximum number of nodes separating any leaves from the root of the tree.
max_bins	The maximum number of bins used for discretizing continuous features and for choosing how to split on features at each node. More bins give higher granularity.
min_instances_per_node	
	Minimum number of instances each child must have after split.
min_info_gain	Minimum information gain for a split to be considered at a tree node. Should be >= 0, defaults to 0.
impurity	Criterion used for information gain calculation. Supported: "entropy" and "gini" (default) for classification and "variance" (default) for regression. For ml_decision_tree, setting "auto" will default to the appropriate criterion based on model type.
seed	Seed for random numbers.
thresholds	Thresholds in multi-class classification to adjust the probability of predicting each class. Array must have length equal to the number of classes, with values > 0 excepting that at most one value may be 0. The class with largest value p/t is predicted, where p is the original probability of that class and t is the class's threshold.
cache_node_ids	If FALSE, the algorithm will pass trees to executors to match instances with nodes. If TRUE, the algorithm will cache node IDs for each instance. Caching can speed up training of deeper trees. Defaults to FALSE.
checkpoint_interval	
	Set checkpoint interval (>= 1) or disable checkpoint (-1). E.g. 10 means that the cache will get checkpointed every 10 iterations, defaults to 10.
max_memory_in_mb	
	Maximum memory in MB allocated to histogram aggregation. If too small, then 1 node will be split per iteration, and its aggregates may exceed this size. Defaults to 256.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by [ft_r_formula]().
label_col	Label column name. The column should be a numeric column. Usually this column is output by [ft_r_formula]().
prediction_col	Prediction column name.
probability_col	
	Column name for predicted class conditional probabilities.
raw_prediction_col	
	Raw prediction (a.k.a. confidence) column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; see Details.
type	The type of model to fit. "regression" treats the response as a continuous variable, while "classification" treats the response as a categorical variable. When "auto" is used, the model type is inferred based on the response variable type – if it is a numeric type, then regression is used; classification otherwise.

variance_col	(Optional) Column name for the biased sample variance of prediction.
response	(Deprecated) The name of the response column (as a length-one character vector.)
features	(Deprecated) The name of features (terms) to use for the model fit.

Details

`ml_decision_tree` is a wrapper around `ml_decision_tree_regressor.tbl_spark` and `ml_decision_tree_classifier` and calls the appropriate method based on model type.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

See Also

Other ml algorithms: `ml_aft_survival_regression()`, `ml_gbt_classifier()`, `ml_generalized_linear_regression()`, `ml_isotonic_regression()`, `ml_linear_regression()`, `ml_linear_svc()`, `ml_logistic_regression()`, `ml_multilayer_perceptron_classifier()`, `ml_naive_bayes()`, `ml_one_vs_rest()`, `ml_random_forest_classifier()`

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

partitions <- iris_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

iris_training <- partitions$training
iris_test <- partitions$test

dt_model <- iris_training %>%
  ml_decision_tree(Species ~ .)

pred <- ml_predict(dt_model, iris_test)

ml_multiclass_classification_evaluator(pred)

## End(Not run)
```

---

ml_default_stop_words   *Default stop words*

Description

Loads the default stop words for the given language.

Usage

```
ml_default_stop_words(
  sc,
  language = c("english", "danish", "dutch", "finnish", "french", "german", "hungarian",
    "italian", "norwegian", "portuguese", "russian", "spanish", "swedish", "turkish"),
  ...
)
```

Arguments

sc	A spark_connection
language	A character string.
...	Optional arguments; currently unused.

Details

Supported languages: danish, dutch, english, finnish, french, german, hungarian, italian, norwegian, portuguese, russian, spanish, swedish, turkish. Defaults to English. See https://anoncvs.postgresql.org/cvsweb.cgi/pgsql/src/backend/snowball/stopwords/ for more details

Value

A list of stop words.

See Also

ft_stop_words_remover

---

ml_evaluate   *Evaluate the Model on a Validation Set*

Description

Compute performance metrics.

Usage

```
ml_evaluate(x, dataset)

## S3 method for class 'ml_model_logistic_regression'
ml_evaluate(x, dataset)

## S3 method for class 'ml_logistic_regression_model'
ml_evaluate(x, dataset)

## S3 method for class 'ml_model_linear_regression'
ml_evaluate(x, dataset)

## S3 method for class 'ml_linear_regression_model'
ml_evaluate(x, dataset)

## S3 method for class 'ml_model_generalized_linear_regression'
ml_evaluate(x, dataset)

## S3 method for class 'ml_generalized_linear_regression_model'
ml_evaluate(x, dataset)

## S3 method for class 'ml_model_clustering'
ml_evaluate(x, dataset)

## S3 method for class 'ml_model_classification'
ml_evaluate(x, dataset)

## S3 method for class 'ml_evaluator'
ml_evaluate(x, dataset)
```

Arguments

x	An ML model object or an evaluator object.
dataset	The dataset to be validate the model on.

Examples

```
## Not run:
sc <- spark_connect(master = "local")
```

```
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

ml_gaussian_mixture(iris_tbl, Species ~ .) %>%
  ml_evaluate(iris_tbl)

ml_kmeans(iris_tbl, Species ~ .) %>%
  ml_evaluate(iris_tbl)

ml_bisecting_kmeans(iris_tbl, Species ~ .) %>%
  ml_evaluate(iris_tbl)

## End(Not run)
```

---

ml_evaluator   *Spark ML - Evaluators*

Description

A set of functions to calculate performance metrics for prediction models. Also see the Spark ML
Documentation https://spark.apache.org/docs/latest/api/scala/index.html#org.apache.spark.ml.evaluation.package

Usage

```
ml_binary_classification_evaluator(
  x,
  label_col = "label",
  raw_prediction_col = "rawPrediction",
  metric_name = "areaUnderROC",
  uid = random_string("binary_classification_evaluator_"),
  ...
)

ml_binary_classification_eval(
  x,
  label_col = "label",
  prediction_col = "prediction",
  metric_name = "areaUnderROC"
)

ml_multiclass_classification_evaluator(
  x,
  label_col = "label",
  prediction_col = "prediction",
  metric_name = "f1",
  uid = random_string("multiclass_classification_evaluator_"),
  ...
)
```

```
ml_classification_eval(
  x,
  label_col = "label",
  prediction_col = "prediction",
  metric_name = "f1"
)

ml_regression_evaluator(
  x,
  label_col = "label",
  prediction_col = "prediction",
  metric_name = "rmse",
  uid = random_string("regression_evaluator_"),
  ...
)
```

Arguments

x	A spark_connection object or a tbl_spark containing label and prediction columns. The latter should be the output of [sdf_predict]().
label_col	Name of column string specifying which column contains the true labels or values.
raw_prediction_col	
	Raw prediction (a.k.a. confidence) column name.
metric_name	The performance metric. See details.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; currently unused.
prediction_col	Name of the column that contains the predicted label or value NOT the scored probability. Column should be of type Double.

Details

The following metrics are supported

- Binary Classification: areaUnderROC (default) or areaUnderPR (not available in Spark 2.X.)

- Multiclass Classification: f1 (default), precision, recall, weightedPrecision, weightedRecall or accuracy; for Spark 2.X: f1 (default), weightedPrecision, weightedRecall or accuracy.

- Regression: rmse (root mean squared error, default), mse (mean squared error), r2, or mae (mean absolute error.)

ml_binary_classification_eval() is an alias for ml_binary_classification_evaluator() for backwards compatibility.

ml_classification_eval() is an alias for ml_multiclass_classification_evaluator() for backwards compatibility.

Value

The calculated performance metric

Examples

```
## Not run:
sc <- spark_connect(master = "local")
mtcars_tbl <- sdf_copy_to(sc, mtcars, name = "mtcars_tbl", overwrite = TRUE)

partitions <- mtcars_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

mtcars_training <- partitions$training
mtcars_test <- partitions$test

# for multiclass classification
rf_model <- mtcars_training %>%
  ml_random_forest(cyl ~ ., type = "classification")

pred <- ml_predict(rf_model, mtcars_test)

ml_multiclass_classification_evaluator(pred)

# for regression
rf_model <- mtcars_training %>%
  ml_random_forest(cyl ~ ., type = "regression")

pred <- ml_predict(rf_model, mtcars_test)

ml_regression_evaluator(pred, label_col = "cyl")

# for binary classification
rf_model <- mtcars_training %>%
  ml_random_forest(am ~ gear + carb, type = "classification")

pred <- ml_predict(rf_model, mtcars_test)

ml_binary_classification_evaluator(pred)

## End(Not run)
```

```
ml_feature_importances
```
Spark ML - Feature Importance for Tree Models

Description

Spark ML - Feature Importance for Tree Models

Usage

```
ml_feature_importances(model, ...)

ml_tree_feature_importance(model, ...)
```

Arguments

model	A decision tree-based model.
...	Optional arguments; currently unused.

Value

For `ml_model`, a sorted data frame with feature labels and their relative importance. For `ml_prediction_model`, a vector of relative importances.

---

ml_fpgrowth	*Frequent Pattern Mining – FPGrowth*

Description

A parallel FP-growth algorithm to mine frequent itemsets.

Usage

```
ml_fpgrowth(
  x,
  items_col = "items",
  min_confidence = 0.8,
  min_support = 0.3,
  prediction_col = "prediction",
  uid = random_string("fpgrowth_"),
  ...
)

ml_association_rules(model)

ml_freq_itemsets(model)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
items_col	Items column name. Default: "items"
min_confidence	Minimal confidence for generating Association Rule. `min_confidence` will not affect the mining for frequent itemsets, but will affect the association rules generation. Default: 0.8

min_support	Minimal support level of the frequent pattern. [0.0, 1.0]. Any pattern that appears more than (min_support * size-of-the-dataset) times will be output in the frequent itemsets. Default: 0.3
prediction_col	Prediction column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; currently unused.
model	A fitted FPGrowth model returned by `ml_fpgrowth()`

---

ml_gaussian_mixture   *Spark ML – Gaussian Mixture clustering.*

Description

This class performs expectation maximization for multivariate Gaussian Mixture Models (GMMs). A GMM represents a composite distribution of independent Gaussian distributions with associated "mixing" weights specifying each's contribution to the composite. Given a set of sample points, this class will maximize the log-likelihood for a mixture of k Gaussians, iterating until the log-likelihood changes by less than `tol`, or until it has reached the max number of iterations. While this process is generally guaranteed to converge, it is not guaranteed to find a global optimum.

Usage

```
ml_gaussian_mixture(
  x,
  formula = NULL,
  k = 2,
  max_iter = 100,
  tol = 0.01,
  seed = NULL,
  features_col = "features",
  prediction_col = "prediction",
  probability_col = "probability",
  uid = random_string("gaussian_mixture_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
formula	Used when x is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
k	The number of clusters to create
max_iter	The maximum number of iterations to use.
tol	Param for the convergence tolerance for iterative algorithms.

seed	A random seed. Set this value if you need your results to be reproducible across repeated calls.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by `ft_r_formula`.
prediction_col	Prediction column name.
probability_col	
	Column name for predicted class conditional probabilities. Note: Not all models output well-calibrated probability estimates! These probabilities should be treated as confidences, not precise probabilities.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments, see Details. #' @return The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

gmm_model <- ml_gaussian_mixture(iris_tbl, Species ~ .)
pred <- sdf_predict(iris_tbl, gmm_model)
ml_clustering_evaluator(pred)

## End(Not run)
```

---

ml_gbt_classifier   *Spark ML – Gradient Boosted Trees*

Description

Perform binary classification and regression using gradient boosted trees. Multiclass classification is not supported yet.

Usage

```
ml_gbt_classifier(
  x,
  formula = NULL,
  max_iter = 20,
  max_depth = 5,
  step_size = 0.1,
  subsampling_rate = 1,
```

```
  feature_subset_strategy = "auto",
  min_instances_per_node = 1L,
  max_bins = 32,
  min_info_gain = 0,
  loss_type = "logistic",
  seed = NULL,
  thresholds = NULL,
  checkpoint_interval = 10,
  cache_node_ids = FALSE,
  max_memory_in_mb = 256,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  probability_col = "probability",
  raw_prediction_col = "rawPrediction",
  uid = random_string("gbt_classifier_"),
  ...
)

ml_gradient_boosted_trees(
  x,
  formula = NULL,
  type = c("auto", "regression", "classification"),
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  probability_col = "probability",
  raw_prediction_col = "rawPrediction",
  checkpoint_interval = 10,
  loss_type = c("auto", "logistic", "squared", "absolute"),
  max_bins = 32,
  max_depth = 5,
  max_iter = 20L,
  min_info_gain = 0,
  min_instances_per_node = 1,
  step_size = 0.1,
  subsampling_rate = 1,
  feature_subset_strategy = "auto",
  seed = NULL,
  thresholds = NULL,
  cache_node_ids = FALSE,
  max_memory_in_mb = 256,
  uid = random_string("gradient_boosted_trees_"),
  response = NULL,
  features = NULL,
  ...
)
```

```
ml_gbt_regressor(
  x,
  formula = NULL,
  max_iter = 20,
  max_depth = 5,
  step_size = 0.1,
  subsampling_rate = 1,
  feature_subset_strategy = "auto",
  min_instances_per_node = 1,
  max_bins = 32,
  min_info_gain = 0,
  loss_type = "squared",
  seed = NULL,
  checkpoint_interval = 10,
  cache_node_ids = FALSE,
  max_memory_in_mb = 256,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  uid = random_string("gbt_regressor_"),
  ...
)
```

Arguments

`x`	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
`formula`	Used when x is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see [ft_r_formula](#) for details.
`max_iter`	Maxmimum number of iterations.
`max_depth`	Maximum depth of the tree (>= 0); that is, the maximum number of nodes separating any leaves from the root of the tree.
`step_size`	Step size (a.k.a. learning rate) in interval (0, 1] for shrinking the contribution of each estimator. (default = 0.1)
`subsampling_rate`	
	Fraction of the training data used for learning each decision tree, in range (0, 1]. (default = 1.0)
`feature_subset_strategy`	
	The number of features to consider for splits at each tree node. See details for options.
`min_instances_per_node`	
	Minimum number of instances each child must have after split.
`max_bins`	The maximum number of bins used for discretizing continuous features and for choosing how to split on features at each node. More bins give higher granularity.
`min_info_gain`	Minimum information gain for a split to be considered at a tree node. Should be >= 0, defaults to 0.

loss_type	Loss function which GBT tries to minimize. Supported: "squared" (L2) and "absolute" (L1) (default = squared) for regression and "logistic" (default) for classification. For ml_gradient_boosted_trees, setting "auto" will default to the appropriate loss type based on model type.
seed	Seed for random numbers.
thresholds	Thresholds in multi-class classification to adjust the probability of predicting each class. Array must have length equal to the number of classes, with values > 0 excepting that at most one value may be 0. The class with largest value p/t is predicted, where p is the original probability of that class and t is the class's threshold.
checkpoint_interval	
	Set checkpoint interval (>= 1) or disable checkpoint (-1). E.g. 10 means that the cache will get checkpointed every 10 iterations, defaults to 10.
cache_node_ids	If FALSE, the algorithm will pass trees to executors to match instances with nodes. If TRUE, the algorithm will cache node IDs for each instance. Caching can speed up training of deeper trees. Defaults to FALSE.
max_memory_in_mb	
	Maximum memory in MB allocated to histogram aggregation. If too small, then 1 node will be split per iteration, and its aggregates may exceed this size. Defaults to 256.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by [ft_r_formula]().
label_col	Label column name. The column should be a numeric column. Usually this column is output by [ft_r_formula]().
prediction_col	Prediction column name.
probability_col	
	Column name for predicted class conditional probabilities.
raw_prediction_col	
	Raw prediction (a.k.a. confidence) column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; see Details.
type	The type of model to fit. "regression" treats the response as a continuous variable, while "classification" treats the response as a categorical variable. When "auto" is used, the model type is inferred based on the response variable type – if it is a numeric type, then regression is used; classification otherwise.
response	(Deprecated) The name of the response column (as a length-one character vector.)
features	(Deprecated) The name of features (terms) to use for the model fit.

Details

The supported options for feature_subset_strategy are

- "auto": Choose automatically for task: If num_trees == 1, set to "all". If num_trees > 1 (forest), set to "sqrt" for classification and to "onethird" for regression.

- "all": use all features

- "onethird": use 1/3 of the features

- "sqrt": use use sqrt(number of features)

- "log2": use log2(number of features)

- "n": when n is in the range (0, 1.0], use n * number of features. When n is in the range (1, number of features), use n features. (default = "auto")

ml_gradient_boosted_trees is a wrapper around ml_gbt_regressor.tbl_spark and ml_gbt_classifier.tbl_spark and calls the appropriate method based on model type.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the predictor appended to it. If a tbl_spark, it will return a tbl_spark with the predictions added to it.

See Also

Other ml algorithms: ml_aft_survival_regression(), ml_decision_tree_classifier(), ml_generalized_linear_r ml_isotonic_regression(), ml_linear_regression(), ml_linear_svc(), ml_logistic_regression(), ml_multilayer_perceptron_classifier(), ml_naive_bayes(), ml_one_vs_rest(), ml_random_forest_classifier(

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

partitions <- iris_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

iris_training <- partitions$training
iris_test <- partitions$test

gbt_model <- iris_training %>%
  ml_gradient_boosted_trees(Sepal_Length ~ Petal_Length + Petal_Width)

pred <- ml_predict(gbt_model, iris_test)

ml_regression_evaluator(pred, label_col = "Sepal_Length")

## End(Not run)
```

```
ml_generalized_linear_regression
```
Spark ML – Generalized Linear Regression

Description

Perform regression using Generalized Linear Model (GLM).

Usage

```
ml_generalized_linear_regression(
  x,
  formula = NULL,
  family = "gaussian",
  link = NULL,
  fit_intercept = TRUE,
  offset_col = NULL,
  link_power = NULL,
  link_prediction_col = NULL,
  reg_param = 0,
  max_iter = 25,
  weight_col = NULL,
  solver = "irls",
  tol = 1e-06,
  variance_power = 0,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  uid = random_string("generalized_linear_regression_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
formula	Used when x is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
family	Name of family which is a description of the error distribution to be used in the model. Supported options: "gaussian", "binomial", "poisson", "gamma" and "tweedie". Default is "gaussian".
link	Name of link function which provides the relationship between the linear predictor and the mean of the distribution function. See for supported link functions.
fit_intercept	Boolean; should the model be fit with an intercept term?
offset_col	Offset column name. If this is not set, we treat all instance offsets as 0.0. The feature specified as offset has a constant coefficient of 1.0.

link_power	Index in the power link function. Only applicable to the Tweedie family. Note that link power 0, 1, -1 or 0.5 corresponds to the Log, Identity, Inverse or Sqrt link, respectively. When not set, this value defaults to 1 - variancePower, which matches the R "statmod" package.
link_prediction_col	
	Link prediction (linear predictor) column name. Default is not set, which means we do not output link prediction.
reg_param	Regularization parameter (aka lambda)
max_iter	The maximum number of iterations to use.
weight_col	The name of the column to use as weights for the model fit.
solver	Solver algorithm for optimization.
tol	Param for the convergence tolerance for iterative algorithms.
variance_power	Power in the variance function of the Tweedie distribution which provides the relationship between the variance and mean of the distribution. Only applicable to the Tweedie family. (see Tweedie Distribution (Wikipedia)) Supported values: 0 and [1, Inf). Note that variance power 0, 1, or 2 corresponds to the Gaussian, Poisson or Gamma family, respectively.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by ft_r_formula.
label_col	Label column name. The column should be a numeric column. Usually this column is output by ft_r_formula.
prediction_col	Prediction column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; see Details.

Details

Valid link functions for each family is listed below. The first link function of each family is the default one.

- gaussian: "identity", "log", "inverse"
- binomial: "logit", "probit", "loglog"
- poisson: "log", "identity", "sqrt"
- gamma: "inverse", "identity", "log"
- tweedie: power link function specified through link_power. The default link power in the tweedie family is 1 - variance_power.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the predictor appended to it. If a tbl_spark, it will return a tbl_spark with the predictions added to it.

See Also

Other ml algorithms: ml_aft_survival_regression(), ml_decision_tree_classifier(), ml_gbt_classifier(), ml_isotonic_regression(), ml_linear_regression(), ml_linear_svc(), ml_logistic_regression(), ml_multilayer_perceptron_classifier(), ml_naive_bayes(), ml_one_vs_rest(), ml_random_forest_classifier(

Examples

```
## Not run:
library(sparklyr)

sc <- spark_connect(master = "local")
mtcars_tbl <- sdf_copy_to(sc, mtcars, name = "mtcars_tbl", overwrite = TRUE)

partitions <- mtcars_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

mtcars_training <- partitions$training
mtcars_test <- partitions$test

# Specify the grid
family <- c("gaussian", "gamma", "poisson")
link <- c("identity", "log")
family_link <- expand.grid(family = family, link = link, stringsAsFactors = FALSE)
family_link <- data.frame(family_link, rmse = 0)

# Train the models
for (i in seq_len(nrow(family_link))) {
  glm_model <- mtcars_training %>%
    ml_generalized_linear_regression(mpg ~ .,
      family = family_link[i, 1],
      link = family_link[i, 2]
    )

  pred <- ml_predict(glm_model, mtcars_test)
  family_link[i, 3] <- ml_regression_evaluator(pred, label_col = "mpg")
}

family_link

## End(Not run)
```

---

ml_glm_tidiers   *Tidying methods for Spark ML linear models*

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_generalized_linear_regression'
tidy(x, exponentiate = FALSE, ...)

## S3 method for class 'ml_model_linear_regression'
tidy(x, ...)

## S3 method for class 'ml_model_generalized_linear_regression'
augment(
  x,
  newdata = NULL,
  type.residuals = c("working", "deviance", "pearson", "response"),
  ...
)

## S3 method for class '`_ml_model_linear_regression`'
augment(
  x,
  new_data = NULL,
  type.residuals = c("working", "deviance", "pearson", "response"),
  ...
)

## S3 method for class 'ml_model_linear_regression'
augment(
  x,
  newdata = NULL,
  type.residuals = c("working", "deviance", "pearson", "response"),
  ...
)

## S3 method for class 'ml_model_generalized_linear_regression'
glance(x, ...)

## S3 method for class 'ml_model_linear_regression'
glance(x, ...)
```

Arguments

x	a Spark ML model.
exponentiate	For GLM, whether to exponentiate the coefficient estimates (typical for logistic regression.)
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.
type.residuals	type of residuals, defaults to "working". Must be set to "working" when newdata is supplied.
new_data	a tbl_spark of new data to use for prediction.

Details

The residuals attached by `augment` are of type "working" by default, which is different from the default of "deviance" for `residuals()` or `sdf_residuals()`.

`ml_isotonic_regression`

Spark ML – Isotonic Regression

Description

Currently implemented using parallelized pool adjacent violators algorithm. Only univariate (single feature) algorithm supported.

Usage

```
ml_isotonic_regression(
  x,
  formula = NULL,
  feature_index = 0,
  isotonic = TRUE,
  weight_col = NULL,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  uid = random_string("isotonic_regression_"),
  ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
formula	Used when x is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see [ft_r_formula](#) for details.
feature_index	Index of the feature if `features_col` is a vector column (default: 0), no effect otherwise.
isotonic	Whether the output sequence should be isotonic/increasing (true) or antitonic/decreasing (false). Default: true
weight_col	The name of the column to use as weights for the model fit.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by [ft_r_formula](#).
label_col	Label column name. The column should be a numeric column. Usually this column is output by [ft_r_formula](#).

prediction_col	Prediction column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; see Details.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

See Also

Other ml algorithms: `ml_aft_survival_regression()`, `ml_decision_tree_classifier()`, `ml_gbt_classifier()`, `ml_generalized_linear_regression()`, `ml_linear_regression()`, `ml_linear_svc()`, `ml_logistic_regression()`, `ml_multilayer_perceptron_classifier()`, `ml_naive_bayes()`, `ml_one_vs_rest()`, `ml_random_forest_classifier(`

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

partitions <- iris_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

iris_training <- partitions$training
iris_test <- partitions$test

iso_res <- iris_tbl %>%
  ml_isotonic_regression(Petal_Length ~ Petal_Width)

pred <- ml_predict(iso_res, iris_test)

pred

## End(Not run)
```

---

ml_isotonic_regression_tidiers
Tidying methods for Spark ML Isotonic Regression

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_isotonic_regression'
tidy(x, ...)

## S3 method for class 'ml_model_isotonic_regression'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_isotonic_regression'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.

---

ml_kmeans   *Spark ML – K-Means Clustering*

Description

K-means clustering with support for k-means‖ initialization proposed by Bahmani et al. Using 'ml_kmeans()' with the formula interface requires Spark 2.0+.

Usage

```
ml_kmeans(
  x,
  formula = NULL,
  k = 2,
  max_iter = 20,
  tol = 1e-04,
  init_steps = 2,
  init_mode = "k-means||",
  seed = NULL,
  features_col = "features",
  prediction_col = "prediction",
  uid = random_string("kmeans_"),
  ...
)

ml_compute_cost(model, dataset)

ml_compute_silhouette_measure(
  model,
  dataset,
  distance_measure = c("squaredEuclidean", "cosine")
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
formula	Used when x is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
k	The number of clusters to create
max_iter	The maximum number of iterations to use.
tol	Param for the convergence tolerance for iterative algorithms.
init_steps	Number of steps for the k-means‖ initialization mode. This is an advanced setting – the default of 2 is almost always enough. Must be > 0. Default: 2.
init_mode	Initialization algorithm. This can be either "random" to choose random points as initial cluster centers, or "k-means‖" to use a parallel variant of k-means++ (Bahmani et al., Scalable K-Means++, VLDB 2012). Default: k-means‖.
seed	A random seed. Set this value if you need your results to be reproducible across repeated calls.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by `ft_r_formula`.
prediction_col	Prediction column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments, see Details. #' @return The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.
model	A fitted K-means model returned by `ml_kmeans()`
dataset	Dataset on which to calculate K-means cost
distance_measure	
	Distance measure to apply when computing the Silhouette measure.

Value

`ml_compute_cost()` returns the K-means cost (sum of squared distances of points to their nearest center) for the model on the given data.

`ml_compute_silhouette_measure()` returns the Silhouette measure of the clustering on the given data.

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)
ml_kmeans(iris_tbl, Species ~ .)

## End(Not run)
```

ml_kmeans_cluster_eval

Evaluate a K-mean clustering

Description

Evaluate a K-mean clustering

Arguments

model   A fitted K-means model returned by ml_kmeans()

dataset   Dataset on which to calculate K-means cost

ml_lda   *Spark ML – Latent Dirichlet Allocation*

Description

Latent Dirichlet Allocation (LDA), a topic model designed for text documents.

Usage

```
ml_lda(
  x,
  formula = NULL,
  k = 10,
  max_iter = 20,
  doc_concentration = NULL,
  topic_concentration = NULL,
  subsampling_rate = 0.05,
  optimizer = "online",
  checkpoint_interval = 10,
  keep_last_checkpoint = TRUE,
  learning_decay = 0.51,
  learning_offset = 1024,
  optimize_doc_concentration = TRUE,
  seed = NULL,
  features_col = "features",
  topic_distribution_col = "topicDistribution",
  uid = random_string("lda_"),
  ...
)

ml_describe_topics(model, max_terms_per_topic = 10)
```

```
ml_log_likelihood(model, dataset)

ml_log_perplexity(model, dataset)

ml_topics_matrix(model)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
formula	Used when x is a tbl_spark. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
k	The number of clusters to create
max_iter	The maximum number of iterations to use.
doc_concentration	
	Concentration parameter (commonly named "alpha") for the prior placed on documents' distributions over topics ("theta"). See details.
topic_concentration	
	Concentration parameter (commonly named "beta" or "eta") for the prior placed on topics' distributions over terms.
subsampling_rate	
	(For Online optimizer only) Fraction of the corpus to be sampled and used in each iteration of mini-batch gradient descent, in range (0, 1]. Note that this should be adjusted in synch with max_iter so the entire corpus is used. Specifically, set both so that maxIterations * miniBatchFraction greater than or equal to 1.
optimizer	Optimizer or inference algorithm used to estimate the LDA model. Supported: "online" for Online Variational Bayes (default) and "em" for Expectation-Maximization.
checkpoint_interval	
	Set checkpoint interval (>= 1) or disable checkpoint (-1). E.g. 10 means that the cache will get checkpointed every 10 iterations, defaults to 10.
keep_last_checkpoint	
	(Spark 2.0.0+) (For EM optimizer only) If using checkpointing, this indicates whether to keep the last checkpoint. If FALSE, then the checkpoint will be deleted. Deleting the checkpoint can cause failures if a data partition is lost, so set this bit with care. Note that checkpoints will be cleaned up via reference counting, regardless.
learning_decay	(For Online optimizer only) Learning rate, set as an exponential decay rate. This should be between (0.5, 1.0] to guarantee asymptotic convergence. This is called "kappa" in the Online LDA paper (Hoffman et al., 2010). Default: 0.51, based on Hoffman et al.
learning_offset	
	(For Online optimizer only) A (positive) learning parameter that downweights early iterations. Larger values make early iterations count less. This is called "tau0" in the Online LDA paper (Hoffman et al., 2010) Default: 1024, following Hoffman et al.

optimize_doc_concentration

        (For Online optimizer only) Indicates whether the `doc_concentration` (Dirichlet parameter for document-topic distribution) will be optimized during training. Setting this to true will make the model more expressive and fit the training data better. Default: `FALSE`

seed   A random seed. Set this value if you need your results to be reproducible across repeated calls.

features_col   Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by [ft_r_formula](#).

topic_distribution_col

        Output column with estimates of the topic mixture distribution for each document (often called "theta" in the literature). Returns a vector of zeros for an empty document.

uid   A character string used to uniquely identify the ML estimator.

...   Optional arguments, see Details. #' @return The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

model   A fitted LDA model returned by `ml_lda()`.

max_terms_per_topic

        Maximum number of terms to collect for each topic. Default value of 10.

dataset   test corpus to use for calculating log likelihood or log perplexity

Details

For 'ml_lda.tbl_spark' with the formula interface, you can specify named arguments in '...' that will be passed 'ft_regex_tokenizer()', 'ft_stop_words_remover()', and 'ft_count_vectorizer()'. For example, to increase the default 'min_token_length', you can use 'ml_lda(dataset, ~ text, min_token_length = 4)'.

Terminology for LDA:

- "term" = "word": an element of the vocabulary

- "token": instance of a term appearing in a document

- "topic": multinomial distribution over terms representing some concept

- "document": one piece of text, corresponding to one row in the input data

Original LDA paper (journal version): Blei, Ng, and Jordan. "Latent Dirichlet Allocation." JMLR, 2003.

Input data (`features_col`): LDA is given a collection of documents as input data, via the `features_col` parameter. Each document is specified as a Vector of length `vocab_size`, where each entry is the count for the corresponding term (word) in the document. Feature transformers such as [ft_tokenizer](#) and [ft_count_vectorizer](#) can be useful for converting text to word count vectors

Value

ml_describe_topics returns a DataFrame with topics and their top-weighted terms.

ml_log_likelihood calculates a lower bound on the log likelihood of the entire corpus

Parameter details

doc_concentration: This is the parameter to a Dirichlet distribution, where larger values mean more smoothing (more regularization). If not set by the user, then doc_concentration is set automatically. If set to singleton vector [alpha], then alpha is replicated to a vector of length k in fitting. Otherwise, the doc_concentration vector must be length k. (default = automatic)

Optimizer-specific parameter settings:

EM

- Currently only supports symmetric distributions, so all values in the vector should be the same.
- Values should be greater than 1.0
- default = uniformly (50 / k) + 1, where 50/k is common in LDA libraries and +1 follows from Asuncion et al. (2009), who recommend a +1 adjustment for EM.

Online

- Values should be greater than or equal to 0
- default = uniformly (1.0 / k), following the implementation from here

topic_concentration:

This is the parameter to a symmetric Dirichlet distribution.

Note: The topics' distributions over terms are called "beta" in the original LDA paper by Blei et al., but are called "phi" in many later papers such as Asuncion et al., 2009.

If not set by the user, then topic_concentration is set automatically. (default = automatic)

Optimizer-specific parameter settings:

EM

- Value should be greater than 1.0
- default = 0.1 + 1, where 0.1 gives a small amount of smoothing and +1 follows Asuncion et al. (2009), who recommend a +1 adjustment for EM.

Online

- Value should be greater than or equal to 0
- default = (1.0 / k), following the implementation from here.

topic_distribution_col: This uses a variational approximation following Hoffman et al. (2010), where the approximate distribution is called "gamma." Technically, this method returns this approximation "gamma" for each document.

Examples

```
## Not run:
library(janeaustenr)
library(dplyr)
sc <- spark_connect(master = "local")
```

```
lines_tbl <- sdf_copy_to(sc,
  austen_books()[c(1:30), ],
  name = "lines_tbl",
  overwrite = TRUE
)

# transform the data in a tidy form
lines_tbl_tidy <- lines_tbl %>%
  ft_tokenizer(
    input_col = "text",
    output_col = "word_list"
  ) %>%
  ft_stop_words_remover(
    input_col = "word_list",
    output_col = "wo_stop_words"
  ) %>%
  mutate(text = explode(wo_stop_words)) %>%
  filter(text != "") %>%
  select(text, book)

lda_model <- lines_tbl_tidy %>%
  ml_lda(~text, k = 4)

# vocabulary and topics
tidy(lda_model)

## End(Not run)
```

---

ml_lda_tidiers   *Tidying methods for Spark ML LDA models*

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_lda'
tidy(x, ...)

## S3 method for class 'ml_model_lda'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_lda'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.

---

ml_linear_regression   *Spark ML – Linear Regression*

Description

Perform regression using linear regression.

Usage

```
ml_linear_regression(
  x,
  formula = NULL,
  fit_intercept = TRUE,
  elastic_net_param = 0,
  reg_param = 0,
  max_iter = 100,
  weight_col = NULL,
  loss = "squaredError",
  solver = "auto",
  standardization = TRUE,
  tol = 1e-06,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  uid = random_string("linear_regression_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
formula	Used when x is a tbl_spark. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
fit_intercept	Boolean; should the model be fit with an intercept term?
elastic_net_param	
	ElasticNet mixing parameter, in range [0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it is an L1 penalty.
reg_param	Regularization parameter (aka lambda)
max_iter	The maximum number of iterations to use.

weight_col	The name of the column to use as weights for the model fit.
loss	The loss function to be optimized. Supported options: "squaredError" and "huber". Default: "squaredError"
solver	Solver algorithm for optimization.
standardization	
	Whether to standardize the training features before fitting the model.
tol	Param for the convergence tolerance for iterative algorithms.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by `ft_r_formula`.
label_col	Label column name. The column should be a numeric column. Usually this column is output by `ft_r_formula`.
prediction_col	Prediction column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; see Details.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

See Also

Other ml algorithms: `ml_aft_survival_regression()`, `ml_decision_tree_classifier()`, `ml_gbt_classifier()`, `ml_generalized_linear_regression()`, `ml_isotonic_regression()`, `ml_linear_svc()`, `ml_logistic_regression()`, `ml_multilayer_perceptron_classifier()`, `ml_naive_bayes()`, `ml_one_vs_rest()`, `ml_random_forest_classifier()`

Examples

```
## Not run:
sc <- spark_connect(master = "local")
mtcars_tbl <- sdf_copy_to(sc, mtcars, name = "mtcars_tbl", overwrite = TRUE)

partitions <- mtcars_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

mtcars_training <- partitions$training
mtcars_test <- partitions$test

lm_model <- mtcars_training %>%
  ml_linear_regression(mpg ~ .)

pred <- ml_predict(lm_model, mtcars_test)

ml_regression_evaluator(pred, label_col = "mpg")

## End(Not run)
```

ml_linear_svc *Spark ML – LinearSVC*

Description

Perform classification using linear support vector machines (SVM). This binary classifier optimizes the Hinge Loss using the OWLQN optimizer. Only supports L2 regularization currently.

Usage

```
ml_linear_svc(
  x,
  formula = NULL,
  fit_intercept = TRUE,
  reg_param = 0,
  max_iter = 100,
  standardization = TRUE,
  weight_col = NULL,
  tol = 1e-06,
  threshold = 0,
  aggregation_depth = 2,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  raw_prediction_col = "rawPrediction",
  uid = random_string("linear_svc_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
formula	Used when x is a tbl_spark. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
fit_intercept	Boolean; should the model be fit with an intercept term?
reg_param	Regularization parameter (aka lambda)
max_iter	The maximum number of iterations to use.
standardization	Whether to standardize the training features before fitting the model.
weight_col	The name of the column to use as weights for the model fit.
tol	Param for the convergence tolerance for iterative algorithms.
threshold	in binary classification prediction, in range [0, 1].
aggregation_depth	(Spark 2.1.0+) Suggested depth for treeAggregate (>= 2).

features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by [ft_r_formula](#).
label_col	Label column name. The column should be a numeric column. Usually this column is output by [ft_r_formula](#).
prediction_col	Prediction column name.
raw_prediction_col	
	Raw prediction (a.k.a. confidence) column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; see Details.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

See Also

Other ml algorithms: [ml_aft_survival_regression()](#), [ml_decision_tree_classifier()](#), [ml_gbt_classifier()](#), [ml_generalized_linear_regression()](#), [ml_isotonic_regression()](#), [ml_linear_regression()](#), [ml_logistic_regression()](#), [ml_multilayer_perceptron_classifier()](#), [ml_naive_bayes()](#), [ml_one_vs_rest()](#), [ml_random_forest_classifier()](#)

Examples

```
## Not run:
library(dplyr)

sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

partitions <- iris_tbl %>%
  filter(Species != "setosa") %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

iris_training <- partitions$training
iris_test <- partitions$test

svc_model <- iris_training %>%
  ml_linear_svc(Species ~ .)

pred <- ml_predict(svc_model, iris_test)

ml_binary_classification_evaluator(pred)

## End(Not run)
```

---

ml_linear_svc_tidiers   *Tidying methods for Spark ML linear svc*

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_linear_svc'
tidy(x, ...)

## S3 method for class 'ml_model_linear_svc'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_linear_svc'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.

---

ml_logistic_regression

Spark ML – Logistic Regression

Description

Perform classification using logistic regression.

Usage

```
ml_logistic_regression(
  x,
  formula = NULL,
  fit_intercept = TRUE,
  elastic_net_param = 0,
  reg_param = 0,
  max_iter = 100,
  threshold = 0.5,
  thresholds = NULL,
  tol = 1e-06,
  weight_col = NULL,
```

```
  aggregation_depth = 2,
  lower_bounds_on_coefficients = NULL,
  lower_bounds_on_intercepts = NULL,
  upper_bounds_on_coefficients = NULL,
  upper_bounds_on_intercepts = NULL,
  features_col = "features",
  label_col = "label",
  family = "auto",
  prediction_col = "prediction",
  probability_col = "probability",
  raw_prediction_col = "rawPrediction",
  uid = random_string("logistic_regression_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
formula	Used when x is a tbl_spark. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
fit_intercept	Boolean; should the model be fit with an intercept term?
elastic_net_param	
	ElasticNet mixing parameter, in range [0, 1]. For alpha = 0, the penalty is an L2 penalty. For alpha = 1, it is an L1 penalty.
reg_param	Regularization parameter (aka lambda)
max_iter	The maximum number of iterations to use.
threshold	in binary classification prediction, in range [0, 1].
thresholds	Thresholds in multi-class classification to adjust the probability of predicting each class. Array must have length equal to the number of classes, with values > 0 excepting that at most one value may be 0. The class with largest value p/t is predicted, where p is the original probability of that class and t is the class's threshold.
tol	Param for the convergence tolerance for iterative algorithms.
weight_col	The name of the column to use as weights for the model fit.
aggregation_depth	
	(Spark 2.1.0+) Suggested depth for treeAggregate (>= 2).
lower_bounds_on_coefficients	
	(Spark 2.2.0+) Lower bounds on coefficients if fitting under bound constrained optimization. The bound matrix must be compatible with the shape (1, number of features) for binomial regression, or (number of classes, number of features) for multinomial regression.
lower_bounds_on_intercepts	
	(Spark 2.2.0+) Lower bounds on intercepts if fitting under bound constrained optimization. The bounds vector size must be equal with 1 for binomial regression, or the number of classes for multinomial regression.

upper_bounds_on_coefficients

(Spark 2.2.0+) Upper bounds on coefficients if fitting under bound constrained optimization. The bound matrix must be compatible with the shape (1, number of features) for binomial regression, or (number of classes, number of features) for multinomial regression.

upper_bounds_on_intercepts

(Spark 2.2.0+) Upper bounds on intercepts if fitting under bound constrained optimization. The bounds vector size must be equal with 1 for binomial regression, or the number of classes for multinomial regression.

features_col   Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by ft_r_formula.

label_col   Label column name. The column should be a numeric column. Usually this column is output by ft_r_formula.

family   (Spark 2.1.0+) Param for the name of family which is a description of the label distribution to be used in the model. Supported options: "auto", "binomial", and "multinomial."

prediction_col   Prediction column name.

probability_col

Column name for predicted class conditional probabilities.

raw_prediction_col

Raw prediction (a.k.a. confidence) column name.

uid   A character string used to uniquely identify the ML estimator.

...   Optional arguments; see Details.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

See Also

Other ml algorithms: [ml_aft_survival_regression()](ml_aft_survival_regression), [ml_decision_tree_classifier()](ml_decision_tree_classifier), [ml_gbt_classifier()](ml_gbt_classifier), [ml_generalized_linear_regression()](ml_generalized_linear_regression), [ml_isotonic_regression()](ml_isotonic_regression), [ml_linear_regression()](ml_linear_regression), [ml_linear_svc()](ml_linear_svc), [ml_multilayer_perceptron_classifier()](ml_multilayer_perceptron_classifier), [ml_naive_bayes()](ml_naive_bayes), [ml_one_vs_rest()](ml_one_vs_rest), [ml_random_forest_classifier()](ml_random_forest_classifier)

Examples

```
## Not run:
sc <- spark_connect(master = "local")
mtcars_tbl <- sdf_copy_to(sc, mtcars, name = "mtcars_tbl", overwrite = TRUE)

partitions <- mtcars_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

mtcars_training <- partitions$training
```

```
mtcars_test <- partitions$test

lr_model <- mtcars_training %>%
  ml_logistic_regression(am ~ gear + carb)

pred <- ml_predict(lr_model, mtcars_test)

ml_binary_classification_evaluator(pred)

## End(Not run)
```

ml_logistic_regression_tidiers

Tidying methods for Spark ML Logistic Regression

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_logistic_regression'
tidy(x, ...)

## S3 method for class 'ml_model_logistic_regression'
augment(x, newdata = NULL, ...)

## S3 method for class '`_ml_model_logistic_regression`'
augment(x, new_data = NULL, ...)

## S3 method for class 'ml_model_logistic_regression'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.
new_data	a tbl_spark of new data to use for prediction.

---

ml_metrics_binary   *Extracts metrics from a fitted table*

Description

The function works best when passed a 'tbl_spark' created by 'ml_predict()'. The output 'tbl_spark' will contain the correct variable types and format that the given Spark model "evaluator" expects.

Usage

```
ml_metrics_binary(
  x,
  truth = label,
  estimate = rawPrediction,
  metrics = c("roc_auc", "pr_auc"),
  ...
)
```

Arguments

x	A 'tbl_spark' containing the estimate (prediction) and the truth (value of what actually happened)
truth	The name of the column from 'x' with an integer field containing the binary response (0 or 1). The 'ml_predict()' function will create a new field named 'label' which contains the expected type and values. 'truth' defaults to 'label'.
estimate	The name of the column from 'x' that contains the prediction. Defaults to 'rawPrediction', since its type and expected values will match 'truth'.
metrics	A character vector with the metrics to calculate. For binary models the possible values are: 'roc_auc' (Area under the Receiver Operator curve), 'pr_auc' (Area under the Precesion Recall curve). Defaults to: 'roc_auc', 'pr_auc'
...	Optional arguments; currently unused.

Details

The 'ml_metrics' family of functions implement Spark's 'evaluate' closer to how the 'yardstick' package works. The functions expect a table containing the truth and estimate, and return a 'tibble' with the results. The 'tibble' has the same format and variable names as the output of the 'yardstick' functions.

Examples

```
## Not run:
sc <- spark_connect("local")
tbl_iris <- copy_to(sc, iris)
prep_iris <- tbl_iris %>%
  mutate(is_setosa = ifelse(Species == "setosa", 1, 0))
iris_split <- sdf_random_split(prep_iris, training = 0.5, test = 0.5)
```

```
model <- ml_logistic_regression(iris_split$training, "is_setosa ~ Sepal_Length")
tbl_predictions <- ml_predict(model, iris_split$test)
ml_metrics_binary(tbl_predictions)

## End(Not run)
```

---

ml_metrics_multiclass   *Extracts metrics from a fitted table*

Description

The function works best when passed a 'tbl_spark' created by 'ml_predict()'. The output 'tbl_spark' will contain the correct variable types and format that the given Spark model "evaluator" expects.

Usage

```
ml_metrics_multiclass(
  x,
  truth = label,
  estimate = prediction,
  metrics = c("accuracy"),
  beta = NULL,
  ...
)
```

Arguments

x	A 'tbl_spark' containing the estimate (prediction) and the truth (value of what actually happened)
truth	The name of the column from 'x' with an integer field containing an the indexed value for each outcome . The 'ml_predict()' function will create a new field named 'label' which contains the expected type and values. 'truth' defaults to 'label'.
estimate	The name of the column from 'x' that contains the prediction. Defaults to 'prediction', since its type and indexed values will match 'truth'.
metrics	A character vector with the metrics to calculate. For multiclass models the possible values are: 'acurracy', 'f_meas' (F-score), 'recall' and 'precision'. This function translates the argument into an acceptable Spark parameter. If no translation is found, then the raw value of the argument is passed to Spark. This makes it possible to request a metric that is not listed here but, depending on version, it is available in Spark. Other metrics form multi-class models are: 'weightedTruePositiveRate', 'weightedFalsePositiveRate', 'weightedFMeasure', 'truePositiveRateByLabel', 'falsePositiveRateByLabel', 'precisionByLabel', 'recallByLabel', 'fMeasureByLabel', 'logLoss', 'hammingLoss'
beta	Numerical value used for precision and recall. Defaults to NULL, but if the Spark session's verion is 3.0 and above, then NULL is changed to 1, unless something different is supplied in this argument.
...	Optional arguments; currently unused.

Details

The 'ml_metrics' family of functions implement Spark's 'evaluate' closer to how the 'yardstick' package works. The functions expect a table containing the truth and estimate, and return a 'tibble' with the results. The 'tibble' has the same format and variable names as the output of the 'yardstick' functions.

Examples

```
## Not run:
sc <- spark_connect("local")
tbl_iris <- copy_to(sc, iris)
iris_split <- sdf_random_split(tbl_iris, training = 0.5, test = 0.5)
model <- ml_random_forest(iris_split$training, "Species ~ .")
tbl_predictions <- ml_predict(model, iris_split$test)

ml_metrics_multiclass(tbl_predictions)

# Request different metrics
ml_metrics_multiclass(tbl_predictions, metrics = c("recall", "precision"))

# Request metrics not translated by the function, but valid in Spark
ml_metrics_multiclass(tbl_predictions, metrics = c("logLoss", "hammingLoss"))

## End(Not run)
```

---

ml_metrics_regression *Extracts metrics from a fitted table*

Description

The function works best when passed a 'tbl_spark' created by 'ml_predict()'. The output 'tbl_spark' will contain the correct variable types and format that the given Spark model "evaluator" expects.

Usage

```
ml_metrics_regression(
  x,
  truth,
  estimate = prediction,
  metrics = c("rmse", "rsq", "mae"),
  ...
)
```

Arguments

x   A 'tbl_spark' containing the estimate (prediction) and the truth (value of what actually happened)

truth	The name of the column from 'x' that contains the value of what actually happened
estimate	The name of the column from 'x' that contains the prediction. Defaults to 'prediction', since it is the default that 'ml_predict()' uses.
metrics	A character vector with the metrics to calculate. For regression models the possible values are: 'rmse' (Root mean squared error), 'mse' (Mean squared error),'rsq' (R squared), 'mae' (Mean absolute error), and 'var' (Explained variance). Defaults to: 'rmse', 'rsq', 'mae'
...	Optional arguments; currently unused.

Details

The 'ml_metrics' family of functions implement Spark's 'evaluate' closer to how the 'yardstick' package works. The functions expect a table containing the truth and estimate, and return a 'tibble' with the results. The 'tibble' has the same format and variable names as the output of the 'yardstick' functions.

Examples

```
## Not run:
sc <- spark_connect("local")
tbl_iris <- copy_to(sc, iris)
iris_split <- sdf_random_split(tbl_iris, training = 0.5, test = 0.5)
training <- iris_split$training
reg_formula <- "Sepal_Length ~ Sepal_Width + Petal_Length + Petal_Width"
model <- ml_generalized_linear_regression(training, reg_formula)
tbl_predictions <- ml_predict(model, iris_split$test)
tbl_predictions %>%
  ml_metrics_regression(Sepal_Length)

## End(Not run)
```

---

ml_model_data   *Extracts data associated with a Spark ML model*

Description

Extracts data associated with a Spark ML model

Usage

```
ml_model_data(object)
```

Arguments

object	a Spark ML model

Value

A tbl_spark

```
ml_multilayer_perceptron_classifier
```
Spark ML – Multilayer Perceptron

Description

Classification model based on the Multilayer Perceptron. Each layer has sigmoid activation function, output layer has softmax.

Usage

```
ml_multilayer_perceptron_classifier(
  x,
  formula = NULL,
  layers = NULL,
  max_iter = 100,
  step_size = 0.03,
  tol = 1e-06,
  block_size = 128,
  solver = "l-bfgs",
  seed = NULL,
  initial_weights = NULL,
  thresholds = NULL,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  probability_col = "probability",
  raw_prediction_col = "rawPrediction",
  uid = random_string("multilayer_perceptron_classifier_"),
  ...
)

ml_multilayer_perceptron(
  x,
  formula = NULL,
  layers,
  max_iter = 100,
  step_size = 0.03,
  tol = 1e-06,
  block_size = 128,
  solver = "l-bfgs",
  seed = NULL,
  initial_weights = NULL,
```

```
    features_col = "features",
    label_col = "label",
    thresholds = NULL,
    prediction_col = "prediction",
    probability_col = "probability",
    raw_prediction_col = "rawPrediction",
    uid = random_string("multilayer_perceptron_classifier_"),
    response = NULL,
    features = NULL,
    ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
formula	Used when `x` is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see [ft_r_formula](#) for details.
layers	A numeric vector describing the layers – each element in the vector gives the size of a layer. For example, `c(4, 5, 2)` would imply three layers, with an input (feature) layer of size 4, an intermediate layer of size 5, and an output (class) layer of size 2.
max_iter	The maximum number of iterations to use.
step_size	Step size to be used for each iteration of optimization (> 0).
tol	Param for the convergence tolerance for iterative algorithms.
block_size	Block size for stacking input data in matrices to speed up the computation. Data is stacked within partitions. If block size is more than remaining data in a partition then it is adjusted to the size of this data. Recommended size is between 10 and 1000. Default: 128
solver	The solver algorithm for optimization. Supported options: "gd" (minibatch gradient descent) or "l-bfgs". Default: "l-bfgs"
seed	A random seed. Set this value if you need your results to be reproducible across repeated calls.
initial_weights	
	The initial weights of the model.
thresholds	Thresholds in multi-class classification to adjust the probability of predicting each class. Array must have length equal to the number of classes, with values > 0 excepting that at most one value may be 0. The class with largest value `p/t` is predicted, where p is the original probability of that class and t is the class's threshold.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by [ft_r_formula](#).
label_col	Label column name. The column should be a numeric column. Usually this column is output by [ft_r_formula](#).

prediction_col	Prediction column name.
probability_col	
	Column name for predicted class conditional probabilities.
raw_prediction_col	
	Raw prediction (a.k.a. confidence) column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; see Details.
response	(Deprecated) The name of the response column (as a length-one character vector.)
features	(Deprecated) The name of features (terms) to use for the model fit.

Details

`ml_multilayer_perceptron()` is an alias for `ml_multilayer_perceptron_classifier()` for backwards compatibility.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

See Also

Other ml algorithms: `ml_aft_survival_regression()`, `ml_decision_tree_classifier()`, `ml_gbt_classifier()`, `ml_generalized_linear_regression()`, `ml_isotonic_regression()`, `ml_linear_regression()`, `ml_linear_svc()`, `ml_logistic_regression()`, `ml_naive_bayes()`, `ml_one_vs_rest()`, `ml_random_forest_classifi`

Examples

```
## Not run:
sc <- spark_connect(master = "local")

iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)
partitions <- iris_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

iris_training <- partitions$training
iris_test <- partitions$test

mlp_model <- iris_training %>%
  ml_multilayer_perceptron_classifier(Species ~ ., layers = c(4, 3, 3))

pred <- ml_predict(mlp_model, iris_test)

ml_multiclass_classification_evaluator(pred)

## End(Not run)
```

```
ml_multilayer_perceptron_tidiers
```
Tidying methods for Spark ML MLP

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_multilayer_perceptron_classification'
tidy(x, ...)

## S3 method for class 'ml_model_multilayer_perceptron_classification'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_multilayer_perceptron_classification'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.

```
ml_naive_bayes          Spark ML – Naive-Bayes
```

Description

Naive Bayes Classifiers. It supports Multinomial NB (see here) which can handle finitely supported discrete data. For example, by converting documents into TF-IDF vectors, it can be used for document classification. By making every vector a binary (0/1) data, it can also be used as Bernoulli NB (see here). The input feature values must be nonnegative.

Usage

```
ml_naive_bayes(
  x,
  formula = NULL,
  model_type = "multinomial",
  smoothing = 1,
  thresholds = NULL,
  weight_col = NULL,
  features_col = "features",
```

```
    label_col = "label",
    prediction_col = "prediction",
    probability_col = "probability",
    raw_prediction_col = "rawPrediction",
    uid = random_string("naive_bayes_"),
    ...
)
```

Arguments

x   A spark_connection, ml_pipeline, or a tbl_spark.

formula   Used when x is a tbl_spark. R formula as a character string or a formula.
                   This is used to transform the input dataframe before fitting, see ft_r_formula for
                   details.

model_type   The model type. Supported options: "multinomial" and "bernoulli". (de-
                   fault = multinomial)

smoothing   The (Laplace) smoothing parameter. Defaults to 1.

thresholds   Thresholds in multi-class classification to adjust the probability of predicting
                   each class. Array must have length equal to the number of classes, with values
                   > 0 excepting that at most one value may be 0. The class with largest value p/t
                   is predicted, where p is the original probability of that class and t is the class's
                   threshold.

weight_col   (Spark 2.1.0+) Weight column name. If this is not set or empty, we treat all
                   instance weights as 1.0.

features_col   Features column name, as a length-one character vector. The column should
                   be single vector column of numeric values. Usually this column is output by
                   ft_r_formula.

label_col   Label column name. The column should be a numeric column. Usually this
                   column is output by ft_r_formula.

prediction_col Prediction column name.

probability_col

                   Column name for predicted class conditional probabilities.

raw_prediction_col

                   Raw prediction (a.k.a. confidence) column name.

uid   A character string used to uniquely identify the ML estimator.

...   Optional arguments; see Details.

Value

The object returned depends on the class of x. If it is a spark_connection, the function returns a
ml_estimator object. If it is a ml_pipeline, it will return a pipeline with the predictor appended
to it. If a tbl_spark, it will return a tbl_spark with the predictions added to it.

See Also

Other ml algorithms: ml_aft_survival_regression(), ml_decision_tree_classifier(), ml_gbt_classifier(), ml_generalized_linear_regression(), ml_isotonic_regression(), ml_linear_regression(), ml_linear_svc(), ml_logistic_regression(), ml_multilayer_perceptron_classifier(), ml_one_vs_rest(), ml_random_forest_classifier()

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

partitions <- iris_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

iris_training <- partitions$training
iris_test <- partitions$test

nb_model <- iris_training %>%
  ml_naive_bayes(Species ~ .)

pred <- ml_predict(nb_model, iris_test)

ml_multiclass_classification_evaluator(pred)

## End(Not run)
```

---

ml_naive_bayes_tidiers

Tidying methods for Spark ML Naive Bayes

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_naive_bayes'
tidy(x, ...)

## S3 method for class 'ml_model_naive_bayes'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_naive_bayes'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.

---

ml_one_vs_rest	*Spark ML – OneVsRest*

Description

Reduction of Multiclass Classification to Binary Classification. Performs reduction using one against all strategy. For a multiclass classification with k classes, train k models (one per class). Each example is scored against all k models and the model with highest score is picked to label the example.

Usage

```
ml_one_vs_rest(
  x,
  formula = NULL,
  classifier = NULL,
  features_col = "features",
  label_col = "label",
  prediction_col = "prediction",
  uid = random_string("one_vs_rest_"),
  ...
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
formula	Used when x is a tbl_spark. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see ft_r_formula for details.
classifier	Object of class ml_estimator. Base binary classifier that we reduce multiclass classification into.
features_col	Features column name, as a length-one character vector. The column should be single vector column of numeric values. Usually this column is output by ft_r_formula.
label_col	Label column name. The column should be a numeric column. Usually this column is output by ft_r_formula.
prediction_col	Prediction column name.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; see Details.

Value

The object returned depends on the class of `x`. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

See Also

Other ml algorithms: `ml_aft_survival_regression()`, `ml_decision_tree_classifier()`, `ml_gbt_classifier()`, `ml_generalized_linear_regression()`, `ml_isotonic_regression()`, `ml_linear_regression()`, `ml_linear_svc()`, `ml_logistic_regression()`, `ml_multilayer_perceptron_classifier()`, `ml_naive_bayes()`, `ml_random_forest_classifier()`

---

ml_pca_tidiers   *Tidying methods for Spark ML Principal Component Analysis*

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_pca'
tidy(x, ...)

## S3 method for class 'ml_model_pca'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_pca'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.

---

ml_pipeline *Spark ML – Pipelines*

Description

Create Spark ML Pipelines

Usage

```
ml_pipeline(x, ..., uid = random_string("pipeline_"))
```

Arguments

x	Either a `spark_connection` or `ml_pipeline_stage` objects
...	`ml_pipeline_stage` objects.
uid	A character string used to uniquely identify the ML estimator.

Value

When x is a `spark_connection`, `ml_pipeline()` returns an empty pipeline object. When x is a `ml_pipeline_stage`, `ml_pipeline()` returns an `ml_pipeline` with the stages set to x and any transformers or estimators given in `...`.

---

ml_power_iteration *Spark ML – Power Iteration Clustering*

Description

Power iteration clustering (PIC) is a scalable and efficient algorithm for clustering vertices of a graph given pairwise similarities as edge properties, described in the paper "Power Iteration Clustering" by Frank Lin and William W. Cohen. It computes a pseudo-eigenvector of the normalized affinity matrix of the graph via power iteration and uses it to cluster vertices. spark.mllib includes an implementation of PIC using GraphX as its backend. It takes an RDD of (srcId, dstId, similarity) tuples and outputs a model with the clustering assignments. The similarities must be nonnegative. PIC assumes that the similarity measure is symmetric. A pair (srcId, dstId) regardless of the ordering should appear at most once in the input data. If a pair is missing from input, their similarity is treated as zero.

Usage

```
ml_power_iteration(
  x,
  k = 4,
  max_iter = 20,
  init_mode = "random",
  src_col = "src",
  dst_col = "dst",
  weight_col = "weight",
  ...
)
```

Arguments

x	A 'spark_connection' or a 'tbl_spark'.
k	The number of clusters to create.
max_iter	The maximum number of iterations to run.
init_mode	This can be either "random", which is the default, to use a random vector as vertex properties, or "degree" to use normalized sum similarities.
src_col	Column in the input Spark dataframe containing 0-based indexes of all source vertices in the affinity matrix described in the PIC paper.
dst_col	Column in the input Spark dataframe containing 0-based indexes of all destination vertices in the affinity matrix described in the PIC paper.
weight_col	Column in the input Spark dataframe containing non-negative edge weights in the affinity matrix described in the PIC paper.
...	Optional arguments. Currently unused.

Value

A 2-column R dataframe with columns named "id" and "cluster" describing the resulting cluster assignments

Examples

```
## Not run:

library(sparklyr)

sc <- spark_connect(master = "local")

r1 <- 1
n1 <- 80L
r2 <- 4
n2 <- 80L

gen_circle <- function(radius, num_pts) {
  # generate evenly distributed points on a circle centered at the origin
  seq(0, num_pts - 1) %>%
```

```r
    lapply(
      function(pt) {
        theta <- 2 * pi * pt / num_pts

        radius * c(cos(theta), sin(theta))
      }
    )
}

guassian_similarity <- function(pt1, pt2) {
  dist2 <- sum((pt2 - pt1)^2)

  exp(-dist2 / 2)
}

gen_pic_data <- function() {
  # generate points on 2 concentric circle centered at the origin and then
  # compute pairwise Gaussian similarity values of all unordered pair of
  # points
  n <- n1 + n2
  pts <- append(gen_circle(r1, n1), gen_circle(r2, n2))
  num_unordered_pairs <- n * (n - 1) / 2

  src <- rep(0L, num_unordered_pairs)
  dst <- rep(0L, num_unordered_pairs)
  sim <- rep(0, num_unordered_pairs)

  idx <- 1
  for (i in seq(2, n)) {
    for (j in seq(i - 1)) {
      src[[idx]] <- i - 1L
      dst[[idx]] <- j - 1L
      sim[[idx]] <- guassian_similarity(pts[[i]], pts[[j]])
      idx <- idx + 1
    }
  }

  dplyr::tibble(src = src, dst = dst, sim = sim)
}

pic_data <- copy_to(sc, gen_pic_data())

clusters <- ml_power_iteration(
  pic_data,
  src_col = "src", dst_col = "dst", weight_col = "sim", k = 2, max_iter = 40
)
print(clusters)

## End(Not run)
```

---

ml_prefixspan   *Frequent Pattern Mining – PrefixSpan*

Description

PrefixSpan algorithm for mining frequent itemsets.

Usage

```
ml_prefixspan(
  x,
  seq_col = "sequence",
  min_support = 0.1,
  max_pattern_length = 10,
  max_local_proj_db_size = 3.2e+07,
  uid = random_string("prefixspan_"),
  ...
)

ml_freq_seq_patterns(model)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
seq_col	The name of the sequence column in dataset (defaults to "sequence"). Rows with nulls in this column are ignored.
min_support	The minimum support required to be considered a frequent sequential pattern.
max_pattern_length	
	The maximum length of a frequent sequential pattern. Any frequent pattern exceeding this length will not be included in the results.
max_local_proj_db_size	
	The maximum number of items allowed in a prefix-projected database before local iterative processing of the projected database begins. This parameter should be tuned with respect to the size of your executors.
uid	A character string used to uniquely identify the ML estimator.
...	Optional arguments; currently unused.
model	A Prefix Span model.

Examples

```
## Not run:
library(sparklyr)
sc <- spark_connect(master = "local", version = "2.4.0")

items_df <- dplyr::tibble(
  seq = list(
```

```
      list(list(1, 2), list(3)),
      list(list(1), list(3, 2), list(1, 2)),
      list(list(1, 2), list(5)),
      list(list(6))
  )
)
items_sdf <- copy_to(sc, items_df, overwrite = TRUE)

prefix_span_model <- ml_prefixspan(
  sc,
  seq_col = "seq",
  min_support = 0.5,
  max_pattern_length = 5,
  max_local_proj_db_size = 32000000
)

frequent_items <- prefix_span_model$frequent_sequential_patterns(items_sdf) %>% collect()

## End(Not run)
```

```
ml_random_forest_classifier
```
Spark ML – Random Forest

Description

Perform classification and regression using random forests.

Usage

```
ml_random_forest_classifier(
  x,
  formula = NULL,
  num_trees = 20,
  subsampling_rate = 1,
  max_depth = 5,
  min_instances_per_node = 1,
  feature_subset_strategy = "auto",
  impurity = "gini",
  min_info_gain = 0,
  max_bins = 32,
  seed = NULL,
  thresholds = NULL,
  checkpoint_interval = 10,
  cache_node_ids = FALSE,
  max_memory_in_mb = 256,
  features_col = "features",
```

```
    label_col = "label",
    prediction_col = "prediction",
    probability_col = "probability",
    raw_prediction_col = "rawPrediction",
    uid = random_string("random_forest_classifier_"),
    ...
  )

  ml_random_forest(
    x,
    formula = NULL,
    type = c("auto", "regression", "classification"),
    features_col = "features",
    label_col = "label",
    prediction_col = "prediction",
    probability_col = "probability",
    raw_prediction_col = "rawPrediction",
    feature_subset_strategy = "auto",
    impurity = "auto",
    checkpoint_interval = 10,
    max_bins = 32,
    max_depth = 5,
    num_trees = 20,
    min_info_gain = 0,
    min_instances_per_node = 1,
    subsampling_rate = 1,
    seed = NULL,
    thresholds = NULL,
    cache_node_ids = FALSE,
    max_memory_in_mb = 256,
    uid = random_string("random_forest_"),
    response = NULL,
    features = NULL,
    ...
  )

  ml_random_forest_regressor(
    x,
    formula = NULL,
    num_trees = 20,
    subsampling_rate = 1,
    max_depth = 5,
    min_instances_per_node = 1,
    feature_subset_strategy = "auto",
    impurity = "variance",
    min_info_gain = 0,
    max_bins = 32,
    seed = NULL,
```

```
    checkpoint_interval = 10,
    cache_node_ids = FALSE,
    max_memory_in_mb = 256,
    features_col = "features",
    label_col = "label",
    prediction_col = "prediction",
    uid = random_string("random_forest_regressor_"),
    ...
)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
formula	Used when x is a `tbl_spark`. R formula as a character string or a formula. This is used to transform the input dataframe before fitting, see [ft_r_formula](#) for details.
num_trees	Number of trees to train (>= 1). If 1, then no bootstrapping is used. If > 1, then bootstrapping is done.
subsampling_rate	Fraction of the training data used for learning each decision tree, in range (0, 1]. (default = 1.0)
max_depth	Maximum depth of the tree (>= 0); that is, the maximum number of nodes separating any leaves from the root of the tree.
min_instances_per_node	Minimum number of instances each child must have after split.
feature_subset_strategy	The number of features to consider for splits at each tree node. See details for options.
impurity	Criterion used for information gain calculation. Supported: "entropy" and "gini" (default) for classification and "variance" (default) for regression. For `ml_decision_tree`, setting `"auto"` will default to the appropriate criterion based on model type.
min_info_gain	Minimum information gain for a split to be considered at a tree node. Should be >= 0, defaults to 0.
max_bins	The maximum number of bins used for discretizing continuous features and for choosing how to split on features at each node. More bins give higher granularity.
seed	Seed for random numbers.
thresholds	Thresholds in multi-class classification to adjust the probability of predicting each class. Array must have length equal to the number of classes, with values > 0 excepting that at most one value may be 0. The class with largest value p/t is predicted, where p is the original probability of that class and t is the class's threshold.
checkpoint_interval	Set checkpoint interval (>= 1) or disable checkpoint (-1). E.g. 10 means that the cache will get checkpointed every 10 iterations, defaults to 10.

cache_node_ids   If FALSE, the algorithm will pass trees to executors to match instances with
                 nodes. If TRUE, the algorithm will cache node IDs for each instance. Caching
                 can speed up training of deeper trees. Defaults to FALSE.

max_memory_in_mb

                 Maximum memory in MB allocated to histogram aggregation. If too small,
                 then 1 node will be split per iteration, and its aggregates may exceed this size.
                 Defaults to 256.

features_col   Features column name, as a length-one character vector. The column should
                 be single vector column of numeric values. Usually this column is output by
                 ft_r_formula.

label_col   Label column name. The column should be a numeric column. Usually this
                 column is output by ft_r_formula.

prediction_col   Prediction column name.

probability_col

                 Column name for predicted class conditional probabilities.

raw_prediction_col

                 Raw prediction (a.k.a. confidence) column name.

uid   A character string used to uniquely identify the ML estimator.

...   Optional arguments; see Details.

type   The type of model to fit. "regression" treats the response as a continuous
                 variable, while "classification" treats the response as a categorical variable.
                 When "auto" is used, the model type is inferred based on the response variable
                 type – if it is a numeric type, then regression is used; classification otherwise.

response   (Deprecated) The name of the response column (as a length-one character vec-
                 tor.)

features   (Deprecated) The name of features (terms) to use for the model fit.

Details

The supported options for feature_subset_strategy are

- "auto": Choose automatically for task: If num_trees == 1, set to "all". If num_trees > 1
  (forest), set to "sqrt" for classification and to "onethird" for regression.

- "all": use all features

- "onethird": use 1/3 of the features

- "sqrt": use use sqrt(number of features)

- "log2": use log2(number of features)

- "n": when n is in the range (0, 1.0], use n * number of features. When n is in the range (1,
  number of features), use n features. (default = "auto")

ml_random_forest is a wrapper around ml_random_forest_regressor.tbl_spark and ml_random_forest_classifier
and calls the appropriate method based on model type.

Value

The object returned depends on the class of x. If it is a `spark_connection`, the function returns a `ml_estimator` object. If it is a `ml_pipeline`, it will return a pipeline with the predictor appended to it. If a `tbl_spark`, it will return a `tbl_spark` with the predictions added to it.

See Also

Other ml algorithms: `ml_aft_survival_regression()`, `ml_decision_tree_classifier()`, `ml_gbt_classifier()`, `ml_generalized_linear_regression()`, `ml_isotonic_regression()`, `ml_linear_regression()`, `ml_linear_svc()`, `ml_logistic_regression()`, `ml_multilayer_perceptron_classifier()`, `ml_naive_bayes()`, `ml_one_vs_rest()`

Examples

```
## Not run:
sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

partitions <- iris_tbl %>%
  sdf_random_split(training = 0.7, test = 0.3, seed = 1111)

iris_training <- partitions$training
iris_test <- partitions$test

rf_model <- iris_training %>%
  ml_random_forest(Species ~ ., type = "classification")

pred <- ml_predict(rf_model, iris_test)

ml_multiclass_classification_evaluator(pred)

## End(Not run)
```

---

ml_stage   *Spark ML – Pipeline stage extraction*

Description

Extraction of stages from a Pipeline or PipelineModel object.

Usage

```
ml_stage(x, stage)

ml_stages(x, stages = NULL)
```

Arguments

x	A `ml_pipeline` or a `ml_pipeline_model` object
stage	The UID of a stage in the pipeline.
stages	The UIDs of stages in the pipeline as a character vector.

Value

For `ml_stage()`: The stage specified.

For `ml_stages()`: A list of stages. If `stages` is not set, the function returns all stages of the pipeline in a list.

---

ml_summary   *Spark ML – Extraction of summary metrics*

Description

Extracts a metric from the summary object of a Spark ML model.

Usage

```
ml_summary(x, metric = NULL, allow_null = FALSE)
```

Arguments

x	A Spark ML model that has a summary.
metric	The name of the metric to extract. If not set, returns the summary object.
allow_null	Whether null results are allowed when the metric is not found in the summary.

---

ml_survival_regression_tidiers
                        Tidying methods for Spark ML Survival Regression

---

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_aft_survival_regression'
tidy(x, ...)

## S3 method for class 'ml_model_aft_survival_regression'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_aft_survival_regression'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.

---

ml_tree_tidiers   *Tidying methods for Spark ML tree models*

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_decision_tree_classification'
tidy(x, ...)

## S3 method for class 'ml_model_decision_tree_regression'
tidy(x, ...)

## S3 method for class 'ml_model_decision_tree_classification'
augment(x, newdata = NULL, ...)

## S3 method for class '`_ml_model_decision_tree_classification`'
augment(x, new_data = NULL, ...)

## S3 method for class 'ml_model_decision_tree_regression'
augment(x, newdata = NULL, ...)

## S3 method for class '`_ml_model_decision_tree_regression`'
augment(x, new_data = NULL, ...)

## S3 method for class 'ml_model_decision_tree_classification'
glance(x, ...)

## S3 method for class 'ml_model_decision_tree_regression'
glance(x, ...)

## S3 method for class 'ml_model_random_forest_classification'
tidy(x, ...)

## S3 method for class 'ml_model_random_forest_regression'
tidy(x, ...)

## S3 method for class 'ml_model_random_forest_classification'
augment(x, newdata = NULL, ...)
```

```
## S3 method for class '`_ml_model_random_forest_classification`'
augment(x, new_data = NULL, ...)

## S3 method for class 'ml_model_random_forest_regression'
augment(x, newdata = NULL, ...)

## S3 method for class '`_ml_model_random_forest_regression`'
augment(x, new_data = NULL, ...)

## S3 method for class 'ml_model_random_forest_classification'
glance(x, ...)

## S3 method for class 'ml_model_random_forest_regression'
glance(x, ...)

## S3 method for class 'ml_model_gbt_classification'
tidy(x, ...)

## S3 method for class 'ml_model_gbt_regression'
tidy(x, ...)

## S3 method for class 'ml_model_gbt_classification'
augment(x, newdata = NULL, ...)

## S3 method for class '`_ml_model_gbt_classification`'
augment(x, new_data = NULL, ...)

## S3 method for class 'ml_model_gbt_regression'
augment(x, newdata = NULL, ...)

## S3 method for class '`_ml_model_gbt_regression`'
augment(x, new_data = NULL, ...)

## S3 method for class 'ml_model_gbt_classification'
glance(x, ...)

## S3 method for class 'ml_model_gbt_regression'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.
new_data	a tbl_spark of new data to use for prediction.

---

ml_uid   *Spark ML – UID*

Description

Extracts the UID of an ML object.

Usage

```
ml_uid(x)
```

Arguments

x   A Spark ML object

---

ml_unsupervised_tidiers
                        Tidying methods for Spark ML unsupervised models

---

Description

These methods summarize the results of Spark ML models into tidy forms.

Usage

```
## S3 method for class 'ml_model_kmeans'
tidy(x, ...)

## S3 method for class 'ml_model_kmeans'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_kmeans'
glance(x, ...)

## S3 method for class 'ml_model_bisecting_kmeans'
tidy(x, ...)

## S3 method for class 'ml_model_bisecting_kmeans'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_bisecting_kmeans'
glance(x, ...)

## S3 method for class 'ml_model_gaussian_mixture'
tidy(x, ...)
```

```
## S3 method for class 'ml_model_gaussian_mixture'
augment(x, newdata = NULL, ...)

## S3 method for class 'ml_model_gaussian_mixture'
glance(x, ...)
```

Arguments

x	a Spark ML model.
...	extra arguments (not used.)
newdata	a tbl_spark of new data to use for prediction.

---

mutate   *Mutate*

Description

See [mutate](#) for more details.

---

na.replace   *Replace Missing Values in Objects*

Description

This S3 generic provides an interface for replacing [NA](#) values within an object.

Usage

```
na.replace(object, ...)
```

Arguments

object	An R object.
...	Arguments passed along to implementing methods.

---

nest   *Nest*

Description

See [nest](#) for more details.

---

pivot_longer   *Pivot longer*

Description

See pivot_longer for more details.

---

pivot_wider   *Pivot wider*

Description

See pivot_wider for more details.

---

random_string   *Random string generation*

Description

Generate a random string with a given prefix.

Usage

```
random_string(prefix = "table")
```

Arguments

prefix   A length-one character vector.

---

reactiveSpark   *Reactive spark reader*

Description

Given a spark object, returns a reactive data source for the contents of the spark object. This function is most useful to read Spark streams.

Usage

```
reactiveSpark(x, intervalMillis = 1000, session = NULL)
```

Arguments

x	An object coercable to a Spark DataFrame.
intervalMillis	Approximate number of milliseconds to wait to retrieve updated data frame. This can be a numeric value, or a function that returns a numeric value.
session	The user session to associate this file reader with, or NULL if none. If non-null, the reader will automatically stop when the session ends.

---

registerDoSpark   *Register a Parallel Backend*

Description

Registers a parallel backend using the foreach package.

Usage

```
registerDoSpark(spark_conn, parallelism = NULL, ...)
```

Arguments

spark_conn	Spark connection to use
parallelism	Level of parallelism to use for task execution (if unspecified, then it will take the value of 'SparkContext.defaultParallelism()' which by default is the number of cores available to the 'sparklyr' application)
...	additional options for sparklyr parallel backend (currently only the only valid option is 'nocompile')

Value

None

Examples

```
## Not run:

sc <- spark_connect(master = "local")
registerDoSpark(sc, nocompile = FALSE)

## End(Not run)
```

register_extension	*Register a Package that Implements a Spark Extension*

Description

Registering an extension package will result in the package being automatically scanned for spark dependencies when a connection to Spark is created.

Usage

```
register_extension(package)

registered_extensions()
```

Arguments

package   The package(s) to register.

Note

Packages should typically register their extensions in their .onLoad hook – this ensures that their extensions are registered when their namespaces are loaded.

replace_na	*Replace NA*

Description

See replace_na for more details.

right_join	*Right join*

Description

See right_join for more details.

---

sdf-saveload   *Save / Load a Spark DataFrame*

Description

Routines for saving and loading Spark DataFrames.

Usage

```
sdf_save_table(x, name, overwrite = FALSE, append = FALSE)

sdf_load_table(sc, name)

sdf_save_parquet(x, path, overwrite = FALSE, append = FALSE)

sdf_load_parquet(sc, path)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
name	The table name to assign to the saved Spark DataFrame.
overwrite	Boolean; overwrite a pre-existing table of the same name?
append	Boolean; append to a pre-existing table of the same name?
sc	A spark_connection object.
path	The path where the Spark DataFrame should be saved.

---

sdf-transform-methods   *Spark ML – Transform, fit, and predict methods (sdf_ interface)*

Description

Deprecated methods for transformation, fit, and prediction. These are mirrors of the corresponding
[ml-transform-methods](#).

Usage

```
sdf_predict(x, model, ...)

sdf_transform(x, transformer, ...)

sdf_fit(x, estimator, ...)

sdf_fit_and_transform(x, estimator, ...)
```

Arguments

x	A `tbl_spark`.
model	A `ml_transformer` or a `ml_model` object.
...	Optional arguments passed to the corresponding `ml_` methods.
transformer	A `ml_transformer` object.
estimator	A `ml_estimator` object.

Value

`sdf_predict()`, `sdf_transform()`, and `sdf_fit_and_transform()` return a transformed dataframe whereas `sdf_fit()` returns a `ml_transformer`.

---

sdf_along	*Create DataFrame for along Object*

Description

Creates a DataFrame along the given object.

Usage

```
sdf_along(sc, along, repartition = NULL, type = c("integer", "integer64"))
```

Arguments

sc	The associated Spark connection.
along	Takes the length from the length of this argument.
repartition	The number of partitions to use when distributing the data across the Spark cluster.
type	The data type to use for the index, either `"integer"` or `"integer64"`.

---

sdf_bind	*Bind multiple Spark DataFrames by row and column*

Description

`sdf_bind_rows()` and `sdf_bind_cols()` are implementation of the common pattern of `do.call(rbind, sdfs)` or `do.call(cbind, sdfs)` for binding many Spark DataFrames into one.

Usage

```
sdf_bind_rows(..., id = NULL)

sdf_bind_cols(...)
```

Arguments

`...`	Spark tbls to combine.
	Each argument can either be a Spark DataFrame or a list of Spark DataFrames
	When row-binding, columns are matched by name, and any missing columns with be filled with NA.
	When column-binding, rows are matched by position, so all data frames must have the same number of rows.
`id`	Data frame identifier.
	When `id` is supplied, a new column of identifiers is created to link each row to its original Spark DataFrame. The labels are taken from the named arguments to `sdf_bind_rows()`. When a list of Spark DataFrames is supplied, the labels are taken from the names of the list. If no names are found a numeric sequence is used instead.

Details

The output of `sdf_bind_rows()` will contain a column if that column appears in any of the inputs.

Value

`sdf_bind_rows()` and `sdf_bind_cols()` return `tbl_spark`

`sdf_broadcast`	*Broadcast hint*

Description

Used to force broadcast hash joins.

Usage

```
sdf_broadcast(x)
```

Arguments

`x`	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.

---

sdf_checkpoint *Checkpoint a Spark DataFrame*

Description

Checkpoint a Spark DataFrame

Usage

```
sdf_checkpoint(x, eager = TRUE)
```

Arguments

x   an object coercible to a Spark DataFrame

eager   whether to truncate the lineage of the DataFrame

---

sdf_coalesce *Coalesces a Spark DataFrame*

Description

Coalesces a Spark DataFrame

Usage

```
sdf_coalesce(x, partitions)
```

Arguments

x   A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.

partitions   number of partitions

---

sdf_collect   *Collect a Spark DataFrame into R.*

Description

Collects a Spark dataframe into R.

Usage

```
sdf_collect(object, impl = c("row-wise", "row-wise-iter", "column-wise"), ...)
```

Arguments

object	Spark dataframe to collect
impl	Which implementation to use while collecting Spark dataframe - row-wise: fetch the entire dataframe into memory and then process it row-by-row - row-wise-iter: iterate through the dataframe using RDD local iterator, processing one row at a time (hence reducing memory footprint) - column-wise: fetch the entire dataframe into memory and then process it column-by-column NOTE: (1) this will not apply to streaming or arrow use cases (2) this parameter will only affect implementation detail, and will not affect result of 'sdf_collect', and should only be set if performance profiling indicates any particular choice will be significantly better than the default choice ("row-wise")
...	Additional options.

---

sdf_copy_to   *Copy an Object into Spark*

Description

Copy an object into Spark, and return an R object wrapping the copied object (typically, a Spark DataFrame).

Usage

```
sdf_copy_to(sc, x, name, memory, repartition, overwrite, struct_columns, ...)

sdf_import(x, sc, name, memory, repartition, overwrite, struct_columns, ...)
```

Arguments

sc	The associated Spark connection.
x	An R object from which a Spark DataFrame can be generated.
name	The name to assign to the copied table in Spark.
memory	Boolean; should the table be cached into memory?
repartition	The number of partitions to use when distributing the table across the Spark cluster. The default (0) can be used to avoid partitioning.
overwrite	Boolean; overwrite a pre-existing table with the name name if one already exists?
struct_columns	(only supported with Spark 2.4.0 or higher) A list of columns from the source data frame that should be converted to Spark SQL StructType columns. The source columns can contain either json strings or nested lists. All rows within each source column should have identical schemas (because otherwise the conversion result will contain unexpected null values or missing values as Spark currently does not support schema discovery on individual rows within a struct column).
...	Optional arguments, passed to implementing methods.

Advanced Usage

sdf_copy_to is an S3 generic that, by default, dispatches to sdf_import. Package authors that would like to implement sdf_copy_to for a custom object type can accomplish this by implementing the associated method on sdf_import.

See Also

Other Spark data frames: [sdf_distinct](), [sdf_random_split](), [sdf_register](), [sdf_sample](), [sdf_sort](), [sdf_weighted_sample]()

Examples

```
## Not run:
sc <- spark_connect(master = "spark://HOST:PORT")
sdf_copy_to(sc, iris)

## End(Not run)
```

---

sdf_crosstab   *Cross Tabulation*

Description

Builds a contingency table at each combination of factor levels.

Usage

```
sdf_crosstab(x, col1, col2)
```

Arguments

x	A Spark DataFrame
col1	The name of the first column. Distinct items will make the first item of each row.
col2	The name of the second column. Distinct items will make the column names of the DataFrame.

Value

A DataFrame containing the contingency table.

---

sdf_debug_string   *Debug Info for Spark DataFrame*

Description

Prints plan of execution to generate x. This plan will, among other things, show the number of partitions in parenthesis at the far left and indicate stages using indentation.

Usage

```
sdf_debug_string(x, print = TRUE)
```

Arguments

x	An R object wrapping, or containing, a Spark DataFrame.
print	Print debug information?

---

sdf_describe   *Compute summary statistics for columns of a data frame*

Description

Compute summary statistics for columns of a data frame

Usage

```
sdf_describe(x, cols = colnames(x))
```

Arguments

x	An object coercible to a Spark DataFrame
cols	Columns to compute statistics for, given as a character vector

---

sdf_dim *Support for Dimension Operations*

Description

sdf_dim(), sdf_nrow() and sdf_ncol() provide similar functionality to dim(), nrow() and ncol().

Usage

```
sdf_dim(x)

sdf_nrow(x)

sdf_ncol(x)
```

Arguments

x   An object (usually a spark_tbl).

---

sdf_distinct *Invoke distinct on a Spark DataFrame*

Description

Invoke distinct on a Spark DataFrame

Usage

```
sdf_distinct(x, ..., name)
```

Arguments

x	A Spark DataFrame.
...	Optional variables to use when determining uniqueness. If there are multiple rows for a given combination of inputs, only the first row will be preserved. If omitted, will use all variables.
name	A name to assign this table. Passed to [sdf_register()].

See Also

Other Spark data frames: [sdf_copy_to()](), [sdf_random_split()](), [sdf_register()](), [sdf_sample()](), [sdf_sort()](), [sdf_weighted_sample()]()

---

sdf_drop_duplicates   *Remove duplicates from a Spark DataFrame*

Description

Remove duplicates from a Spark DataFrame

Usage

```
sdf_drop_duplicates(x, cols = NULL)
```

Arguments

x	An object coercible to a Spark DataFrame
cols	Subset of Columns to consider, given as a character vector

---

sdf_expand_grid   *Create a Spark dataframe containing all combinations of inputs*

Description

Given one or more R vectors/factors or single-column Spark dataframes, perform an expand.grid operation on all of them and store the result in a Spark dataframe

Usage

```
sdf_expand_grid(
  sc,
  ...,
  broadcast_vars = NULL,
  memory = TRUE,
  repartition = NULL,
  partition_by = NULL
)
```

Arguments

sc	The associated Spark connection.
...	Each input variable can be either a R vector/factor or a Spark dataframe. Unnamed inputs will assume the default names of 'Var1', 'Var2', etc in the result, similar to what 'expand.grid' does for unnamed inputs.
broadcast_vars	Indicates which input(s) should be broadcasted to all nodes of the Spark cluster during the join process (default: none).
memory	Boolean; whether the resulting Spark dataframe should be cached into memory (default: TRUE)

repartition	Number of partitions the resulting Spark dataframe should have
partition_by	Vector of column names used for partitioning the resulting Spark dataframe, only supported for Spark 2.0+

Examples

```
## Not run:
sc <- spark_connect(master = "local")
grid_sdf <- sdf_expand_grid(sc, seq(5), rnorm(10), letters)

## End(Not run)
```

---

sdf_from_avro   *Convert column(s) from avro format*

Description

Convert column(s) from avro format

Usage

```
sdf_from_avro(x, cols)
```

Arguments

x	An object coercible to a Spark DataFrame
cols	Named list of columns to transform from Avro format plus a valid Avro schema string for each column, where column names are keys and column schema strings are values (e.g., c(example_primitive_col = "string", example_complex_col = "{\"type\":\"record\",\"name\":\"person\",\"fields\":[{\"name\":\"person_name\",\" {\"name\":\"person_id\",\"type\":\"long\"}]}")

---

sdf_is_streaming   *Spark DataFrame is Streaming*

Description

Is the given Spark DataFrame a streaming data?

Usage

```
sdf_is_streaming(x)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.

---

sdf_last_index   *Returns the last index of a Spark DataFrame*

Description

Returns the last index of a Spark DataFrame. The Spark `mapPartitionsWithIndex` function is used to iterate through the last nonempty partition of the RDD to find the last record.

Usage

```
sdf_last_index(x, id = "id")
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
id	The name of the index column.

---

sdf_len   *Create DataFrame for Length*

Description

Creates a DataFrame for the given length.

Usage

```
sdf_len(sc, length, repartition = NULL, type = c("integer", "integer64"))
```

Arguments

sc	The associated Spark connection.
length	The desired length of the sequence.
repartition	The number of partitions to use when distributing the data across the Spark cluster.
type	The data type to use for the index, either `"integer"` or `"integer64"`.

---

sdf_num_partitions *Gets number of partitions of a Spark DataFrame*

Description

Gets number of partitions of a Spark DataFrame

Usage

```
sdf_num_partitions(x)
```

Arguments

x   A spark_connection, ml_pipeline, or a tbl_spark.

---

sdf_partition_sizes *Compute the number of records within each partition of a Spark DataFrame*

Description

Compute the number of records within each partition of a Spark DataFrame

Usage

```
sdf_partition_sizes(x)
```

Arguments

x   A spark_connection, ml_pipeline, or a tbl_spark.

Examples

```
## Not run:
library(sparklyr)
sc <- spark_connect(master = "spark://HOST:PORT")
example_sdf <- sdf_len(sc, 100L, repartition = 10L)
example_sdf %>%
  sdf_partition_sizes() %>%
  print()

## End(Not run)
```

---

sdf_persist   *Persist a Spark DataFrame*

Description

Persist a Spark DataFrame, forcing any pending computations and (optionally) serializing the results to disk.

Usage

```
sdf_persist(x, storage.level = "MEMORY_AND_DISK", name = NULL)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
storage.level	The storage level to be used. Please view the Spark Documentation for information on what storage levels are accepted.
name	A name to assign this table. Passed to [sdf_register()].

Details

Spark DataFrames invoke their operations lazily – pending operations are deferred until their results are actually needed. Persisting a Spark DataFrame effectively 'forces' any pending computations, and then persists the generated Spark DataFrame as requested (to memory, to disk, or otherwise).

Users of Spark should be careful to persist the results of any computations which are non-deterministic – otherwise, one might see that the values within a column seem to 'change' as new operations are performed on that data set.

---

sdf_pivot   *Pivot a Spark DataFrame*

Description

Construct a pivot table over a Spark Dataframe, using a syntax similar to that from `reshape2::dcast`.

Usage

```
sdf_pivot(x, formula, fun.aggregate = "count")
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
formula	A two-sided R formula of the form x_1 + x_2 + ... ~ y_1. The left-hand side of the formula indicates which variables are used for grouping, and the right-hand side indicates which variable is used for pivoting. Currently, only a single pivot column is supported.
fun.aggregate	How should the grouped dataset be aggregated? Can be a length-one character vector, giving the name of a Spark aggregation function to be called; a named R list mapping column names to an aggregation method, or an R function that is invoked on the grouped dataset.

Examples

```
## Not run:
library(sparklyr)
library(dplyr)

sc <- spark_connect(master = "local")
iris_tbl <- sdf_copy_to(sc, iris, name = "iris_tbl", overwrite = TRUE)

# aggregating by mean
iris_tbl %>%
  mutate(Petal_Width = ifelse(Petal_Width > 1.5, "High", "Low")) %>%
  sdf_pivot(Petal_Width ~ Species,
    fun.aggregate = list(Petal_Length = "mean")
  )

# aggregating all observations in a list
iris_tbl %>%
  mutate(Petal_Width = ifelse(Petal_Width > 1.5, "High", "Low")) %>%
  sdf_pivot(Petal_Width ~ Species,
    fun.aggregate = list(Petal_Length = "collect_list")
  )

## End(Not run)
```

---

sdf_project	*Project features onto principal components*

Description

Project features onto principal components

Usage

```
sdf_project(
  object,
  newdata,
  features = dimnames(object$pc)[[1]],
  feature_prefix = NULL,
  ...
)
```

Arguments

object	A Spark PCA model object
newdata	An object coercible to a Spark DataFrame
features	A vector of names of columns to be projected
feature_prefix	The prefix used in naming the output features
...	Optional arguments; currently unused.

---

sdf_quantile   *Compute (Approximate) Quantiles with a Spark DataFrame*

Description

Given a numeric column within a Spark DataFrame, compute approximate quantiles.

Usage

```
sdf_quantile(
  x,
  column,
  probabilities = c(0, 0.25, 0.5, 0.75, 1),
  relative.error = 1e-05,
  weight.column = NULL
)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
column	The column(s) for which quantiles should be computed. Multiple columns are only supported in Spark 2.0+.
probabilities	A numeric vector of probabilities, for which quantiles should be computed.
relative.error	The maximal possible difference between the actual percentile of a result and its expected percentile (e.g., if 'relative.error' is 0.01 and 'probabilities' is 0.95, then any value between the 94th and 96th percentile will be considered an acceptable approximation).

weight.column	If not NULL, then a generalized version of the Greenwald- Khanna algorithm will be run to compute weighted percentiles, with each sample from 'column' having a relative weight specified by the corresponding value in 'weight.column'. The weights can be considered as relative frequencies of sample data points.

---

sdf_random_split   *Partition a Spark Dataframe*

Description

Partition a Spark DataFrame into multiple groups. This routine is useful for splitting a DataFrame into, for example, training and test datasets.

Usage

```
sdf_random_split(
  x,
  ...,
  weights = NULL,
  seed = sample(.Machine$integer.max, 1)
)

sdf_partition(x, ..., weights = NULL, seed = sample(.Machine$integer.max, 1))
```

Arguments

x	An object coercable to a Spark DataFrame.
...	Named parameters, mapping table names to weights. The weights will be normalized such that they sum to 1.
weights	An alternate mechanism for supplying weights – when specified, this takes precedence over the . . . arguments.
seed	Random seed to use for randomly partitioning the dataset. Set this if you want your partitioning to be reproducible on repeated runs.

Details

The sampling weights define the probability that a particular observation will be assigned to a particular partition, not the resulting size of the partition. This implies that partitioning a DataFrame with, for example,

```
sdf_random_split(x, training = 0.5, test = 0.5)
```

is not guaranteed to produce training and test partitions of equal size.

Value

An R list of tbl_sparks.

See Also

Other Spark data frames: `sdf_copy_to()`, `sdf_distinct()`, `sdf_register()`, `sdf_sample()`,
`sdf_sort()`, `sdf_weighted_sample()`

Examples

```
## Not run:
# randomly partition data into a 'training' and 'test'
# dataset, with 60% of the observations assigned to the
# 'training' dataset, and 40% assigned to the 'test' dataset
data(diamonds, package = "ggplot2")
diamonds_tbl <- copy_to(sc, diamonds, "diamonds")
partitions <- diamonds_tbl %>%
  sdf_random_split(training = 0.6, test = 0.4)
print(partitions)

# alternate way of specifying weights
weights <- c(training = 0.6, test = 0.4)
diamonds_tbl %>% sdf_random_split(weights = weights)

## End(Not run)
```

---

sdf_rbeta   *Generate random samples from a Beta distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from
a Betal distribution.

Usage

```
sdf_rbeta(
  sc,
  n,
  shape1,
  shape2,
  num_partitions = NULL,
  seed = NULL,
  output_col = "x"
)
```

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
shape1	Non-negative parameter (alpha) of the Beta distribution.

shape2	Non-negative parameter (beta) of the Beta distribution.
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: [sdf_rbinom()](), [sdf_rcauchy()](), [sdf_rchisq()](), [sdf_rexp()](),
[sdf_rgamma()](), [sdf_rgeom()](), [sdf_rhyper()](), [sdf_rlnorm()](), [sdf_rnorm()](), [sdf_rpois()](), [sdf_rt()](),
[sdf_runif()](), [sdf_rweibull()]()

---

sdf_rbinom	*Generate random samples from a binomial distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a binomial distribution.

Usage

```
sdf_rbinom(
  sc,
  n,
  size,
  prob,
  num_partitions = NULL,
  seed = NULL,
  output_col = "x"
)
```

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
size	Number of trials (zero or more).
prob	Probability of success on each trial.
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rcauchy()`, `sdf_rchisq()`, `sdf_rexp()`, `sdf_rgamma()`, `sdf_rgeom()`, `sdf_rhyper()`, `sdf_rlnorm()`, `sdf_rnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_runif()`, `sdf_rweibull()`

---

sdf_rcauchy   *Generate random samples from a Cauchy distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a Cauchy distribution.

Usage

```
sdf_rcauchy(
  sc,
  n,
  location = 0,
  scale = 1,
  num_partitions = NULL,
  seed = NULL,
  output_col = "x"
)
```

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
location	Location parameter of the distribution.
scale	Scale parameter of the distribution.
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rbinom()`, `sdf_rchisq()`, `sdf_rexp()`, `sdf_rgamma()`, `sdf_rgeom()`, `sdf_rhyper()`, `sdf_rlnorm()`, `sdf_rnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_runif()`, `sdf_rweibull()`

---

sdf_rchisq   *Generate random samples from a chi-squared distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a chi-squared distribution.

Usage

```
sdf_rchisq(sc, n, df, num_partitions = NULL, seed = NULL, output_col = "x")
```

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
df	Degrees of freedom (non-negative, but can be non-integer).
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rbinom()`, `sdf_rcauchy()`, `sdf_rexp()`, `sdf_rgamma()`, `sdf_rgeom()`, `sdf_rhyper()`, `sdf_rlnorm()`, `sdf_rnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_runif()`, `sdf_rweibull()`

---

sdf_read_column   *Read a Column from a Spark DataFrame*

Description

Read a single column from a Spark DataFrame, and return the contents of that column back to R.

Usage

```
sdf_read_column(x, column)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
column	The name of a column within `x`.

Details

It is expected for this operation to preserve row order.

---

sdf_register   *Register a Spark DataFrame*

Description

Registers a Spark DataFrame (giving it a table name for the Spark SQL context), and returns a
`tbl_spark`.

Usage

```
sdf_register(x, name = NULL)
```

Arguments

x	A Spark DataFrame.
name	A name to assign this table.

See Also

Other Spark data frames: `sdf_copy_to()`, `sdf_distinct()`, `sdf_random_split()`, `sdf_sample()`,
`sdf_sort()`, `sdf_weighted_sample()`

---

sdf_repartition   *Repartition a Spark DataFrame*

Description

Repartition a Spark DataFrame

Usage

```
sdf_repartition(x, partitions = NULL, partition_by = NULL)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
partitions	number of partitions
partition_by	vector of column names used for partitioning, only supported for Spark 2.0+

sdf_residuals.ml_model_generalized_linear_regression
Model Residuals

Description

This generic method returns a Spark DataFrame with model residuals added as a column to the model training data.

Usage

```
## S3 method for class 'ml_model_generalized_linear_regression'
sdf_residuals(
  object,
  type = c("deviance", "pearson", "working", "response"),
  ...
)

## S3 method for class 'ml_model_linear_regression'
sdf_residuals(object, ...)

sdf_residuals(object, ...)
```

Arguments

object	Spark ML model object.
type	type of residuals which should be returned.
...	additional arguments

---

sdf_rexp   *Generate random samples from an exponential distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from an exponential distribution.

Usage

```
sdf_rexp(sc, n, rate = 1, num_partitions = NULL, seed = NULL, output_col = "x")
```

Arguments

`sc`	A Spark connection.
`n`	Sample Size (default: 1000).
`rate`	Rate of the exponential distribution (default: 1). The exponential distribution with rate lambda has mean 1 / lambda and density f(x) = lambda e ^ - lambda x.
`num_partitions`	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
`seed`	Random seed (default: a random long integer).
`output_col`	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rbinom()`, `sdf_rcauchy()`, `sdf_rchisq()`, `sdf_rgamma()`, `sdf_rgeom()`, `sdf_rhyper()`, `sdf_rlnorm()`, `sdf_rnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_runif()`, `sdf_rweibull()`

---

sdf_rgamma   *Generate random samples from a Gamma distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a Gamma distribution.

Usage

```
sdf_rgamma(
  sc,
  n,
  shape,
  rate = 1,
  num_partitions = NULL,
  seed = NULL,
  output_col = "x"
)
```

Arguments

`sc`	A Spark connection.
`n`	Sample Size (default: 1000).
`shape`	Shape parameter (greater than 0) for the Gamma distribution.
`rate`	Rate parameter (greater than 0) for the Gamma distribution (scale is 1/rate).
`num_partitions`	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
`seed`	Random seed (default: a random long integer).
`output_col`	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rbinom()`, `sdf_rcauchy()`, `sdf_rchisq()`, `sdf_rexp()`, `sdf_rgeom()`, `sdf_rhyper()`, `sdf_rlnorm()`, `sdf_rnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_runif()`, `sdf_rweibull()`

---

sdf_rgeom   *Generate random samples from a geometric distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a geometric distribution.

Usage

```
sdf_rgeom(sc, n, prob, num_partitions = NULL, seed = NULL, output_col = "x")
```

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
prob	Probability of success in each trial.
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rbinom()`, `sdf_rcauchy()`, `sdf_rchisq()`, `sdf_rexp()`, `sdf_rgamma()`, `sdf_rhyper()`, `sdf_rlnorm()`, `sdf_rnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_runif()`, `sdf_rweibull()`

---

sdf_rhyper　　　　　　　　　　　*Generate random samples from a hypergeometric distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a hypergeometric distribution.

Usage

```
sdf_rhyper(
  sc,
  nn,
  m,
  n,
  k,
  num_partitions = NULL,
  seed = NULL,
  output_col = "x"
)
```

Arguments

sc	A Spark connection.
nn	Sample Size.
m	The number of successes among the population.
n	The number of failures among the population.
k	The number of draws.
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: sdf_rbeta(), sdf_rbinom(), sdf_rcauchy(), sdf_rchisq(), sdf_rexp(), sdf_rgamma(), sdf_rgeom(), sdf_rlnorm(), sdf_rnorm(), sdf_rpois(), sdf_rt(), sdf_runif(), sdf_rweibull()

---

sdf_rlnorm *Generate random samples from a log normal distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a log normal distribution.

Usage

```
sdf_rlnorm(
  sc,
  n,
  meanlog = 0,
  sdlog = 1,
  num_partitions = NULL,
  seed = NULL,
  output_col = "x"
)
```

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
meanlog	The mean of the normally distributed natural logarithm of this distribution.
sdlog	The Standard deviation of the normally distributed natural logarithm of this distribution.
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rbinom()`, `sdf_rcauchy()`, `sdf_rchisq()`, `sdf_rexp()`, `sdf_rgamma()`, `sdf_rgeom()`, `sdf_rhyper()`, `sdf_rnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_runif()`, `sdf_rweibull()`

---

sdf_rnorm   *Generate random samples from the standard normal distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from the standard normal distribution.

Usage

```
sdf_rnorm(
  sc,
  n,
  mean = 0,
  sd = 1,
  num_partitions = NULL,
  seed = NULL,
  output_col = "x"
)
```

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
mean	The mean value of the normal distribution.
sd	The standard deviation of the normal distribution.
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rbinom()`, `sdf_rcauchy()`, `sdf_rchisq()`, `sdf_rexp()`, `sdf_rgamma()`, `sdf_rgeom()`, `sdf_rhyper()`, `sdf_rlnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_runif()`, `sdf_rweibull()`

---

sdf_rpois   *Generate random samples from a Poisson distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a Poisson distribution.

Usage

```
sdf_rpois(sc, n, lambda, num_partitions = NULL, seed = NULL, output_col = "x")
```

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
lambda	Mean, or lambda, of the Poisson distribution.
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: sdf_rbeta(), sdf_rbinom(), sdf_rcauchy(), sdf_rchisq(), sdf_rexp(), sdf_rgamma(), sdf_rgeom(), sdf_rhyper(), sdf_rlnorm(), sdf_rnorm(), sdf_rt(), sdf_runif(), sdf_rweibull()

---

sdf_rt   *Generate random samples from a t-distribution*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a t-distribution.

Usage

```
sdf_rt(sc, n, df, num_partitions = NULL, seed = NULL, output_col = "x")
```

Arguments

`sc`	A Spark connection.
`n`	Sample Size (default: 1000).
`df`	Degrees of freedom (> 0, maybe non-integer).
`num_partitions`	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
`seed`	Random seed (default: a random long integer).
`output_col`	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: [sdf_rbeta](), [sdf_rbinom](), [sdf_rcauchy](), [sdf_rchisq](),
[sdf_rexp](), [sdf_rgamma](), [sdf_rgeom](), [sdf_rhyper](), [sdf_rlnorm](), [sdf_rnorm](), [sdf_rpois](),
[sdf_runif](), [sdf_rweibull]()

---

sdf_runif	*Generate random samples from the uniform distribution U(0, 1).*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from
the uniform distribution U(0, 1).

Usage

```
sdf_runif(
  sc,
  n,
  min = 0,
  max = 1,
  num_partitions = NULL,
  seed = NULL,
  output_col = "x"
)
```

Arguments

`sc`	A Spark connection.
`n`	Sample Size (default: 1000).
`min`	The lower limit of the distribution.
`max`	The upper limit of the distribution.
`num_partitions`	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
`seed`	Random seed (default: a random long integer).
`output_col`	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rbinom()`, `sdf_rcauchy()`, `sdf_rchisq()`, `sdf_rexp()`, `sdf_rgamma()`, `sdf_rgeom()`, `sdf_rhyper()`, `sdf_rlnorm()`, `sdf_rnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_rweibull()`

---

sdf_rweibull	*Generate random samples from a Weibull distribution.*

Description

Generator method for creating a single-column Spark dataframes comprised of i.i.d. samples from a Weibull distribution.

Usage

```
sdf_rweibull(
  sc,
  n,
  shape,
  scale = 1,
  num_partitions = NULL,
  seed = NULL,
  output_col = "x"
)
```

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
shape	The shape of the Weibull distribution.
scale	The scale of the Weibull distribution (default: 1).
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

See Also

Other Spark statistical routines: `sdf_rbeta()`, `sdf_rbinom()`, `sdf_rcauchy()`, `sdf_rchisq()`, `sdf_rexp()`, `sdf_rgamma()`, `sdf_rgeom()`, `sdf_rhyper()`, `sdf_rlnorm()`, `sdf_rnorm()`, `sdf_rpois()`, `sdf_rt()`, `sdf_runif()`

sdf_sample   *Randomly Sample Rows from a Spark DataFrame*

Description

Draw a random sample of rows (with or without replacement) from a Spark DataFrame.

Usage

```
sdf_sample(x, fraction = 1, replacement = TRUE, seed = NULL)
```

Arguments

x	An object coercable to a Spark DataFrame.
fraction	The fraction to sample.
replacement	Boolean; sample with replacement?
seed	An (optional) integer seed.

See Also

Other Spark data frames: `sdf_copy_to()`, `sdf_distinct()`, `sdf_random_split()`, `sdf_register()`, `sdf_sort()`, `sdf_weighted_sample()`

sdf_schema   *Read the Schema of a Spark DataFrame*

Description

Read the schema of a Spark DataFrame.

Usage

```
sdf_schema(x, expand_nested_cols = FALSE, expand_struct_cols = FALSE)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
expand_nested_cols	Whether to expand columns containing nested array of structs (which are usually created by tidyr::nest on a Spark data frame)
expand_struct_cols	Whether to expand columns containing structs

Details

The type column returned gives the string representation of the underlying Spark type for that column; for example, a vector of numeric values would be returned with the type "DoubleType". Please see the Spark Scala API Documentation for information on what types are available and exposed by Spark.

Value

An R list, with each list element describing the name and type of a column.

---

sdf_separate_column   *Separate a Vector Column into Scalar Columns*

Description

Given a vector column in a Spark DataFrame, split that into n separate columns, each column made up of the different elements in the column column.

Usage

```
sdf_separate_column(x, column, into = NULL)
```

Arguments

x	A spark_connection, ml_pipeline, or a tbl_spark.
column	The name of a (vector-typed) column.
into	A specification of the columns that should be generated from column. This can either be a vector of column names, or an R list mapping column names to the (1-based) index at which a particular vector element should be extracted.

---

sdf_seq   *Create DataFrame for Range*

Description

Creates a DataFrame for the given range

Usage

```
sdf_seq(
  sc,
  from = 1L,
  to = 1L,
  by = 1L,
  repartition = NULL,
  type = c("integer", "integer64")
)
```

Arguments

sc	The associated Spark connection.
from, to	The start and end to use as a range
by	The increment of the sequence.
repartition	The number of partitions to use when distributing the data across the Spark cluster. Defaults to the minimum number of partitions.
type	The data type to use for the index, either "integer" or "integer64".

---

sdf_sort   *Sort a Spark DataFrame*

Description

Sort a Spark DataFrame by one or more columns, with each column sorted in ascending order.

Usage

```
sdf_sort(x, columns)
```

Arguments

x	An object coercable to a Spark DataFrame.
columns	The column(s) to sort by.

See Also

Other Spark data frames: sdf_copy_to(), sdf_distinct(), sdf_random_split(), sdf_register(), sdf_sample(), sdf_weighted_sample()

---

sdf_sql   *Spark DataFrame from SQL*

Description

Defines a Spark DataFrame from a SQL query, useful to create Spark DataFrames without collecting the results immediately.

Usage

```
sdf_sql(sc, sql)
```

Arguments

sc	A spark_connection.
sql	a 'SQL' query used to generate a Spark DataFrame.

---

sdf_to_avro *Convert column(s) to avro format*

Description

Convert column(s) to avro format

Usage

```
sdf_to_avro(x, cols = colnames(x))
```

Arguments

x	An object coercible to a Spark DataFrame
cols	Subset of Columns to convert into avro format

---

sdf_unnest_longer *Unnest longer*

Description

Expand a struct column or an array column within a Spark dataframe into one or more rows, similar
what to tidyr::unnest_longer does to an R dataframe. An index column, if included, will be 1-based
if 'col' is an array column.

Usage

```
sdf_unnest_longer(
  data,
  col,
  values_to = NULL,
  indices_to = NULL,
  include_indices = NULL,
  names_repair = "check_unique",
  ptype = list(),
  transform = list()
)
```

Arguments

data	The Spark dataframe to be unnested
col	The struct column to extract components from
values_to	Name of column to store vector values. Defaults to 'col'.
indices_to	A string giving the name of column which will contain the inner names or position (if not named) of the values. Defaults to 'col' with '_id' suffix

include_indices

>   Whether to include an index column. An index column will be included by
>   default if 'col' is a struct column. It will also be included if 'indices_to' is not
>   'NULL'.

names_repair   Strategy for fixing duplicate column names (the semantic will be exactly identi-
>   cal to that of '.name_repair' option in [tibble](#))

ptype   Optionally, supply an R data frame prototype for the output. Each column of the
>   unnested result will be casted based on the Spark equivalent of the type of the
>   column with the same name within 'ptype', e.g., if 'ptype' has a column 'x' of
>   type 'character', then column 'x' of the unnested result will be casted from its
>   original SQL type to StringType.

transform   Optionally, a named list of transformation functions applied

Examples

```
## Not run:
library(sparklyr)
sc <- spark_connect(master = "local", version = "2.4.0")

# unnesting a struct column
sdf <- copy_to(
  sc,
  dplyr::tibble(
    x = 1:3,
    y = list(list(a = 1, b = 2), list(a = 3, b = 4), list(a = 5, b = 6))
  )
)

unnested <- sdf %>% sdf_unnest_longer(y, indices_to = "attr")

# unnesting an array column
sdf <- copy_to(
  sc,
  dplyr::tibble(
    x = 1:3,
    y = list(1:10, 1:5, 1:2)
  )
)

unnested <- sdf %>% sdf_unnest_longer(y, indices_to = "array_idx")

## End(Not run)
```

---

sdf_unnest_wider   *Unnest wider*

Description

Flatten a struct column within a Spark dataframe into one or more columns, similar what to tidyr::unnest_wider does to an R dataframe

Usage

```
sdf_unnest_wider(
  data,
  col,
  names_sep = NULL,
  names_repair = "check_unique",
  ptype = list(),
  transform = list()
)
```

Arguments

data	The Spark dataframe to be unnested
col	The struct column to extract components from
names_sep	If 'NULL', the default, the names will be left as is. If a string, the inner and outer names will be pasted together using 'names_sep' as the delimiter.
names_repair	Strategy for fixing duplicate column names (the semantic will be exactly identical to that of '.name_repair' option in [tibble](#))
ptype	Optionally, supply an R data frame prototype for the output. Each column of the unnested result will be casted based on the Spark equivalent of the type of the column with the same name within 'ptype', e.g., if 'ptype' has a column 'x' of type 'character', then column 'x' of the unnested result will be casted from its original SQL type to StringType.
transform	Optionally, a named list of transformation functions applied to each component (e.g., list('x = as.character') to cast column 'x' to String).

Examples

```
## Not run:
library(sparklyr)
sc <- spark_connect(master = "local", version = "2.4.0")

sdf <- copy_to(
  sc,
  dplyr::tibble(
    x = 1:3,
    y = list(list(a = 1, b = 2), list(a = 3, b = 4), list(a = 5, b = 6))
  )
)

# flatten struct column 'y' into two separate columns 'y_a' and 'y_b'
unnested <- sdf %>% sdf_unnest_wider(y, names_sep = "_")
```

```
## End(Not run)
```

---

sdf_weighted_sample   *Perform Weighted Random Sampling on a Spark DataFrame*

Description

Draw a random sample of rows (with or without replacement) from a Spark DataFrame If the sampling is done without replacement, then it will be conceptually equivalent to an iterative process such that in each step the probability of adding a row to the sample set is equal to its weight divided by summation of weights of all rows that are not in the sample set yet in that step.

Usage

```
sdf_weighted_sample(x, weight_col, k, replacement = TRUE, seed = NULL)
```

Arguments

x	An object coercable to a Spark DataFrame.
weight_col	Name of the weight column
k	Sample set size
replacement	Whether to sample with replacement
seed	An (optional) integer seed

See Also

Other Spark data frames: [sdf_copy_to](), [sdf_distinct](), [sdf_random_split](), [sdf_register](), [sdf_sample](), [sdf_sort]()

---

sdf_with_sequential_id
                    Add a Sequential ID Column to a Spark DataFrame

---

Description

Add a sequential ID column to a Spark DataFrame. The Spark zipWithIndex function is used to produce these. This differs from sdf_with_unique_id in that the IDs generated are independent of partitioning.

Usage

```
sdf_with_sequential_id(x, id = "id", from = 1L)
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
id	The name of the column to host the generated IDs.
from	The starting value of the id column

`sdf_with_unique_id`	*Add a Unique ID Column to a Spark DataFrame*

Description

Add a unique ID column to a Spark DataFrame. The Spark `monotonicallyIncreasingId` function is used to produce these and is guaranteed to produce unique, monotonically increasing ids; however, there is no guarantee that these IDs will be sequential. The table is persisted immediately after the column is generated, to ensure that the column is stable – otherwise, it can differ across new computations.

Usage

```
sdf_with_unique_id(x, id = "id")
```

Arguments

x	A `spark_connection`, `ml_pipeline`, or a `tbl_spark`.
id	The name of the column to host the generated IDs.

---

select	*Select*

Description

See [select](#) for more details.

---

separate	*Separate*

Description

See [separate](#) for more details.

---

spark-api   *Access the Spark API*

Description

Access the commonly-used Spark objects associated with a Spark instance. These objects provide access to different facets of the Spark API.

Usage

```
spark_context(sc)

java_context(sc)

hive_context(sc)

spark_session(sc)
```

Arguments

sc   A spark_connection.

Details

The Scala API documentation is useful for discovering what methods are available for each of these objects. Use invoke to call methods on these objects.

Spark Context

The main entry point for Spark functionality. The **Spark Context** represents the connection to a Spark cluster, and can be used to create RDDs, accumulators and broadcast variables on that cluster.

Java Spark Context

A Java-friendly version of the aforementioned **Spark Context**.

Hive Context

An instance of the Spark SQL execution engine that integrates with data stored in Hive. Configuration for Hive is read from hive-site.xml on the classpath.

Starting with Spark >= 2.0.0, the **Hive Context** class has been deprecated – it is superceded by the **Spark Session** class, and hive_context will return a **Spark Session** object instead. Note that both classes share a SQL interface, and therefore one can invoke SQL through these objects.

Spark Session

Available since Spark 2.0.0, the **Spark Session** unifies the **Spark Context** and **Hive Context** classes into a single interface. Its use is recommended over the older APIs for code targeting Spark 2.0.0 and above.

---

spark-connections   *Manage Spark Connections*

Description

These routines allow you to manage your connections to Spark.

Call 'spark_disconnect()' on each open Spark connection

Usage

```
spark_connect(
  master,
  spark_home = Sys.getenv("SPARK_HOME"),
  method = c("shell", "livy", "databricks", "test", "qubole", "synapse"),
  app_name = "sparklyr",
  version = NULL,
  config = spark_config(),
  extensions = sparklyr::registered_extensions(),
  packages = NULL,
  scala_version = NULL,
  ...
)

spark_connection_is_open(sc)

spark_disconnect(sc, ...)

spark_disconnect_all(...)

spark_submit(
  master,
  file,
  spark_home = Sys.getenv("SPARK_HOME"),
  app_name = "sparklyr",
  version = NULL,
  config = spark_config(),
  extensions = sparklyr::registered_extensions(),
  scala_version = NULL,
  ...
)
```

Arguments

master	Spark cluster url to connect to. Use `"local"` to connect to a local instance of Spark installed via `spark_install`.
spark_home	The path to a Spark installation. Defaults to the path provided by the SPARK_HOME environment variable. If SPARK_HOME is defined, it will always be used unless the version parameter is specified to force the use of a locally installed version.
method	The method used to connect to Spark. Default connection method is `"shell"` to connect using spark-submit, use `"livy"` to perform remote connections using HTTP, or `"databricks"` when using a Databricks clusters.
app_name	The application name to be used while running in the Spark cluster.
version	The version of Spark to use. Required for `"local"` Spark connections, optional otherwise.
config	Custom configuration for the generated Spark connection. See `spark_config` for details.
extensions	Extension R packages to enable for this connection. By default, all packages enabled through the use of `sparklyr::register_extension` will be passed here.
packages	A list of Spark packages to load. For example, `"delta"` or `"kafka"` to enable Delta Lake or Kafka. Also supports full versions like `"io.delta:delta-core_2.11:0.4.0"`. This is similar to adding packages into the `sparklyr.shell.packages` configuration option. Notice that the version parameter is used to choose the correct package, otherwise assumes the latest version is being used.
scala_version	Load the sparklyr jar file that is built with the version of Scala specified (this currently only makes sense for Spark 2.4, where sparklyr will by default assume Spark 2.4 on current host is built with Scala 2.11, and therefore 'scala_version = '2.12'' is needed if sparklyr is connecting to Spark 2.4 built with Scala 2.12)
...	Additional params to be passed to each 'spark_disconnect()' call (e.g., 'terminate = TRUE')
sc	A spark_connection.
file	Path to R source file to submit for batch execution.

Details

By default, when using `method = "livy"`, jars are downloaded from GitHub. But an alternative path (local to Livy server or on HDFS or HTTP(s)) to sparklyr JAR can also be specified through the `sparklyr.livy.jar` setting.

Examples

```
conf <- spark_config()
conf$`sparklyr.shell.conf` <- c(
  "spark.executor.extraJavaOptions=-Duser.timezone='UTC'",
  "spark.driver.extraJavaOptions=-Duser.timezone='UTC'",
  "spark.sql.session.timeZone='UTC'"
)
```

```
sc <- spark_connect(
  master = "spark://HOST:PORT", config = conf
)
connection_is_open(sc)

spark_disconnect(sc)
```

---

sparklyr_get_backend_port

Return the port number of a 'sparklyr' backend.

Description

Retrieve the port number of the 'sparklyr' backend associated with a Spark connection.

Usage

```
sparklyr_get_backend_port(sc)
```

Arguments

sc   A spark_connection.

Value

The port number of the 'sparklyr' backend associated with sc.

---

spark_adaptive_query_execution

Retrieves or sets status of Spark AQE

Description

Retrieves or sets whether Spark adaptive query execution is enabled

Usage

```
spark_adaptive_query_execution(sc, enable = NULL)
```

Arguments

sc   A spark_connection.

enable   Whether to enable Spark adaptive query execution. Defaults to NULL to retrieve
                configuration entries.

See Also

Other Spark runtime configuration: `spark_advisory_shuffle_partition_size()`, `spark_auto_broadcast_join_thres`
`spark_coalesce_initial_num_partitions()`, `spark_coalesce_min_num_partitions()`, `spark_coalesce_shuffle_`
`spark_session_config()`

---

spark_advisory_shuffle_partition_size
Retrieves or sets advisory size of the shuffle partition

Description

Retrieves or sets advisory size in bytes of the shuffle partition during adaptive optimization

Usage

```
spark_advisory_shuffle_partition_size(sc, size = NULL)
```

Arguments

sc   A spark_connection.

size   Advisory size in bytes of the shuffle partition. Defaults to NULL to retrieve con-
            figuration entries.

See Also

Other Spark runtime configuration: `spark_adaptive_query_execution()`, `spark_auto_broadcast_join_threshold()`,
`spark_coalesce_initial_num_partitions()`, `spark_coalesce_min_num_partitions()`, `spark_coalesce_shuffle_`
`spark_session_config()`

---

spark_apply   *Apply an R Function in Spark*

Description

Applies an R function to a Spark object (typically, a Spark DataFrame).

Usage

```
spark_apply(
  x,
  f,
  columns = NULL,
  memory = TRUE,
  group_by = NULL,
  packages = NULL,
  context = NULL,
  name = NULL,
  barrier = NULL,
  fetch_result_as_sdf = TRUE,
  partition_index_param = "",
  arrow_max_records_per_batch = NULL,
  auto_deps = FALSE,
  ...
)
```

Arguments

x	An object (usually a `spark_tbl`) coercable to a Spark DataFrame.
f	A function that transforms a data frame partition into a data frame. The function f has signature `f(df, context, group1, group2, ...)` where df is a data frame with the data to be processed, `context` is an optional object passed as the `context` parameter and `group1` to `groupN` contain the values of the `group_by` values. When `group_by` is not specified, f takes only one argument. Can also be an rlang anonymous function. For example, as `~ .x + 1` to define an expression that adds one to the given `.x` data frame.
columns	A vector of column names or a named vector of column types for the transformed object. When not specified, a sample of 10 rows is taken to infer out the output columns automatically, to avoid this performance penalty, specify the column types. The sample size is configurable using the `sparklyr.apply.schema.infer` configuration option.
memory	Boolean; should the table be cached into memory?
group_by	Column name used to group by data frame partitions.
packages	Boolean to distribute `.libPaths()` packages to each node, a list of packages to distribute, or a package bundle created with `spark_apply_bundle()`. Defaults to `TRUE` or the `sparklyr.apply.packages` value set in `spark_config()`. For clusters using Yarn cluster mode, `packages` can point to a package bundle created using `spark_apply_bundle()` and made available as a Spark file using `config$sparklyr.shell.files`. For clusters using Livy, packages can be manually installed on the driver node. For offline clusters where `available.packages()` is not available, manually download the packages database from https://cran.r-project.org/web/packages/packages.rds and set `Sys.setenv(sparklyr.apply.packagesdb = "<pathl-to-rds>")`. Otherwise, all packages will be used by default.

For clusters where R packages already installed in every worker node, the `spark.r.libpaths` config entry can be set in `spark_config()` to the local packages library. To specify multiple paths collapse them (without spaces) with a comma delimiter (e.g., `"/lib/path/one,/lib/path/two"`).

context   Optional object to be serialized and passed back to `f()`.

name   Optional table name while registering the resulting data frame.

barrier   Optional to support Barrier Execution Mode in the scheduler.

fetch_result_as_sdf

Whether to return the transformed results in a Spark Dataframe (defaults to `TRUE`). When set to `FALSE`, results will be returned as a list of R objects instead.

NOTE: `fetch_result_as_sdf` must be set to `FALSE` when the transformation function being applied is returning R objects that cannot be stored in a Spark Dataframe (e.g., complex numbers or any other R data type that does not have an equivalent representation among Spark SQL data types).

partition_index_param

Optional if non-empty, then `f` also receives the index of the partition being processed as a named argument with this name, in addition to all positional argument(s) it will receive

NOTE: when `fetch_result_as_sdf` is set to `FALSE`, object returned from the transformation function also must be serializable by the `base::serialize` function in R.

arrow_max_records_per_batch

Maximum size of each Arrow record batch, ignored if Arrow serialization is not enabled.

auto_deps   [Experimental] Whether to infer all required R packages by examining the closure `f()` and only distribute required R and their transitive dependencies to Spark worker nodes (default: FALSE). NOTE: this option will only take effect if `packages` is set to `TRUE` or is a character vector of R package names. If `packages` is a character vector of R package names, then both the set of packages specified by `packages` and the set of inferred packages will be distributed to Spark workers.

...   Optional arguments; currently unused.

Configuration

`spark_config()` settings can be specified to change the workers environment.

For instance, to set additional environment variables to each worker node use the `sparklyr.apply.env.*` config, to launch workers without `--vanilla` use `sparklyr.apply.options.vanilla` set to `FALSE`, to run a custom script before launching Rscript use `sparklyr.apply.options.rscript.before`.

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(master = "local[3]")
```

```
# creates an Spark data frame with 10 elements then multiply times 10 in R
sdf_len(sc, 10) %>% spark_apply(function(df) df * 10)

# using barrier mode
sdf_len(sc, 3, repartition = 3) %>%
  spark_apply(nrow, barrier = TRUE, columns = c(id = "integer")) %>%
  collect()

## End(Not run)
```

spark_apply_bundle   *Create Bundle for Spark Apply*

Description

Creates a bundle of packages for spark_apply().

Usage

```
spark_apply_bundle(packages = TRUE, base_path = getwd(), session_id = NULL)
```

Arguments

packages	List of packages to pack or TRUE to pack all.
base_path	Base path used to store the resulting bundle.
session_id	An optional ID string to include in the bundle file name to allow the bundle to be session-specific

spark_apply_log   *Log Writer for Spark Apply*

Description

Writes data to log under spark_apply().

Usage

```
spark_apply_log(..., level = "INFO")
```

Arguments

...	Arguments to write to log.
level	Severity level for this entry; recommended values: INFO, ERROR or WARN.

spark_auto_broadcast_join_threshold
                    Retrieves or sets the auto broadcast join threshold

Description

Configures the maximum size in bytes for a table that will be broadcast to all worker nodes when performing a join. By setting this value to -1 broadcasting can be disabled. Note that currently statistics are only supported for Hive Metastore tables where the command 'ANALYZE TABLE <tableName> COMPUTE STATISTICS noscan' has been run, and file-based data source tables where the statistics are computed directly on the files of data.

Usage

```
spark_auto_broadcast_join_threshold(sc, threshold = NULL)
```

Arguments

sc   A spark_connection.

threshold   Maximum size in bytes for a table that will be broadcast to all worker nodes when performing a join. Defaults to NULL to retrieve configuration entries.

See Also

Other Spark runtime configuration: spark_adaptive_query_execution(), spark_advisory_shuffle_partition_size(
spark_coalesce_initial_num_partitions(), spark_coalesce_min_num_partitions(), spark_coalesce_shuffle_p
spark_session_config()

spark_coalesce_initial_num_partitions
                    Retrieves or sets initial number of shuffle partitions before coalescing

Description

Retrieves or sets initial number of shuffle partitions before coalescing

Usage

```
spark_coalesce_initial_num_partitions(sc, num_partitions = NULL)
```

Arguments

sc   A spark_connection.

num_partitions Initial number of shuffle partitions before coalescing. Defaults to NULL to retrieve configuration entries.

See Also

Other Spark runtime configuration: `spark_adaptive_query_execution`(), `spark_advisory_shuffle_partition_size`(
`spark_auto_broadcast_join_threshold`(), `spark_coalesce_min_num_partitions`(), `spark_coalesce_shuffle_pa
`spark_session_config`()

---

spark_coalesce_min_num_partitions

Retrieves or sets the minimum number of shuffle partitions after coalescing

---

Description

Retrieves or sets the minimum number of shuffle partitions after coalescing

Usage

```
spark_coalesce_min_num_partitions(sc, num_partitions = NULL)
```

Arguments

sc   A spark_connection.

num_partitions Minimum number of shuffle partitions after coalescing. Defaults to `NULL` to
                retrieve configuration entries.

See Also

Other Spark runtime configuration: `spark_adaptive_query_execution`(), `spark_advisory_shuffle_partition_size`(
`spark_auto_broadcast_join_threshold`(), `spark_coalesce_initial_num_partitions`(), `spark_coalesce_shuffle`
`spark_session_config`()

---

spark_coalesce_shuffle_partitions

Retrieves or sets whether coalescing contiguous shuffle partitions is enabled

---

Description

Retrieves or sets whether coalescing contiguous shuffle partitions is enabled

Usage

```
spark_coalesce_shuffle_partitions(sc, enable = NULL)
```

Arguments

sc	A spark_connection.
enable	Whether to enable coalescing of contiguous shuffle partitions. Defaults to NULL to retrieve configuration entries.

See Also

Other Spark runtime configuration: spark_adaptive_query_execution(), spark_advisory_shuffle_partition_size(
spark_auto_broadcast_join_threshold(), spark_coalesce_initial_num_partitions(), spark_coalesce_min_nu
spark_session_config()

---

spark_compilation_spec

Define a Spark Compilation Specification

Description

For use with compile_package_jars. The Spark compilation specification is used when compiling
Spark extension Java Archives, and defines which versions of Spark, as well as which versions of
Scala, should be used for compilation.

Usage

```
spark_compilation_spec(
  spark_version = NULL,
  spark_home = NULL,
  scalac_path = NULL,
  scala_filter = NULL,
  jar_name = NULL,
  jar_path = NULL,
  jar_dep = NULL,
  embedded_srcs = "embedded_sources.R"
)
```

Arguments

spark_version	The Spark version to build against. This can be left unset if the path to a suitable Spark home is supplied.
spark_home	The path to a Spark home installation. This can be left unset if spark_version is supplied; in such a case, sparklyr will attempt to discover the associated Spark installation using spark_home_dir.
scalac_path	The path to the scalac compiler to be used during compilation of your Spark extension. Note that you should ensure the version of scalac selected matches the version of scalac used with the version of Spark you are compiling against.

scala_filter	An optional R function that can be used to filter which scala files are used during compilation. This can be useful if you have auxiliary files that should only be included with certain versions of Spark.
jar_name	The name to be assigned to the generated jar.
jar_path	The path to the jar tool to be used during compilation of your Spark extension.
jar_dep	An optional list of additional jar dependencies.
embedded_srcs	Embedded source file(s) under <R package root>/java to be included in the root of the resulting jar file as resources

Details

Most Spark extensions won't need to define their own compilation specification, and can instead rely on the default behavior of compile_package_jars.

---

spark_config   *Read Spark Configuration*

Description

Read Spark Configuration

Usage

```
spark_config(file = "config.yml", use_default = TRUE)
```

Arguments

file	Name of the configuration file
use_default	TRUE to use the built-in defaults provided in this package

Details

Read Spark configuration using the **config** package.

Value

Named list with configuration data

---

spark_config_kubernetes

Kubernetes Configuration

Description

Convenience function to initialize a Kubernetes configuration instead of spark_config(), exposes
common properties to set in Kubernetes clusters.

Usage

```
spark_config_kubernetes(
  master,
  version = "3.2.3",
  image = "spark:sparklyr",
  driver = random_string("sparklyr-"),
  account = "spark",
  jars = "local:///opt/sparklyr",
  forward = TRUE,
  executors = NULL,
  conf = NULL,
  timeout = 120,
  ports = c(8880, 8881, 4040),
  fix_config = identical(.Platform$OS.type, "windows"),
  ...
)
```

Arguments

master	Kubernetes url to connect to, found by running kubectl cluster-info.
version	The version of Spark being used.
image	Container image to use to launch Spark and sparklyr. Also known as spark.kubernetes.container.ima
driver	Name of the driver pod. If not set, the driver pod name is set to "sparklyr" suffixed by id to avoid name conflicts. Also known as spark.kubernetes.driver.pod.name.
account	Service account that is used when running the driver pod. The driver pod uses this service account when requesting executor pods from the API server. Also known as spark.kubernetes.authenticate.driver.serviceAccountName.
jars	Path to the sparklyr jars; either, a local path inside the container image with the sparklyr jars copied when the image was created or, a path accesible by the container where the sparklyr jars were copied. You can find a path to the sparklyr jars by running system.file("java/", package = "sparklyr").
forward	Should ports used in sparklyr be forwarded automatically through Kubernetes? Default to TRUE which runs kubectl port-forward and pkill kubectl on disconnection.
executors	Number of executors to request while connecting.

conf	A named list of additional entries to add to `sparklyr.shell.conf`.
timeout	Total seconds to wait before giving up on connection.
ports	Ports to forward using kubectl.
fix_config	Should the spark-defaults.conf get fixed? `TRUE` for Windows.
...	Additional parameters, currently not in use.

---

spark_config_settings   *Retrieve Available Settings*

Description

Retrieves available sparklyr settings that can be used in configuration files or `spark_config()`.

Usage

```
spark_config_settings()
```

---

spark_connection   *Retrieve the Spark Connection Associated with an R Object*

Description

Retrieve the `spark_connection` associated with an R object.

Usage

```
spark_connection(x, ...)
```

Arguments

| x | An R object from which a `spark_connection` can be obtained. |
| ... | Optional arguments; currently unused. |

---

spark_connection-class

spark_connection class

Description

spark_connection class

spark_connection_find   *Find Spark Connection*

Description

Finds an active spark connection in the environment given the connection parameters.

Usage

```
spark_connection_find(master = NULL, app_name = NULL, method = NULL)
```

Arguments

master	The Spark master parameter.
app_name	The Spark application name.
method	The method used to connect to Spark.

---

spark_connect_method   *Function that negotiates the connection with the Spark back-end*

Description

Function that negotiates the connection with the Spark back-end

Usage

```
spark_connect_method(
  x,
  method,
  master,
  spark_home,
  config,
  app_name,
  version,
  hadoop_version,
  extensions,
  scala_version,
  ...
)
```

Arguments

x	A dummy method object to determine which code to use to connect
method	The method used to connect to Spark. Default connection method is "shell" to connect using spark-submit, use "livy" to perform remote connections using HTTP, or "databricks" when using a Databricks clusters.
master	Spark cluster url to connect to. Use "local" to connect to a local instance of Spark installed via [spark_install](#).
spark_home	The path to a Spark installation. Defaults to the path provided by the SPARK_HOME environment variable. If SPARK_HOME is defined, it will always be used unless the version parameter is specified to force the use of a locally installed version.
config	Custom configuration for the generated Spark connection. See [spark_config](#) for details.
app_name	The application name to be used while running in the Spark cluster.
version	The version of Spark to use. Required for "local" Spark connections, optional otherwise.
hadoop_version	Version of Hadoop to use
extensions	Extension R packages to enable for this connection. By default, all packages enabled through the use of [sparklyr::register_extension](#) will be passed here.
scala_version	Load the sparklyr jar file that is built with the version of Scala specified (this currently only makes sense for Spark 2.4, where sparklyr will by default assume Spark 2.4 on current host is built with Scala 2.11, and therefore 'scala_version = '2.12'' is needed if sparklyr is connecting to Spark 2.4 built with Scala 2.12)
...	Additional params to be passed to each 'spark_disconnect()' call (e.g., 'terminate = TRUE')

---

spark_context_config   *Runtime configuration interface for the Spark Context.*

Description

Retrieves the runtime configuration interface for the Spark Context.

Usage

```
spark_context_config(sc)
```

Arguments

sc	A spark_connection.

---

spark_dataframe *Retrieve a Spark DataFrame*

Description

This S3 generic is used to access a Spark DataFrame object (as a Java object reference) from an R object.

Usage

```
spark_dataframe(x, ...)
```

Arguments

x   An R object wrapping, or containing, a Spark DataFrame.

...   Optional arguments; currently unused.

Value

A [spark_jobj](#) representing a Java object reference to a Spark DataFrame.

---

spark_default_compilation_spec
                    Default Compilation Specification for Spark Extensions

---

Description

This is the default compilation specification used for Spark extensions, when used with [compile_package_jars](#).

Usage

```
spark_default_compilation_spec(
  pkg = infer_active_package_name(),
  locations = NULL
)
```

Arguments

pkg   The package containing Spark extensions to be compiled.

locations   Additional locations to scan. By default, the directories /opt/scala and /usr/local/scala
            will be scanned.

---

spark_dependency *Define a Spark dependency*

Description

Define a Spark dependency consisting of a set of custom JARs, Spark packages, and customized dbplyr SQL translation env.

Usage

```
spark_dependency(
  jars = NULL,
  packages = NULL,
  initializer = NULL,
  catalog = NULL,
  repositories = NULL,
  dbplyr_sql_variant = NULL,
  ...
)
```

Arguments

jars	Character vector of full paths to JAR files.
packages	Character vector of Spark packages names.
initializer	Optional callback function called when initializing a connection.
catalog	Optional location where extension JAR files can be downloaded for Livy.
repositories	Character vector of Spark package repositories.
dbplyr_sql_variant	

Customization of dbplyr SQL translation env. Must be a named list of the following form: list(scalar = list(scalar_fn1 = ..., scalar_fn2 = ..., <etc>), aggregate = list(agg_fn1 = ..., agg_fn2 = ..., <etc>), window = list(wnd_fn1 = ..., wnd_fn2 = ..., <etc>)) See sql_variant for details.

...	Additional optional arguments.

Value

An object of type 'spark_dependency'

---

spark_dependency_fallback

Fallback to Spark Dependency

Description

Helper function to assist falling back to previous Spark versions.

Usage

```
spark_dependency_fallback(spark_version, supported_versions)
```

Arguments

spark_version   The Spark version being requested in `spark_dependencies`.

supported_versions

The Spark versions that are supported by this extension.

Value

A Spark version to use.

---

spark_extension   *Create Spark Extension*

Description

Creates an R package ready to be used as an Spark extension.

Usage

```
spark_extension(path)
```

Arguments

path   Location where the extension will be created.

---

spark_home_set *Set the SPARK_HOME environment variable*

Description

Set the SPARK_HOME environment variable. This slightly speeds up some operations, including the connection time.

Usage

```
spark_home_set(path = NULL, ...)
```

Arguments

path   A string containing the path to the installation location of Spark. If NULL, the
                path to the most latest Spark/Hadoop versions is used.

...   Additional parameters not currently used.

Value

The function is mostly invoked for the side-effect of setting the SPARK_HOME environment variable. It also returns TRUE if the environment was successfully set, and FALSE otherwise.

Examples

```
## Not run:
# Not run due to side-effects
spark_home_set()

## End(Not run)
```

---

spark_ide_connection_open
                Set of functions to provide integration with the RStudio IDE

Description

Set of functions to provide integration with the RStudio IDE

Usage

```
spark_ide_connection_open(con, env, connect_call)

spark_ide_connection_closed(con)

spark_ide_connection_updated(con, hint)

spark_ide_connection_actions(con)

spark_ide_objects(con, catalog, schema, name, type)

spark_ide_columns(
  con,
  table = NULL,
  view = NULL,
  catalog = NULL,
  schema = NULL
)

spark_ide_preview(
  con,
  rowLimit,
  table = NULL,
  view = NULL,
  catalog = NULL,
  schema = NULL
)
```

Arguments

con	Valid Spark connection
env	R environment of the interactive R session
connect_call	R code that can be used to re-connect to the Spark connection
hint	Name of the Spark connection that the RStudio IDE can use as reference.
catalog	Name of the top level of the requested table or view
schema	Name of the second most top level of the requested level or view
name	The new of the view or table being requested
type	Type of the object being requested, 'view' or 'table'
table	Name of the requested table
view	Name of the requested view
rowLimit	The number of rows to show in the 'Preview' pane of the RStudio IDE

Details

These function are meant for downstream packages, that provide additional backends to 'spark-lyr', to override the opening, closing, update, and preview functionality. The arguments are driven

by what the RStudio IDE API expects them to be, so this is the reason why some use 'type' to designated views or tables, and others have one argument for 'table', and another for 'view'.

---

spark_insert_table   *Inserts a Spark DataFrame into a Spark table*

Description

Inserts a Spark DataFrame into a Spark table

Usage

```
spark_insert_table(
  x,
  name,
  mode = NULL,
  overwrite = FALSE,
  options = list(),
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
name	The name to assign to the newly generated table.
mode	A `character` element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also `https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes` for your version of Spark.
overwrite	Boolean; overwrite the table with the given name if it already exists?
options	A list of strings with additional options.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: `collect_from_rds()`, `spark_load_table()`, `spark_read()`, `spark_read_avro()`, `spark_read_binary()`, `spark_read_csv()`, `spark_read_delta()`, `spark_read_image()`, `spark_read_jdbc()`, `spark_read_json()`, `spark_read_libsvm()`, `spark_read_orc()`, `spark_read_parquet()`, `spark_read_source()`, `spark_read_table()`, `spark_read_text()`, `spark_save_table()`, `spark_write_avro()`, `spark_write_csv()`, `spark_write_delta()`, `spark_write_jdbc()`, `spark_write_json()`, `spark_write_orc()`, `spark_write_parquet()`, `spark_write_source()`, `spark_write_table()`, `spark_write_text()`

---

spark_install	*Download and install various versions of Spark*

Description

Install versions of Spark for use with local Spark connections (i.e. `spark_connect(master = "local")`)

Usage

```
spark_install(
  version = NULL,
  hadoop_version = NULL,
  reset = TRUE,
  logging = "INFO",
  verbose = interactive()
)

spark_uninstall(version, hadoop_version)

spark_install_dir()

spark_install_tar(tarfile)

spark_installed_versions()

spark_available_versions(
  show_hadoop = FALSE,
  show_minor = FALSE,
  show_future = FALSE
)
```

Arguments

version	Version of Spark to install. See `spark_available_versions` for a list of supported versions
hadoop_version	Version of Hadoop to install. See `spark_available_versions` for a list of supported versions
reset	Attempts to reset settings to defaults.
logging	Logging level to configure install. Supported options: "WARN", "INFO"
verbose	Report information as Spark is downloaded / installed
tarfile	Path to TAR file conforming to the pattern spark-###-bin-(hadoop)?### where ### reference spark and hadoop versions respectively.
show_hadoop	Show Hadoop distributions?
show_minor	Show minor Spark versions?
show_future	Should future versions which have not been released be shown?

Value

List with information about the installed version.

---

spark_integ_test_skip *It lets the package know if it should test a particular functionality or not*

Description

It lets the package know if it should test a particular functionality or not

Usage

```
spark_integ_test_skip(sc, test_name)
```

Arguments

sc	Spark connection
test_name	The name of the test

Details

It expects a boolean to be returned. If TRUE, the corresponding test will be skipped. If FALSE the test will be conducted.

---

spark_jobj *Retrieve a Spark JVM Object Reference*

Description

This S3 generic is used for accessing the underlying Java Virtual Machine (JVM) Spark objects associated with R objects. These objects act as references to Spark objects living in the JVM. Methods on these objects can be called with the [invoke](#) family of functions.

Usage

```
spark_jobj(x, ...)
```

Arguments

x	An R object containing, or wrapping, a spark_jobj.
...	Optional arguments; currently unused.

See Also

[invoke](#), for calling methods on Java object references.

spark_jobj-class   *spark_jobj class*

Description

spark_jobj class

spark_last_error   *Surfaces the last error from Spark captured by internal 'spark_error'*
                          function

Description

Surfaces the last error from Spark captured by internal 'spark_error' function

Usage

```
spark_last_error()
```

spark_load_table   *Reads from a Spark Table into a Spark DataFrame.*

Description

Reads from a Spark Table into a Spark DataFrame.

Usage

```
spark_load_table(
  sc,
  name,
  path,
  options = list(),
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
options	A list of strings with additional options. See https://spark.apache.org/docs/latest/sql-programming-guide.html#configuration.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_log   *View Entries in the Spark Log*

Description

View the most recent entries in the Spark log. This can be useful when inspecting output / errors produced by Spark during the invocation of various commands.

Usage

```
spark_log(sc, n = 100, filter = NULL, ...)
```

Arguments

sc	A spark_connection.
n	The max number of log entries to retrieve. Use NULL to retrieve all entries within the log.
filter	Character string to filter log entries.
...	Optional arguments; currently unused.

| spark_read | *Read file(s) into a Spark DataFrame using a custom reader* |

Description

Run a custom R function on Spark workers to ingest data from one or more files into a Spark DataFrame, assuming all files follow the same schema.

Usage

```
spark_read(sc, paths, reader, columns, packages = TRUE, ...)
```

Arguments

sc	A spark_connection.
paths	A character vector of one or more file URIs (e.g., c("hdfs://localhost:9000/file.txt", "hdfs://localhost:9000/file2.txt"))
reader	A self-contained R function that takes a single file URI as argument and returns the data read from that file as a data frame.
columns	a named list of column names and column types of the resulting data frame (e.g., list(column_1 = "integer", column_2 = "character")), or a list of column names only if column types should be inferred from the data (e.g., list("column_1", "column_2"), or NULL if column types should be inferred and resulting data frame can have arbitrary column names
packages	A list of R packages to distribute to Spark workers
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

Examples

```
## Not run:

library(sparklyr)
sc <- spark_connect(
  master = "yarn",
  spark_home = "~/spark/spark-2.4.5-bin-hadoop2.7"
)
```

```
# This is a contrived example to show reader tasks will be distributed across
# all Spark worker nodes
spark_read(
  sc,
  rep("/dev/null", 10),
  reader = function(path) system("hostname", intern = TRUE),
  columns = c(hostname = "string")
) %>% sdf_collect()

## End(Not run)
```

---

spark_read_avro   *Read Apache Avro data into a Spark DataFrame.*

Description

Notice this functionality requires the Spark connection sc to be instantiated with either an explicitly specified Spark version (i.e., spark_connect(..., version = <version>, packages = c("avro", <other package(s)>), ...)) or a specific version of Spark avro package to use (e.g., spark_connect(..., packages = c("org.apache.spark:spark-avro_2.12:3.0.0", <other package(s)>), ...)).

Usage

```
spark_read_avro(
  sc,
  name = NULL,
  path = name,
  avro_schema = NULL,
  ignore_extension = TRUE,
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
avro_schema	Optional Avro schema in JSON format
ignore_extension	
	If enabled, all files with and without .avro extension are loaded (default: TRUE)
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.

memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(),
spark_read(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(),
spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(),
spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(),
spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(),
spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_read_binary	*Read binary data into a Spark DataFrame.*

Description

Read binary files within a directory and convert each file into a record within the resulting Spark
dataframe. The output will be a Spark dataframe with the following columns and possibly partition
columns:

- path: StringType
- modificationTime: TimestampType
- length: LongType
- content: BinaryType

Usage

```
spark_read_binary(
  sc,
  name = NULL,
  dir = name,
  path_glob_filter = "*",
  recursive_file_lookup = FALSE,
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
dir	Directory to read binary files from.

path_glob_filter

> Glob pattern of binary files to be loaded (e.g., "*.jpg").

recursive_file_lookup

> If FALSE (default), then partition discovery will be enabled (i.e., if a partition naming scheme is present, then partitions specified by subdirectory names such as "date=2019-07-01" will be created and files outside subdirectories following a partition naming scheme will be ignored). If TRUE, then all nested directories will be searched even if their names do not follow a partition naming scheme.

repartition   The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.

memory   Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)

overwrite   Boolean; overwrite the table with the given name if it already exists?

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_read_csv   *Read a CSV file into a Spark DataFrame*

Description

Read a tabular data file into a Spark DataFrame.

Usage

```
spark_read_csv(
  sc,
  name = NULL,
  path = name,
  header = TRUE,
  columns = NULL,
  infer_schema = is.null(columns),
  delimiter = ",",
  quote = "\"",
  escape = "\\",
  charset = "UTF-8",
  null_value = NULL,
  options = list(),
  repartition = 0,
```

```
  memory = TRUE,
  overwrite = TRUE,
  ...
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
header	Boolean; should the first row of data be used as a header? Defaults to TRUE.
columns	A vector of column names or a named vector of column types. If specified, the elements can be "binary" for BinaryType, "boolean" for BooleanType, "byte" for ByteType, "integer" for IntegerType, "integer64" for LongType, "double" for DoubleType, "character" for StringType, "timestamp" for TimestampType and "date" for DateType.
infer_schema	Boolean; should column types be automatically inferred? Requires one extra pass over the data. Defaults to is.null(columns).
delimiter	The character used to delimit each column. Defaults to ','.
quote	The character used as a quote. Defaults to '"'.
escape	The character used to escape other characters. Defaults to '\'.
charset	The character set. Defaults to '"UTF-8"'.
null_value	The character to use for null, or missing, values. Defaults to NULL.
options	A list of strings with additional options.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?
...	Optional arguments; currently unused.

Details

You can read data from HDFS (hdfs://), S3 (s3a://), as well as the local file system (file://).

When header is FALSE, the column names are generated with a V prefix; e.g. V1, V2, ...

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_read_delta   *Read from Delta Lake into a Spark DataFrame.*

Description

Read from Delta Lake into a Spark DataFrame.

Usage

```
spark_read_delta(
  sc,
  path,
  name = NULL,
  version = NULL,
  timestamp = NULL,
  options = list(),
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE,
  ...
)
```

Arguments

sc	A spark_connection.
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
name	The name to assign to the newly generated table.
version	The version of the delta table to read.
timestamp	The timestamp of the delta table to read. For example, "2019-01-01" or "2019-01-01'T'00:00:00.000Z
options	A list of strings with additional options.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

spark_read_image   *Read image data into a Spark DataFrame.*

Description

Read image files within a directory and convert each file into a record within the resulting Spark dataframe. The output will be a Spark dataframe consisting of struct types containing the following attributes:

- origin: StringType

- height: IntegerType

- width: IntegerType

- nChannels: IntegerType

- mode: IntegerType

- data: BinaryType

Usage

```
spark_read_image(
  sc,
  name = NULL,
  dir = name,
  drop_invalid = TRUE,
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
dir	Directory to read binary files from.
drop_invalid	Whether to drop files that are not valid images from the result (default: TRUE).
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?

See Also

Other Spark serialization routines: [collect_from_rds](), [spark_insert_table](), [spark_load_table](),
[spark_read](), [spark_read_avro](), [spark_read_binary](), [spark_read_csv](), [spark_read_delta](),
[spark_read_jdbc](), [spark_read_json](), [spark_read_libsvm](), [spark_read_orc](), [spark_read_parquet](),
[spark_read_source](), [spark_read_table](), [spark_read_text](), [spark_save_table](), [spark_write_avro](),
[spark_write_csv](), [spark_write_delta](), [spark_write_jdbc](), [spark_write_json](), [spark_write_orc](),
[spark_write_parquet](), [spark_write_source](), [spark_write_table](), [spark_write_text]()

| spark_read_jdbc | *Read from JDBC connection into a Spark DataFrame.* |

Description

Read from JDBC connection into a Spark DataFrame.

Usage

```
spark_read_jdbc(
  sc,
  name,
  options = list(),
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE,
  columns = NULL,
  ...
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
options	A list of strings with additional options. See https://spark.apache.org/docs/latest/sql-programming-guide.html#configuration.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?
columns	A vector of column names or a named vector of column types. If specified, the elements can be "binary" for BinaryType, "boolean" for BooleanType, "byte" for ByteType, "integer" for IntegerType, "integer64" for LongType, "double" for DoubleType, "character" for StringType, "timestamp" for TimestampType and "date" for DateType.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(),
spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(),
spark_read_image(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(),
spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(),
spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(),
spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

Examples

```
## Not run:
sc <- spark_connect(
  master = "local",
  config = list(
   `sparklyr.shell.driver-class-path` = "/usr/share/java/mysql-connector-java-8.0.25.jar"
  )
)
spark_read_jdbc(
  sc,
  name = "my_sql_table",
  options = list(
    url = "jdbc:mysql://localhost:3306/my_sql_schema",
    driver = "com.mysql.jdbc.Driver",
    user = "me",
    password = "******",
    dbtable = "my_sql_table"
  )
)

## End(Not run)
```

---

spark_read_json   *Read a JSON file into a Spark DataFrame*

Description

Read a table serialized in the JavaScript Object Notation format into a Spark DataFrame.

Usage

```
spark_read_json(
  sc,
  name = NULL,
  path = name,
  options = list(),
  repartition = 0,
  memory = TRUE,
```

```
  overwrite = TRUE,
  columns = NULL,
  ...
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
options	A list of strings with additional options.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?
columns	A vector of column names or a named vector of column types. If specified, the elements can be "binary" for BinaryType, "boolean" for BooleanType, "byte" for ByteType, "integer" for IntegerType, "integer64" for LongType, "double" for DoubleType, "character" for StringType, "timestamp" for TimestampType and "date" for DateType.
...	Optional arguments; currently unused.

Details

You can read data from HDFS (hdfs://), S3 (s3a://), as well as the local file system (file://).

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_read_libsvm   *Read libsvm file into a Spark DataFrame.*

Description

Read libsvm file into a Spark DataFrame.

Usage

```
spark_read_libsvm(
  sc,
  name = NULL,
  path = name,
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE,
  options = list(),
  ...
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?
options	A list of strings with additional options.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_read_orc   *Read a ORC file into a Spark DataFrame*

Description

Read a ORC file into a Spark DataFrame.

Usage

```
spark_read_orc(
  sc,
  name = NULL,
  path = name,
  options = list(),
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE,
  columns = NULL,
  schema = NULL,
  ...
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
options	A list of strings with additional options. See https://spark.apache.org/docs/latest/sql-programming-guide.html#configuration.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?
columns	A vector of column names or a named vector of column types. If specified, the elements can be "binary" for BinaryType, "boolean" for BooleanType, "byte" for ByteType, "integer" for IntegerType, "integer64" for LongType, "double" for DoubleType, "character" for StringType, "timestamp" for TimestampType and "date" for DateType.
schema	A (java) read schema. Useful for optimizing read operation on nested data.
...	Optional arguments; currently unused.

Details

You can read data from HDFS (hdfs://), S3 (s3a://), as well as the local file system (file://).

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_read_parquet   *Read a Parquet file into a Spark DataFrame*

Description

Read a Parquet file into a Spark DataFrame.

Usage

```
spark_read_parquet(
  sc,
  name = NULL,
  path = name,
  options = list(),
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE,
  columns = NULL,
  schema = NULL,
  ...
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
options	A list of strings with additional options. See https://spark.apache.org/docs/latest/sql-programming-guide.html#configuration.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?
columns	A vector of column names or a named vector of column types. If specified, the elements can be "binary" for BinaryType, "boolean" for BooleanType, "byte" for ByteType, "integer" for IntegerType, "integer64" for LongType, "double" for DoubleType, "character" for StringType, "timestamp" for TimestampType and "date" for DateType.
schema	A (java) read schema. Useful for optimizing read operation on nested data.
...	Optional arguments; currently unused.

Details

You can read data from HDFS (`hdfs://`), S3 (`s3a://`), as well as the local file system (`file://`).

See Also

Other Spark serialization routines: [collect_from_rds()](), [spark_insert_table()](), [spark_load_table()](),
[spark_read()](), [spark_read_avro()](), [spark_read_binary()](), [spark_read_csv()](), [spark_read_delta()](),
[spark_read_image()](), [spark_read_jdbc()](), [spark_read_json()](), [spark_read_libsvm()](), [spark_read_orc()](),
[spark_read_source()](), [spark_read_table()](), [spark_read_text()](), [spark_save_table()](), [spark_write_avro()](),
[spark_write_csv()](), [spark_write_delta()](), [spark_write_jdbc()](), [spark_write_json()](), [spark_write_orc()](),
[spark_write_parquet()](), [spark_write_source()](), [spark_write_table()](), [spark_write_text()]()

| spark_read_source | *Read from a generic source into a Spark DataFrame.* |

Description

Read from a generic source into a Spark DataFrame.

Usage

```
spark_read_source(
  sc,
  name = NULL,
  path = name,
  source,
  options = list(),
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE,
  columns = NULL,
  ...
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
source	A data source capable of reading data.
options	A list of strings with additional options. See https://spark.apache.org/docs/latest/sql-programming-guide.html#configuration.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.

memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?
columns	A vector of column names or a named vector of column types. If specified, the elements can be ″binary″ for BinaryType, ″boolean″ for BooleanType, ″byte″ for ByteType, ″integer″ for IntegerType, ″integer64″ for LongType, ″double″ for DoubleType, ″character″ for StringType, ″timestamp″ for TimestampType and ″date″ for DateType.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: [collect_from_rds](), [spark_insert_table](), [spark_load_table](),
[spark_read](), [spark_read_avro](), [spark_read_binary](), [spark_read_csv](), [spark_read_delta](),
[spark_read_image](), [spark_read_jdbc](), [spark_read_json](), [spark_read_libsvm](), [spark_read_orc](),
[spark_read_parquet](), [spark_read_table](), [spark_read_text](), [spark_save_table](), [spark_write_avro](),
[spark_write_csv](), [spark_write_delta](), [spark_write_jdbc](), [spark_write_json](), [spark_write_orc](),
[spark_write_parquet](), [spark_write_source](), [spark_write_table](), [spark_write_text]()

---

spark_read_table	*Reads from a Spark Table into a Spark DataFrame.*

Description

Reads from a Spark Table into a Spark DataFrame.

Usage

```
spark_read_table(
  sc,
  name,
  options = list(),
  repartition = 0,
  memory = TRUE,
  columns = NULL,
  ...
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
options	A list of strings with additional options. See https://spark.apache.org/docs/latest/sql-programming-guide.html#configuration.
repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.

memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
columns	A vector of column names or a named vector of column types. If specified, the elements can be ″binary″ for BinaryType, ″boolean″ for BooleanType, ″byte″ for ByteType, ″integer″ for IntegerType, ″integer64″ for LongType, ″double″ for DoubleType, ″character″ for StringType, ″timestamp″ for TimestampType and ″date″ for DateType.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: [collect_from_rds](), [spark_insert_table](), [spark_load_table](),
[spark_read](), [spark_read_avro](), [spark_read_binary](), [spark_read_csv](), [spark_read_delta](),
[spark_read_image](), [spark_read_jdbc](), [spark_read_json](), [spark_read_libsvm](), [spark_read_orc](),
[spark_read_parquet](), [spark_read_source](), [spark_read_text](), [spark_save_table](), [spark_write_avro](),
[spark_write_csv](), [spark_write_delta](), [spark_write_jdbc](), [spark_write_json](), [spark_write_orc](),
[spark_write_parquet](), [spark_write_source](), [spark_write_table](), [spark_write_text]()

---

spark_read_text	*Read a Text file into a Spark DataFrame*

Description

Read a Text file into a Spark DataFrame

Usage

```
spark_read_text(
  sc,
  name = NULL,
  path = name,
  repartition = 0,
  memory = TRUE,
  overwrite = TRUE,
  options = list(),
  whole = FALSE,
  ...
)
```

Arguments

sc	A spark_connection.
name	The name to assign to the newly generated table.
path	The path to the file. Needs to be accessible from the cluster. Supports the '″hdfs://″', '″s3a://″' and '″file://″' protocols.

repartition	The number of partitions used to distribute the generated table. Use 0 (the default) to avoid partitioning.
memory	Boolean; should the data be loaded eagerly into memory? (That is, should the table be cached?)
overwrite	Boolean; overwrite the table with the given name if it already exists?
options	A list of strings with additional options.
whole	Read the entire text file as a single entry? Defaults to `FALSE`.
...	Optional arguments; currently unused.

Details

You can read data from HDFS (`hdfs://`), S3 (`s3a://`), as well as the local file system (`file://`).

See Also

Other Spark serialization routines: `collect_from_rds()`, `spark_insert_table()`, `spark_load_table()`, `spark_read()`, `spark_read_avro()`, `spark_read_binary()`, `spark_read_csv()`, `spark_read_delta()`, `spark_read_image()`, `spark_read_jdbc()`, `spark_read_json()`, `spark_read_libsvm()`, `spark_read_orc()`, `spark_read_parquet()`, `spark_read_source()`, `spark_read_table()`, `spark_save_table()`, `spark_write_avro()`, `spark_write_csv()`, `spark_write_delta()`, `spark_write_jdbc()`, `spark_write_json()`, `spark_write_orc()`, `spark_write_parquet()`, `spark_write_source()`, `spark_write_table()`, `spark_write_text()`

---

spark_save_table	*Saves a Spark DataFrame as a Spark table*

Description

Saves a Spark DataFrame and as a Spark table.

Usage

```
spark_save_table(x, path, mode = NULL, options = list())
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
mode	A `character` element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also `https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes` for your version of Spark.
options	A list of strings with additional options.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(),
spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(),
spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(),
spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_write_avro(),
spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(),
spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_session_config   *Runtime configuration interface for the Spark Session*

Description

Retrieves or sets runtime configuration entries for the Spark Session

Usage

```
spark_session_config(sc, config = TRUE, value = NULL)
```

Arguments

sc	A spark_connection.
config	The configuration entry name(s) (e.g., "spark.sql.shuffle.partitions"). Defaults to NULL to retrieve all configuration entries.
value	The configuration value to be set. Defaults to NULL to retrieve configuration entries.

See Also

Other Spark runtime configuration: spark_adaptive_query_execution(), spark_advisory_shuffle_partition_size(
spark_auto_broadcast_join_threshold(), spark_coalesce_initial_num_partitions(), spark_coalesce_min_num
spark_coalesce_shuffle_partitions()

---

spark_statistical_routines

Generate random samples from some distribution

Description

Generator methods for creating single-column Spark dataframes comprised of i.i.d. samples from
some distribution.

Arguments

sc	A Spark connection.
n	Sample Size (default: 1000).
num_partitions	Number of partitions in the resulting Spark dataframe (default: default parallelism of the Spark cluster).
seed	Random seed (default: a random long integer).
output_col	Name of the output column containing sample values (default: "x").

---

spark_table_name   *Generate a Table Name from Expression*

Description

Attempts to generate a table name from an expression; otherwise, assigns an auto-generated generic name with "sparklyr_" prefix.

Usage

```
spark_table_name(expr)
```

Arguments

expr	The expression to attempt to use as name

---

spark_version   *Get the Spark Version Associated with a Spark Connection*

Description

Retrieve the version of Spark associated with a Spark connection.

Usage

```
spark_version(sc)
```

Arguments

sc	A spark_connection.

Details

Suffixes for e.g. preview versions, or snapshotted versions, are trimmed – if you require the full Spark version, you can retrieve it with invoke(spark_context(sc), "version").

Value

The Spark version as a numeric_version.

---

spark_version_from_home

Get the Spark Version Associated with a Spark Installation

Description

Retrieve the version of Spark associated with a Spark installation.

Usage

```
spark_version_from_home(spark_home, default = NULL)
```

Arguments

spark_home   The path to a Spark installation.

default   The default version to be inferred, in case version lookup failed, e.g. no Spark installation was found at `spark_home`.

---

spark_web   *Open the Spark web interface*

Description

Open the Spark web interface

Usage

```
spark_web(sc, ...)
```

Arguments

sc   A `spark_connection`.

...   Optional arguments; currently unused.

---

spark_write   *Write Spark DataFrame to file using a custom writer*

Description

Run a custom R function on Spark worker to write a Spark DataFrame into file(s). If Spark's
speculative execution feature is enabled (i.e., 'spark.speculation' is true), then each write task may
be executed more than once and the user-defined writer function will need to ensure no concurrent
writes happen to the same file path (e.g., by appending UUID to each file name).

Usage

```
spark_write(x, writer, paths, packages = NULL)
```

Arguments

x	A Spark Dataframe to be saved into file(s)
writer	A writer function with the signature function(partition, path) where `partition` is a R dataframe containing all rows from one partition of the original Spark Dataframe x and path is a string specifying the file to write `partition` to
paths	A single destination path or a list of destination paths, each one specifying a location for a partition from x to be written to. If number of partition(s) in x is not equal to `length(paths)` then x will be re-partitioned to contain `length(paths)` partition(s)
packages	Boolean to distribute `.libPaths()` packages to each node, a list of packages to distribute, or a package bundle created with

Examples

```
## Not run:

library(sparklyr)

sc <- spark_connect(master = "local[3]")

# copy some test data into a Spark Dataframe
sdf <- sdf_copy_to(sc, iris, overwrite = TRUE)

# create a writer function
writer <- function(df, path) {
  write.csv(df, path)
}

spark_write(
  sdf,
  writer,
  # re-partition sdf into 3 partitions and write them to 3 separate files
  paths = list("file:///tmp/file1", "file:///tmp/file2", "file:///tmp/file3"),
```

```
)

spark_write(
  sdf,
  writer,
  # save all rows into a single file
  paths = list("file:///tmp/all_rows")
)

## End(Not run)
```

---

spark_write_avro   *Serialize a Spark DataFrame into Apache Avro format*

Description

Notice this functionality requires the Spark connection sc to be instantiated with either an explicitly specified Spark version (i.e., spark_connect(..., version = <version>, packages = c("avro", <other package(s)>), ...)) or a specific version of Spark avro package to use (e.g., spark_connect(..., packages = c("org.apache.spark:spark-avro_2.12:3.0.0", <other package(s)>), ...)).

Usage

```
spark_write_avro(
  x,
  path,
  avro_schema = NULL,
  record_name = "topLevelRecord",
  record_namespace = "",
  compression = "snappy",
  partition_by = NULL
)
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
avro_schema	Optional Avro schema in JSON format
record_name	Optional top level record name in write result (default: "topLevelRecord")
record_namespace	
	Record namespace in write result (default: "")
compression	Compression codec to use (default: "snappy")
partition_by	A character vector. Partitions the output by the given columns on the file system.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(),
spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(),
spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(),
spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(),
spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(),
spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_write_csv   *Write a Spark DataFrame to a CSV*

Description

Write a Spark DataFrame to a tabular (typically, comma-separated) file.

Usage

```
spark_write_csv(
  x,
  path,
  header = TRUE,
  delimiter = ",",
  quote = "\"",
  escape = "\\",
  charset = "UTF-8",
  null_value = NULL,
  options = list(),
  mode = NULL,
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
header	Should the first row of data be used as a header? Defaults to TRUE.
delimiter	The character used to delimit each column, defaults to ,.
quote	The character used as a quote. Defaults to '' " ''.
escape	The character used to escape other characters, defaults to \.
charset	The character set, defaults to "UTF-8".
null_value	The character to use for default values, defaults to NULL.
options	A list of strings with additional options.

mode	A character element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes for your version of Spark.
partition_by	A character vector. Partitions the output by the given columns on the file system.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: `collect_from_rds()`, `spark_insert_table()`, `spark_load_table()`, `spark_read()`, `spark_read_avro()`, `spark_read_binary()`, `spark_read_csv()`, `spark_read_delta()`, `spark_read_image()`, `spark_read_jdbc()`, `spark_read_json()`, `spark_read_libsvm()`, `spark_read_orc()`, `spark_read_parquet()`, `spark_read_source()`, `spark_read_table()`, `spark_read_text()`, `spark_save_table()`, `spark_write_avro()`, `spark_write_delta()`, `spark_write_jdbc()`, `spark_write_json()`, `spark_write_orc()`, `spark_write_parquet()`, `spark_write_source()`, `spark_write_table()`, `spark_write_text()`

| spark_write_delta | *Writes a Spark DataFrame into Delta Lake* |

Description

Writes a Spark DataFrame into Delta Lake.

Usage

```
spark_write_delta(
  x,
  path,
  mode = NULL,
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
mode	A character element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes for your version of Spark.

options	A list of strings with additional options.
partition_by	A character vector. Partitions the output by the given columns on the file system.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_write_jdbc	*Writes a Spark DataFrame into a JDBC table*

Description

Writes a Spark DataFrame into a JDBC table

Usage

```
spark_write_jdbc(
  x,
  name,
  mode = NULL,
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
name	The name to assign to the newly generated table.
mode	A character element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes for your version of Spark.
options	A list of strings with additional options.
partition_by	A character vector. Partitions the output by the given columns on the file system.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(),
spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(),
spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(),
spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(),
spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_json(), spark_write_orc(),
spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

Examples

```
## Not run:
sc <- spark_connect(
  master = "local",
  config = list(
   `sparklyr.shell.driver-class-path` = "/usr/share/java/mysql-connector-java-8.0.25.jar"
  )
)
spark_write_jdbc(
  sdf_len(sc, 10),
  name = "my_sql_table",
  options = list(
    url = "jdbc:mysql://localhost:3306/my_sql_schema",
    driver = "com.mysql.jdbc.Driver",
    user = "me",
    password = "******",
    dbtable = "my_sql_table"
  )
)

## End(Not run)
```

| spark_write_json | *Write a Spark DataFrame to a JSON file* |

Description

Serialize a Spark DataFrame to the JavaScript Object Notation format.

Usage

```
spark_write_json(
  x,
  path,
  mode = NULL,
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
mode	A `character` element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes for your version of Spark.
options	A list of strings with additional options.
partition_by	A `character` vector. Partitions the output by the given columns on the file system.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: [collect_from_rds()](), [spark_insert_table()](), [spark_load_table()](), [spark_read()](), [spark_read_avro()](), [spark_read_binary()](), [spark_read_csv()](), [spark_read_delta()](), [spark_read_image()](), [spark_read_jdbc()](), [spark_read_json()](), [spark_read_libsvm()](), [spark_read_orc()](), [spark_read_parquet()](), [spark_read_source()](), [spark_read_table()](), [spark_read_text()](), [spark_save_table()](), [spark_write_avro()](), [spark_write_csv()](), [spark_write_delta()](), [spark_write_jdbc()](), [spark_write_orc()](), [spark_write_parquet()](), [spark_write_source()](), [spark_write_table()](), [spark_write_text()]()

---

spark_write_orc   *Write a Spark DataFrame to a ORC file*

Description

Serialize a Spark DataFrame to the [ORC]() format.

Usage

```
spark_write_orc(
  x,
  path,
  mode = NULL,
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
mode	A character element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure. For more details see also https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes for your version of Spark.
options	A list of strings with additional options. See https://spark.apache.org/docs/latest/sql-programming-guide.html#configuration.
partition_by	A character vector. Partitions the output by the given columns on the file system.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_parquet(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_write_parquet   *Write a Spark DataFrame to a Parquet file*

Description

Serialize a Spark DataFrame to the Parquet format.

Usage

```
spark_write_parquet(
  x,
  path,
  mode = NULL,
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
mode	A `character` element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes for your version of Spark.
options	A list of strings with additional options. See https://spark.apache.org/docs/latest/sql-programming-guide.html#configuration.
partition_by	A `character` vector. Partitions the output by the given columns on the file system.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_source(), spark_write_table(), spark_write_text()

---

spark_write_rds   *Write Spark DataFrame to RDS files*

Description

Write Spark dataframe to RDS files. Each partition of the dataframe will be exported to a separate RDS file so that all partitions can be processed in parallel.

Usage

```
spark_write_rds(x, dest_uri)
```

Arguments

x	A Spark DataFrame to be exported
dest_uri	Can be a URI template containing 'partitionId' (e.g., "hdfs://my_data_part_{partitionId}.rds") where 'partitionId' will be substituted with ID of each partition using 'glue', or a list of URIs to be assigned to RDS output from all partitions (e.g., "hdfs://my_data_part_0.rds", "hdfs://my_data_part_1.rds", and so on) If working with a Spark instance running locally, then all URIs should be in "file://<local file path>" form. Otherwise the scheme of the URI should reflect the underlying file system the

                    Spark instance is working with (e.g., "hdfs://"). If the resulting list of URI(s) does not contain unique values, then it will be post-processed with 'make.unique()' to ensure uniqueness.

Value

A tibble containing partition ID and RDS file location for each partition of the input Spark dataframe.

---

spark_write_source   *Writes a Spark DataFrame into a generic source*

Description

Writes a Spark DataFrame into a generic source.

Usage

```
spark_write_source(
  x,
  source,
  mode = NULL,
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
source	A data source capable of reading data.
mode	A `character` element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes for your version of Spark.
options	A list of strings with additional options.
partition_by	A `character` vector. Partitions the output by the given columns on the file system.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_table(), spark_write_text()

| spark_write_table | *Writes a Spark DataFrame into a Spark table* |

Description

Writes a Spark DataFrame into a Spark table

Usage

```
spark_write_table(
  x,
  name,
  mode = NULL,
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
name	The name to assign to the newly generated table.
mode	A `character` element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes for your version of Spark.
options	A list of strings with additional options.
partition_by	A `character` vector. Partitions the output by the given columns on the file system.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_text()

---

spark_write_text   *Write a Spark DataFrame to a Text file*

Description

Serialize a Spark DataFrame to the plain text format.

Usage

```
spark_write_text(
  x,
  path,
  mode = NULL,
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
mode	A character element. Specifies the behavior when data or table already exists. Supported values include: 'error', 'append', 'overwrite' and ignore. Notice that 'overwrite' will also change the column structure.
	For more details see also https://spark.apache.org/docs/latest/sql-programming-guide.html#save-modes for your version of Spark.
options	A list of strings with additional options.
partition_by	A character vector. Partitions the output by the given columns on the file system.
...	Optional arguments; currently unused.

See Also

Other Spark serialization routines: collect_from_rds(), spark_insert_table(), spark_load_table(), spark_read(), spark_read_avro(), spark_read_binary(), spark_read_csv(), spark_read_delta(), spark_read_image(), spark_read_jdbc(), spark_read_json(), spark_read_libsvm(), spark_read_orc(), spark_read_parquet(), spark_read_source(), spark_read_table(), spark_read_text(), spark_save_table(), spark_write_avro(), spark_write_csv(), spark_write_delta(), spark_write_jdbc(), spark_write_json(), spark_write_orc(), spark_write_parquet(), spark_write_source(), spark_write_table()

---

src_databases   *Show database list*

Description

Show database list

Usage

```
src_databases(sc, ...)
```

Arguments

sc	A spark_connection.
...	Optional arguments; currently unused.

---

stream_find   *Find Stream*

Description

Finds and returns a stream based on the stream's identifier.

Usage

```
stream_find(sc, id)
```

Arguments

sc	The associated Spark connection.
id	The stream identifier to find.

Examples

```
## Not run:
sc <- spark_connect(master = "local")
sdf_len(sc, 10) %>%
  spark_write_parquet(path = "parquet-in")

stream <- stream_read_parquet(sc, "parquet-in") %>%
  stream_write_parquet("parquet-out")

stream_id <- stream_id(stream)
stream_find(sc, stream_id)

## End(Not run)
```

---

stream_generate_test   *Generate Test Stream*

Description

Generates a local test stream, useful when testing streams locally.

Usage

```
stream_generate_test(
  df = rep(1:1000),
  path = "source",
  distribution = floor(10 + 1e+05 * stats::dbinom(1:20, 20, 0.5)),
  iterations = 50,
  interval = 1
)
```

Arguments

df	The data frame used as a source of rows to the stream, will be cast to data frame if needed. Defaults to a sequence of one thousand entries.
path	Path to save stream of files to, defaults to "source".
distribution	The distribution of rows to use over each iteration, defaults to a binomial distribution. The stream will cycle through the distribution if needed.
iterations	Number of iterations to execute before stopping, defaults to fifty.
interval	The inverval in seconds use to write the stream, defaults to one second.

Details

This function requires the `callr` package to be installed.

---

stream_id   *Spark Stream's Identifier*

Description

Retrieves the identifier of the Spark stream.

Usage

```
stream_id(stream)
```

Arguments

stream	The spark stream object.

---

stream_lag   *Apply lag function to columns of a Spark Streaming DataFrame*

Description

Given a streaming Spark dataframe as input, this function will return another streaming dataframe
that contains all columns in the input and column(s) that are shifted behind by the offset(s) specified
in '...' (see example)

Usage

```
stream_lag(x, cols, thresholds = NULL)
```

Arguments

x	An object coercable to a Spark Streaming DataFrame.
cols	A list of expressions for a single or multiple variables to create that will contain the value of a previous entry.
thresholds	Optional named list of timestamp column(s) and corresponding time duration(s) for deterimining whether a previous record is sufficiently recent relative to the current record. If the any of the time difference(s) between the current and a previous record is greater than the maximal duration allowed, then the previous record is discarded and will not be part of the query result. The durations can be specified with numeric types (which will be interpreted as max difference allowed in number of milliseconds between 2 UNIX timestamps) or time duration strings such as "5s", "5sec", "5min", "5hour", etc. Any timestamp column in 'x' that is not of timestamp of date Spark SQL types will be interepreted as number of milliseconds since the UNIX epoch.

Examples

```
## Not run:

library(sparklyr)

sc <- spark_connect(master = "local", version = "2.2.0")

streaming_path <- tempfile("days_df_")
days_df <- dplyr::tibble(
  today = weekdays(as.Date(seq(7), origin = "1970-01-01"))
)
num_iters <- 7
stream_generate_test(
  df = days_df,
  path = streaming_path,
  distribution = rep(nrow(days_df), num_iters),
  iterations = num_iters
)
```

```
stream_read_csv(sc, streaming_path) %>%
  stream_lag(cols = c(yesterday = today ~ 1, two_days_ago = today ~ 2)) %>%
  collect() %>%
  print(n = 10L)

## End(Not run)
```

---

stream_name · *Spark Stream's Name*

Description

Retrieves the name of the Spark stream if available.

Usage

```
stream_name(stream)
```

Arguments

stream   The spark stream object.

---

stream_read_csv · *Read files created by the stream*

Description

Read files created by the stream

Usage

```
stream_read_csv(
  sc,
  path,
  name = NULL,
  header = TRUE,
  columns = NULL,
  delimiter = ",",
  quote = "\"",
  escape = "\\",
  charset = "UTF-8",
  null_value = NULL,
  options = list(),
  ...
```

```
)

stream_read_text(sc, path, name = NULL, options = list(), ...)

stream_read_json(sc, path, name = NULL, columns = NULL, options = list(), ...)

stream_read_parquet(
  sc,
  path,
  name = NULL,
  columns = NULL,
  options = list(),
  ...
)

stream_read_orc(sc, path, name = NULL, columns = NULL, options = list(), ...)

stream_read_kafka(sc, name = NULL, options = list(), ...)

stream_read_socket(sc, name = NULL, columns = NULL, options = list(), ...)

stream_read_delta(sc, path, name = NULL, options = list(), ...)

stream_read_cloudfiles(sc, path, name = NULL, options = list(), ...)

stream_read_table(sc, path, name = NULL, options = list(), ...)
```

Arguments

`sc`	A `spark_connection`.
`path`	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
`name`	The name to assign to the newly generated stream.
`header`	Boolean; should the first row of data be used as a header? Defaults to `TRUE`.
`columns`	A vector of column names or a named vector of column types. If specified, the elements can be `"binary"` for BinaryType, `"boolean"` for BooleanType, `"byte"` for ByteType, `"integer"` for IntegerType, `"integer64"` for LongType, `"double"` for DoubleType, `"character"` for StringType, `"timestamp"` for TimestampType and `"date"` for DateType.
`delimiter`	The character used to delimit each column. Defaults to '`,`'.
`quote`	The character used as a quote. Defaults to '`"`'.
`escape`	The character used to escape other characters. Defaults to '`\`'.
`charset`	The character set. Defaults to '`"UTF-8"`'.
`null_value`	The character to use for null, or missing, values. Defaults to `NULL`.
`options`	A list of strings with additional options.
`...`	Optional arguments; currently unused.

Examples

```
## Not run:

sc <- spark_connect(master = "local")

dir.create("csv-in")
write.csv(iris, "csv-in/data.csv", row.names = FALSE)

csv_path <- file.path("file://", getwd(), "csv-in")

stream <- stream_read_csv(sc, csv_path) %>% stream_write_csv("csv-out")

stream_stop(stream)

## End(Not run)
```

---

stream_render   *Render Stream*

Description

Collects streaming statistics to render the stream as an 'htmlwidget'.

Usage

```
stream_render(stream = NULL, collect = 10, stats = NULL, ...)
```

Arguments

stream	The stream to render
collect	The interval in seconds to collect data before rendering the 'htmlwidget'.
stats	Optional stream statistics collected using stream_stats(), when specified, stream should be omitted.
...	Additional optional arguments.

Examples

```
## Not run:
library(sparklyr)
sc <- spark_connect(master = "local")

dir.create("iris-in")
write.csv(iris, "iris-in/iris.csv", row.names = FALSE)

stream <- stream_read_csv(sc, "iris-in/") %>%
  stream_write_csv("iris-out/")
```

```
stream_render(stream)
stream_stop(stream)

## End(Not run)
```

---

stream_stats   *Stream Statistics*

Description

Collects streaming statistics, usually, to be used with stream_render() to render streaming statistics.

Usage

```
stream_stats(stream, stats = list())
```

Arguments

stream	The stream to collect statistics from.
stats	An optional stats object generated using stream_stats().

Value

A stats object containing streaming statistics that can be passed back to the stats parameter to continue aggregating streaming stats.

Examples

```
## Not run:
sc <- spark_connect(master = "local")
sdf_len(sc, 10) %>%
  spark_write_parquet(path = "parquet-in")

stream <- stream_read_parquet(sc, "parquet-in") %>%
  stream_write_parquet("parquet-out")

stream_stats(stream)

## End(Not run)
```

---

stream_stop *Stops a Spark Stream*

Description

Stops processing data from a Spark stream.

Usage

```
stream_stop(stream)
```

Arguments

stream   The spark stream object to be stopped.

---

stream_trigger_continuous

Spark Stream Continuous Trigger

Description

Creates a Spark structured streaming trigger to execute continuously. This mode is the most performant but not all operations are supported.

Usage

```
stream_trigger_continuous(checkpoint = 5000)
```

Arguments

checkpoint   The checkpoint interval specified in milliseconds.

See Also

[stream_trigger_interval](#)

---

stream_trigger_interval

Spark Stream Interval Trigger

Description

Creates a Spark structured streaming trigger to execute over the specified interval.

Usage

```
stream_trigger_interval(interval = 1000)
```

Arguments

interval   The execution interval specified in milliseconds.

See Also

stream_trigger_continuous

---

stream_view   *View Stream*

Description

Opens a Shiny gadget to visualize the given stream.

Usage

```
stream_view(stream, ...)
```

Arguments

stream   The stream to visualize.
...   Additional optional arguments.

Examples

```
## Not run:
library(sparklyr)
sc <- spark_connect(master = "local")

dir.create("iris-in")
write.csv(iris, "iris-in/iris.csv", row.names = FALSE)

stream_read_csv(sc, "iris-in/") %>%
```

```
    stream_write_csv("iris-out/") %>%
    stream_view() %>%
    stream_stop()

## End(Not run)
```

`stream_watermark`   *Watermark Stream*

Description

Ensures a stream has a watermark defined, which is required for some operations over streams.

Usage

```
stream_watermark(x, column = "timestamp", threshold = "10 minutes")
```

Arguments

x	An object coercable to a Spark Streaming DataFrame.
column	The name of the column that contains the event time of the row, if the column is missing, a column with the current time will be added.
threshold	The minimum delay to wait to data to arrive late, defaults to ten minutes.

`stream_write_csv`   *Write files to the stream*

Description

Write files to the stream

Usage

```
stream_write_csv(
  x,
  path,
  mode = c("append", "complete", "update"),
  trigger = stream_trigger_interval(),
  checkpoint = file.path(path, "checkpoint"),
  header = TRUE,
  delimiter = ",",
  quote = "\"",
  escape = "\\",
  charset = "UTF-8",
  null_value = NULL,
```

```
  options = list(),
  partition_by = NULL,
  ...
)

stream_write_text(
  x,
  path,
  mode = c("append", "complete", "update"),
  trigger = stream_trigger_interval(),
  checkpoint = file.path(path, "checkpoints", random_string("")),
  options = list(),
  partition_by = NULL,
  ...
)

stream_write_json(
  x,
  path,
  mode = c("append", "complete", "update"),
  trigger = stream_trigger_interval(),
  checkpoint = file.path(path, "checkpoints", random_string("")),
  options = list(),
  partition_by = NULL,
  ...
)

stream_write_parquet(
  x,
  path,
  mode = c("append", "complete", "update"),
  trigger = stream_trigger_interval(),
  checkpoint = file.path(path, "checkpoints", random_string("")),
  options = list(),
  partition_by = NULL,
  ...
)

stream_write_orc(
  x,
  path,
  mode = c("append", "complete", "update"),
  trigger = stream_trigger_interval(),
  checkpoint = file.path(path, "checkpoints", random_string("")),
  options = list(),
  partition_by = NULL,
  ...
)
```

```
stream_write_kafka(
  x,
  mode = c("append", "complete", "update"),
  trigger = stream_trigger_interval(),
  checkpoint = file.path("checkpoints", random_string("")),
  options = list(),
  partition_by = NULL,
  ...
)

stream_write_console(
  x,
  mode = c("append", "complete", "update"),
  options = list(),
  trigger = stream_trigger_interval(),
  partition_by = NULL,
  ...
)

stream_write_delta(
  x,
  path,
  mode = c("append", "complete", "update"),
  checkpoint = file.path("checkpoints", random_string("")),
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '"hdfs://"', '"s3a://"' and '"file://"' protocols.
mode	Specifies how data is written to a streaming sink. Valid values are "append", "complete" or "update".
trigger	The trigger for the stream query, defaults to micro-batches running every 5 seconds. See stream_trigger_interval and stream_trigger_continuous.
checkpoint	The location where the system will write all the checkpoint information to guarantee end-to-end fault-tolerance.
header	Should the first row of data be used as a header? Defaults to TRUE.
delimiter	The character used to delimit each column, defaults to ,.
quote	The character used as a quote. Defaults to '"'.
escape	The character used to escape other characters, defaults to \.
charset	The character set, defaults to "UTF-8".

`null_value`	The character to use for default values, defaults to NULL.
`options`	A list of strings with additional options.
`partition_by`	Partitions the output by the given list of columns.
`...`	Optional arguments; currently unused.

See Also

Other Spark stream serialization: `stream_write_memory()`, `stream_write_table()`

Examples

```
## Not run:

sc <- spark_connect(master = "local")

dir.create("csv-in")
write.csv(iris, "csv-in/data.csv", row.names = FALSE)

csv_path <- file.path("file://", getwd(), "csv-in")

stream <- stream_read_csv(sc, csv_path) %>% stream_write_csv("csv-out")

stream_stop(stream)

## End(Not run)
```

---

stream_write_memory   *Write Memory Stream*

Description

Writes a Spark dataframe stream into a memory stream.

Usage

```
stream_write_memory(
  x,
  name = random_string("sparklyr_tmp_"),
  mode = c("append", "complete", "update"),
  trigger = stream_trigger_interval(),
  checkpoint = file.path("checkpoints", name, random_string("")),
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
name	The name to assign to the newly generated stream.
mode	Specifies how data is written to a streaming sink. Valid values are `"append"`, `"complete"` or `"update"`.
trigger	The trigger for the stream query, defaults to micro-batches running every 5 seconds. See `stream_trigger_interval` and `stream_trigger_continuous`.
checkpoint	The location where the system will write all the checkpoint information to guarantee end-to-end fault-tolerance.
options	A list of strings with additional options.
partition_by	Partitions the output by the given list of columns.
...	Optional arguments; currently unused.

See Also

Other Spark stream serialization: `stream_write_csv()`, `stream_write_table()`

---

stream_write_table	*Write Stream to Table*

Description

Writes a Spark dataframe stream into a table.

Usage

```
stream_write_table(
  x,
  path,
  format = NULL,
  mode = c("append", "complete", "update"),
  checkpoint = file.path("checkpoints", random_string("")),
  options = list(),
  partition_by = NULL,
  ...
)
```

Arguments

x	A Spark DataFrame or dplyr operation
path	The path to the file. Needs to be accessible from the cluster. Supports the '`"hdfs://"`', '`"s3a://"`' and '`"file://"`' protocols.
format	Specifies format of data written to table E.g. `"delta"`, `"parquet"`. Defaults to `NULL` which will use system default format.

mode	Specifies how data is written to a streaming sink. Valid values are `"append"`, `"complete"` or `"update"`.
checkpoint	The location where the system will write all the checkpoint information to guarantee end-to-end fault-tolerance.
options	A list of strings with additional options.
partition_by	Partitions the output by the given list of columns.
...	Optional arguments; currently unused.

See Also

Other Spark stream serialization: `stream_write_csv()`, `stream_write_memory()`

---

tbl_cache   *Cache a Spark Table*

Description

Force a Spark table with name `name` to be loaded into memory. Operations on cached tables should normally (although not always) be more performant than the same operation performed on an uncached table.

Usage

```
tbl_cache(sc, name, force = TRUE)
```

Arguments

sc	A `spark_connection`.
name	The table name.
force	Force the data to be loaded into memory? This is accomplished by calling the `count` API on the associated Spark DataFrame.

---

tbl_change_db   *Use specific database*

Description

Use specific database

Usage

```
tbl_change_db(sc, name)
```

Arguments

sc	A `spark_connection`.
name	The database name.

---

tbl_uncache *Uncache a Spark Table*

Description

Force a Spark table with name name to be unloaded from memory.

Usage

```
tbl_uncache(sc, name)
```

Arguments

sc	A spark_connection.
name	The table name.

---

transform_sdf *transform a subset of column(s) in a Spark Dataframe*

Description

transform a subset of column(s) in a Spark Dataframe

Usage

```
transform_sdf(x, cols, fn)
```

Arguments

x	An object coercible to a Spark DataFrame
cols	Subset of columns to apply transformation to
fn	Transformation function taking column name as the 1st parameter, the corresponding org.apache.spark.sql.Column object as the 2nd parameter, and returning a transformed org.apache.spark.sql.Column object

---

unite *Unite*

Description

See [unite](#unite) for more details.

---

unnest *Unnest*

Description

See unnest for more details.

[.tbl_spark *Subsetting operator for Spark dataframe*

Description

Susetting operator for Spark dataframe allowing a subset of column(s) to be selected using syntaxes similar to those supported by R dataframes

Usage

```
## S3 method for class 'tbl_spark'
x[i]
```

Arguments

x	The Spark dataframe
i	Expression specifying subset of column(s) to include or exclude from the result (e.g., '["col1"]', '[c("col1", "col2")]', '[1:10]', '[-1]', '[NULL]', or '[]')

%->% *Infix operator for composing a lambda expression*

Description

Infix operator that allows a lambda expression to be composed in R and be translated to Spark SQL equivalent using ' dbplyr::translate_sql functionalities

Usage

```
params %->% ...
```

Arguments

params	Parameter(s) of the lambda expression, can be either a single parameter or a comma separated listed of parameters in the form of .(param1, param2, ...) (see examples)
...	Body of the lambda expression, *must be within parentheses*

Details

Notice when composing a lambda expression in R, the body of the lambda expression *must always be surrounded with parentheses*, otherwise a parsing error will occur.

Examples

```
## Not run:

a %->% (mean(a) + 1) # translates to <SQL> `a` -> (AVG(`a`) OVER () + 1.0)

.(a, b) %->% (a < 1 && b > 1) # translates to <SQL> `a`,`b` -> (`a` < 1.0 AND `b` > 1.0)

## End(Not run)
```

Index